

Welcome to Programming for Data Science

Welcome to the course manual for CSC310 at URI with Professor Brown.

This class meets MWF 3-3:50pm in Chafee Social Sci Center 235.

This website will contain the syllabus, class notes and other reference material for the class.

[Course Calendar on BrightSpace](#)



Tip

[subscribe to that calendar](#) in your favorite calendar application

Basic Facts

About this course

Data science exists at the intersection of computer science, statistics, and machine learning. That means writing programs to access and manipulate data so that it becomes available for analysis using statistical and machine learning techniques is at the core of data science. Data scientists use their data and analytical ability to find and interpret rich data sources; manage large amounts of data despite hardware, software, and bandwidth constraints; merge data sources; ensure consistency of datasets; create visualizations to aid in understanding data; build mathematical models using the data; and present and communicate the data insights/findings.

This course provides a survey of data science. Topics include data driven programming in Python; data sets, file formats and meta-data; descriptive statistics, data visualization, and foundations of predictive data modeling and machine learning; accessing web data and databases; distributed data management. You will work on weekly substantial programming problems such as accessing data in database and visualize it or build machine learning models of a given data set.

Basic programming skills (CSC201 or CSC211) are a prerequisite to this course. This course is a prerequisite course to machine learning, where you learn how machine learning algorithms work. In this course, we will start with a very fast review of basic programming ideas, since you've already done that before. We will learn how to *use* machine learning algorithms to do data science, but not how to *build* machine learning algorithms, we'll use packages that implement the algorithms for us.

About this syllabus

This syllabus is a *living* document and accessible from BrightSpace, as a pdf for download directly online at rhodyprog4ds.github.io/BrownFall21/syllabus. If you choose to download a copy of it, note that it is only a copy. You can get notification of changes from GitHub by "watching" the You can view the date of changes and exactly what changes were made on the Github [commit history](#) page.

Creating an [issue](#) is also a good way to ask questions about anything in the course it will prompt additions and expand the FAQ section.

About your instructor

Name: Dr. Sarah M Brown Office hours: TBA via zoom, link on BrightSpace

Dr. Sarah M Brown is a second year Assistant Professor of Computer Science, who does research on how social context changes machine learning. Dr. Brown earned a PhD in Electrical Engineering from Northeastern University, completed a postdoctoral fellowship at University of California Berkeley, and worked as a postdoctoral research associate at Brown University before joining URI. At Brown University, Dr. Brown taught the Data and Society course for the Master's in Data Science Program. You can learn more about me at my [website](#) or my research on my [lab site](#).

You can call me Professor Brown or Dr. Brown, I use she/her pronouns.

The best way to contact me is e-mail or an issue on an assignment repo. For more details, see the [Communication Section](#)

Tools and Resources

We will use a variety of tools to conduct class and to facilitate your programming. You will need a computer with Linux, MacOS, or Windows. It is unlikely that a tablet will be able to do all of the things required in this course. A Chromebook may work, especially with developer tools turned on. Ask Dr. Brown if you need help getting access to an adequate computer.

All of the tools and resources below are either:

- paid for by URI **OR**
- freely available online.

BrightSpace

This will be the central location from which you can access all other materials. Any links that are for private discussion among those enrolled in the course will be available only from our course [Brightspace site](#).

This is also where your grades will appear and how I will post announcements.

For announcements, you can [customize](#) how you receive them.

! Important

TL;DR [\[1\]](#)

- check Brightspace
- Log in to Prismia Chat
- Make a GitHub Account
- Install Python
- Install Git

Prismia chat

Our class link for [Prismia chat](#) is available on Brightspace. We will use this for chatting and in-class understanding checks.

On Prismia, all students see the instructor's messages, but only the Instructor and TA see student responses.

i Note

Seeing the BrightSpace site requires logging in with your URI SSO and being enrolled in the course

Course Manual

The course manual will have content including the class policies, scheduling, class notes, assignment information, and additional resources. This will be linked from Brightspace and available publicly online at rhodyprog4ds.github.io/BrownFall21/. Links to the course reference text and code documentation will also be included here in the assignments and class notes.

GitHub Classroom

You will need a [GitHub](#) Account. If you do not already have one, please [create one](#) by the first day of class. If you have one, but have not used it recently, you may need to update your password and login credentials as the [Authentication rules](#) changed over the summer. In order to use the command line with https, you will need to [create a Personal Access Token](#) for each device you use. In order to use the command line with SSH, set up your public key.

Programming Environment

This is a programming course, so you will need a programming environment. In order to complete assignments you need the items listed in the requirements list. The easiest way to meet these requirements is to follow the recommendations below. I will provide instruction assuming that you have followed the recommendations.

Requirements:

- Python with scientific computing packages (numpy, scipy, jupyter, pandas, seaborn, sklearn)
- [Git](#)
- A web browser compatible with [Jupyter Notebooks](#)

⚠ Warning

Everything in this class will be tested with the up to date (or otherwise specified) version of Jupyter Notebooks. Google Colab is similar, but not the same, and some things may not work there. It is an okay backup, but should not be your primary work environment.

📌 Note

all Git instructions will be given as instructions for the command line interface and GitHub specific instructions via the web interface. You may choose to use GitHub desktop or built in IDE tools, but the instructional team may not be able to help.

Recommendation:

- Install python via [Anaconda](#)
- if you use Windows, install Git with [GitBash \(video instructions\)](#).
- if you use MacOS, install Git with the Xcode Command Line Tools. On Mavericks (10.9) or above you can do this by trying to run git from the Terminal the very first time.`git --version`

Optional:

- Text Editor: you may want a text editor outside of the Jupyter environment. Jupyter can edit markdown files (that you'll need for your portfolio), in browser, but it is more common to use a text editor like Atom or Sublime for this purpose.

Video install instructions for Anaconda:

- [Windows](#)
- [Mac](#)
- I don't have a video for linux, but it's a little more straight forward.

💡 A tip from Dr. Brown

I use [atom](#), but I decided to use it by downloading both Atom and Sublime and trying different things in each for a week. I liked Atom better after that and I've stuck with it since. I used Atom to write all of the content in this syllabus. VScode will also work, if needed

Textbook

The text for this class is a reference book and will not be a source of assignments. It will be a helpful reference and you may be directed there for answers to questions or alternate explanations of topics.

Python for Data Science is available free [online](#):

Zoom (backup only, Fall 2021 is in person)

This is where we will meet if for any reason we cannot be in person. You will find the link to class zoom sessions on Brightspace.

URI provides all faculty, staff, and students with a paid Zoom account. It *can* run in your browser or on a mobile device, but you will be able to participate in class best if you download the [Zoom client](#) on your computer. Please [log in](#) and [configure your account](#). Please add a photo of yourself to your account so that we can still see your likeness in some form when your camera is off. You may also wish to use a virtual background and you are welcome to do so.

Class will be interactive, so if you cannot be in a quiet place at class time, headphones with a built in microphone are strongly recommended.

For help, you can access the [instructions provided by IT](#).

[1] Too long; didn't read.

Data Science Achievements

In this course there are 5 learning outcomes that I expect you to achieve by the end of the semester. To get there, you'll focus on 15 smaller achievements that will be the basis of your grade. This section will describe how the topics covered, the learning outcomes, and the achievements are covered over time. In the next section, you'll see how these achievements turn into grades.

Learning Outcomes

By the end of the semester

1. (process) Describe the process of data science, define each phase, and identify standard tools
2. (data) Access and combine data in multiple formats for analysis
3. (exploratory) Perform exploratory data analyses including descriptive statistics and visualization
4. (modeling) Select models for data by applying and evaluating multiple models to a single dataset
5. (communicate) Communicate solutions to problems with data in common industry formats

We will build your skill in the **process** and **communicate** outcomes over the whole semester. The middle three skills will correspond roughly to the content taught for each of the first three portfolio checks.

Schedule

The course will meet MWF 3-3:50pm in Chafee Social Sci Center 235. Every class will include participatory live coding (instructor types code while explaining, students follow along)) instruction and small exercises for you to progress toward level 1 achievements of the new skills introduced in class that day.

Programming assignments that will be due each week Tuesday by 11:59pm.

week	topics	skills
1	[admin, python review]	process
2	Loading data, Python review	[access, prepare, summarize]
3	Exploratory Data Analysis	[summarize, visualize]
4	Data Cleaning	[prepare, summarize, visualize]
5	Databases, Merging DataFrames	[access, construct, summarize]
6	Modeling, Naive Bayes, classification performance metrics	[classification, evaluate]
7	decision trees, cross validation	[classification, evaluate]
8	Regression	[regression, evaluate]
9	Clustering	[clustering, evaluate]
10	SVM, parameter tuning	[optimize, tools]
11	KNN, Model comparison	[compare, tools]
12	Text Analysis	[unstructured]
13	Images Analysis	[unstructured, tools]
14	Deep Learning	[tools, compare]

Note

On the [Course Calendar on BrightSpace](#) page you can get a feed link to add to the calendar of your choice by clicking on the subscribe (star) button on the top right of the page. Class is for 1 hour there because of Brightspace/zoom integration limitations, but that calendar includes the zoom link.

Achievement Definitions

The table below describes how your participation, assignments, and portfolios will be assessed to earn each achievement. The keyword for each skill is a short name that will be used to refer to skills throughout the course materials; the full description of the skill is in this table.

	skill	Level 1	Level 2	Level 3
keyword				
python	pythonic code writing	python code that mostly runs, occasional pep8 adherence	python code that reliably runs, frequent pep8 adherence	reliable, efficient, pythonic code that consistently adheres to pep8
process	describe data science as a process	Identify basic components of data science	Describe and define each stage of the data science process	Compare different ways that data science can facilitate decision making
access	access data in multiple formats	load data from at least one format; identify the most common data formats	Load data for processing from the most common formats; Compare and contrast most common formats	access data from both common and uncommon formats and identify best practices for formats in different contexts
construct	construct datasets from multiple sources	identify what should happen to merge datasets or when they can be merged	apply basic merges	merge data that is not automatically aligned
summarize	Summarize and describe data	Describe the shape and structure of a dataset in basic terms	compute summary standard statistics of a whole dataset and grouped data	Compute and interpret various summary statistics of subsets of data
visualize	Visualize data	identify plot types, generate basic plots from pandas	generate multiple plot types with complete labeling with pandas and seaborn	generate complex plots with pandas and plotting libraries and customize with matplotlib or additional parameters
prepare	prepare data for analysis	identify if data is or is not ready for analysis, potential problems with data	apply data reshaping, cleaning, and filtering as directed	apply data reshaping, cleaning, and filtering manipulations reliably and correctly by assessing data as received
classification	Apply classification	identify and describe what classification is, apply pre-fit classification models	fit preselected classification model to a dataset	fit and apply classification models and select appropriate classification models for different contexts
regression	Apply Regression	identify what data that can be used for regression looks like	can fit linear regression models	can fit and explain regularized or nonlinear regression
clustering	Clustering	describe what clustering is	apply basic clustering	apply multiple clustering techniques, and interpret results
evaluate	Evaluate model performance	Explain basic performance metrics for different data science tasks	Apply basic model evaluation metrics to a held out test set	Evaluate a model with multiple metrics and cross validation
optimize	Optimize model parameters	Identify when model parameters need to be optimized	Manually optimize basic model parameters such as model order	Select optimal parameters based of multiple quantitative criteria and automate parameter tuning
compare	compare models	Qualitatively compare model classes	Compare model classes in specific terms and fit models in terms of traditional model performance metrics	Evaluate tradeoffs between different model comparison types

	skill	Level 1	Level 2	Level 3
keyword				
unstructured	model	Identify options for representing text data and use them once data is tranformed	Apply at least one representation to transform unstructured data for model fitting or summarizing	apply multiple representations and compare and contrast them for different end results
	unstructured data			
workflow	use industry standard data science tools and workflows to solve data science problems	Solve well strucutred problems with a single tool pipeline	Solve semi-structured, completely specified problems, apply common structure to learn new features of standard tools	Scope, choose an appropriate tool pipeline and solve data science problems, describe strengths and weakensses of common tools

Assignments and Skills

Using the keywords from the table above, this table shows which assignments you will be able to demonstrate which skills and the total number of assignments that assess each skill. This is the number of opportunities you have to earn Level 2 and still preserve 2 chances to earn Level 3 for each skill.

	A1	A2	A3	A4	A5	A6	A7	A8	A9	A10	A11	A12	A13	Assignment
keyword														
python	1	1	1	1	0	0	0	0	0	0	0	0	0	
process	1	1	0	0	0	0	0	0	0	0	0	0	0	
access	0	1	1	1	0	0	0	0	0	0	0	0	0	
construct	0	0	0	0	1	1	0	0	0	0	0	0	0	
summarize	0	0	1	1	1	1	1	1	1	1	1	1	1	1
visualize	0	0	1	1	0	1	1	1	1	1	1	1	1	1
prepare	0	0	0	1	1	0	0	0	0	0	0	0	0	
classification	0	0	0	0	0	1	1	0	0	1	0	0	0	
regression	0	0	0	0	0	0	0	1	0	0	1	0	0	
clustering	0	0	0	0	0	0	0	0	1	0	1	0	0	
evaluate	0	0	0	0	0	0	0	0	0	1	1	0	0	
optimize	0	0	0	0	0	0	0	0	0	1	1	0	0	
compare	0	0	0	0	0	0	0	0	0	0	1	0	1	
unstructured	0	0	0	0	0	0	0	0	0	0	0	1	1	
workflow	0	0	0	0	0	0	0	0	0	1	1	1	1	

Portfolios and Skills

The objective of your portfolio submissions is to earn Level 3 achievements. The following table shows what Level 3 looks like for each skill and identifies which portfolio submissions you can earn that Level 3 in that skill.

		Level 3	P1	P2	P3	P4
keyword						
python	reliable, efficient, pythonic code that consistently adheres to pep8		1	1	0	0
process	Compare different ways that data science can facilitate decision making		0	1	1	0
access	access data from both common and uncommon formats and identify best practices for formats in different contexts		1	1	0	0
construct	merge data that is not automatically aligned		1	1	0	0
summarize	Compute and interpret various summary statistics of subsets of data		1	1	0	0
visualize	generate complex plots with pandas and plotting libraries and customize with matplotlib or additional parameters		1	1	0	0
prepare	apply data reshaping, cleaning, and filtering manipulations reliably and correctly by assessing data as received		1	1	0	0
classification	fit and apply classification models and select appropriate classification models for different contexts		0	1	1	0
regression	can fit and explain regularized or nonlinear regression		0	1	1	0
clustering	apply multiple clustering techniques, and interpret results		0	1	1	0
evaluate	Evaluate a model with multiple metrics and cross validation		0	1	1	0
optimize	Select optimal parameters based of mutiple quantiatieve criteria and automate parameter tuning		0	0	1	1
compare	Evaluate tradeoffs between different model comparison types		0	0	1	1
unstructured	apply multiple representations and compare and contrast them for different end results		0	0	1	1
workflow	Scope, choose an appropriate tool pipeline and solve data science problems, describe strengths and weaknesses of common tools		0	0	1	1

Grading

This section of the syllabus describes the principles and mechanics of the grading for the course. This course will be graded on a basis of a set of *skills* (described in detail the next section of the syllabus). This is in contrast to more common grading on a basis of points earned through assignments.

Principles of Grading

Learning happens through practice and feedback. My goal as a teacher is for you to learn. The grading in this course is based on your learning of the material, rather than your completion of the activities that are assigned.

This course is designed to encourage you to work steadily at learning the material and demonstrating your new knowledge. There are no single points of failure, where you lose points that cannot be recovered. Also, you cannot cram anything one time and then forget it. The material will build and you have to demonstrate that you retained things.

- Earning a C in this class means you have a general understanding of Data Science and could participate in a basic conversation about all of the topics we cover. I expect everyone to reach this level.
- Earning a B means that you could solve simple data science problems on your own and complete parts of more complex problems as instructed by, for example, a supervisor in an internship or entry level job. This is a very accessible goal, it does not require you to get anything on the first try or to explore topics on your own. I expect most students to reach this level.
- Earning an A means that you could solve moderately complex problems independently and discuss the quality of others' data science solutions. This class will be challenging, it requires you to explore topics a little deeper than we cover them in class, but unlike typical grading it does not require all of your assignments to be near perfect.

Grading this way also is more amenable to the fact that there are correct and incorrect ways to do things, but there is not always a single correct answer to a realistic data science problem. Your work will be assessed on whether or not it demonstrates your learning of the targeted skills. You will also receive feedback on how to improve.

How it works

There are 15 skills that you will be graded on in this course. While learning these skills, you will work through a progression of learning. Your grade will be based on earning 45 achievements that are organized into 15 skill groups with 3 levels for each.

These map onto letter grades roughly as follows:

- If you achieve level 1 in all of the skills, you will earn at least a C in the course.
- To earn a B, you must earn all of the level 1 and level 2 achievements.
- To earn an A, you must earn all of the achievements.

You will have at least three opportunities to earn every level 2 achievement. You will have at least two opportunities to earn every level 3 achievement. You will have three *types* of opportunities to demonstrate your current skill level: participation, assignments, and a portfolio.

Each level of achievement corresponds to a phase in your learning of the skill:

- To earn level 1 achievements, you will need to demonstrate basic awareness of the required concepts and know approximately what to do, but you may need specific instructions of which things to do or to look up examples to modify every step of the way. You can earn level 1 achievements in class, assignments, or portfolio submissions.
- To earn level 2 achievements you will need to demonstrate understanding of the concepts and the ability to apply them with instruction after earning the level 1 achievement for that skill. You can earn level 2 achievements in assignments or portfolio submissions.
- To earn level 3 achievements you will be required to consistently execute each skill and demonstrate deep understanding of the course material, after achieving level 2 in that skill. You can earn level 3 achievements only through your portfolio submissions.

For each skill these are defined in the [Achievement Definition Table](#)

Participation

While attending synchronous class sessions, there will be understanding checks and in class exercises. Completing in class exercises and correctly answering questions in class can earn level 1 achievements. In class questions will be administered through the classroom chat platform Prismia.chat; these records will be used to update your skill progression. You can also earn level 1 achievements from adding annotation to a section of the class notes.

Assignments

For your learning to progress and earn level 2 achievements, you must practice with the skills outside of class time.

Assignments will each evaluate certain skills. After your assignment is reviewed, you will get qualitative feedback on your work, and an assessment of your demonstration of the targeted skills.

Portfolio Checks

To earn level 3 achievements, you will build a portfolio consisting of reflections, challenge problems, and longer analyses over the course of the semester. You will submit your portfolio for review 4 times. The first two will cover the skills taught up until 1 week before the submission deadline.

The third and fourth portfolio checks will cover all of the skills. The fourth will be due during finals. This means that, if you have achieved mastery of all of the skills by the 3rd portfolio check, you do not need to submit the fourth one.

Portfolio prompts will be given throughout the class, some will be structured questions, others may be questions that arise in class, for which there is not time to answer.

TLDR

You *could* earn a C through in class participation alone, if you make nearly zero mistakes. To earn a B, you must complete assignments and participate in class. To earn an A you must participate, complete assignments, and build a portfolio.

Detailed mechanics

Warning

If you will skip an assignment, please accept the GitHub assignment and then close the Feedback pull request with a comment. This way we can make sure that you have support you need.

On Brightspace there are 45 Grade items that you will get a 0 or a 1 grade for. These will be revealed, so that you can view them as you have an opportunity to demonstrate each one. The table below shows the minimum number of skills at each level to earn each letter grade.

	Level 3	Level 2	Level 1
letter grade			
A	15	15	15
A-	10	15	15
B+	5	15	15
B	0	15	15
B-	0	10	15
C+	0	5	15
C	0	0	15
C-	0	0	10
D+	0	0	5
D	0	0	3

For example, if you achieve level 2 on all of the skills and level 3 on 7 skills, that will be a B+.

If you achieve level 3 on 14 of the skills, but only level 1 on one of the skills, that will be a B-, because the minimum number of level 2 achievements for a B is 15. In this scenario the total number of achievements is 14 at level 3, 14 at level 2 and 15 at level 3, because you have to earn achievements within a skill in sequence.

The letter grade can be computed as follows

Note

In this example, you will have also achieved level 1 on all of the skills, because it is a prerequisite to level 2.

```
def compute_grade(num_level1,num_level2,num_level3):
    """
    Computes a grade for CSC/DSP310 from numbers of achievements at each level

    Parameters:
    -----
    num_level1 : int
        number of level 1 achievements earned
    num_level2 : int
        number of level 2 achievements earned
    num_level3 : int
        number of level 3 achievements earned

    Returns:
    -----
    letter_grade : string
        letter grade with modifier (+/-)
    """
    if num_level1 == 15:
        if num_level2 == 15:
            if num_level3 == 15:
                grade = 'A'
            elif num_level3 >= 10:
                grade = 'A-'
            elif num_level3 >= 5:
                grade = 'B+'
            else:
                grade = 'B'
        elif num_level2 >= 10:
            grade = 'B-'
        elif num_level2 >= 5:
            grade = 'C+'
        else:
            grade = 'C'
    elif num_level1 >= 10:
        grade = 'C-'
    elif num_level1 >= 5:
        grade = 'D+'
    elif num_level1 >= 3:
        grade = 'D'
    else:
        grade = 'F'

    return grade
```

For example you can run the code like this in a cell to see the output

```
compute_grade(15,15,15)
```

```
'A'
```

```
compute_grade(14,14,14)
```

```
'C-'
```

Or use `assert` to test it formally

```
assert compute_grade(14,14,14) == 'C-'
```

```
assert compute_grade(15,15,15) == 'A'
```

```
assert compute_grade(15,15,11) == 'A-'
```

Late work

Late assignments will not be graded. Every skill will be assessed through more than one assignment, so missing assignments occasionally not necessarily hurt your grade. If you do not submit any assignments that cover a given skill, you may earn the level 2 achievement in that skill through a portfolio check, but you will not be able to earn the level 3 achievement in that skill. If you submit work that is not complete, however, it will be assessed and receive feedback. Submitting pseudocode or code with errors and comments about what you have tried could earn a level 1 achievement. Additionally, most assignments cover multiple skills, so partially completing the assignment may earn level 2 for one, but not all. Submitting *something* even if it is not perfect is important to keeping conversation open and getting feedback and help continuously.

Building your Data Science Portfolio should be an ongoing process, where you commit work to your portfolio frequently. If something comes up and you cannot finish all that you would like assessed by the deadline, open an [Extension Request](#) issue on your repository.

In this issue, include:

1. A new deadline proposal
2. What additional work you plan to add
3. Why the extension is important to your learning
4. Why the extension will not hinder your ability to complete the next assignment on time.

This request should be no more than 7 sentences.

Portfolio due dates will be announced well in advance and prompts for it will be released weekly. You should spend some time working on it each week, applying what you've learned so far, from the feedback on previous assignments.

Examples

If you always attend and get everything correct, you will earn an A and you won't need to submit the 4th portfolio check or assignment 13.

Getting A Without Perfection

Note

You may visit office hours to discuss assignments that you did not complete on time to get feedback and check your own understanding, but they will not count toward skill demonstration.

Map to an A

How Achievements were earned

	Level 1	Level 2	Level 3
python	A1	A3	P1
process	A1	P1	P2
access	2	A2	P1
construct	5	A5	P1
summarize	3	A3	P1
visualize	3	A3	P2
prepare	4	A5	P2
classification	A10	P2	P3
regression	8	A11	P2
clustering	9	A9	P3
evaluate	7	A11	P3
optimize	10	A11	P4
compare	11	A13	P3
unstructured	12	A13	P4
tools	11	A13	P3

Activity Legend

In class	Assignment	Portfolio Check












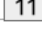

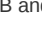

Other Activities

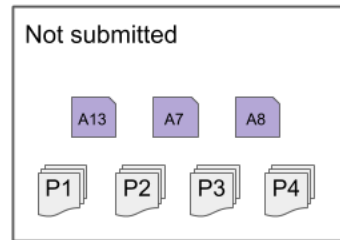
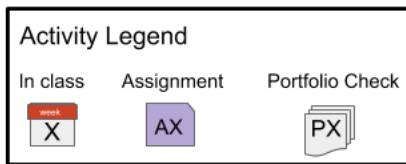
	Attended, but did not understand
	Submitted, but incorrect
	Missed class
	Not submitted
	Submitted, but incorrect
	Not submitted
	Not submitted
	Attended, but all level 1 complete
	Attended, but all level 1 complete

In this example the student made several mistakes, but still earned an A. This is the advantage to this grading scheme. For the **python**, **process**, and **classification** skills, the level 1 achievements were earned on assignments, not in class. For the **process** and **classification** skills, the level 2 achievements were not earned on assignments, only on portfolio checks, but they were earned on the first portfolio of those skills, so the level 3 achievements were earned on the second portfolio check for that skill. This student's fourth portfolio only demonstrated two skills: **optimize** and **unstructured**. It included only 1 analysis, a text analysis with optimizing the parameters of the model. Assignments 4 and 7 were both submitted, but didn't earn any achievements, the student got feedback though, that they were able to apply in later assignments to earn the achievements. The student missed class week 6 and chose to not submit assignment 6 and use week 7 to catch up. The student had too much work in another class and chose to skip assignment 8. The student tried assignment 12, but didn't finish it on time, so it was not graded, but the student visited office hours to understand and be sure to earn the level 2 **unstructured** achievement on assignment 13.

Getting a B with minimal work

Map to a B easily

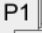
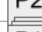

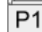
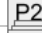
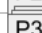
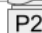
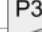






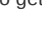
	Level 1	Level 2	Level 3
python	 1	A3	
process	 1	A1	
access	 2	A2	
construct	 5	A5	
summarize	 3	A3	
visualize	 3	A3	
prepare	 4	A4	
classification	 10	A6	
regression	 8	A11	
clustering	 9	A9	
evaluate	 7	A10	
optimize	 10	A10	
compare	 11	A11	
unstructured	 12	A12	
tools	 11	A12	

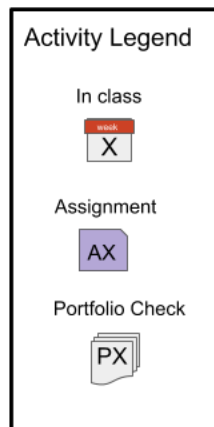


In this example, the student earned all level 1 achievements in class and all level 2 on assignments. This student was content with getting a B and chose to not submit a portfolio.

Getting a B while having trouble

Map to a B, having trouble

	Level 1	Level 2	Level 3
python	A1	 P1	
process	A1	 P2	
access	A2	 P1	
construct	A5	 P1	
summarize	A3	 P1	
visualize	A3	 P2	
prepare	A5	 P2	
classification	A10	 P3	
regression	A11	 P2	
clustering	A9	 P3	
evaluate	A11	 P3	
optimize	A11	 P4	
compare	A13	 P3	
unstructured	A13	 P4	
tools	A13	 P3	



In this example, the student struggled to understand in class and on assignments. Assignments were submitted that showed some understanding, but all had some serious mistakes, so only level 1 achievements were earned from assignments. The student wanted to get a B and worked hard to get the level 2 achievements on the portfolio checks.

Ram Tokens

Ram Tokens in this course will be used as a currency for extra effort. You can earn Ram Tokens by doing work that supports your learning or class activities, but do not directly demonstrate achievements. You can spend Ram Tokens to get extra grading. This will be mostly applicable to Portfolio Checks. In Checks 3 & 4, some achievements will not be eligible for grading as per the [table](#). However, you can exchange Ram Tokens to make more achievements eligible for assessment. This system rewards you for putting in consistent effort, even if it takes you many tries to understand a concept.

To accumulate Ram Tokens, you submit a 'Deposit' to the [Ram Token Bank: http://drsmb.co/ramtoken](http://drsmb.co/ramtoken) with a link to what you did to earn a token. To apply Ram tokens for extra grading, submit the same form, with a link to the assignment and add the

Support

Academic Enhancement Center

Academic Enhancement Center (for undergraduate courses): Located in Roosevelt Hall, the AEC offers free face-to-face and web-based services to undergraduate students seeking academic support. Peer tutoring is available for STEM-related courses by appointment online and in-person. The Writing Center offers peer tutoring focused on supporting undergraduate writers at any stage of a writing assignment. The UCS160 course and academic skills consultations offer students strategies and activities aimed at improving their studying and test-taking skills. Complete details about each of these programs, up-to-date schedules, contact information and self-service study resources are all available on the AEC website.

- **STEM Tutoring** helps students navigate 100 and 200 level math, chemistry, physics, biology, and other select STEM courses. The STEM Tutoring program offers free online and limited in-person peer-tutoring this fall. Undergraduates in introductory STEM courses have a variety of small group times to choose from and can select occasional or weekly appointments. Appointments and locations will be visible in the TutorTrac system on September 14th, 2020. The TutorTrac application is available through [URI Microsoft 365 single sign-on](#) and by visiting aec.uri.edu. More detailed information and instructions can be found on the AEC tutoring page.
- **Academic Skills Development** resources helps students plan work, manage time, and study more effectively. In Fall 2020, all Academic Skills and Strategies programming are offered both online and in-person. UCS160: Success in Higher Education is a one-credit course on developing a more effective approach to studying. Academic Consultations are 30-minute, 1 to 1 appointments that students can schedule on Starfish with Dr. David Hayes to address individual academic issues. Study Your Way to Success is a self-guided web portal connecting students to tips and strategies on studying and time management related topics. For more information on these programs, visit the Academic Skills Page or contact Dr. Hayes directly at davidhayes@uri.edu.
- The **Undergraduate Writing Center** provides free writing support to students in any class, at any stage of the writing process: from understanding an assignment and brainstorming ideas, to developing, organizing, and revising a draft. Fall 2020 services are offered through two online options: 1) real-time synchronous appointments with a peer consultant (25- and 50-minute slots, available Sunday - Friday), and 2) written asynchronous consultations with a 24-hour turn-around response time (available Monday - Friday). Synchronous appointments are video-based, with audio, chat, document-sharing, and live captioning capabilities, to meet a range of accessibility needs. View the synchronous and asynchronous schedules and book online, visit uri.mywconline.com.

Policies

Anti-Bias Statement:

We respect the rights and dignity of each individual and group. We reject prejudice and intolerance, and we work to understand differences. We believe that equity and inclusion are critical components for campus community members to thrive. If you are a target or a witness of a bias incident, you are encouraged to submit a report to the URI Bias Response Team at www.uri.edu/brt. There you will also find people and resources to help.

Disability Services for Students Statement:

Your access in this course is important. Please send me your Disability Services for Students (DSS) accommodation letter early in the semester so that we have adequate time to discuss and arrange your approved academic accommodations. If you have not yet established services through DSS, please contact them to engage in a confidential

conversation about the process for requesting reasonable accommodations in the classroom. DSS can be reached by calling: 401-874-2098, visiting: web.uri.edu/disability, or emailing: dss@etal.uri.edu. We are available to meet with students enrolled in Kingston as well as Providence courses.

Academic Honesty

Students are expected to be honest in all academic work. A student's name on any written work, quiz or exam shall be regarded as assurance that the work is the result of the student's own independent thought and study. Work should be stated in the student's own words, properly attributed to its source. Students have an obligation to know how to quote, paraphrase, summarize, cite and reference the work of others with integrity. The following are examples of academic dishonesty.

- Using material, directly or paraphrasing, from published sources (print or electronic) without appropriate citation
- Claiming disproportionate credit for work not done independently
- Unauthorized possession or access to exams
- Unauthorized communication during exams
- Unauthorized use of another's work or preparing work for another student
- Taking an exam for another student
- Altering or attempting to alter grades
- The use of notes or electronic devices to gain an unauthorized advantage during exams
- Fabricating or falsifying facts, data or references
- Facilitating or aiding another's academic dishonesty
- Submitting the same paper for more than one course without prior approval from the instructors

URI COVID-19 Statement

The University is committed to delivering its educational mission while protecting the health and safety of our community. While the university has worked to create a healthy learning environment for all, it is up to all of us to ensure our campus stays that way.

As members of the URI community, students are required to comply with standards of conduct and take precautions to keep themselves and others safe. Visit web.uri.edu/coronavirus/ for the latest information about the URI COVID-19 response.

- [Universal indoor masking](#) is required by all community members, on all campuses, regardless of vaccination status. If the universal mask mandate is discontinued during the semester, students who have an approved exemption and are not fully vaccinated will need to continue to wear a mask indoors and maintain physical distance.
- Students who are experiencing symptoms of illness should not come to class. Please stay in your home/room and notify URI Health Services via phone at 401-874-2246.
- If you are already on campus and start to feel ill, go home/back to your room and self-isolate. Notify URI Health Services via phone immediately at 401-874-2246.

If you are unable to attend class, please notify me at brownsarahm@uri.edu. We will work together to ensure that course instruction and work is completed for the semester.

Course Communications

Help Hours

Day	Time	Location	Host
Monday	1:00:00 PM-2:30	inperson roomtbd	Chamundi
Wednesday	4:00:00 PM	inperson roomtbd	Chamundi
Wednesday	2:00:00 PM-3	inperson roomtbd	Chamundi
Wednesday	7:00:00 PM-8:30	gather.town	Sarah
Friday	5:00:00 PM-6:30pm	gather.town	Chamundi
By appointment	TBD	in person Tyler 134	Sarah

We have several different ways to communicate in this course. This section summarizes them

To reach out, By usage

usage	platform	area	note
in class	prismia	chat	outside of class time this is not monitored closely
any time	prismia	message board	for discussion with peers
any time	prismia	download transcript	use after class to get preliminary notes eg if you miss a class
private questions to your assignment	github	issue on assignment repo	eg bugs in your code"
for general questions that can help others	github	issue on course website	eg what the instructions of an assignment mean or questions about the syllabus
to share resources	github	pull request on website	remember to request ram tokens if applicable
matters that don't fit into another category	e-mail	to brownsarahm@uri.edu	remember to include `[CSC310]` or `[DSP310]` (note `verbatim` no space)

Note

e-mail is last because it's not collaborative; other platforms allow us (Proessor + TA) to collaborate on who responds to things more easily.

By Platform

Use e-mail for

usage	area	note
matters that don't fit into another category	to brownsarahm@uri.edu	remember to include `[CSC310]` or `[DSP310]` (note `verbatim` no space)

Use github for

usage	area	note
private questions to your assignment	issue on assignment repo	eg bugs in your code"
for general questions that can help others	issue on course website	eg what the instructions of an assignment mean or questions about the syllabus
to share resources	pull request on website	remember to request ram tokens if applicable

Use prismia for


usage	area	note
in class	chat	outside of class time this is not monitored closely
any time	message board	for discussion with peers
any time	download transcript	use after class to get preliminary notes eg if you miss a class

Tips

For assignment help

- **send in advance, leave time for a response** I check e-mail/github a small number of times per day, during work hours, almost exclusively. You might see me post to this site, post to BrightSpace, or comment on your assignments outside of my normal working hours, but I will not reliably see emails that arrive during those hours. This means that it is important to start assignments early.

Using issues

- use issues for content directly related to assignments. If you push your code to the repository and then open an issue, I can see your code and your question at the same time and download it to run it if I need to debug it
- use issues for questions about this syllabus or class notes. At the top right there's a GitHub logo  that allows you to open a issue (for a question) or suggest an edit (eg if you think there's a typo or you find an additional helpful resource related to something)

For E-mail

- use e-mail for general inquiries or notifications
- Please include `[CSC310]` or `[DSP310]` in the subject line of your email along with the topic of your message. This is important, because your messages are important, but I also get a lot of e-mail. Consider these a cheat code to my inbox: I have setup a filter that will flag your e-mail if you use one of those in the subject to ensure that I see it.

Note

Whether you use CSC or DSP does not matter.

1. Welcome to Programming to Data Science

Today's goals:

1. Operate tools for in-class participation
 2. Understand what Data Science is, in broad terms
 3. Understand the syllabus (grading, topics covered, schedule, etc)
 4. Understand how to learn in this course
-

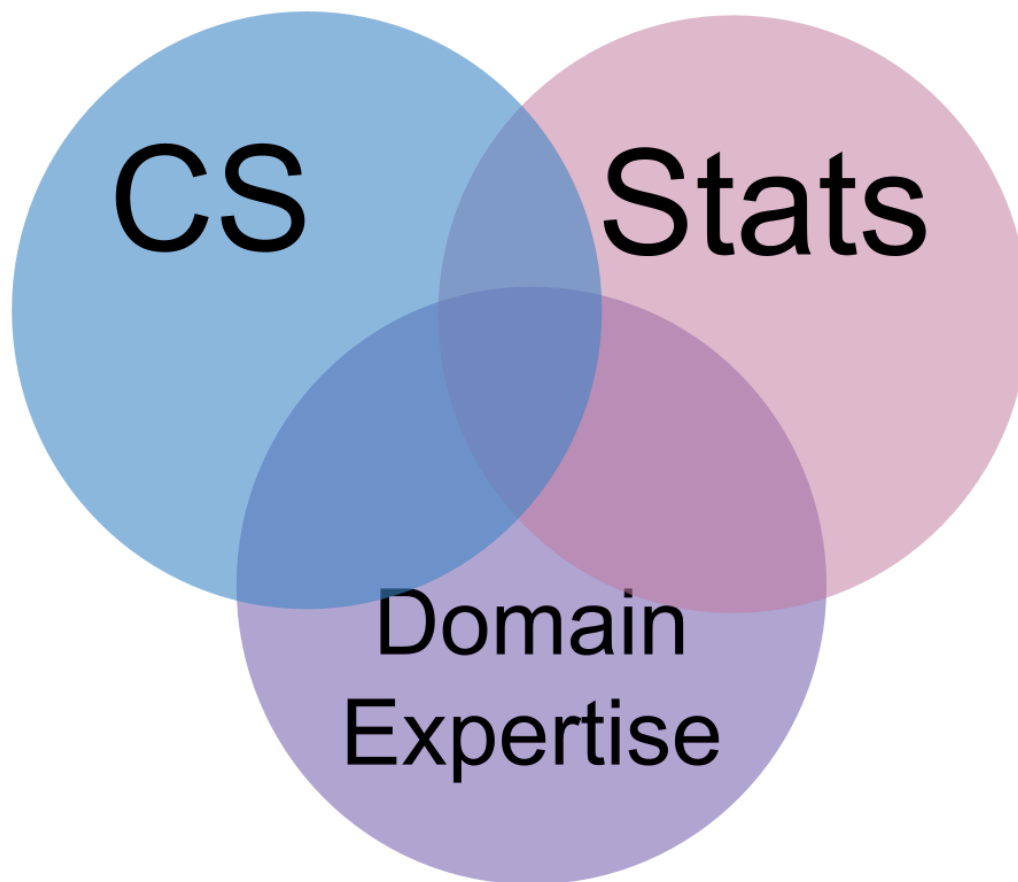
1.1. Prismia Chat

We will use these to monitor your participation in class and to gather information. Features:

- instructor only
- reply to you directly
- share responses for all

1.2. What is Data Science

In general:




statistics is the type of math we use to make sense of data. Formally, a statistic is just a function of data.

computer science is so that we can manipulate visualize and automate the inferences we make.

domain expertise helps us have the intuition to know if what we did worked right. A statistic must be interpreted in context; the relevant context determines what they mean and which are valid. The context will say whether automating something is safe or not, it can help us tell whether our code actually worked right or not.

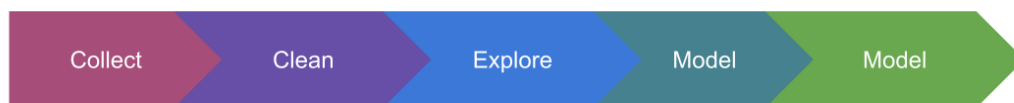
For this class



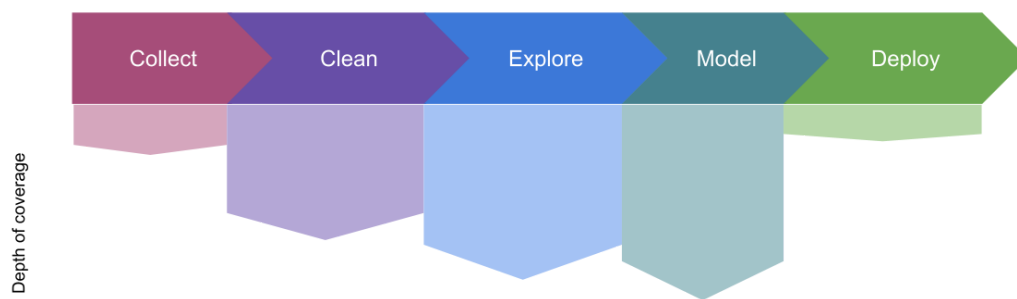
We'll focus on the programming as our main means of studying data science, but we will use bits of the other parts. In particular, you're encouraged to choose datasets that you have domain expertise about, or that you want to learn about.

But there are many definitions. We'll use this one, but you may come across others.

1.2.1. How does data science happen?

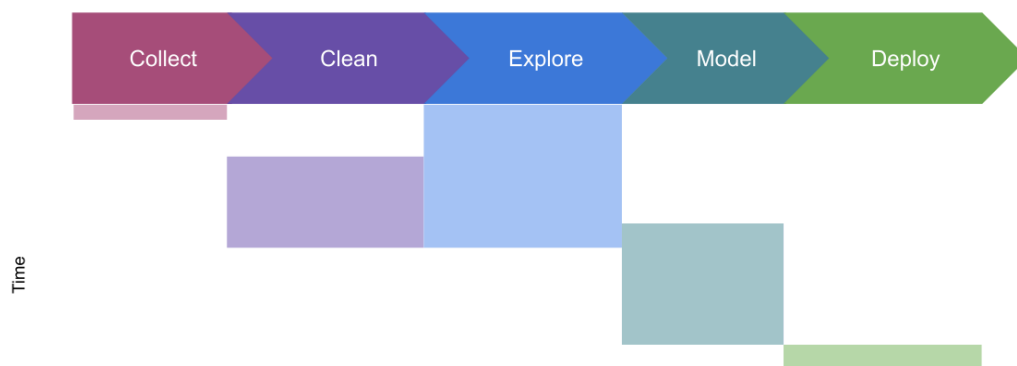


1.2.2. how we'll cover it, in depth



- *collect*: Discuss only a little; Minimal programming involved
- *clean*: Cover the main programming techniques; Some requires domain knowledge beyond scope of course
- *explore*: Cover the main programming techniques; Some requires domain knowledge beyond scope of course
- *model*: Cover the main programming, basic idea of models; How to use models, not how learning algorithms work
- *deploy*: A little bit at the end, but a lot of preparation for decision making around deployment

1.2.2.1. how we'll cover it in, time



We'll cover exploratory data analysis before cleaning because those tools will help us check how we've cleaned the data.

1.3. How this class will work

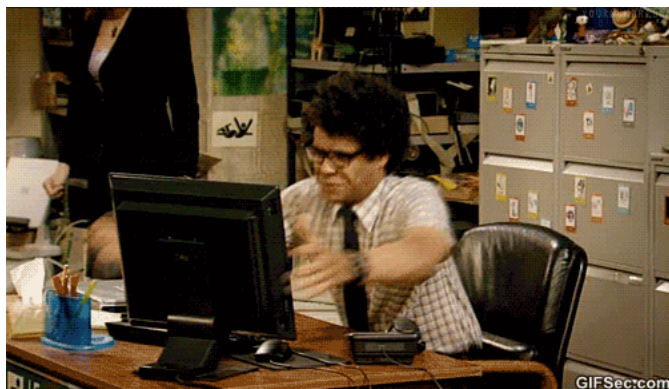
- today is an exception
- in general we'll be live coding

Let's look at the [syllabus](#)

Read carefully to make sure you understand the grading; it's not typical points and an average.

Class is designed to avoid this:

1.4.



1.5. Learning Cycle



Read more about how I'm designing this course to help you learn on the [how to learn](#) page.

1.6. Check your understanding of the syllabus

It's easy when reading something long to lose track of it. Your eyes can go over each word, without actually retaining the information, but it's important to understand the syllabus for the course.

You can find the answers to the following questions on the syllabus. If you've already read it, try answering them to check your understanding. If you haven't read it yet, use these to guide you to get familiar with finding key facts about the course on the syllabus.

1. What do you need to bring to class each day?
2. What is the basis of grading for this course?
3. How do you reference the course text?
4. What is the penalty for missing an assignment?

More information about the course is available throughout the site, the next few questions will help you self-check that you've found the important things. Remember, the goal is not necessarily to memorize all of this, but to be able to find it.

1. When & what are you expected to read for this class?

- ☐ read the text book before class
- ☐ review notes & documentation after class
- ☐ preview the notes & documentation before class
- ☐ read documentation and text book after class

1. Your assignment says to find a dataset that has variables of a specific type, which website can you use?
2. Your assignment says to find a dataset of any type about something you're interested in, which resource would you use?

2. Jupyter Notebook Tour & Python Review

2.1. A jupyter notebook tour

Launch a [jupyter notebook](#):

- on Windows, use anaconda terminal
- on Mac/Linux, use terminal

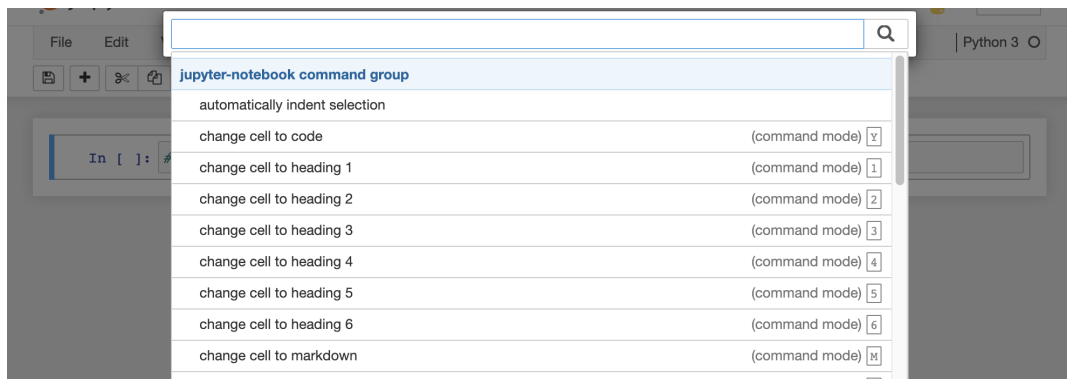
```
cd path/to/where/you/save/notes
jupyter notebook
```

A Jupyter notebook has two modes. When you first open, it is in command mode. The border is blue in command mode.

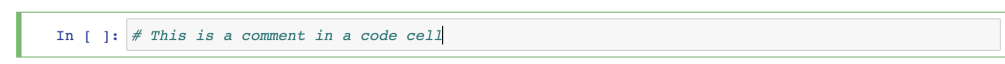


```
In [ ]: # This is a comment in a code cell
```

When you press a key in command mode it works like a shortcut. For example **p** shows the command search menu.

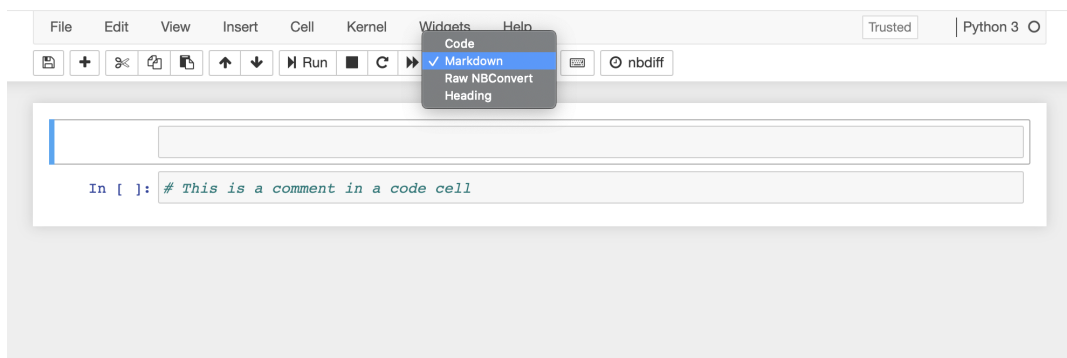


If you press **enter** (or **return**) or click on the highlighted cell, which is the boxes we can type in, it changes to edit mode. The border is green in edit mode



```
In [ ]: # This is a comment in a code cell
```

There are two type of cells that we will used: code and markdown. You can change that in command mode with **y** for code and **m** for markdown or on the cell type menu at the top of the notebook.



++

This is a markdown cell

- we can make
- itemized lists of
- bullet points

1. and we can make nubmered
2. lists, and not have to worry
3. about renumbering them
4. if we add a step in the middle later

2.1.1. Notebook Reminders

Blue border is command mode, green border is edit mode

use Escape to get to command mode

Common command mode actions:

- m: switch cell to markdown
- y: switch cell to code
- a: add a cell above
- b: add a cell below
- c: copy cell
- v: paste the cell
- O + O: restart kernel
- p: command menu

use enter/return to get to edit mode

In code cells, we can use a python interpreter, for example as a calculator.

```
4+6
```

```
10
```

It prints out the last line of code that it ran, even though it executes all of them

```
name = 'sarah'  
4+5  
name *3
```

```
'sarahsarahsarah'
```

Note

For a little more python review, see my [2020 CSC310 notes](#) this is just enough for this assignment.

2.2. Just enough Git for Assignment 1

2.2.1. Assignment 1:

Goals for this assignment

- setup your portfolio
- check that you understand the grading
- review Python basics
- practice with git and GitHub

2.2.2. Why Version control

We often want to keep track of the different versions in case we want to go back, but this can be painful:

"FINAL".doc



FINAL.doc!



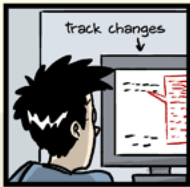
FINAL_rev.2.doc



FINAL_rev.6.COMMENTS.doc



FINAL_rev.8.comments5.
CORRECTIONS.doc



FINAL_rev.18.comments7.
corrections9.MORE.30.doc

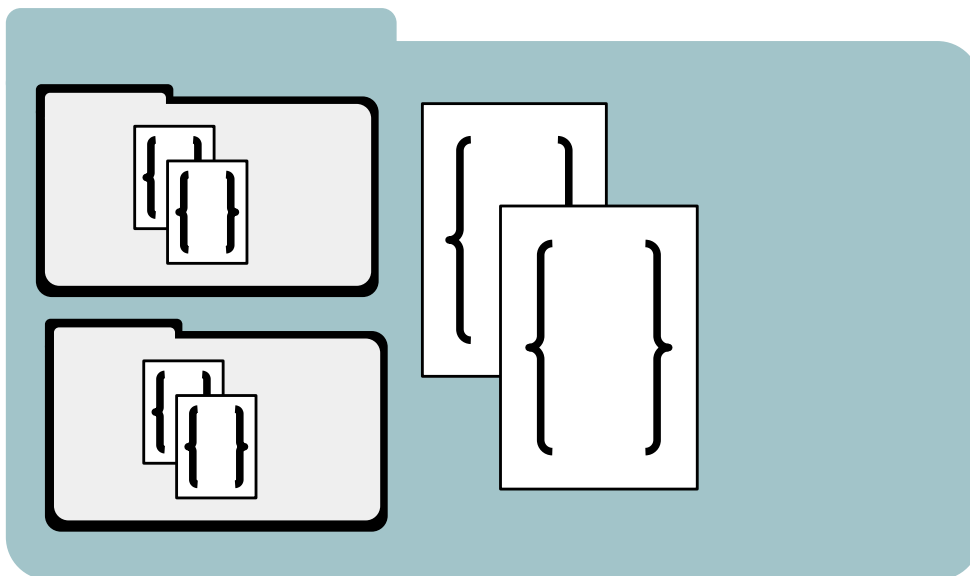


FINAL_rev.22.comments49.
corrections.10. #@\$%WHYDID
ICOMETOGRADSCHOOL?????.doc

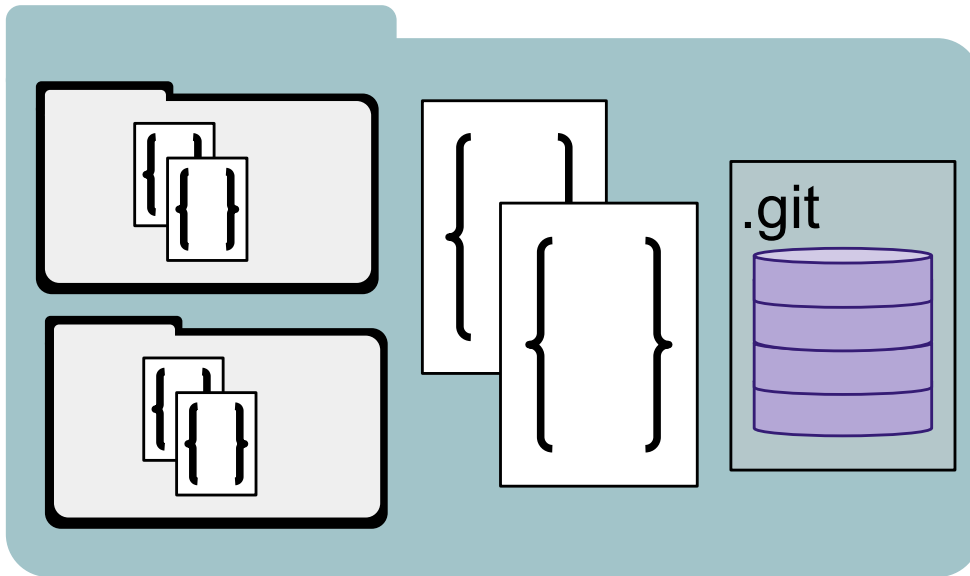
WWW.PHDCOMICS.COM

JORGE CHAM © 2012

We typically organize projects in folder

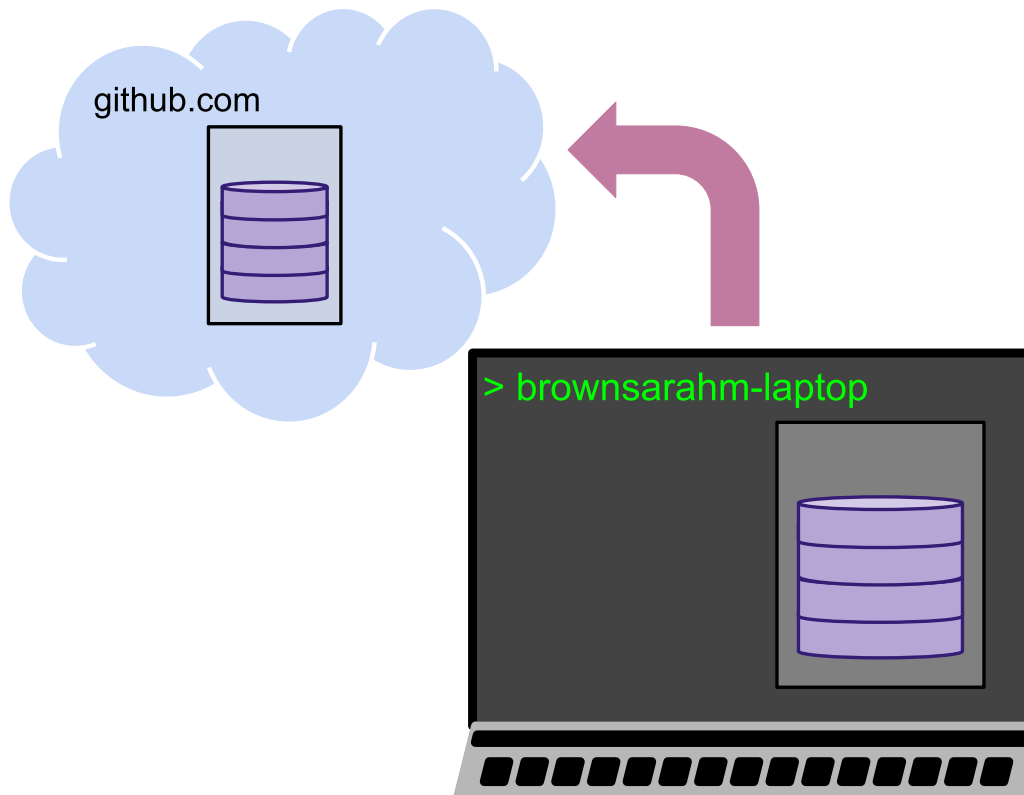


A [repository](#) is a folder with a hidden directory named `.git`



The `git` application manages that hidden directory, we don't write to it directly, which is why we keep it hidden.

Git is a distributed system, you have a local version and a remote version.



Once a repository exists on GitHub, we get a local copy by cloning it after we get its address from the GitHub interface, by clicking on the green code button that is below the menu area to the right. It's at the top right corner of the list of files in the repository.

rhodyprog4ds / portfolio-brownsarahm (Private)

generated from rhodyprog4ds/portfolio

<> Code Issues Pull requests Actions Projects Wiki Security Insights Settings

feedback had recent pushes 1 minute ago Compare & pull request

main 5 branches 1 tag

Go to file Add file Code

brownsarahm update toc to include notebook

.github	correct path for jupyter conversion
about	move notebook
template_files	convert notebooks to md
.gitignore	merge gh changes and ignore
README.md	Initial commit

Clone with HTTPS Use SSH

Use Git or checkout with SVN using the web URL.

https://github.com/rhodyprog4ds/por

Open with GitHub Desktop

Download ZIP

For this part, use GitBash on windows or terminal otherwise: If you set up a Personal Access Token you can use the https version

After `cd/to/where/you/want/your/repo/locally`:

```
git clone https://github.com/rhodyprog4ds/portfolio-example
```

If you set up ssh keys you use that instead

```
git clone git@github.com:rhodyprog4ds/portfolio-example.git
```

Once it's cloned, then you can navigate into the new folder:

```
cd portfolio-example
```

Then you can change files, for example adding to the intro.

Some common actions in Git, you'll want.

Check on the status of your repository:

```
git status
```

Add files to the staging area:

```
git add filename
```

Add all changes to the staging area:

```
git add .
```

Commit your changes to the repository:

```
git commit -m 'a message that will help your future self know what this part is'
```

Push your changes to GitHub

```
git push
```

Pull changes from GitHub

Note

These notes can be downloaded as an actual notebook, click the GitHub logo at the top of the page and choose .ipynb. The following is not runnabel in the notebook as is.


```
git pull
```

You can also go through these same basic steps: add, commit, push

2.3. More on git

- [GitHub Hello World](#)
- [Software Carpentry Git Novice Lesson](#)

Also, in Spring 2022, I'm teaching a section of CSC392: Topics in Computing, Introduction to Computer Systems, that will cover tools of the trade (git, bash, etc) and how they all work in great detail.

2.4. More on Python

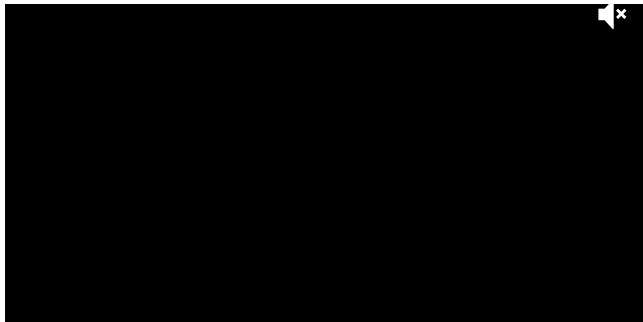
Read [Pep 8](#) to see what good style in Python is.

3. Getting help, object inspection, loading data

3.1. First, Don't Worry members

Class Response Summary:





USES CLASSIC MEME FORMAT



PEOPLE LIKE IT



3.2. Getting Help in Jupyter

Python has a `print` function and we can use the help in jupyter to learn about how to use it in different ways.

Given this code excerpt, how could you print out "Sarah_Brown"?

```
first = 'Sarah'
last = 'Brown'
```

We can use jupyter popup help with shift +tab or ?

```
print?
```

Or the base python `help` function

```
help(print)
```

```
Help on built-in function print in module builtins:

print(...)
    print(value, ..., sep=' ', end='\n', file=sys.stdout, flush=False)

    Prints the values to a stream, or to sys.stdout by default.
    Optional keyword arguments:
    file: a file-like object (stream); defaults to the current sys.stdout.
    sep:   string inserted between values, default a space.
    end:   string appended after the last value, default a newline.
    flush: whether to forcibly flush the stream.
```

Notice that function can take multiple arguments and has a keyword argument (must be used like `argument=value`) described as `sep=' '`. This means that by default it adds a space

```
print(first,last)
```

```
Sarah Brown
```

But we can change the separator.

```
Sarah_Brown
```

Note that it also defaults to end to use `\n`

```
print(first,last)
print('hello')
```

```
Sarah Brown
hello
```

Where does this help information come from?

Note

You can copy code from the notes, try hovering over this

```
def compute_grade(num_level1,num_level2,num_level3):
    '''
    Computes a grade for CSC/DSP310 from numbers of achievements at each level

    Parameters:
    -----
    num_level1 : int
        number of level 1 achievements earned
    num_level2 : int
        number of level 2 achievements earned
    num_level3 : int
        number of level 3 achievements earned

    Returns:
    -----
    letter_grade : string
        letter grade with modifier (+/-)
    '''
    if num_level1 == 15:
        if num_level2 == 15:
            if num_level3 == 15:
                grade = 'A'
            elif num_level3 >= 10:
                grade = 'A-'
            elif num_level3 >= 5:
                grade = 'B+'
            else:
                grade = 'B'
        elif num_level2 >= 10:
            grade = 'B-'
        elif num_level2 >= 5:
            grade = 'C+'
        else:
            grade = 'C'
    elif num_level1 >= 10:
        grade = 'C-'
    elif num_level1 >= 5:
        grade = 'D+'
    elif num_level1 >= 3:
        grade = 'D'
    else:
        grade = 'F'

    return grade
```

We can apply `help` on the function we wrote

```
help(compute_grade)
```

```
Help on function compute_grade in module __main__:

compute_grade(num_level1, num_level2, num_level3)
    Computes a grade for CSC/DSP310 from numbers of achievements at each level

    Parameters:
    -----
    num_level1 : int
        number of level 1 achievements earned
    num_level2 : int
        number of level 2 achievements earned
    num_level3 : int
        number of level 3 achievements earned

    Returns:
    -----
    letter_grade : string
        letter grade with modifier (+/-)
```

It gets the docstring

3.3. Everything is an Object in Python

we can use the builtin function `type` to inspect them, and get attributes with `.`

```
type(compute_grade)
```

```
function
```

```
compute_grade.__name__
```

```
'compute_grade'
```

```
c = 4.5
```

```
type(c)
```

```
float
```

```
c= 'hello'
```

```
type(c)
```

```
str
```

When do we use single vs double quotes?

- You can use either, unless you need to put one inside the string then use the other.

```
my_sentence = "The professor's name is Dr. Brown"
```

```
my_sentence = 'The professor's name is Dr. Brown'
```

```
File "/tmp/ipykernel_1651/607286316.py", line 1
  my_sentence = 'The professor's name is Dr. Brown'
                                     ^
SyntaxError: invalid syntax
```

Yes we can escape special characters:

```
my_sentence = 'The professor\'s name is Dr. Brown'
```

but, it's less readable and not recommended.

3.4. Good Code is always relative

In programming for data science, we are often trying to tell a story.

💡 Try it yourself

How might this goal change your code for this class relative to other code you have written or could imagine writing?

Python is a fully [open source project](#) and as such is governed by [community standards](#) and [conventions](#).

💡 Try it yourself

Find PEP8 (note that following it is part of earning python achievements)

The [documentation](#) for the full language is online too.

Guido van Rossum was the first main developer and wrote [essays](#) about python too.

it's [pretty popular](#)

3.5. Coffee Data

We're going to use a dataset about [coffee quality](#) today.

How was this dataset collected?

- reviews added to DB
- then scraped

Where did it come from?

- coffee Quality Institute's trained reviewers.

what format is it provided in?

- csv (Comma Separated Values)

what other information is in this repository?

- the code to scrape

Get raw url for the dataset click on the raw button on the [csv page](#), then copy the url.

 a screenshot from github of the data file page with the raw button circled in pink

We'll save that url as a variable to work with it.

```
data_url = 'https://raw.githubusercontent.com/jldbc/coffee-quality-  
database/master/data/robusta_data_cleaned.csv'
```

We will use a library called Pandas

```
import pandas as pd  
# import library and give it an alias (nickname) pd
```

```
pd.read_csv(data_url)
```

Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	
0	1 Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ai c produ
1	2 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethurā e
2	3 Robusta	andrew hetzel	India	sethuraman estate	NaN	
3	4 Robusta	ugacof	Uganda	ugacof project area	NaN	uq
4	5 Robusta	katuka development trust ltd	Uganda	katikamu capca farmers association	NaN	kā develop
5	6 Robusta	andrew hetzel	India	NaN	NaN	
6	7 Robusta	andrew hetzel	India	sethuraman estates	NaN	
7	8 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	7	sethurā e
8	9 Robusta	nishant gurjer	India	sethuraman estate	RKR	sethurā e
9	10 Robusta	ugacof	Uganda	ishaka	NaN	nsu
10	11 Robusta	ugacof	Uganda	ugacof project area	NaN	uq
11	12 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	RC AB	sethurā e
12	13 Robusta	andrew hetzel	India	sethuraman estates	NaN	
13	14 Robusta	kasoz coffee farmers association	Uganda	kasoz coffee farmers	NaN	
14	15 Robusta	ankole coffee producers coop	Uganda	kyangundu coop society	NaN	ai c produ coop t
15	16 Robusta	andrew hetzel	India	sethuraman estate	NaN	
16	17 Robusta	andrew hetzel	India	sethuraman estates	NaN	sethurā es
17	18 Robusta	kawacom uganda ltd	Uganda	bushenyi	NaN	kawā
18	19 Robusta	nitubaasa ltd	Uganda	kigezi coffee farmers association	NaN	nitub
19	20 Robusta	mannya coffee project	Uganda	mannya coffee project	NaN	ma c pr
20	21 Robusta	andrew hetzel	India	sethuraman estates	NaN	
21	22 Robusta	andrew hetzel	India	sethuraman estates	NaN	sethurā es
22	23 Robusta	andrew hetzel	United States	sethuraman estates	NaN	sethurā es

Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number		
23	24	Robusta	luis robles	Ecuador	robustasa	Lavado 1	our ow
24	25	Robusta	luis robles	Ecuador	robustasa	Lavado 3	labor
25	26	Robusta	james moore	United States	fazenda cazengo	NaN	caz
26	27	Robusta	cafe politico	India	NaN	NaN	
27	28	Robusta	cafe politico	Vietnam	NaN	NaN	

28 rows × 44 columns

Try it yourself

Read the data in again, but with the index correct and save it to a variable.

Once we read it in, we can view the first 5 rows with the `head` method.

```
coffee_df.head()
```

	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICC
1	Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ankole coffee producers	
2	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethuraman estate	14/114
3	Robusta	andrew hetzel	India	sethuraman estate	NaN	NaN	
4	Robusta	ugacof	Uganda	ugacof project area	NaN	ugacof	
5	Robusta	katuka development trust ltd	Uganda	katikamu capca farmers association	NaN	katuka development trust	

5 rows × 43 columns

Important

Remember to comment & annotate your code

3.6. Follow Up questions

3.6.1. General Questions

How do you create code to scrape data from a website and compile it into a csv file?



Will we be using pandas a lot during the semester?



3.6.2. Clarifying

How do you auto finish your directories



How do you properly shut down Jupyter Notebook



Is pd some sort of variable we set or was it built in?



How should I be organized for this class? Keep it all in a single folder? Keep it on GitHub?



I'm still not sure how to keep everything together in a portfolio for the semester?



I am still wondering if I am using anaconda or just normal terminal



Can I push this code into my portfolio using the anaconda terminal



3.6.3. Grading Questions

How do we keep track of which achievements we've earned?



I don't really have many questions from today, but I was wondering if office hours were posted.



Will we always submit homework through the portfolio folder in github?



I'm just confused as how to view my feedback from the assignment



3.6.4. Questions we'll answer later this week

- does each column have a number assigned to it in data frames?
- Can other data types be imported into a notebook and edited the same way as .csv files?

3.7. Try it yourself

- How could you check if `pd` is built in or if we defined it?
- If we wanted to see more than 5 rows when printing the head of the dataset how would we do so?

4. Pandas DataFrames

Today, we're going to explore [DataFrames](#) in greater detail. We'll continue using that same coffee dataset.

```
coffee_data_url = 'https://raw.githubusercontent.com/jldbc/coffee-quality-database/master/data/robusta_data_cleaned.csv'
```

4.1. More about loading libraries

We can import pandas without the alias `pd` if we want, but then we have to use the full name everywhere

```
import pandas
```

```
pandas.read_csv(coffee_data_url)
```

Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	
0	1 Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ai c produ
1	2 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethurā e
2	3 Robusta	andrew hetzel	India	sethuraman estate	NaN	
3	4 Robusta	ugacof	Uganda	ugacof project area	NaN	uq
4	5 Robusta	katuka development trust ltd	Uganda	katikamu capca farmers association	NaN	kā develop
5	6 Robusta	andrew hetzel	India	NaN	NaN	
6	7 Robusta	andrew hetzel	India	sethuraman estates	NaN	
7	8 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	7	sethurā e
8	9 Robusta	nishant gurjer	India	sethuraman estate	RKR	sethurā e
9	10 Robusta	ugacof	Uganda	ishaka	NaN	nsu
10	11 Robusta	ugacof	Uganda	ugacof project area	NaN	uq
11	12 Robusta	nishant gurjer	India	sethuraman estate kaapi royale	RC AB	sethurā e
12	13 Robusta	andrew hetzel	India	sethuraman estates	NaN	
13	14 Robusta	kasoz coffee farmers association	Uganda	kasoz coffee farmers	NaN	
14	15 Robusta	ankole coffee producers coop	Uganda	kyangundu coop society	NaN	ai c produ coop t
15	16 Robusta	andrew hetzel	India	sethuraman estate	NaN	
16	17 Robusta	andrew hetzel	India	sethuraman estates	NaN	sethurā es
17	18 Robusta	kawacom uganda ltd	Uganda	bushenyi	NaN	kawā
18	19 Robusta	nitubaasa ltd	Uganda	kigezi coffee farmers association	NaN	nitub
19	20 Robusta	mannya coffee project	Uganda	mannya coffee project	NaN	ma c pr
20	21 Robusta	andrew hetzel	India	sethuraman estates	NaN	
21	22 Robusta	andrew hetzel	India	sethuraman estates	NaN	sethurā es
22	23 Robusta	andrew hetzel	United States	sethuraman estates	NaN	sethurā es

Unnamed: 0	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number		
23	24	Robusta	luis robles	Ecuador	robustasa	Lavado 1	our ow
24	25	Robusta	luis robles	Ecuador	robustasa	Lavado 3	labor
25	26	Robusta	james moore	United States	fazenda cazengo	NaN	caz
26	27	Robusta	cafe politico	India	NaN	NaN	
27	28	Robusta	cafe politico	Vietnam	NaN	NaN	

28 rows × 44 columns

We'll use `pd` because that's the more common convention and so that we can type fewer characters throughout our code

```
import pandas as pd
```

4.2. Examining DataFrames

```
df = pd.read_csv(coffee_data_url, index_col=0)
```

We can look at the first 5 rows with `head`

```
df.head()
```

	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICC
1	Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ankole coffee producers	
2	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethuraman estate	14/1148
3	Robusta	andrew hetzel	India	sethuraman estate	NaN	NaN	
4	Robusta	ugacof	Uganda	ugacof project area	NaN	ugacof	
5	Robusta	katuka development trust ltd	Uganda	katikamu capca farmers association	NaN	katuka development trust	

5 rows × 43 columns

Using `help`, we can see that `head` takes one parameter and has a default value of 5, which is why we got 5 rows, but we can get 2 instead

```
df.head(2)
```

	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICO.Number
1	Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ankole coffee producers	
2	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethuraman estate	14/1148/20

2 rows × 43 columns

We can look at the last rows with `tail`

```
df.tail(3)
```

	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	ICO.Number	
26	Robusta	james moore	United States	fazenda cazengo	NaN	cafe cazengo	NaN	c
27	Robusta	cafe politico	India	NaN	NaN	NaN	14-1118-2014-0087	
28	Robusta	cafe politico	Vietnam	NaN	NaN	NaN	NaN	

3 rows × 43 columns

I told you this was a DataFrame, but we can check with type.

```
type(df)
```

```
pandas.core.frame.DataFrame
```

We can also examine its parts. It consists of several; first the column headings

```
df.columns
```

```
Index(['Species', 'Owner', 'Country.of.Origin', 'Farm.Name', 'Lot.Number',  
      'Mill', 'ICO.Number', 'Company', 'Altitude', 'Region', 'Producer',  
      'Number.of.Bags', 'Bag.Weight', 'In.Country.Partner', 'Harvest.Year',  
      'Grading.Date', 'Owner.1', 'Variety', 'Processing.Method',  
      'Fragrance...Aroma', 'Flavor', 'Aftertaste', 'Salt...Acid',  
      'Bitter...Sweet', 'Mouthfeel', 'Uniform.Cup', 'Clean.Cup', 'Balance',  
      'Copper.Points', 'Total.Cup.Points', 'Moisture', 'Category.One.Defects',  
      'Quakers', 'Color', 'Category.Two.Defects', 'Expiration',  
      'Certification.Body', 'Certification.Address', 'Certification.Contact',  
      'unit_of_measurement', 'altitude_low_meters', 'altitude_high_meters',  
      'altitude_mean_meters'],  
      dtype='object')
```

These are a special type called Index

```
type(df.columns)
```

```
pandas.core.indexes.base.Index
```

It also has an index

```
df.index
```

```
Int64Index([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16, 17,  
            18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28],  
            dtype='int64')
```

and values

```
df.values
```

```
array([[ 'Robusta', 'ankole coffee producers coop', 'Uganda', ..., 1488.0,  
        1488.0, 1488.0],  
       [ 'Robusta', 'nishant gurjer', 'India', ..., 3170.0, 3170.0,  
        3170.0],  
       [ 'Robusta', 'andrew hetzel', 'India', ..., 1000.0, 1000.0, 1000.0],  
       ...,  
       [ 'Robusta', 'james moore', 'United States', ..., 795.0, 795.0,  
        795.0],  
       [ 'Robusta', 'cafe politico', 'India', ..., nan, nan, nan],  
       [ 'Robusta', 'cafe politico', 'Vietnam', ..., nan, nan, nan]],  
      dtype=object)
```

it also knows its own shape

```
df.shape
```

```
(28, 43)
```

we can use builtin functions on our DataFrame too not just its own methods and attributes.

```
len(df)
```

```
28
```

Why does `len` turn green? it's a python reserve word

4.3. Building a Data Frame programmatically

One way to build a data frame is from a dictionary:

```
people = {'names': ['Sarah', 'Connor', 'Kenza'],  
          'username': ['brownsarahm', 'sudoPsych', 'kddlh']}
```

```
people
```

```
{'names': ['Sarah', 'Connor', 'Kenza'],  
 'username': ['brownsarahm', 'sudoPsych', 'kddlh']}
```

```
type(people)
```

```
dict
```

```
people_df = pd.DataFrame(people)  
people_df
```

	names	username
0	Sarah	brownsarahm
1	Connor	sudoPsych
2	Kenza	kddlh

```
type(people['names'])
```

```
list
```

```
type(people)
```

```
dict
```

```
type({4,5,5})
```

```
set
```

```
{4,5,5}
```

```
{4, 5}
```

```
people['names']
```

```
['Sarah', 'Connor', 'Kenza']
```

```
type(set(people['names']))
```

```
set
```

```
unique_people = set(people['names'])  
type(unique_people)
```

```
set
```

```
df.columns
```

```
Index(['Species', 'Owner', 'Country.of.Origin', 'Farm.Name', 'Lot.Number',  
      'Mill', 'ICO.Number', 'Company', 'Altitude', 'Region', 'Producer',  
      'Number.of.Bags', 'Bag.Weight', 'In.Country.Partner', 'Harvest.Year',  
      'Grading.Date', 'Owner.1', 'Variety', 'Processing.Method',  
      'Fragrance...Aroma', 'Flavor', 'Aftertaste', 'Salt...Acid',  
      'Bitter...Sweet', 'Mouthfeel', 'Uniform.Cup', 'Clean.Cup', 'Balance',  
      'Cupper.Points', 'Total.Cup.Points', 'Moisture', 'Category.One.Defects',  
      'Quakers', 'Color', 'Category.Two.Defects', 'Expiration',  
      'Certification.Body', 'Certification.Address', 'Certification.Contact',  
      'unit_of_measurement', 'altitude_low_meters', 'altitude_high_meters',  
      'altitude_mean_meters'],  
      dtype='object')
```

```
for col in df.columns:  
    print(col.split('.'))
```

```
['Species']  
['Owner']  
['Country', 'of', 'Origin']  
['Farm', 'Name']  
['Lot', 'Number']  
['Mill']  
['ICO', 'Number']  
['Company']  
['Altitude']  
['Region']  
['Producer']  
['Number', 'of', 'Bags']  
['Bag', 'Weight']  
['In', 'Country', 'Partner']  
['Harvest', 'Year']  
['Grading', 'Date']  
['Owner', '1']  
['Variety']  
['Processing', 'Method']  
['Fragrance', '', '', 'Aroma']  
['Flavor']  
['Aftertaste']  
['Salt', '', '', 'Acid']  
['Bitter', '', '', 'Sweet']  
['Mouthfeel']  
['Uniform', 'Cup']  
['Clean', 'Cup']  
['Balance']  
['Cupper', 'Points']  
['Total', 'Cup', 'Points']  
['Moisture']  
['Category', 'One', 'Defects']  
['Quakers']  
['Color']  
['Category', 'Two', 'Defects']  
['Expiration']  
['Certification', 'Body']  
['Certification', 'Address']  
['Certification', 'Contact']  
['unit_of_measurement']  
['altitude_low_meters']  
['altitude_high_meters']  
['altitude_mean_meters']
```

```
for key,value in people.items():  
    print(key,':',value)
```

```
names : ['Sarah', 'Connor', 'Kenza']
username : ['brownsarahm', 'sudoPsych', 'kddlh']
```

```
df['Owner']
```

```
1      ankole coffee producers coop
2              nishant gurjer
3              andrew hetzel
4              ugacof
5      katuka development trust ltd
6              andrew hetzel
7              andrew hetzel
8              nishant gurjer
9              nishant gurjer
10             ugacof
11             ugacof
12             nishant gurjer
13             andrew hetzel
14      kasozi coffee farmers association
15      ankole coffee producers coop
16              andrew hetzel
17              andrew hetzel
18      kawacom uganda ltd
19      nitubaasa ltd
20      mannya coffee project
21              andrew hetzel
22              andrew hetzel
23              andrew hetzel
24              luis robles
25              luis robles
26              james moore
27              cafe politico
28              cafe politico
Name: Owner, dtype: object
```

```
df.Owner
```

```
1      ankole coffee producers coop
2              nishant gurjer
3              andrew hetzel
4              ugacof
5      katuka development trust ltd
6              andrew hetzel
7              andrew hetzel
8              nishant gurjer
9              nishant gurjer
10             ugacof
11             ugacof
12             nishant gurjer
13             andrew hetzel
14      kasozi coffee farmers association
15      ankole coffee producers coop
16              andrew hetzel
17              andrew hetzel
18      kawacom uganda ltd
19      nitubaasa ltd
20      mannya coffee project
21              andrew hetzel
22              andrew hetzel
23              andrew hetzel
24              luis robles
25              luis robles
26              james moore
27              cafe politico
28              cafe politico
Name: Owner, dtype: object
```

```
df
```


	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	IC
1	Robusta	ankole coffee producers coop	Uganda	kyangundu cooperative society	NaN	ankole coffee producers	
2	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	25	sethuraman estate	14/114
3	Robusta	andrew hetzel	India	sethuraman estate	NaN	NaN	
4	Robusta	ugacof	Uganda	ugacof project area	NaN	ugacof	
5	Robusta	katuka development trust ltd	Uganda	katikamu capca farmers association	NaN	katuka development trust	
6	Robusta	andrew hetzel	India	NaN	NaN	(self)	
7	Robusta	andrew hetzel	India	sethuraman estates	NaN	NaN	
8	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	7	sethuraman estate	14/114
9	Robusta	nishant gurjer	India	sethuraman estate	RKR	sethuraman estate	14/114
10	Robusta	ugacof	Uganda	ishaka	NaN	nsubuga umar	
11	Robusta	ugacof	Uganda	ugacof project area	NaN	ugacof	
12	Robusta	nishant gurjer	India	sethuraman estate kaapi royale	RC AB	sethuraman estate	14/114
13	Robusta	andrew hetzel	India	sethuraman estates	NaN	NaN	
14	Robusta	kasoz coffee farmers association	Uganda	kasoz coffee farmers	NaN	NaN	
15	Robusta	ankole coffee producers coop	Uganda	kyangundu coop society	NaN	ankole coffee producers coop union ltd	
16	Robusta	andrew hetzel	India	sethuraman estate	NaN	NaN	
17	Robusta	andrew hetzel	India	sethuraman estates	NaN	sethuraman estates	
18	Robusta	kawacom uganda ltd	Uganda	bushenyi	NaN	kawacom	
19	Robusta	nitubaasa ltd	Uganda	kigezi coffee farmers association	NaN	nitubaasa	
20	Robusta	mannya coffee project	Uganda	mannya coffee project	NaN	mannya coffee project	
21	Robusta	andrew hetzel	India	sethuraman estates	NaN	NaN	
22	Robusta	andrew hetzel	India	sethuraman estates	NaN	sethuraman estates	
23	Robusta	andrew hetzel	United States	sethuraman estates	NaN	sethuraman estates	

	Species	Owner	Country.of.Origin	Farm.Name	Lot.Number	Mill	IC
24	Robusta	luis robles	Ecuador	robustasa	Lavado 1	our own lab	
25	Robusta	luis robles	Ecuador	robustasa	Lavado 3	own laboratory	
26	Robusta	james moore	United States	fazenda cazengo	NaN	cafe cazengo	
27	Robusta	cafe politico	India	NaN	NaN	NaN	14-1
28	Robusta	cafe politico	Vietnam	NaN	NaN	NaN	

28 rows × 43 columns

Key points:

write three things to remember from today's class

4.4. Questions After Classroom

many overlapping questions today

4.5. General

How to know which function to use in certain problems or situations



4.6. Clarifying

Is there a way to have a set show the duplicates that get discarded?



being able to access the code somewhere without asking to scroll would be nice



4.7. Course Admin

When will homeworks be posted/due typically?



4.8. Questions we'll answer later

can you use cast a pandas dataframe into a set?



4.9. Try it yourself

- Create variables of three different types with facts about yourself. Use descriptive variable names relative to the contents, not their types.
- Create a list, again with a descriptive name, and print out the types

```
<class 'str'>
<class 'int'>
<class 'list'>
```

- Write a function, `type_extractor` that takes a list and a type and returns the item of that type from the list
- Test your function on all three items from your dictionary.
- Use one type of jupyter help on your function, what does it display? If it doesn't display anything modify your function so that help will work.
- Make yourself notes in the most memorable way for you about what a DataFrame is.

5. More Loading Data, Indexing, and Iterables

As always, we'll start with loading pandas.

```
import pandas as pd
```

5.1. Checking in on help hours

if you missed class, check over the office hours schedule and e-mail if you can or cannot attend at least one time

5.2. Portfolio Preparation and Maintenance

We'll spend a little time today getting your portfolio ready for the first check.

5.2.1. Access your portfolio

Go to your portfllo

- from the [course organization](#)
- from the list of your recent repositories on the left hand side of the [GitHub home page](#)

optionally, open it locally as well (we're going to update content and)

5.2.2. Start your Know, Want to Know, Learned Table

In each portfolio submission introduction, you'll reflect on what you've learned. To get ready for that, we'll first make note of what you already know and what you want to know.

1. edit `submission_1_intro` in your portfolio locally or on GitHub:
2. In the KWL section in the first two bullets after each skill with what you know and want to know. You can edit these in more detail later.

This will render as a table in your built portfolio, for reference on the syntax, [refer to the Tables section of the jupyterbook Myst Markdown cheatsheet](#).

⚠ Warning

If you work on this in the GitHub website, be sure to pull these changes locally before you start working offline next

5.2.3. Merge the setup work

Once you're done, Go to your pull request tab, and select the feedback Pull Reques. Commit any suggestions if you'd like and then merge the PR.

⚠ Warning

only do this after grading

Note

To view the feedback, after merging the PR, remove `is:open` from the search bar on the PR page

5.3. Indexing

```
topics = ['what is data science', 'jupyter',  
          'conditional', 'functions', 'lists',  
          'dictionaries', 'pandas' ]
```

What will `topics[-1]` return?

```
topics[-1]
```

```
'pandas'
```

Using negative indices starts from the right. The last element is `-1`. The first is `0`.

5.4. Reading DataFrames from Websites

We'll first read from the course website.

```
course_comms_url =  
'https://rhodyprog4ds.github.io/BrownFall21/syllabus/communication.html'
```

So far, we've read data in from a .csv file with `pd.read_csv` and created a DataFrame with the constructor `pd.DataFrame` using a dictionary. Pandas provides many interfaces for reading in data. They're described on the [Pandas IO page](#).

We can use the `read_html` method to read from this page. We know that it has multiple tables on the page, so let's see what it does:

```
pd.read_html(course_comms_url)
```

Note

Using the documentation for a library (and the base language) is totally expected and normal part of programming. That's what you should use as your primary source for questions in this class. Other sources can become outdated pretty quickly as the language changes, but most of the libraries we'll use have processes in place to ensure that their own documentation gets updated at the same time the code does.

Warning

If you use other sources and get advised to solutions that are deprecated you may not earn achievements for that work.

```

[      Day      Time      Location      Host
0      Monday  1:00:00 PM-2:30  inperson roomtbd  Chamundi
1      Wednesday  4:00:00 PM  inperson roomtbd  Chamundi
2      Wednesday  2:00:00 PM-3  inperson roomtbd  Chamundi
3      Wednesday  7:00:00 PM-8:30  gather.town  Sarah
4      Friday  5:00:00 PM-6:30pm  gather.town  Chamundi
5  By appointment      TBD  in person Tyler 134  Sarah,
      usage platform \
0      in class  prismia
1      any time  prismia
2      any time  prismia
3      private questions to your assignment  github
4      for general questions that can help others  github
5      to share resources  github
6  matters that don't fit into another category  e-mail

      area      note
0      chat  outside of class time this is not monitored cl...
1      message board  for discussion with peers
2      download transcript  use after class to get preliminary notes eg if...
3  issue on assignment repo  eg bugs in your code"
4  issue on course website  eg what the instructions of an assignment mean...
5  pull request on website  remember to request ram tokens if applicable
6  to brownsarahm@uri.edu  remember to include `[CSC310]` or `[DSP310]` (... ,
      usage      area \
0  matters that don't fit into another category  to brownsarahm@uri.edu

      note
0  remember to include `[CSC310]` or `[DSP310]` (... ,
      usage      area \
0      private questions to your assignment  issue on assignment repo
1  for general questions that can help others  issue on course website
2      to share resources  pull request on website

      note
0      eg bugs in your code"
1  eg what the instructions of an assignment mean...
2      remember to request ram tokens if applicable ,
      usage      area \
0  in class      chat
1  any time      message board
2  any time  download transcript

      note
0  outside of class time this is not monitored cl...
1      for discussion with peers
2  use after class to get preliminary notes eg if... ]

```

It appears to have read all of them, lets check the type:

```
type(pd.read_html(course_comms_url))
```

```
list
```

Since we know it's a list, we'll save it to a variable that indicates that.

```
comms_list = pd.read_html(course_comms_url)
```

If we get just the first element,

```
type(comms_list[0])
```

```
pandas.core.frame.DataFrame
```

it's a DataFrame and prints accordingly.

```
comms_list[0]
```

	Day	Time	Location	Host
0	Monday	1:00:00 PM-2:30	inperson roomtbd	Chamundi
1	Wednesday	4:00:00 PM	inperson roomtbd	Chamundi
2	Wednesday	2:00:00 PM-3	inperson roomtbd	Chamundi
3	Wednesday	7:00:00 PM-8:30	gather.town	Sarah
4	Friday	5:00:00 PM-6:30pm	gather.town	Chamundi
5	By appointment	TBD	in person Tyler 134	Sarah

Since it's a list, we can use base python's `len` function to check how many tables there are

```
len(comms_list)
```

```
5
```

We've seen the first table and know it's the help hours, so we can save that to a separate variable and use it

```
help_df = comms_list[0]
```

We've inspected the dataframe some before, but we can also check the type of each column.

```
help_df.dtypes
```

```
Day          object
Time         object
Location     object
Host         object
dtype: object
```

Note

You can read more about the [details of data types](#) in Pandas

Question: Why does it have `dtype:object` after the type for each row?

First to understand this, let's save the thing we're curious to a variable so we can examine it multiple ways more easily.

```
help_df_types = help_df.dtypes
```

Next we'll check the type of this object and its shape

```
type(help_df_types)
```

```
pandas.core.series.Series
```

a `Series` is like a dataframe, but just one row with headings, and then rotated.

```
help_df_types.shape
```

```
(4,)
```

This means that it's length is 4 and it's a 1 dimensional object; the column headers have converted to an index and are treated as metadata, but not a part of the actual data.

So, the line we're interested in is not a part of the object, because it's length 4 and the thing we're curious about is the fifth line.

We'll pick one variable from the DataFrame and check its type

```
type(help_df['Day'])
```

```
pandas.core.series.Series
```

This is also a Series, so let's check its output

```
help_df['Day']
```

```

0      Monday
1      Wednesday
2      Wednesday
3      Wednesday
4      Friday
5      By appointment
Name: Day, dtype: object

```

The last line of this one is information about the Series, its name, and its dtype.

Let's make another series, and see how it prints

```
pd.Series([5,4,5])
```

```

0    5
1    4
2    5
dtype: int64

```

The last line is the dtype of the Series; so in our original object, that last line is because the list of dtypes is the of type object.

```
help_df_types
```

```

Day      object
Time     object
Location object
Host     object
dtype: object

```

5.5. How do we know what to check?

we examined the DataFrame so far by (me) knowing what to look for.

In python objects you can programmatically find what to look for with the `__dict__` attribute or we can rely on the [online documentation](#) or use it via help.

In ipython (what we use in jupyter, by default) we can use the `?` for help

```
pd.DataFrame?
```

```
help(pd.DataFrame)
```

💡 Everything is Data

writing good documentation lets people who use your code get help for free Not only do help tools use the docs, but that website is generated programmatically using a tool called Sphinx from the documentation inside the code. You can access the docstring of a python using the `.__doc__` attribute

[ipython help](#) read about how it works, if it doesn't work for you to try to figure out why

Help on class DataFrame in module pandas.core.frame:

```
class DataFrame(pandas.core.generic.NDFrame, pandas.core.arraylike.OpsMixin)
| DataFrame(data=None, index: 'Axes | None' = None, columns: 'Axes | None' = None,
| dtype: 'Dtype | None' = None, copy: 'bool | None' = None)
|
| Two-dimensional, size-mutable, potentially heterogeneous tabular data.
|
| Data structure also contains labeled axes (rows and columns).
| Arithmetic operations align on both row and column labels. Can be
| thought of as a dict-like container for Series objects. The primary
| pandas data structure.
|
| Parameters
| -----
| data : ndarray (structured or homogeneous), Iterable, dict, or DataFrame
|       Dict can contain Series, arrays, constants, dataclass or list-like objects.
If
|       data is a dict, column order follows insertion-order.
|
|       .. versionchanged:: 0.25.0
|          If data is a list of dicts, column order follows insertion-order.
|
| index : Index or array-like
|       Index to use for resulting frame. Will default to RangeIndex if
|       no indexing information part of input data and no index provided.
| columns : Index or array-like
|       Column labels to use for resulting frame when data does not have them,
|       defaulting to RangeIndex(0, 1, 2, ..., n). If data contains column labels,
|       will perform column selection instead.
```



```

dtype : dtype, default None
    Data type to force. Only a single dtype is allowed. If None, infer.
copy : bool or None, default None
    Copy data from inputs.
    For dict data, the default of None behaves like ``copy=True``. For DataFrame
    or 2d ndarray input, the default of None behaves like ``copy=False``.

.. versionchanged:: 1.3.0

See Also
-----
DataFrame.from_records : Constructor from tuples, also record arrays.
DataFrame.from_dict : From dicts of Series, arrays, or dicts.
read_csv : Read a comma-separated values (csv) file into DataFrame.
read_table : Read general delimited file into DataFrame.
read_clipboard : Read text from clipboard into DataFrame.

Examples
-----
Constructing DataFrame from a dictionary.

>>> d = {'col1': [1, 2], 'col2': [3, 4]}
>>> df = pd.DataFrame(data=d)
>>> df
   col1  col2
0      1      3
1      2      4

Notice that the inferred dtype is int64.

>>> df.dtypes
col1    int64
col2    int64
dtype: object

To enforce a single dtype:

>>> df = pd.DataFrame(data=d, dtype=np.int8)
>>> df.dtypes
col1    int8
col2    int8
dtype: object

Constructing DataFrame from numpy ndarray:

>>> df2 = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6], [7, 8, 9]]),
...                     columns=['a', 'b', 'c'])
>>> df2
   a  b  c
0  1  2  3
1  4  5  6
2  7  8  9

Constructing DataFrame from a numpy ndarray that has labeled columns:

>>> data = np.array([(1, 2, 3), (4, 5, 6), (7, 8, 9)],
...                  dtype=[("a", "i4"), ("b", "i4"), ("c", "i4")])
>>> df3 = pd.DataFrame(data, columns=['c', 'a'])
...
>>> df3
   c  a
0  3  1
1  6  4
2  9  7

Constructing DataFrame from dataclass:

>>> from dataclasses import make_dataclass
>>> Point = make_dataclass("Point", [("x", int), ("y", int)])
>>> pd.DataFrame([Point(0, 0), Point(0, 3), Point(2, 3)])
   x  y
0  0  0
1  0  3
2  2  3

Method resolution order:
    DataFrame
    pandas.core.generic.NDFrame
    pandas.core.base.PandasObject
    pandas.core.accessor.DirNamesMixin
    pandas.core.indexing.IndexingMixin
    pandas.core.arraylike.OpsMixin
    builtins.object

Methods defined here:

__divmod__(self, other) -> 'tuple[DataFrame, DataFrame]'

```

```

__getitem__(self, key)

__init__(self, data=None, index: 'Axes | None' = None, columns: 'Axes | None' =
None, dtype: 'Dtype | None' = None, copy: 'bool | None' = None)
    Initialize self. See help(type(self)) for accurate signature.

__len__(self) -> 'int'
    Returns length of info axis, but here we use the index.

__matmul__(self, other: 'AnyArrayLike | FrameOrSeriesUnion') ->
'FrameOrSeriesUnion'
    Matrix multiplication using binary `@` operator in Python>=3.5.

__rdivmod__(self, other) -> 'tuple[DataFrame, DataFrame]'

__repr__(self) -> 'str'
    Return a string representation for a particular DataFrame.

__rmatmul__(self, other)
    Matrix multiplication using binary `@` operator in Python>=3.5.

__setitem__(self, key, value)

add(self, other, axis='columns', level=None, fill_value=None)
    Get Addition of dataframe and other, element-wise (binary operator `add`).

    Equivalent to ``dataframe + other``, but with support to substitute a
fill_value
    for missing data in one of the inputs. With reverse version, `radd`.

    Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
    arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters
-----
other : scalar, sequence, Series, or DataFrame
    Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
    Whether to compare by the index (0 or 'index') or columns
    (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
    Broadcast across a level, matching Index values on the
    passed MultiIndex level.
fill_value : float or None, default None
    Fill existing missing (NaN) values, and any new element needed for
    successful DataFrame alignment, with this value before computation.
    If data in both corresponding DataFrame locations is missing
    the result will be missing.

Returns
-----
DataFrame
    Result of the arithmetic operation.

See Also
-----
DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes
-----
Mismatched indices will be unioned together.

Examples
-----
>>> df = pd.DataFrame({'angles': [0, 3, 4],
...                     'degrees': [360, 180, 360]},
...                     index=['circle', 'triangle', 'rectangle'])
>>> df
   angles  degrees
circle      0     360
triangle    3     180
rectangle   4     360

Add a scalar with operator version which return the same
results.

>>> df + 1
   angles  degrees
circle      1     361

```

triangle	4	181
rectangle	5	361

```
>>> df.add(1)
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

Divide by constant with reverse version.

```
>>> df.div(10)
      angles  degrees
circle    0.0     36.0
triangle  0.3     18.0
rectangle 0.4     36.0
```

```
>>> df.rdiv(10)
      angles  degrees
circle      inf  0.027778
triangle  3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
      angles  degrees
circle     -1     358
triangle    2     178
rectangle   3     358
```

```
>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle     -1     358
triangle    2     178
rectangle   3     358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
      angles  degrees
circle     -1     359
triangle    2     179
rectangle   3     359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                       index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle     3
rectangle    4
```

```
>>> df * other
      angles  degrees
circle      0      NaN
triangle     9      NaN
rectangle   16      NaN
```

```
>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0      0.0
triangle     9      0.0
rectangle   16      0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                               'degrees': [360, 180, 360, 360, 540, 720]},
...                              index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                      ['circle', 'triangle', 'rectangle',
...                                       'square', 'pentagon', 'hexagon']])
>>> df_multindex
      angles  degrees
A circle      0     360
  triangle    3     180
  rectangle    4     360
B square      4     360
  pentagon    5     540
  hexagon     6     720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
      angles  degrees
A circle      NaN      1.0
  triangle    1.0      1.0
  rectangle    1.0      1.0
```

B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

```
agg = aggregate(self, func=None, axis: 'Axis' = 0, *args, **kwargs)
```

```
aggregate(self, func=None, axis: 'Axis' = 0, *args, **kwargs)
```

Aggregate using one or more operations over the specified axis.

Parameters

func : function, str, list or dict

Function to use for aggregating the data. If a function, must either work when passed a DataFrame or when passed to DataFrame.apply.

Accepted combinations are:

- function
- string function name
- list of functions and/or function names, e.g. ``[np.sum, 'mean']``
- dict of axis labels -> functions, function names or list of such.

axis : {0 or 'index', 1 or 'columns'}, default 0

If 0 or 'index': apply function to each column.

If 1 or 'columns': apply function to each row.

***args**

Positional arguments to pass to `func`.

****kwargs**

Keyword arguments to pass to `func`.

Returns

scalar, Series or DataFrame

The return can be:

- * scalar : when Series.agg is called with single function
- * Series : when DataFrame.agg is called with a single function
- * DataFrame : when DataFrame.agg is called with several functions

Return scalar, Series or DataFrame.

The aggregation operations are always performed over an axis, either the index (default) or the column axis. This behavior is different from `numpy` aggregation functions (`mean`, `median`, `prod`, `sum`, `std`, `var`), where the default is to compute the aggregation of the flattened array, e.g., ``numpy.mean(arr_2d)`` as opposed to ``numpy.mean(arr_2d, axis=0)``.

`agg` is an alias for `aggregate`. Use the alias.

See Also

DataFrame.apply : Perform any type of operations.

DataFrame.transform : Perform transformation type operations.

core.groupby.GroupBy : Perform operations over groups.

core.resample.Resampler : Perform operations over resampled bins.

core.window.Rolling : Perform operations over rolling window.

core.window.Expanding : Perform operations over expanding window.

core.window.ExponentialMovingWindow : Perform operation over exponential

weighted

window.

Notes

`agg` is an alias for `aggregate`. Use the alias.

Functions that mutate the passed object can produce unexpected behavior or errors and are not supported. See :ref:`gotchas.udf-mutation` for more details.

A passed user-defined-function will be passed a Series for evaluation.

Examples

```
>>> df = pd.DataFrame([[1, 2, 3],
...                     [4, 5, 6],
...                     [7, 8, 9],
...                     [np.nan, np.nan, np.nan]],
...                    columns=['A', 'B', 'C'])
```

Aggregate these functions over the rows.

```
>>> df.agg(['sum', 'min'])
```

	A	B	C
sum	12.0	15.0	18.0
min	1.0	2.0	3.0

Different aggregations per column.

```
>>> df.agg({'A' : ['sum', 'min'], 'B' : ['min', 'max']})
      A      B
sum  12.0  NaN
min   1.0   2.0
max   NaN   8.0
```

Aggregate different functions over the columns and rename the index of the resulting DataFrame.

```
>>> df.agg(x=('A', max), y=('B', 'min'), z=('C', np.mean))
      A      B      C
x  7.0  NaN  NaN
y  NaN  2.0  NaN
z  NaN  NaN  6.0
```

Aggregate over the columns.

```
>>> df.agg("mean", axis="columns")
0      2.0
1      5.0
2      8.0
3      NaN
dtype: float64
```

align(self, other, join: 'str' = 'outer', axis: 'Axis | None' = None, level: 'Level | None' = None, copy: 'bool' = True, fill_value=None, method: 'str | None' = None, limit=None, fill_axis: 'Axis' = 0, broadcast_axis: 'Axis | None' = None) -> 'DataFrame'

Align two objects on their axes with the specified join method.

Join method is specified for each axis Index.

Parameters

other : DataFrame or Series

join : {'outer', 'inner', 'left', 'right'}, default 'outer'

axis : allowed axis of the other object, default None

Align on index (0), columns (1), or both (None).

level : int or level name, default None

Broadcast across a level, matching Index values on the passed MultiIndex level.

copy : bool, default True

Always returns new objects. If copy=False and no reindexing is required then original objects are returned.

fill_value : scalar, default np.NaN

Value to use for missing values. Defaults to NaN, but can be any "compatible" value.

method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None

Method to use for filling holes in reindexed Series:

- pad / ffill: propagate last valid observation forward to next valid.
- backfill / bfill: use NEXT valid observation to fill gap.

limit : int, default None

If method is specified, this is the maximum number of consecutive NaN values to forward/backward fill. In other words, if there is a gap with more than this number of consecutive NaNs, it will only be partially filled. If method is not specified, this is the maximum number of entries along the entire axis where NaNs will be filled. Must be greater than 0 if not None.

fill_axis : {0 or 'index', 1 or 'columns'}, default 0

Filling axis, method and limit.

broadcast_axis : {0 or 'index', 1 or 'columns'}, default None

Broadcast values along this axis, if aligning two objects of different dimensions.

Returns

(left, right) : (DataFrame, type of other)

Aligned objects.

all(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
Return whether all elements are True, potentially over an axis.

Returns True unless there at least one element within a series or along a Dataframe axis that is False or equivalent (e.g. zero or empty).

Parameters

axis : {0 or 'index', 1 or 'columns', None}, default 0

Indicate which axis or axes should be reduced.

* 0 / 'index' : reduce the index, return a Series whose index is the

```

    original column labels.
    * 1 / 'columns' : reduce the columns, return a Series whose index is the
    original index.
    * None : reduce all axes, return a scalar.

bool_only : bool, default None
    Include only boolean columns. If None, will attempt to use everything,
    then use only boolean data. Not implemented for Series.
skipna : bool, default True
    Exclude NA/null values. If the entire row/column is NA and skipna is
    True, then the result will be True, as for an empty row/column.
    If skipna is False, then NA are treated as True, because these are not
    equal to zero.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
**kwargs : any, default None
    Additional keywords have no effect but might be accepted for
    compatibility with NumPy.

Returns
-----
Series or DataFrame
    If level is specified, then, DataFrame is returned; otherwise, Series
    is returned.

See Also
-----
Series.all : Return True if all elements are True.
DataFrame.any : Return True if one (or more) elements are True.

Examples
-----
**Series**

>>> pd.Series([True, True]).all()
True
>>> pd.Series([True, False]).all()
False
>>> pd.Series([], dtype="float64").all()
True
>>> pd.Series([np.nan]).all()
True
>>> pd.Series([np.nan]).all(skipna=False)
True

**DataFrames**

Create a dataframe from a dictionary.

>>> df = pd.DataFrame({'col1': [True, True], 'col2': [True, False]})
>>> df
   col1  col2
0   True   True
1   True  False

Default behaviour checks if column-wise values all return True.

>>> df.all()
col1    True
col2   False
dtype: bool

Specify ``axis='columns'`` to check if row-wise values all return True.

>>> df.all(axis='columns')
0     True
1    False
dtype: bool

Or ``axis=None`` for whether every value is True.

>>> df.all(axis=None)
False

any(self, axis=0, bool_only=None, skipna=True, level=None, **kwargs)
    Return whether any element is True, potentially over an axis.

Returns False unless there is at least one element within a series or
along a DataFrame axis that is True or equivalent (e.g. non-zero or
non-empty).

Parameters
-----
axis : {0 or 'index', 1 or 'columns', None}, default 0
    Indicate which axis or axes should be reduced.

```

```

    * 0 / 'index' : reduce the index, return a Series whose index is the
        original column labels.
    * 1 / 'columns' : reduce the columns, return a Series whose index is the
        original index.
    * None : reduce all axes, return a scalar.

bool_only : bool, default None
    Include only boolean columns. If None, will attempt to use everything,
    then use only boolean data. Not implemented for Series.
skipna : bool, default True
    Exclude NA/null values. If the entire row/column is NA and skipna is
    True, then the result will be False, as for an empty row/column.
    If skipna is False, then NA are treated as True, because these are not
    equal to zero.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
**kwargs : any, default None
    Additional keywords have no effect but might be accepted for
    compatibility with NumPy.

Returns
-----
Series or DataFrame
    If level is specified, then, DataFrame is returned; otherwise, Series
    is returned.

See Also
-----
numpy.any : Numpy version of this method.
Series.any : Return whether any element is True.
Series.all : Return whether all elements are True.
DataFrame.any : Return whether any element is True over requested axis.
DataFrame.all : Return whether all elements are True over requested axis.

Examples
-----
**Series**

For Series input, the output is a scalar indicating whether any element
is True.

>>> pd.Series([False, False]).any()
False
>>> pd.Series([True, False]).any()
True
>>> pd.Series([], dtype="float64").any()
False
>>> pd.Series([np.nan]).any()
False
>>> pd.Series([np.nan]).any(skipna=False)
True

**DataFrame**

Whether each column contains at least one True element (the default).

>>> df = pd.DataFrame({"A": [1, 2], "B": [0, 2], "C": [0, 0]})
>>> df
   A  B  C
0  1  0  0
1  2  2  0

>>> df.any()
A    True
B    True
C   False
dtype: bool

Aggregating over the columns.

>>> df = pd.DataFrame({"A": [True, False], "B": [1, 2]})
>>> df
   A  B
0  True  1
1 False  2

>>> df.any(axis='columns')
0    True
1    True
dtype: bool

>>> df = pd.DataFrame({"A": [True, False], "B": [1, 0]})
>>> df
   A  B
0  True  1
1 False  0

```

```

>>> df.any(axis='columns')
0    True
1    False
dtype: bool

Aggregating over the entire DataFrame with ``axis=None``.

>>> df.any(axis=None)
True

`any` for an empty DataFrame is an empty Series.

>>> pd.DataFrame([]).any()
Series([], dtype: bool)

```

append(self, other, ignore_index: 'bool' = False, verify_integrity: 'bool' = False, sort: 'bool' = False) -> 'DataFrame'

Append rows of `other` to the end of caller, returning a new object.

Columns in `other` that are not in the caller are added as new columns.

Parameters

other : DataFrame or Series/dict-like object, or list of these
The data to append.

ignore_index : bool, default False
If True, the resulting axis will be labeled 0, 1, ..., n - 1.

verify_integrity : bool, default False
If True, raise ValueError on creating index with duplicates.

sort : bool, default False
Sort columns if the columns of `self` and `other` are not aligned.

.. versionchanged:: 1.0.0

Changed to not sort by default.

Returns

DataFrame

A new DataFrame consisting of the rows of caller and the rows of `other`.

See Also

concat : General function to concatenate DataFrame or Series objects.

Notes

If a list of dict/series is passed and the keys are all contained in the DataFrame's index, the order of the columns in the resulting DataFrame will be unchanged.

Iteratively appending rows to a DataFrame can be more computationally intensive than a single concatenate. A better solution is to append those rows to a list and then concatenate the list with the original DataFrame all at once.

Examples

```

>>> df = pd.DataFrame([[1, 2], [3, 4]], columns=list('AB'), index=['x', 'y'])
>>> df
   A  B
x  1  2
y  3  4
>>> df2 = pd.DataFrame([[5, 6], [7, 8]], columns=list('AB'), index=['x',
'y'])
>>> df.append(df2)
   A  B
x  1  2
y  3  4
x  5  6
y  7  8

```

With `ignore_index` set to True:

```

>>> df.append(df2, ignore_index=True)
   A  B
0  1  2
1  3  4
2  5  6
3  7  8

```

The following, while not recommended methods for generating DataFrames, show two ways to generate a DataFrame from multiple data sources.

Less efficient:


```
>>> df = pd.DataFrame(columns=['A'])
>>> for i in range(5):
...     df = df.append({'A': i}, ignore_index=True)
>>> df
   A
0  0
1  1
2  2
3  3
4  4
```

More efficient:

```
>>> pd.concat([pd.DataFrame([i], columns=['A']) for i in range(5)],
...            ignore_index=True)
   A
0  0
1  1
2  2
3  3
4  4
```

```
apply(self, func: 'AggFuncType', axis: 'Axis' = 0, raw: 'bool' = False,
result_type=None, args=(), **kwargs)
```

Apply a function along an axis of the DataFrame.

Objects passed to the function are Series objects whose index is either the DataFrame's index (`axis=0`) or the DataFrame's columns (`axis=1`). By default (`result_type=None`), the final return type is inferred from the return type of the applied function. Otherwise, it depends on the `result_type` argument.

Parameters

`func` : function

Function to apply to each column or row.

`axis` : {0 or 'index', 1 or 'columns'}, default 0

Axis along which the function is applied:

* 0 or 'index': apply function to each column.

* 1 or 'columns': apply function to each row.

`raw` : bool, default False

Determines if row or column is passed as a Series or ndarray object:

* `False` : passes each row or column as a Series to the function.

* `True` : the passed function will receive ndarray objects instead.

If you are just applying a NumPy reduction function this will achieve much better performance.

`result_type` : {'expand', 'reduce', 'broadcast', None}, default None

These only act when `axis=1` (columns):

* 'expand' : list-like results will be turned into columns.

* 'reduce' : returns a Series if possible rather than expanding list-like results. This is the opposite of 'expand'.

* 'broadcast' : results will be broadcast to the original shape of the DataFrame, the original index and columns will be retained.

The default behaviour (None) depends on the return value of the applied function: list-like results will be returned as a Series of those. However if the apply function returns a Series these are expanded to columns.

`args` : tuple

Positional arguments to pass to `func` in addition to the array/series.

`**kwargs`

Additional keyword arguments to pass as keywords arguments to `func`.

Returns

Series or DataFrame

Result of applying `func` along the given axis of the DataFrame.

See Also

`DataFrame.applymap`: For elementwise operations.

`DataFrame.aggregate`: Only perform aggregating type operations.

`DataFrame.transform`: Only perform transforming type operations.

Notes

Functions that mutate the passed object can produce unexpected behavior or errors and are not supported. See :ref:`gotchas.udf-mutation` for more details.

Examples

```
-----
>>> df = pd.DataFrame([[4, 9]] * 3, columns=['A', 'B'])
>>> df
   A  B
0  4  9
1  4  9
2  4  9
```

Using a numpy universal function (in this case the same as ``np.sqrt(df)``):

```
>>> df.apply(np.sqrt)
   A  B
0  2.0  3.0
1  2.0  3.0
2  2.0  3.0
```

Using a reducing function on either axis

```
>>> df.apply(np.sum, axis=0)
A    12
B     27
dtype: int64
```

```
>>> df.apply(np.sum, axis=1)
0     13
1     13
2     13
dtype: int64
```

Returning a list-like will result in a Series

```
>>> df.apply(lambda x: [1, 2], axis=1)
0    [1, 2]
1    [1, 2]
2    [1, 2]
dtype: object
```

Passing ``result_type='expand'`` will expand list-like results to columns of a Dataframe

```
>>> df.apply(lambda x: [1, 2], axis=1, result_type='expand')
   0  1
0  1  2
1  1  2
2  1  2
```

Returning a Series inside the function is similar to passing ``result_type='expand'``. The resulting column names will be the Series index.

```
>>> df.apply(lambda x: pd.Series([1, 2], index=['foo', 'bar']), axis=1)
   foo  bar
0    1    2
1    1    2
2    1    2
```

Passing ``result_type='broadcast'`` will ensure the same shape result, whether list-like or scalar is returned by the function, and broadcast it along the axis. The resulting column names will be the originals.

```
>>> df.apply(lambda x: [1, 2], axis=1, result_type='broadcast')
   A  B
0  1  2
1  1  2
2  1  2
```

`applymap(self, func: 'PythonFuncType', na_action: 'str | None' = None, **kwargs)`
-> 'DataFrame'

Apply a function to a Dataframe elementwise.

This method applies a function that accepts and returns a scalar to every element of a DataFrame.

Parameters

func : callable

Python function, returns a single value from a single value.

na_action : {None, 'ignore'}, default None

If 'ignore', propagate NaN values, without passing them to func.

```

.. versionadded:: 1.2

**kwargs
    Additional keyword arguments to pass as keywords arguments to
    `func`.

.. versionadded:: 1.3.0

Returns
-----
DataFrame
    Transformed DataFrame.

See Also
-----
DataFrame.apply : Apply a function along input axis of DataFrame.

Examples
-----
>>> df = pd.DataFrame([[1, 2.12], [3.356, 4.567]])
>>> df
   0      1
0  1.000  2.120
1  3.356  4.567

>>> df.applymap(lambda x: len(str(x)))
   0  1
0  3  4
1  5  5

Like Series.map, NA values can be ignored:

>>> df_copy = df.copy()
>>> df_copy.iloc[0, 0] = pd.NA
>>> df_copy.applymap(lambda x: len(str(x)), na_action='ignore')
   0  1
0 <NA> 4
1  5  5

Note that a vectorized version of `func` often exists, which will
be much faster. You could square each number elementwise.

>>> df.applymap(lambda x: x**2)
   0      1
0  1.000000  4.494400
1 11.262736 20.857489

But it's better to avoid applymap in that case.

>>> df ** 2
   0      1
0  1.000000  4.494400
1 11.262736 20.857489

asfreq(self, freq: 'Frequency', method=None, how: 'str | None' = None, normalize:
'bool' = False, fill_value=None) -> 'DataFrame'
    Convert time series to specified frequency.

Returns the original data conformed to a new index with the specified
frequency.

If the index of this DataFrame is a :class:`~pandas.PeriodIndex`, the new
index
is the result of transforming the original index with
:meth:`~PeriodIndex.asfreq` <pandas.PeriodIndex.asfreq>` (so the original index
will map one-to-one to the new index).

Otherwise, the new index will be equivalent to ``pd.date_range(start, end,
freq=freq)`` where ``start`` and ``end`` are, respectively, the first and
last entries in the original index (see :func:`~pandas.date_range`). The
values corresponding to any timesteps in the new index which were not present
in the original index will be null (``NaN``), unless a method for filling
such unknowns is provided (see the ``method`` parameter below).

The :meth:`~resample` method is more appropriate if an operation on each group
of
timesteps (such as an aggregate) is necessary to represent the data at the
new
frequency.

Parameters
-----
freq : DateOffset or str
    Frequency DateOffset or string.
method : {'backfill', 'bfill', 'pad', 'ffill'}, default None
    Method to use for filling holes in reindexed Series (note this
    does not fill NaNs that already were present):

```

```

    * 'pad' / 'ffill': propagate last valid observation forward to next
      valid
    * 'backfill' / 'bfill': use NEXT valid observation to fill.
how : {'start', 'end'}, default end
      For PeriodIndex only (see PeriodIndex.asfreq).
normalize : bool, default False
      Whether to reset output index to midnight.
fill_value : scalar, optional
      Value to use for missing values, applied during upsampling (note
      this does not fill NaNs that already were present).

Returns
-----
DataFrame
    DataFrame object reindexed to the specified frequency.

See Also
-----
reindex : Conform DataFrame to new index with optional filling logic.

Notes
-----
To learn more about the frequency strings, please see `this link
<https://pandas.pydata.org/pandas-
docs/stable/user\_guide/timeseries.html#offset-aliases>`_.

```

Examples

Start by creating a series with 4 one minute timestamps.

```

>>> index = pd.date_range('1/1/2000', periods=4, freq='T')
>>> series = pd.Series([0.0, None, 2.0, 3.0], index=index)
>>> df = pd.DataFrame({'s': series})
>>> df

```

```

      s
2000-01-01 00:00:00    0.0
2000-01-01 00:01:00   NaN
2000-01-01 00:02:00    2.0
2000-01-01 00:03:00    3.0

```

Upsample the series into 30 second bins.

```

>>> df.asfreq(freq='30S')
      s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30   NaN
2000-01-01 00:01:00   NaN
2000-01-01 00:01:30   NaN
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30   NaN
2000-01-01 00:03:00    3.0

```

Upsample again, providing a ``fill value``.

```

>>> df.asfreq(freq='30S', fill_value=9.0)
      s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30    9.0
2000-01-01 00:01:00   NaN
2000-01-01 00:01:30    9.0
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30    9.0
2000-01-01 00:03:00    3.0

```

Upsample again, providing a ``method``.

```

>>> df.asfreq(freq='30S', method='bfill')
      s
2000-01-01 00:00:00    0.0
2000-01-01 00:00:30   NaN
2000-01-01 00:01:00   NaN
2000-01-01 00:01:30    2.0
2000-01-01 00:02:00    2.0
2000-01-01 00:02:30    3.0
2000-01-01 00:03:00    3.0

```

```

assign(self, **kwargs) -> 'DataFrame'
    Assign new columns to a DataFrame.

```

Returns a new object with all original columns in addition to new ones.
Existing columns that are re-assigned will be overwritten.

Parameters

```

**kwargs : dict of {str: callable or Series}
    The column names are keywords. If the values are

```

callable, they are computed on the DataFrame and assigned to the new columns. The callable must not change input DataFrame (though pandas doesn't check it). If the values are not callable, (e.g. a Series, scalar, or array), they are simply assigned.

Returns

DataFrame

A new DataFrame with the new columns in addition to all the existing columns.

Notes

Assigning multiple columns within the same ``assign`` is possible. Later items in ``**kwargs`` may refer to newly created or modified columns in 'df'; items are computed and assigned into 'df' in order.

Examples

```
>>> df = pd.DataFrame({'temp_c': [17.0, 25.0]},
...                    index=['Portland', 'Berkeley'])
>>> df
```

```
      temp_c
Portland  17.0
Berkeley  25.0
```

Where the value is a callable, evaluated on `df`:

```
>>> df.assign(temp_f=lambda x: x.temp_c * 9 / 5 + 32)
      temp_c  temp_f
Portland   17.0    62.6
Berkeley   25.0    77.0
```

Alternatively, the same behavior can be achieved by directly referencing an existing Series or sequence:

```
>>> df.assign(temp_f=df['temp_c'] * 9 / 5 + 32)
      temp_c  temp_f
Portland   17.0    62.6
Berkeley   25.0    77.0
```

You can create multiple columns within the same assign where one of the columns depends on another one defined within the same assign:

```
>>> df.assign(temp_f=lambda x: x['temp_c'] * 9 / 5 + 32,
...           temp_k=lambda x: (x['temp_f'] + 459.67) * 5 / 9)
      temp_c  temp_f  temp_k
Portland   17.0    62.6   290.15
Berkeley   25.0    77.0   298.15
```

bfill(self: 'DataFrame', axis: 'None | Axis' = None, inplace: 'bool' = False, limit: 'None | int' = None, downcast=None) -> 'DataFrame | None'
Synonym for :meth:`DataFrame.fillna` with ``method='bfill'``.

Returns

Series/DataFrame or None

Object with missing values filled or None if ``inplace=True``.

boxplot = boxplot_frame(self, column=None, by=None, ax=None, fontsize=None, rot=0, grid=True, figsize=None, layout=None, return_type=None, backend=None, **kwargs)

Make a box plot from DataFrame columns.

Make a box-and-whisker plot from DataFrame columns, optionally grouped by some other columns. A box plot is a method for graphically depicting groups of numerical data through their quartiles.

The box extends from the Q1 to Q3 quartile values of the data, with a line at the median (Q2). The whiskers extend from the edges of box to show the range of the data. By default, they extend no more than $1.5 * IQR$ ($IQR = Q3 - Q1$) from the edges of the box, ending at the farthest data point within that interval. Outliers are plotted as separate dots.

For further details see

Wikipedia's entry for `boxplot` <https://en.wikipedia.org/wiki/Box_plot>`_.

Parameters

column : str or list of str, optional

Column name or list of names, or vector.

Can be any valid input to :meth:`pandas.DataFrame.groupby`.

by : str or array-like, optional

Column in the DataFrame to :meth:`pandas.DataFrame.groupby`.

One box-plot will be done per value of columns in `by`.

ax : object of class matplotlib.axes.Axes, optional

The matplotlib axes to be used by boxplot.

```

fontsize : float or str
    Tick label font size in points or as a string (e.g., 'large').
rot : int or float, default 0
    The rotation angle of labels (in degrees)
    with respect to the screen coordinate system.
grid : bool, default True
    Setting this to True will show the grid.
figsize : A tuple (width, height) in inches
    The size of the figure to create in matplotlib.
layout : tuple (rows, columns), optional
    For example, (3, 5) will display the subplots
    using 3 columns and 5 rows, starting from the top-left.
return_type : {'axes', 'dict', 'both'} or None, default 'axes'
    The kind of object to return. The default is ``axes``.

    * 'axes' returns the matplotlib axes the boxplot is drawn on.
    * 'dict' returns a dictionary whose values are the matplotlib
      Lines of the boxplot.
    * 'both' returns a namedtuple with the axes and dict.
    * when grouping with ``by``, a Series mapping columns to
      ``return_type`` is returned.

    If ``return_type`` is `None`, a NumPy array
    of axes with the same shape as ``layout`` is returned.
backend : str, default None
    Backend to use instead of the backend specified in the option
    ``plotting.backend``. For instance, 'matplotlib'. Alternatively, to
    specify the ``plotting.backend`` for the whole session, set
    ``pd.options.plotting.backend``.

.. versionadded:: 1.0.0

**kwargs
    All other plotting keyword arguments to be passed to
    :func:`matplotlib.pyplot.boxplot`.

Returns
-----
result
    See Notes.

See Also
-----
Series.plot.hist: Make a histogram.
matplotlib.pyplot.boxplot : Matplotlib equivalent plot.

Notes
-----
The return type depends on the `return_type` parameter:

* 'axes' : object of class matplotlib.axes.Axes
* 'dict' : dict of matplotlib.lines.Line2D objects
* 'both' : a namedtuple with structure (ax, lines)

For data grouped with ``by``, return a Series of the above or a numpy
array:

* :class:`~pandas.Series`
* :class:`~numpy.array` (for ``return_type = None``)

Use ``return_type='dict'`` when you want to tweak the appearance
of the lines after plotting. In this case a dict containing the Lines
making up the boxes, caps, fliers, medians, and whiskers is returned.

Examples
-----

Boxplots can be created for every column in the dataframe
by ``df.boxplot()`` or indicating the columns to be used:

.. plot::
   :context: close-figs

   >>> np.random.seed(1234)
   >>> df = pd.DataFrame(np.random.randn(10, 4),
   ...                   columns=['Col1', 'Col2', 'Col3', 'Col4'])
   >>> boxplot = df.boxplot(column=['Col1', 'Col2', 'Col3'])

Boxplots of variables distributions grouped by the values of a third
variable can be created using the option ``by``. For instance:

.. plot::
   :context: close-figs

   >>> df = pd.DataFrame(np.random.randn(10, 2),
   ...                   columns=['Col1', 'Col2'])
   >>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',

```

```
... 'B', 'B', 'B', 'B', 'B'])
>>> boxplot = df.boxplot(by='X')
```

A list of strings (i.e. ``['X', 'Y']``) can be passed to boxplot in order to group the data by combination of the variables in the x-axis:

```
.. plot::
:context: close-figs

>>> df = pd.DataFrame(np.random.randn(10, 3),
...                     columns=['Col1', 'Col2', 'Col3'])
>>> df['X'] = pd.Series(['A', 'A', 'A', 'A', 'A',
...                     'B', 'B', 'B', 'B', 'B'])
>>> df['Y'] = pd.Series(['A', 'B', 'A', 'B', 'A',
...                     'B', 'A', 'B', 'A', 'B'])
>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by=['X', 'Y'])
```

The layout of boxplot can be adjusted giving a tuple to ``layout``:

```
.. plot::
:context: close-figs

>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       layout=(2, 1))
```

Additional formatting can be done to the boxplot, like suppressing the grid (``grid=False``), rotating the labels in the x-axis (i.e. ``rot=45``) or changing the fontsize (i.e. ``fontsize=15``):

```
.. plot::
:context: close-figs

>>> boxplot = df.boxplot(grid=False, rot=45, fontsize=15)
```

The parameter ``return_type`` can be used to select the type of element returned by ``boxplot``. When ``return_type='axes'`` is selected, the matplotlib axes on which the boxplot is drawn are returned:

```
>>> boxplot = df.boxplot(column=['Col1', 'Col2'], return_type='axes')
>>> type(boxplot)
<class 'matplotlib.axes._subplots.AxesSubplot'>
```

When grouping with ``by``, a Series mapping columns to ``return_type`` is returned:

```
>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       return_type='axes')
>>> type(boxplot)
<class 'pandas.core.series.Series'>
```

If ``return_type`` is ``None``, a NumPy array of axes with the same shape as ``layout`` is returned:

```
>>> boxplot = df.boxplot(column=['Col1', 'Col2'], by='X',
...                       return_type=None)
>>> type(boxplot)
<class 'numpy.ndarray'>
```

```
clip(self: 'DataFrame', lower=None, upper=None, axis: 'Axis | None' = None,
inplace: 'bool' = False, *args, **kwargs) -> 'DataFrame | None'
Trim values at input threshold(s).
```

Assigns values outside boundary to boundary values. Thresholds can be singular values or array like, and in the latter case the clipping is performed element-wise in the specified axis.

Parameters

```
-----
lower : float or array-like, default None
    Minimum threshold value. All values below this
    threshold will be set to it. A missing
    threshold (e.g. 'NA') will not clip the value.
upper : float or array-like, default None
    Maximum threshold value. All values above this
    threshold will be set to it. A missing
    threshold (e.g. 'NA') will not clip the value.
axis : int or str axis name, optional
    Align object with lower and upper along the given axis.
inplace : bool, default False
    Whether to perform the operation in place on the data.
*args, **kwargs
    Additional keywords have no effect but might be accepted
    for compatibility with numpy.
```

Returns

```
-----
Series or DataFrame or None
```

Same type as calling object with the values outside the clip boundaries replaced or None if ``inplace=True``.

See Also

Series.clip : Trim values at input threshold in series.
DataFrame.clip : Trim values at input threshold in dataframe.
numpy.clip : Clip (limit) the values in an array.

Examples

>>> data = {'col_0': [9, -3, 0, -1, 5], 'col_1': [-2, -7, 6, 8, -5]}
>>> df = pd.DataFrame(data)
>>> df
 col_0 col_1
0 9 -2
1 -3 -7
2 0 6
3 -1 8
4 5 -5

Clips per column using lower and upper thresholds:

>>> df.clip(-4, 6)
 col_0 col_1
0 6 -2
1 -3 -4
2 0 6
3 -1 6
4 5 -4

Clips using specific lower and upper thresholds per column element:

>>> t = pd.Series([2, -4, -1, 6, 3])
>>> t
0 2
1 -4
2 -1
3 6
4 3
dtype: int64

>>> df.clip(t, t + 4, axis=0)
 col_0 col_1
0 6 2
1 -3 -4
2 0 3
3 6 8
4 5 3

Clips using specific lower threshold per column element, with missing values:

>>> t = pd.Series([2, -4, np.NaN, 6, 3])
>>> t
0 2.0
1 -4.0
2 NaN
3 6.0
4 3.0
dtype: float64

>>> df.clip(t, axis=0)
 col_0 col_1
0 9 2
1 -3 -4
2 0 6
3 6 8
4 5 3

combine(self, other: 'DataFrame', func, fill_value=None, overwrite: 'bool' = True) -> 'DataFrame'

Perform column-wise combine with another DataFrame.

Combines a DataFrame with `other` DataFrame using `func` to element-wise combine columns. The row and column indexes of the resulting DataFrame will be the union of the two.

Parameters

other : DataFrame

The DataFrame to merge column-wise.

func : function

Function that takes two series as inputs and return a Series or a scalar. Used to merge the two dataframes column by columns.

fill_value : scalar value, default None

The value to fill NaNs with prior to passing any column to the merge func.


```

overwrite : bool, default True
    If True, columns in `self` that do not exist in `other` will be
    overwritten with NaNs.

Returns
-----
DataFrame
    Combination of the provided DataFrames.

See Also
-----
DataFrame.combine_first : Combine two DataFrame objects and default to
    non-null values in frame calling the method.

Examples
-----
Combine using a simple function that chooses the smaller column.

>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> take_smaller = lambda s1, s2: s1 if s1.sum() < s2.sum() else s2
>>> df1.combine(df2, take_smaller)
   A  B
0  0  3
1  0  3

Example using a true element-wise combine function.

>>> df1 = pd.DataFrame({'A': [5, 0], 'B': [2, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine(df2, np.minimum)
   A  B
0  1  2
1  0  3

Using `fill_value` fills Nones prior to passing the column to the
merge function.

>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine(df2, take_smaller, fill_value=-5)
   A  B
0  0 -5.0
1  0  4.0

However, if the same element in both dataframes is None, that None
is preserved

>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [None, 3]})
>>> df1.combine(df2, take_smaller, fill_value=-5)
   A  B
0  0 -5.0
1  0  3.0

Example that demonstrates the use of `overwrite` and behavior when
the axis differ between the dataframes.

>>> df1 = pd.DataFrame({'A': [0, 0], 'B': [4, 4]})
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [-10, 1], }, index=[1, 2])
>>> df1.combine(df2, take_smaller)
   A  B  C
0 NaN NaN NaN
1 NaN 3.0 -10.0
2 NaN 3.0  1.0

>>> df1.combine(df2, take_smaller, overwrite=False)
   A  B  C
0 0.0 NaN NaN
1 0.0 3.0 -10.0
2 NaN 3.0  1.0

Demonstrating the preference of the passed in dataframe.

>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1], }, index=[1, 2])
>>> df2.combine(df1, take_smaller)
   A  B  C
0 0.0 NaN NaN
1 0.0 3.0 NaN
2 NaN 3.0 NaN

>>> df2.combine(df1, take_smaller, overwrite=False)
   A  B  C
0 0.0 NaN NaN
1 0.0 3.0 1.0
2 NaN 3.0 1.0

```

```

combine_first(self, other: 'DataFrame') -> 'DataFrame'
    Update null elements with value in the same location in `other`.

Combine two DataFrame objects by filling null values in one DataFrame
with non-null values from other DataFrame. The row and column indexes
of the resulting DataFrame will be the union of the two.

Parameters
-----
other : DataFrame
    Provided DataFrame to use to fill null values.

Returns
-----
DataFrame
    The result of combining the provided DataFrame with the other object.

See Also
-----
DataFrame.combine : Perform series-wise operation on two DataFrames
    using a given function.

Examples
-----
>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [None, 4]})
>>> df2 = pd.DataFrame({'A': [1, 1], 'B': [3, 3]})
>>> df1.combine_first(df2)
   A  B
0  1.0 3.0
1  0.0 4.0

Null values still persist if the location of that null value
does not exist in `other`

>>> df1 = pd.DataFrame({'A': [None, 0], 'B': [4, None]})
>>> df2 = pd.DataFrame({'B': [3, 3], 'C': [1, 1]}, index=[1, 2])
>>> df1.combine_first(df2)
   A  B  C
0 NaN 4.0 NaN
1 0.0 3.0 1.0
2 NaN 3.0 1.0

compare(self, other: 'DataFrame', align_axis: 'Axis' = 1, keep_shape: 'bool' =
False, keep_equal: 'bool' = False) -> 'DataFrame'
    Compare to another DataFrame and show the differences.

    .. versionadded:: 1.1.0

Parameters
-----
other : DataFrame
    Object to compare with.

align_axis : {0 or 'index', 1 or 'columns'}, default 1
    Determine which axis to align the comparison on.

    * 0, or 'index' : Resulting differences are stacked vertically
      with rows drawn alternately from self and other.
    * 1, or 'columns' : Resulting differences are aligned horizontally
      with columns drawn alternately from self and other.

keep_shape : bool, default False
    If true, all rows and columns are kept.
    Otherwise, only the ones with different values are kept.

keep_equal : bool, default False
    If true, the result keeps values that are equal.
    Otherwise, equal values are shown as NaNs.

Returns
-----
DataFrame
    DataFrame that shows the differences stacked side by side.

    The resulting index will be a MultiIndex with 'self' and 'other'
    stacked alternately at the inner level.

Raises
-----
ValueError
    When the two DataFrames don't have identical labels or shape.

See Also
-----
Series.compare : Compare with another Series and show differences.
DataFrame.equals : Test whether two objects contain the same elements.

```

Notes

Matching NaNs will not appear as a difference.

Can only compare identically-labeled
(i.e. same shape, identical row and column labels) DataFrames

Examples

```
>>> df = pd.DataFrame(  
...     {  
...         "col1": ["a", "a", "b", "b", "a"],  
...         "col2": [1.0, 2.0, 3.0, np.nan, 5.0],  
...         "col3": [1.0, 2.0, 3.0, 4.0, 5.0]  
...     },  
...     columns=["col1", "col2", "col3"],  
... )  
>>> df  
   col1  col2  col3  
0     a   1.0   1.0  
1     a   2.0   2.0  
2     b   3.0   3.0  
3     b   NaN   4.0  
4     a   5.0   5.0
```

```
>>> df2 = df.copy()  
>>> df2.loc[0, 'col1'] = 'c'  
>>> df2.loc[2, 'col3'] = 4.0  
>>> df2  
   col1  col2  col3  
0     c   1.0   1.0  
1     a   2.0   2.0  
2     b   3.0   4.0  
3     b   NaN   4.0  
4     a   5.0   5.0
```

Align the differences on columns

```
>>> df.compare(df2)  
   col1  col3  
self other self other  
0     a     c   NaN   NaN  
2  NaN   NaN   3.0   4.0
```

Stack the differences on rows

```
>>> df.compare(df2, align_axis=0)  
   col1  col3  
0 self    a   NaN  
  other    c   NaN  
2 self  NaN   3.0  
  other  NaN   4.0
```

Keep the equal values

```
>>> df.compare(df2, keep_equal=True)  
   col1  col3  
self other self other  
0     a     c   1.0   1.0  
2     b     b   3.0   4.0
```

Keep all original rows and columns

```
>>> df.compare(df2, keep_shape=True)  
   col1  col2  col3  
self other self other self other  
0     a     c   NaN   NaN   NaN   NaN  
1  NaN   NaN   NaN   NaN   NaN   NaN  
2  NaN   NaN   NaN   NaN   3.0   4.0  
3  NaN   NaN   NaN   NaN   NaN   NaN  
4  NaN   NaN   NaN   NaN   NaN   NaN
```

Keep all original rows and columns and also all original values

```
>>> df.compare(df2, keep_shape=True, keep_equal=True)  
   col1  col2  col3  
self other self other self other  
0     a     c   1.0   1.0   1.0   1.0  
1     a     a   2.0   2.0   2.0   2.0  
2     b     b   3.0   3.0   3.0   4.0  
3     b     b   NaN   NaN   4.0   4.0  
4     a     a   5.0   5.0   5.0   5.0
```

```
corr(self, method: 'str | Callable[[np.ndarray, np.ndarray], float]' = 'pearson',  
min_periods: 'int' = 1) -> 'DataFrame'  
Compute pairwise correlation of columns, excluding NA/null values.
```

Parameters

method : {'pearson', 'kendall', 'spearman'} or callable
Method of correlation:

- * pearson : standard correlation coefficient
- * kendall : Kendall Tau correlation coefficient
- * spearman : Spearman rank correlation
- * callable: callable with input two 1d ndarrays and returning a float. Note that the returned matrix from corr will have 1 along the diagonals and will be symmetric regardless of the callable's behavior.

min_periods : int, optional

Minimum number of observations required per pair of columns to have a valid result. Currently only available for Pearson and Spearman correlation.

Returns

DataFrame

Correlation matrix.

See Also

DataFrame.corrwith : Compute pairwise correlation with another DataFrame or Series.

Series.corr : Compute the correlation between two Series.

Examples

```
>>> def histogram_intersection(a, b):  
...     v = np.minimum(a, b).sum().round(decimals=1)  
...     return v  
>>> df = pd.DataFrame([(0.2, .3), (.0, .6), (.6, .0), (.2, .1)],  
...                   columns=['dogs', 'cats'])  
>>> df.corr(method=histogram_intersection)  
      dogs  cats  
dogs    1.0   0.3  
cats    0.3   1.0
```

corrwith(self, other, axis: 'Axis' = 0, drop=False, method='pearson') -> 'Series'
Compute pairwise correlation.

Pairwise correlation is computed between rows or columns of DataFrame with rows or columns of Series or DataFrame. DataFrames are first aligned along both axes before computing the correlations.

Parameters

other : DataFrame, Series

Object with which to compute correlations.

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to use. 0 or 'index' to compute column-wise, 1 or 'columns' for row-wise.

drop : bool, default False

Drop missing indices from result.

method : {'pearson', 'kendall', 'spearman'} or callable

Method of correlation:

- * pearson : standard correlation coefficient
- * kendall : Kendall Tau correlation coefficient
- * spearman : Spearman rank correlation
- * callable: callable with input two 1d ndarrays and returning a float.

Returns

Series

Pairwise correlations.

See Also

DataFrame.corr : Compute pairwise correlation of columns.

count(self, axis: 'Axis' = 0, level: 'Level | None' = None, numeric_only: 'bool' = False)

Count non-NA cells for each column or row.

The values 'None', 'NaN', 'NaT', and optionally 'numpy.inf' (depending on 'pandas.options.mode.use_inf_as_na') are considered NA.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

If 0 or 'index' counts are generated for each column.

If 1 or 'columns' counts are generated for each row.

level : int or str, optional
If the axis is a `MultiIndex` (hierarchical), count along a particular `level`, collapsing into a `DataFrame`.
A `str` specifies the level name.
numeric_only : bool, default False
Include only `float`, `int` or `boolean` data.

Returns

Series or DataFrame
For each column/row the number of non-NA/null entries.
If `level` is specified returns a `DataFrame`.

See Also

Series.count: Number of non-NA elements in a Series.
DataFrame.value_counts: Count unique combinations of columns.
DataFrame.shape: Number of DataFrame rows and columns (including NA elements).
DataFrame.isna: Boolean same-sized DataFrame showing places of NA elements.

Examples

Constructing DataFrame from a dictionary:

```
>>> df = pd.DataFrame({"Person":  
...                     ["John", "Myla", "Lewis", "John", "Myla"],  
...                     "Age": [24., np.nan, 21., 33, 26],  
...                     "Single": [False, True, True, True, False]})  
>>> df  
   Person  Age  Single  
0   John  24.0   False  
1   Myla   NaN    True  
2  Lewis  21.0    True  
3   John  33.0    True  
4   Myla  26.0   False
```

Notice the uncounted NA values:

```
>>> df.count()  
Person    5  
Age        4  
Single    5  
dtype: int64
```

Counts for each **row**:

```
>>> df.count(axis='columns')  
0    3  
1    2  
2    3  
3    3  
4    3  
dtype: int64
```

cov(self, min_periods: 'int | None' = None, ddof: 'int | None' = 1) ->
'DataFrame'

Compute pairwise covariance of columns, excluding NA/null values.

Compute the pairwise covariance among the series of a DataFrame.
The returned data frame is the `covariance matrix`
<https://en.wikipedia.org/wiki/Covariance_matrix>`__` of the columns
of the DataFrame.

Both NA and null values are automatically excluded from the
calculation. (See the note below about bias from missing values.)
A threshold can be set for the minimum number of
observations for each value created. Comparisons with observations
below this threshold will be returned as ``NaN``.

This method is generally used for the analysis of time series data to
understand the relationship between different measures
across time.

Parameters

min_periods : int, optional
Minimum number of observations required per pair of columns
to have a valid result.

ddof : int, default 1
Delta degrees of freedom. The divisor used in calculations
is ``N - ddof``, where ``N`` represents the number of elements.

.. versionadded:: 1.1.0

Returns

DataFrame

The covariance matrix of the series of the DataFrame.

See Also

`Series.cov` : Compute covariance with another Series.

`core.window.ExponentialMovingWindow.cov`: Exponential weighted sample covariance.

`core.window.Expanding.cov` : Expanding sample covariance.

`core.window.Rolling.cov` : Rolling sample covariance.

Notes

Returns the covariance matrix of the DataFrame's time series.

The covariance is normalized by N-ddof.

For DataFrames that have Series that are missing data (assuming that data is 'missing at random

<https://en.wikipedia.org/wiki/Missing_data#Missing_at_random>`__)

the returned covariance matrix will be an unbiased estimate of the variance and covariance between the member Series.

However, for many applications this estimate may not be acceptable because the estimate covariance matrix is not guaranteed to be positive semi-definite. This could lead to estimate correlations having absolute values which are greater than one, and/or a non-invertible covariance matrix. See 'Estimation of covariance matrices' <https://en.wikipedia.org/w/index.php?title=Estimation_of_covariance_matrices>`__ for more details.

Examples

```
>>> df = pd.DataFrame([(1, 2), (0, 3), (2, 0), (1, 1)],
```

```
...                      columns=['dogs', 'cats'])
```

```
>>> df.cov()
```

```
          dogs      cats
dogs  0.666667 -1.000000
cats -1.000000  1.666667
```

```
>>> np.random.seed(42)
```

```
>>> df = pd.DataFrame(np.random.randn(1000, 5),
```

```
...                      columns=['a', 'b', 'c', 'd', 'e'])
```

```
>>> df.cov()
```

```
          a          b          c          d          e
a  0.998438 -0.020161  0.059277 -0.008943  0.014144
b -0.020161  1.059352 -0.008543 -0.024738  0.009826
c  0.059277 -0.008543  1.010670 -0.001486 -0.000271
d -0.008943 -0.024738 -0.001486  0.921297 -0.013692
e  0.014144  0.009826 -0.000271 -0.013692  0.977795
```

****Minimum number of periods****

This method also supports an optional ``min_periods`` keyword that specifies the required minimum number of non-NA observations for each column pair in order to have a valid result:

```
>>> np.random.seed(42)
```

```
>>> df = pd.DataFrame(np.random.randn(20, 3),
```

```
...                      columns=['a', 'b', 'c'])
```

```
>>> df.loc[df.index[:5], 'a'] = np.nan
```

```
>>> df.loc[df.index[5:10], 'b'] = np.nan
```

```
>>> df.cov(min_periods=12)
```

```
          a          b          c
a  0.316741      NaN -0.150812
b      NaN  1.248003  0.191417
c -0.150812  0.191417  0.895202
```

`cummax(self, axis=None, skipna=True, *args, **kwargs)`

Return cumulative maximum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative maximum.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.

`skipna` : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

`*args, **kwargs`

Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame

Return cumulative maximum of Series or DataFrame.

See Also

core.window.Expanding.max : Similar functionality
but ignores ``NaN`` values.

DataFrame.max : Return the maximum over
DataFrame axis.

DataFrame.cummax : Return cumulative maximum over DataFrame axis.

DataFrame.cummin : Return cumulative minimum over DataFrame axis.

DataFrame.cumsum : Return cumulative sum over DataFrame axis.

DataFrame.cumprod : Return cumulative product over DataFrame axis.

Examples

****Series****

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
```

```
>>> s
```

```
0    2.0
```

```
1    NaN
```

```
2    5.0
```

```
3   -1.0
```

```
4    0.0
```

```
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cummax()
```

```
0    2.0
```

```
1    NaN
```

```
2    5.0
```

```
3    5.0
```

```
4    5.0
```

```
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummax(skipna=False)
```

```
0    2.0
```

```
1    NaN
```

```
2    NaN
```

```
3    NaN
```

```
4    NaN
```

```
dtype: float64
```

****DataFrame****

```
>>> df = pd.DataFrame([[2.0, 1.0],  
...                    [3.0, np.nan],  
...                    [1.0, 0.0]],  
...                    columns=list('AB'))
```

```
>>> df
```

```
   A    B
```

```
0  2.0  1.0
```

```
1  3.0  NaN
```

```
2  1.0  0.0
```

By default, iterates over rows and finds the maximum
in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummax()
```

```
   A    B
```

```
0  2.0  1.0
```

```
1  3.0  NaN
```

```
2  3.0  1.0
```

To iterate over columns and find the maximum in each row,
use ``axis=1``

```
>>> df.cummax(axis=1)
```

```
   A    B
```

```
0  2.0  2.0
```

```
1  3.0  NaN
```

```
2  1.0  1.0
```

cummin(self, axis=None, skipna=True, *args, **kwargs)

Return cumulative minimum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative
minimum.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.
skipna : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.
*args, **kwargs
Additional keywords have no effect but might be accepted for compatibility with NumPy.

Returns

Series or DataFrame
Return cumulative minimum of Series or DataFrame.

See Also

core.window.Expanding.min : Similar functionality
but ignores ``NaN`` values.
DataFrame.min : Return the minimum over
DataFrame axis.
DataFrame.cummax : Return cumulative maximum over DataFrame axis.
DataFrame.cummin : Return cumulative minimum over DataFrame axis.
DataFrame.cumsum : Return cumulative sum over DataFrame axis.
DataFrame.cumprod : Return cumulative product over DataFrame axis.

Examples

Series

>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0 2.0
1 NaN
2 5.0
3 -1.0
4 0.0
dtype: float64

By default, NA values are ignored.

```
>>> s.cummin()
0    2.0
1    NaN
2    2.0
3   -1.0
4   -1.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cummin(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

DataFrame

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A  B
0  2.0  1.0
1  3.0 NaN
2  1.0  0.0
```

By default, iterates over rows and finds the minimum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cummin()
   A  B
0  2.0  1.0
1  2.0 NaN
2  1.0  0.0
```

To iterate over columns and find the minimum in each row, use ``axis=1``

```
>>> df.cummin(axis=1)
   A  B
0  2.0  1.0
1  3.0 NaN
2  1.0  0.0
```



```

cumprod(self, axis=None, skipna=True, *args, **kwargs)
    Return cumulative product over a DataFrame or Series axis.

    Returns a DataFrame or Series of the same size containing the cumulative
    product.

    Parameters
    -----
    axis : {0 or 'index', 1 or 'columns'}, default 0
        The index or the name of the axis. 0 is equivalent to None or 'index'.
    skipna : bool, default True
        Exclude NA/null values. If an entire row/column is NA, the result
        will be NA.
    *args, **kwargs
        Additional keywords have no effect but might be accepted for
        compatibility with NumPy.

    Returns
    -----
    Series or DataFrame
        Return cumulative product of Series or DataFrame.

    See Also
    -----
    core.window.Expanding.prod : Similar functionality
        but ignores ``NaN`` values.
    DataFrame.prod : Return the product over
        DataFrame axis.
    DataFrame.cummax : Return cumulative maximum over DataFrame axis.
    DataFrame.cummin : Return cumulative minimum over DataFrame axis.
    DataFrame.cumsum : Return cumulative sum over DataFrame axis.
    DataFrame.cumprod : Return cumulative product over DataFrame axis.

    Examples
    -----
    **Series**

    >>> s = pd.Series([2, np.nan, 5, -1, 0])
    >>> s
    0    2.0
    1    NaN
    2    5.0
    3   -1.0
    4    0.0
    dtype: float64

    By default, NA values are ignored.

    >>> s.cumprod()
    0    2.0
    1    NaN
    2   10.0
    3  -10.0
    4   -0.0
    dtype: float64

    To include NA values in the operation, use ``skipna=False``

    >>> s.cumprod(skipna=False)
    0    2.0
    1    NaN
    2    NaN
    3    NaN
    4    NaN
    dtype: float64

    **DataFrame**

    >>> df = pd.DataFrame([[2.0, 1.0],
    ...                    [3.0, np.nan],
    ...                    [1.0, 0.0]],
    ...                    columns=list('AB'))
    >>> df
       A    B
    0  2.0  1.0
    1  3.0  NaN
    2  1.0  0.0

    By default, iterates over rows and finds the product
    in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

    >>> df.cumprod()
       A    B
    0  2.0  1.0
    1  6.0  NaN
    2  6.0  0.0

```

To iterate over columns and find the product in each row,
use ``axis=1``

```
>>> df.cumprod(axis=1)
   A    B
0  2.0  2.0
1  3.0  NaN
2  1.0  0.0
```

`cumsum(self, axis=None, skipna=True, *args, **kwargs)`

Return cumulative sum over a DataFrame or Series axis.

Returns a DataFrame or Series of the same size containing the cumulative sum.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

The index or the name of the axis. 0 is equivalent to None or 'index'.
skipna : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result
will be NA.

*args, **kwargs
Additional keywords have no effect but might be accepted for
compatibility with NumPy.

Returns

Series or DataFrame

Return cumulative sum of Series or DataFrame.

See Also

core.window.Expanding.sum : Similar functionality
but ignores ``NaN`` values.

DataFrame.sum : Return the sum over
DataFrame axis.

DataFrame.cummax : Return cumulative maximum over DataFrame axis.

DataFrame.cummin : Return cumulative minimum over DataFrame axis.

DataFrame.cumsum : Return cumulative sum over DataFrame axis.

DataFrame.cumprod : Return cumulative product over DataFrame axis.

Examples

Series

```
>>> s = pd.Series([2, np.nan, 5, -1, 0])
>>> s
0    2.0
1    NaN
2    5.0
3   -1.0
4    0.0
dtype: float64
```

By default, NA values are ignored.

```
>>> s.cumsum()
0    2.0
1    NaN
2    7.0
3    6.0
4    6.0
dtype: float64
```

To include NA values in the operation, use ``skipna=False``

```
>>> s.cumsum(skipna=False)
0    2.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64
```

DataFrame

```
>>> df = pd.DataFrame([[2.0, 1.0],
...                    [3.0, np.nan],
...                    [1.0, 0.0]],
...                    columns=list('AB'))
>>> df
   A    B
0  2.0  1.0
1  3.0  NaN
2  1.0  0.0
```

By default, iterates over rows and finds the sum in each column. This is equivalent to ``axis=None`` or ``axis='index'``.

```
>>> df.cumsum()
   A    B
0  2.0  1.0
1  5.0  NaN
2  6.0  1.0
```

To iterate over columns and find the sum in each row, use ``axis=1``

```
>>> df.cumsum(axis=1)
   A    B
0  2.0  3.0
1  3.0  NaN
2  1.0  1.0
```

`diff(self, periods: 'int' = 1, axis: 'Axis' = 0) -> 'DataFrame'`
First discrete difference of element.

Calculates the difference of a Dataframe element compared with another element in the Dataframe (default is element in previous row).

Parameters

periods : int, default 1
Periods to shift for calculating difference, accepts negative values.
axis : {0 or 'index', 1 or 'columns'}, default 0
Take difference over rows (0) or columns (1).

Returns

Dataframe
First differences of the Series.

See Also

Dataframe.pct_change: Percent change over given number of periods.
Dataframe.shift: Shift index by desired number of periods with an optional time freq.
Series.diff: First discrete difference of object.

Notes

For boolean dtypes, this uses :meth:`operator.xor` rather than :meth:`operator.sub`.
The result is calculated according to current dtype in Dataframe, however dtype of the result is always float64.

Examples

Difference with previous row

```
>>> df = pd.DataFrame({'a': [1, 2, 3, 4, 5, 6],
...                     'b': [1, 1, 2, 3, 5, 8],
...                     'c': [1, 4, 9, 16, 25, 36]})
>>> df
   a  b  c
0  1  1  1
1  2  1  4
2  3  2  9
3  4  3  16
4  5  5  25
5  6  8  36
```

```
>>> df.diff()
   a  b  c
0 NaN NaN NaN
1  1.0 0.0 3.0
2  1.0 1.0 5.0
3  1.0 1.0 7.0
4  1.0 2.0 9.0
5  1.0 3.0 11.0
```

Difference with previous column

```
>>> df.diff(axis=1)
   a  b  c
0 NaN 0  0
1 NaN -1 3
2 NaN -1 7
3 NaN -1 13
4 NaN 0 20
5 NaN 2 28
```

Difference with 3rd previous row

```
>>> df.diff(periods=3)
      a    b    c
0  NaN  NaN  NaN
1  NaN  NaN  NaN
2  NaN  NaN  NaN
3  3.0  2.0  15.0
4  3.0  4.0  21.0
5  3.0  6.0  27.0
```

Difference with following row

```
>>> df.diff(periods=-1)
      a    b    c
0 -1.0  0.0 -3.0
1 -1.0 -1.0 -5.0
2 -1.0 -1.0 -7.0
3 -1.0 -2.0 -9.0
4 -1.0 -3.0 -11.0
5  NaN  NaN  NaN
```

Overflow in input dtype

```
>>> df = pd.DataFrame({'a': [1, 0]}, dtype=np.uint8)
>>> df.diff()
      a
0  NaN
1 255.0
```

`div = truediv(self, other, axis='columns', level=None, fill_value=None)`

`divide = truediv(self, other, axis='columns', level=None, fill_value=None)`

`dot(self, other: 'AnyArrayLike | FrameOrSeriesUnion') -> 'FrameOrSeriesUnion'`
Compute the matrix multiplication between the DataFrame and other.

This method computes the matrix product between the DataFrame and the values of an other Series, DataFrame or a numpy array.

It can also be called using `self @ other` in Python `>= 3.5`.

Parameters

`other` : Series, DataFrame or array-like
The other object to compute the matrix product with.

Returns

Series or DataFrame

If other is a Series, return the matrix product between self and other as a Series. If other is a DataFrame or a numpy.array, return the matrix product of self and other in a DataFrame of a np.array.

See Also

`Series.dot`: Similar method for Series.

Notes

The dimensions of DataFrame and other must be compatible in order to compute the matrix multiplication. In addition, the column names of DataFrame and the index of other must contain the same values, as they will be aligned prior to the multiplication.

The dot method for Series computes the inner product, instead of the matrix product here.

Examples

Here we multiply a DataFrame with a Series.

```
>>> df = pd.DataFrame([[0, 1, -2, -1], [1, 1, 1, 1]])
>>> s = pd.Series([1, 1, 2, 1])
>>> df.dot(s)
0    -4
1     5
dtype: int64
```

Here we multiply a DataFrame with another DataFrame.

```
>>> other = pd.DataFrame([[0, 1], [1, 2], [-1, -1], [2, 0]])
>>> df.dot(other)
0    1
0    1    4
1    2    2
```

Note that the dot method give the same result as @

```
>>> df @ other
   0  1
0  1  4
1  2  2
```

The dot method works also if other is an np.array.

```
>>> arr = np.array([[0, 1], [1, 2], [-1, -1], [2, 0]])
>>> df.dot(arr)
   0  1
0  1  4
1  2  2
```

Note how shuffling of the objects does not change the result.

```
>>> s2 = s.reindex([1, 0, 2, 3])
>>> df.dot(s2)
   0  -4
1  1  5
dtype: int64
```

```
drop(self, labels=None, axis: 'Axis' = 0, index=None, columns=None, level: 'Level
None' = None, inplace: 'bool' = False, errors: 'str' = 'raise')
Drop specified labels from rows or columns.
```

Remove rows or columns by specifying label names and corresponding axis, or by specifying directly index or column names. When using a multi-index, labels on different levels can be removed by specifying the level. See the `'user guide <advanced.shown_levels>'` for more information about the now unused levels.

Parameters

labels : single label or list-like
Index or column labels to drop.
axis : {0 or 'index', 1 or 'columns'}, default 0
Whether to drop labels from the index (0 or 'index') or columns (1 or 'columns').
index : single label or list-like
Alternative to specifying axis ('`labels, axis=0`' is equivalent to '`index=labels`').
columns : single label or list-like
Alternative to specifying axis ('`labels, axis=1`' is equivalent to '`columns=labels`').
level : int or level name, optional
For MultiIndex, level from which the labels will be removed.
inplace : bool, default False
If False, return a copy. Otherwise, do operation inplace and return None.
errors : {'ignore', 'raise'}, default 'raise'
If 'ignore', suppress error and only existing labels are dropped.

Returns

DataFrame or None
DataFrame without the removed index or column labels or None if '`inplace=True`'.

Raises

KeyError
If any of the labels is not found in the selected axis.

See Also

DataFrame.loc : Label-location based indexer for selection by label.
DataFrame.dropna : Return DataFrame with labels on given axis omitted where (all or any) data are missing.
DataFrame.drop_duplicates : Return DataFrame with duplicate rows removed, optionally only considering certain columns.
Series.drop : Return Series with specified index labels removed.

Examples

```
>>> df = pd.DataFrame(np.arange(12).reshape(3, 4),
...                    columns=['A', 'B', 'C', 'D'])
>>> df
   A  B  C  D
0  0  1  2  3
1  4  5  6  7
2  8  9 10 11
```

Drop columns

```
>>> df.drop(['B', 'C'], axis=1)
   A  D
0  0  3
1  4  7
2  8 11
```

```
>>> df.drop(columns=['B', 'C'])
   A  D
0  0  3
1  4  7
2  8 11
```

Drop a row by index

```
>>> df.drop([0, 1])
   A  B  C  D
2  8  9 10 11
```

Drop columns and/or rows of MultiIndex DataFrame

```
>>> midx = pd.MultiIndex(levels=[['lama', 'cow', 'falcon'],
...                               ['speed', 'weight', 'length']],
...                       codes=[[0, 0, 0, 1, 1, 1, 2, 2, 2],
...                               [0, 1, 2, 0, 1, 2, 0, 1, 2]])
>>> df = pd.DataFrame(index=midx, columns=['big', 'small'],
...                   data=[[45, 30], [200, 100], [1.5, 1], [30, 20],
...                           [250, 150], [1.5, 0.8], [320, 250],
...                           [1, 0.8], [0.3, 0.2]])
>>> df
```

		big	small
lama	speed	45.0	30.0
	weight	200.0	100.0
	length	1.5	1.0
cow	speed	30.0	20.0
	weight	250.0	150.0
	length	1.5	0.8
falcon	speed	320.0	250.0
	weight	1.0	0.8
	length	0.3	0.2

```
>>> df.drop(index='cow', columns='small')
           big
lama  speed  45.0
      weight 200.0
      length  1.5
falcon speed 320.0
      weight  1.0
      length  0.3
```

```
>>> df.drop(index='length', level=1)
           big  small
lama  speed  45.0   30.0
      weight 200.0  100.0
cow    speed  30.0   20.0
      weight 250.0  150.0
falcon speed 320.0  250.0
      weight  1.0   0.8
```

```
drop_duplicates(self, subset: 'Hashable | Sequence[Hashable] | None' = None,
keep: "Literal['first'] | Literal['last'] | Literal[False]" = 'first', inplace:
'bool' = False, ignore_index: 'bool' = False) -> 'DataFrame | None'
Return DataFrame with duplicate rows removed.
```

Considering certain columns is optional. Indexes, including time indexes are ignored.

Parameters

```
subset : column label or sequence of labels, optional
    Only consider certain columns for identifying duplicates, by
    default use all of the columns.
keep : {'first', 'last', False}, default 'first'
    Determines which duplicates (if any) to keep.
    - ``first`` : Drop duplicates except for the first occurrence.
    - ``last`` : Drop duplicates except for the last occurrence.
    - False : Drop all duplicates.
inplace : bool, default False
    Whether to drop duplicates in place or to return a copy.
ignore_index : bool, default False
    If True, the resulting axis will be labeled 0, 1, ..., n - 1.
```

.. versionadded:: 1.0.0

Returns

DataFrame or None

DataFrame with duplicates removed or None if ``inplace=True``.

See Also

DataFrame.value_counts: Count unique combinations of columns.

Examples

Consider dataset containing ramen rating.

```
>>> df = pd.DataFrame({
...     'brand': ['Yum Yum', 'Yum Yum', 'Indomie', 'Indomie', 'Indomie'],
...     'style': ['cup', 'cup', 'cup', 'pack', 'pack'],
...     'rating': [4, 4, 3.5, 15, 5]
... })
>>> df
   brand style  rating
0  Yum Yum  cup     4.0
1  Yum Yum  cup     4.0
2  Indomie  cup     3.5
3  Indomie pack    15.0
4  Indomie pack     5.0
```

By default, it removes duplicate rows based on all columns.

```
>>> df.drop_duplicates()
   brand style  rating
0  Yum Yum  cup     4.0
2  Indomie  cup     3.5
3  Indomie pack    15.0
4  Indomie pack     5.0
```

To remove duplicates on specific column(s), use ``subset``.

```
>>> df.drop_duplicates(subset=['brand'])
   brand style  rating
0  Yum Yum  cup     4.0
2  Indomie  cup     3.5
```

To remove duplicates and keep last occurrences, use ``keep``.

```
>>> df.drop_duplicates(subset=['brand', 'style'], keep='last')
   brand style  rating
1  Yum Yum  cup     4.0
2  Indomie  cup     3.5
4  Indomie pack     5.0
```

```
dropna(self, axis: 'Axis' = 0, how: 'str' = 'any', thresh=None, subset=None,
inplace: 'bool' = False)
    Remove missing values.
```

See the :ref:`User Guide <missing_data>` for more on which values are considered missing, and how to work with missing data.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

Determine if rows or columns which contain missing values are removed.

* 0, or 'index' : Drop rows which contain missing values.

* 1, or 'columns' : Drop columns which contain missing value.

.. versionchanged:: 1.0.0

Pass tuple or list to drop on multiple axes.

Only a single axis is allowed.

how : {'any', 'all'}, default 'any'

Determine if row or column is removed from DataFrame, when we have at least one NA or all NA.

* 'any' : If any NA values are present, drop that row or column.

* 'all' : If all values are NA, drop that row or column.

thresh : int, optional

Require that many non-NA values.

subset : array-like, optional

Labels along other axis to consider, e.g. if you are dropping rows these would be a list of columns to include.

inplace : bool, default False

If True, do operation inplace and return None.

Returns

DataFrame or None

DataFrame with NA entries dropped from it or None if ``inplace=True``.

See Also

DataFrame.isna: Indicate missing values.
DataFrame.notna : Indicate existing (non-missing) values.
DataFrame.fillna : Replace missing values.
Series.dropna : Drop missing values.
Index.dropna : Drop missing indices.

Examples

```
>>> df = pd.DataFrame({"name": ['Alfred', 'Batman', 'Catwoman'],
...                    "toy": [np.nan, 'Batmobile', 'Bullwhip'],
...                    "born": [pd.NaT, pd.Timestamp("1940-04-25"),
...                             pd.NaT]})
```

```
>>> df
   name      toy      born
0  Alfred     NaN      NaT
1  Batman  Batmobile 1940-04-25
2 Catwoman  Bullwhip      NaT
```

Drop the rows where at least one element is missing.

```
>>> df.dropna()
   name      toy      born
1  Batman  Batmobile 1940-04-25
```

Drop the columns where at least one element is missing.

```
>>> df.dropna(axis='columns')
   name
0  Alfred
1  Batman
2 Catwoman
```

Drop the rows where all elements are missing.

```
>>> df.dropna(how='all')
   name      toy      born
0  Alfred     NaN      NaT
1  Batman  Batmobile 1940-04-25
2 Catwoman  Bullwhip      NaT
```

Keep only the rows with at least 2 non-NA values.

```
>>> df.dropna(thresh=2)
   name      toy      born
1  Batman  Batmobile 1940-04-25
2 Catwoman  Bullwhip      NaT
```

Define in which columns to look for missing values.

```
>>> df.dropna(subset=['name', 'toy'])
   name      toy      born
1  Batman  Batmobile 1940-04-25
2 Catwoman  Bullwhip      NaT
```

Keep the DataFrame with valid entries in the same variable.

```
>>> df.dropna(inplace=True)
>>> df
   name      toy      born
1  Batman  Batmobile 1940-04-25
```

```
    duplicated(self, subset: 'Hashable | Sequence[Hashable] | None' = None, keep:
Literal['first'] | Literal['last'] | Literal[False]" = 'first') -> 'Series'
    Return boolean Series denoting duplicate rows.
```

Considering certain columns is optional.

Parameters

subset : column label or sequence of labels, optional
Only consider certain columns for identifying duplicates, by
default use all of the columns.
keep : {'first', 'last', False}, default 'first'
Determines which duplicates (if any) to mark.

- ``first`` : Mark duplicates as ``True`` except for the first
occurrence.
- ``last`` : Mark duplicates as ``True`` except for the last occurrence.
- False : Mark all duplicates as ``True``.

Returns

Series

Boolean series for each duplicated rows.

See Also

Index.duplicated : Equivalent method on index.

Series.duplicated : Equivalent method on Series.

Series.drop_duplicates : Remove duplicate values from Series.

DataFrame.drop_duplicates : Remove duplicate values from DataFrame.

Examples

Consider dataset containing ramen rating.

```
>>> df = pd.DataFrame({
...     'brand': ['Yum Yum', 'Yum Yum', 'Indomie', 'Indomie', 'Indomie'],
...     'style': ['cup', 'cup', 'cup', 'pack', 'pack'],
...     'rating': [4, 4, 3.5, 15, 5]
... })
>>> df
   brand style  rating
0  Yum Yum  cup     4.0
1  Yum Yum  cup     4.0
2  Indomie  cup     3.5
3  Indomie pack    15.0
4  Indomie pack     5.0
```

By default, for each set of duplicated values, the first occurrence is set on False and all others on True.

```
>>> df.duplicated()
0    False
1     True
2    False
3    False
4    False
dtype: bool
```

By using 'last', the last occurrence of each set of duplicated values is set on False and all others on True.

```
>>> df.duplicated(keep='last')
0     True
1    False
2    False
3    False
4    False
dtype: bool
```

By setting 'keep' on False, all duplicates are True.

```
>>> df.duplicated(keep=False)
0     True
1     True
2    False
3    False
4    False
dtype: bool
```

To find duplicates on specific column(s), use 'subset'.

```
>>> df.duplicated(subset=['brand'])
0    False
1     True
2    False
3     True
4     True
dtype: bool
```

eq(self, other, axis='columns', level=None)

Get Equal to of dataframe and other, element-wise (binary operator 'eq').

Among flexible wrappers ('eq', 'ne', 'le', 'lt', 'ge', 'gt') to comparison operators.

Equivalent to '==', '!=', '<=', '<', '>=', '>' with support to choose axis (rows or columns) and level for comparison.

Parameters

other : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').

level : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

DataFrame of bool

Result of the comparison.

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.

DataFrame.ne : Compare DataFrames for inequality elementwise.

DataFrame.le : Compare DataFrames for less than inequality
or equality elementwise.

DataFrame.lt : Compare DataFrames for strictly less than
inequality elementwise.

DataFrame.ge : Compare DataFrames for greater than inequality
or equality elementwise.

DataFrame.gt : Compare DataFrames for strictly greater than
inequality elementwise.

Notes

Mismatched indices will be unioned together.

`NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                    'revenue': [100, 250, 300]},
...                    index=['A', 'B', 'C'])
```

```
>>> df
   cost  revenue
A   250     100
B   150     250
C   100     300
```

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
   cost  revenue
A  False     True
B  False     False
C   True     False
```

```
>>> df.eq(100)
   cost  revenue
A  False     True
B  False     False
C   True     False
```

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
   cost  revenue
A   True     True
B   True     False
C  False     True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
   cost  revenue
A   True     False
B   True     True
C   True     True
D   True     True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
   cost  revenue
A   True     True
B  False     False
C  False     False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
   cost  revenue
A   True     False
B  False     True
C   True     False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                       index=['A', 'B', 'C', 'D'])
```

```
>>> other
      revenue
A         300
B         250
C         100
D         150

>>> df.gt(other)
      cost revenue
A  False   False
B  False   False
C  False    True
D  False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                       ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225

>>> df.le(df_multindex, level=1)
      cost revenue
Q1 A    True     True
   B    True     True
   C    True     True
Q2 A   False     True
   B    True    False
   C    True    False
```

`eval(self, expr: 'str', inplace: 'bool' = False, **kwargs)`
 Evaluate a string describing operations on DataFrame columns.

Operates on columns only, not specific rows or elements. This allows ``eval`` to run arbitrary code, which can make you vulnerable to code injection if you pass user input to this function.

Parameters

`expr` : str

The expression string to evaluate.

`inplace` : bool, default False

If the expression contains an assignment, whether to perform the operation inplace and mutate the existing DataFrame. Otherwise, a new DataFrame is returned.

`**kwargs`

See the documentation for `:func:`eval`` for complete details on the keyword arguments accepted by `:meth:`~pandas.DataFrame.query``.

Returns

ndarray, scalar, pandas object, or None

The result of the evaluation or None if ``inplace=True``.

See Also

`DataFrame.query` : Evaluates a boolean expression to query the columns of a frame.

`DataFrame.assign` : Can evaluate an expression or function to create new values for a column.

`eval` : Evaluate a Python expression as a string using various backends.

Notes

For more details see the API documentation for `:func:`~eval``. For detailed examples see `:ref:`enhancing performance with eval` <enhancingperf.eval>`.

Examples

```
>>> df = pd.DataFrame({'A': range(1, 6), 'B': range(10, 0, -2)})
>>> df
   A  B
0  1 10
1  2  8
2  3  6
3  4  4
```

```

4 5 2
>>> df.eval('A + B')
0 11
1 10
2 9
3 8
4 7
dtype: int64

```

Assignment is allowed though by default the original DataFrame is not modified.

```

>>> df.eval('C = A + B')
   A  B  C
0  1 10 11
1  2  8 10
2  3  6  9
3  4  4  8
4  5  2  7
>>> df
   A  B
0  1 10
1  2  8
2  3  6
3  4  4
4  5  2

```

Use ``inplace=True`` to modify the original DataFrame.

```

>>> df.eval('C = A + B', inplace=True)
>>> df
   A  B  C
0  1 10 11
1  2  8 10
2  3  6  9
3  4  4  8
4  5  2  7

```

Multiple columns can be assigned to using multi-line expressions:

```

>>> df.eval(
...     '''
...     C = A + B
...     D = A - B
...     '''
... )
   A  B  C  D
0  1 10 11 -9
1  2  8 10 -6
2  3  6  9 -3
3  4  4  8  0
4  5  2  7  3

```

`explode(self, column: 'str | tuple | list[str | tuple]', ignore_index: 'bool' = False) -> 'DataFrame'`

Transform each element of a list-like to a row, replicating index values.

.. versionadded:: 0.25.0

Parameters

`column` : str or tuple or list thereof

Column(s) to explode.

For multiple columns, specify a non-empty list with each element be str or tuple, and all specified columns their list-like data on same row of the frame must have matching length.

.. versionadded:: 1.3.0

Multi-column explode

`ignore_index` : bool, default False

If True, the resulting index will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.1.0

Returns

DataFrame

Exploded lists to rows of the subset columns;
index will be duplicated for these rows.

Raises

ValueError :

- * If columns of the frame are not unique.
- * If specified columns to explode is empty list.
- * If specified columns to explode have not matching count of

elements rowwise in the frame.

See Also

DataFrame.unstack : Pivot a level of the (necessarily hierarchical)
index labels.
DataFrame.melt : Unpivot a DataFrame from wide format to long format.
Series.explode : Explode a DataFrame from list-like columns to long format.

Notes

This routine will explode list-likes including lists, tuples, sets,
Series, and np.ndarray. The result dtype of the subset rows will
be object. Scalars will be returned unchanged, and empty list-likes will
result in a np.nan for that row. In addition, the ordering of rows in the
output will be non-deterministic when exploding sets.

Examples

>>> df = pd.DataFrame({'A': [[0, 1, 2], 'foo', []], [3, 4]],
... 'B': 1,
... 'C': [['a', 'b', 'c'], np.nan, [], ['d', 'e']]})
>>> df

	A	B	C
0	[0, 1, 2]	1	[a, b, c]
1	foo	1	NaN
2	[]	1	[]
3	[3, 4]	1	[d, e]

Single-column explode.

```
>>> df.explode('A')
   A  B  C
0  0  1  [a, b, c]
0  1  1  [a, b, c]
0  2  1  [a, b, c]
1  foo 1      NaN
2  NaN 1      []
3   3  1  [d, e]
3   4  1  [d, e]
```

Multi-column explode.

```
>>> df.explode(list('AC'))
   A  B  C
0  0  1  a
0  1  1  b
0  2  1  c
1  foo 1  NaN
2  NaN 1  NaN
3   3  1  d
3   4  1  e
```

ffill(self: 'DataFrame', axis: 'None | Axis' = None, inplace: 'bool' = False,
limit: 'None | int' = None, downcast=None) -> 'DataFrame | None'

Synonym for :meth:`DataFrame.fillna` with ``method='ffill'``.

Returns

Series/DataFrame or None
Object with missing values filled or None if ``inplace=True``.

fillna(self, value: 'object | ArrayLike | None' = None, method: 'FillnaOptions |
None' = None, axis: 'Axis | None' = None, inplace: 'bool' = False, limit=None,
downcast=None) -> 'DataFrame | None'

Fill NA/NaN values using the specified method.

Parameters

value : scalar, dict, Series, or DataFrame
Value to use to fill holes (e.g. 0), alternately a
dict/Series/DataFrame of values specifying which value to use for
each index (for a Series) or column (for a DataFrame). Values not
in the dict/Series/DataFrame will not be filled. This value cannot
be a list.
method : {'backfill', 'bfill', 'pad', 'ffill', None}, default None
Method to use for filling holes in reindexed Series
pad / ffill: propagate last valid observation forward to next valid
backfill / bfill: use next valid observation to fill gap.
axis : {0 or 'index', 1 or 'columns'}
Axis along which to fill missing values.
inplace : bool, default False
If True, fill in-place. Note: this will modify any
other views on this object (e.g., a no-copy slice for a column in a
DataFrame).
limit : int, default None
If method is specified, this is the maximum number of consecutive

NaN values to forward/backward fill. In other words, if there is a gap with more than this number of consecutive NaNs, it will only be partially filled. If method is not specified, this is the maximum number of entries along the entire axis where NaNs will be filled. Must be greater than 0 if not None.

downcast : dict, default is None

A dict of item->dtype of what to downcast if possible, or the string 'infer' which will try to downcast to an appropriate equal type (e.g. float64 to int64 if possible).

Returns

DataFrame or None

Object with missing values filled or None if ``inplace=True``.

See Also

interpolate : Fill NaN values using interpolation.
reindex : Conform object to new index.
asfreq : Convert TimeSeries to specified frequency.

Examples

```
>>> df = pd.DataFrame([[np.nan, 2, np.nan, 0],
...                     [3, 4, np.nan, 1],
...                     [np.nan, np.nan, np.nan, 5],
...                     [np.nan, 3, np.nan, 4]],
...                     columns=list("ABCD"))
>>> df
```

	A	B	C	D
0	NaN	2.0	NaN	0
1	3.0	4.0	NaN	1
2	NaN	NaN	NaN	5
3	NaN	3.0	NaN	4

Replace all NaN elements with 0s.

```
>>> df.fillna(0)
```

	A	B	C	D
0	0.0	2.0	0.0	0
1	3.0	4.0	0.0	1
2	0.0	0.0	0.0	5
3	0.0	3.0	0.0	4

We can also propagate non-null values forward or backward.

```
>>> df.fillna(method="ffill")
```

	A	B	C	D
0	NaN	2.0	NaN	0
1	3.0	4.0	NaN	1
2	3.0	4.0	NaN	5
3	3.0	3.0	NaN	4

Replace all NaN elements in column 'A', 'B', 'C', and 'D', with 0, 1, 2, and 3 respectively.

```
>>> values = {"A": 0, "B": 1, "C": 2, "D": 3}
>>> df.fillna(value=values)
```

	A	B	C	D
0	0.0	2.0	2.0	0
1	3.0	4.0	2.0	1
2	0.0	1.0	2.0	5
3	0.0	3.0	2.0	4

Only replace the first NaN element.

```
>>> df.fillna(value=values, limit=1)
```

	A	B	C	D
0	0.0	2.0	2.0	0
1	3.0	4.0	NaN	1
2	NaN	1.0	NaN	5
3	NaN	3.0	NaN	4

When filling using a DataFrame, replacement happens along the same column names and same indices

```
>>> df2 = pd.DataFrame(np.zeros((4, 4)), columns=list("ABCE"))
>>> df.fillna(df2)
```

	A	B	C	D
0	0.0	2.0	0.0	0
1	3.0	4.0	0.0	1
2	0.0	0.0	0.0	5
3	0.0	3.0	0.0	4

floordiv(self, other, axis='columns', level=None, fill_value=None)

Get Integer division of dataframe and other, element-wise (binary operator 'floordiv').

Equivalent to ``dataframe // other``, but with support to substitute a
fill_value
for missing data in one of the inputs. With reverse version, `rfloordiv`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
Whether to compare by the index (0 or 'index') or columns
(1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
Broadcast across a level, matching Index values on the
passed MultiIndex level.
fill_value : float or None, default None
Fill existing missing (NaN) values, and any new element needed for
successful DataFrame alignment, with this value before computation.
If data in both corresponding DataFrame locations is missing
the result will be missing.

Returns

DataFrame
Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df

	angles	degrees
circle	0	360
triangle	3	180
rectangle	4	360

Add a scalar with operator version which return the same
results.

```
>>> df + 1
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
```

	angles	degrees
circle	-1	359
triangle	2	179
rectangle	3	359

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                        index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                               'degrees': [360, 180, 360, 360, 540, 720]},
...                              index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                     ['circle', 'triangle', 'rectangle',
...                                      'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

	angles	degrees
A circle	0	360
triangle	3	180
rectangle	4	360
B square	4	360
pentagon	5	540
hexagon	6	720

```
>>> df.div(df_multindex, level=1, fill_value=0)
```

	angles	degrees
A circle	NaN	1.0
triangle	1.0	1.0
rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

```
ge(self, other, axis='columns', level=None)
```

Get Greater than or equal to of dataframe and other, element-wise (binary operator `ge`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

other : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').

level : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns


```
-----
DataFrame of bool
    Result of the comparison.
```

See Also

```
-----
DataFrame.eq : Compare DataFrames for equality elementwise.
DataFrame.ne : Compare DataFrames for inequality elementwise.
DataFrame.le : Compare DataFrames for less than inequality
               or equality elementwise.
DataFrame.lt : Compare DataFrames for strictly less than
               inequality elementwise.
DataFrame.ge : Compare DataFrames for greater than inequality
               or equality elementwise.
DataFrame.gt : Compare DataFrames for strictly greater than
               inequality elementwise.
```

Notes

```
-----
Mismatched indices will be unioned together.
`NaN` values are considered different (i.e. `NaN` != `NaN`).
```

Examples

```
-----
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                    'revenue': [100, 250, 300]},
...                    index=['A', 'B', 'C'])
>>> df
   cost  revenue
A   250     100
B   150     250
C   100     300
```

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
   cost  revenue
A  False     True
B  False     False
C   True     False
```

```
>>> df.eq(100)
   cost  revenue
A  False     True
B  False     False
C   True     False
```

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
   cost  revenue
A   True     True
B   True     False
C  False     True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
   cost  revenue
A   True     False
B   True     True
C   True     True
D   True     True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
   cost  revenue
A   True     True
B  False     False
C  False     False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
   cost  revenue
A   True     False
B  False     True
C   True     False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                       index=['A', 'B', 'C', 'D'])
>>> other
```

```

revenue
A    300
B    250
C    100
D    150

>>> df.gt(other)
      cost  revenue
A  False    False
B  False    False
C  False     True
D  False    False

```

Compare to a MultiIndex by level.

```

>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                       ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225

>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A    True     True
   B    True     True
   C    True     True
Q2 A   False     True
   B    True    False
   C    True    False

```

```

groupby(self, by=None, axis: 'Axis' = 0, level: 'Level | None' = None, as_index:
'bool' = True, sort: 'bool' = True, group_keys: 'bool' = True, squeeze: 'bool |
lib.NoDefault' = <no_default>, observed: 'bool' = False, dropna: 'bool' = True) ->
'DataFrameGroupBy'

```

Group DataFrame using a mapper or by a Series of columns.

A groupby operation involves some combination of splitting the object, applying a function, and combining the results. This can be used to group large amounts of data and compute operations on these groups.

Parameters

```

-----
by : mapping, function, label, or list of labels
    Used to determine the groups for the groupby.
    If ``by`` is a function, it's called on each value of the object's
    index. If a dict or Series is passed, the Series or dict VALUES
    will be used to determine the groups (the Series' values are first
    aligned; see ``.align()`` method). If an ndarray is passed, the
    values are used as-is to determine the groups. A label or list of
    labels may be passed to group by the columns in ``self``. Notice
    that a tuple is interpreted as a (single) key.
axis : {0 or 'index', 1 or 'columns'}, default 0
    Split along rows (0) or columns (1).
level : int, level name, or sequence of such, default None
    If the axis is a MultiIndex (hierarchical), group by a particular
    level or levels.
as_index : bool, default True
    For aggregated output, return object with group labels as the
    index. Only relevant for DataFrame input. as_index=False is
    effectively "SQL-style" grouped output.
sort : bool, default True
    Sort group keys. Get better performance by turning this off.
    Note this does not influence the order of observations within each
    group. Groupby preserves the order of rows within each group.
group_keys : bool, default True
    When calling apply, add group keys to index to identify pieces.
squeeze : bool, default False
    Reduce the dimensionality of the return type if possible,
    otherwise return a consistent type.

.. deprecated:: 1.1.0

observed : bool, default False
    This only applies if any of the groupers are Categoricals.
    If True: only show observed values for categorical groupers.
    If False: show all values for categorical groupers.
dropna : bool, default True
    If True, and if group keys contain NA values, NA values together
    with row/column will be dropped.

```

If False, NA values will also be treated as the key in groups

.. versionadded:: 1.1.0

Returns

DataFrameGroupBy

Returns a groupby object that contains information about the groups.

See Also

`resample` : Convenience method for frequency conversion and resampling of time series.

Notes

See the ``user guide`

<<https://pandas.pydata.org/pandas-docs/stable/groupby.html>>`__ for more.

Examples

```
>>> df = pd.DataFrame({'Animal': ['Falcon', 'Falcon',  
...                               'Parrot', 'Parrot'],  
...                   'Max Speed': [380., 370., 24., 26.]})
```

```
>>> df
   Animal  Max Speed
0  Falcon    380.0
1  Falcon    370.0
2  Parrot     24.0
3  Parrot     26.0
```

```
>>> df.groupby(['Animal']).mean()
           Max Speed
Animal
Falcon    375.0
Parrot    25.0
```

****Hierarchical Indexes****

We can groupby different levels of a hierarchical index using the ``level`` parameter:

```
>>> arrays = [['Falcon', 'Falcon', 'Parrot', 'Parrot'],  
...           ['Captive', 'Wild', 'Captive', 'Wild']]
>>> index = pd.MultiIndex.from_arrays(arrays, names=('Animal', 'Type'))
>>> df = pd.DataFrame({'Max Speed': [390., 350., 30., 20.]},  
...                   index=index)
```

```
>>> df
           Max Speed
Animal Type
Falcon Captive    390.0
       Wild      350.0
Parrot Captive     30.0
       Wild       20.0
```

```
>>> df.groupby(level=0).mean()
           Max Speed
Animal
Falcon    370.0
Parrot     25.0
```

```
>>> df.groupby(level="Type").mean()
           Max Speed
Type
Captive    210.0
Wild      185.0
```

We can also choose to include NA in group keys or not by setting ``dropna`` parameter, the default setting is ``True``:

```
>>> l = [[1, 2, 3], [1, None, 4], [2, 1, 3], [1, 2, 2]]
>>> df = pd.DataFrame(l, columns=["a", "b", "c"])
```

```
>>> df.groupby(by=["b"]).sum()
   a  c
b
1.0 2  3
2.0 2  5
```

```
>>> df.groupby(by=["b"], dropna=False).sum()
   a  c
b
1.0 2  3
2.0 2  5
NaN 1  4
```

```
>>> l = [{"a": 12, 12}, [None, 12.3, 33.], [{"b": 12.3, 123}, [{"a": 1, 1}]]
>>> df = pd.DataFrame(l, columns=["a", "b", "c"])
```

```
>>> df.groupby(by="a").sum()
```

```

      b      c
a
a  13.0   13.0
b  12.3  123.0

```

```

>>> df.groupby(by="a", dropna=False).sum()
      b      c
a
a  13.0   13.0
b  12.3  123.0
NaN 12.3   33.0

```

gt(self, other, axis='columns', level=None)
Get Greater than of dataframe and other, element-wise (binary operator `gt`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

other : scalar, sequence, Series, or DataFrame
Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}, default 'columns'
Whether to compare by the index (0 or 'index') or columns (1 or 'columns').
level : int or label
Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

DataFrame of bool
Result of the comparison.

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.
DataFrame.ne : Compare DataFrames for inequality elementwise.
DataFrame.le : Compare DataFrames for less than inequality or equality elementwise.
DataFrame.lt : Compare DataFrames for strictly less than inequality elementwise.
DataFrame.ge : Compare DataFrames for greater than inequality or equality elementwise.
DataFrame.gt : Compare DataFrames for strictly greater than inequality elementwise.

Notes

Mismatched indices will be unioned together.
`NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```

>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                    'revenue': [100, 250, 300]},
...                    index=['A', 'B', 'C'])
>>> df
   cost  revenue
A   250      100
B   150      250
C   100      300

```

Comparison with a scalar, using either the operator or method:

```

>>> df == 100
   cost  revenue
A  False     True
B  False     False
C   True     False

>>> df.eq(100)
   cost  revenue
A  False     True
B  False     False
C   True     False

```

When `other` is a :class:`Series`, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```

>>> df != pd.Series([100, 250], index=["cost", "revenue"])
   cost  revenue
A   True     True
B   True     False

```

C False True

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
   cost  revenue
A  True   False
B  True    True
C  True    True
D  True    True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
   cost  revenue
A  True    True
B  False  False
C  False  False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
   cost  revenue
A  True   False
B  False    True
C  True   False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                       index=['A', 'B', 'C', 'D'])
>>> other
   revenue
A       300
B       250
C       100
D       150
```

```
>>> df.gt(other)
   cost  revenue
A  False  False
B  False  False
C  False    True
D  False  False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                              'revenue': [100, 250, 300, 200, 175, 225]},
...                              index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                      ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
   cost  revenue
Q1 A   250     100
   B   150     250
   C   100     300
Q2 A   150     200
   B   300     175
   C   220     225

>>> df.le(df_multindex, level=1)
   cost  revenue
Q1 A   True    True
   B   True    True
   C   True    True
Q2 A  False    True
   B   True   False
   C   True   False
```

```
hist = hist_frame(data: 'DataFrame', column: 'IndexLabel' = None, by=None, grid:
'bool' = True, xlabelsize: 'int | None' = None, xrot: 'float | None' = None,
ylabelsize: 'int | None' = None, yrot: 'float | None' = None, ax=None, sharex: 'bool'
= False, sharey: 'bool' = False, figsize: 'tuple[int, int] | None' = None, layout:
'tuple[int, int] | None' = None, bins: 'int | Sequence[int]' = 10, backend: 'str |
None' = None, legend: 'bool' = False, **kwargs)
```

Make a histogram of the DataFrame's columns.

A `histogram`_ is a representation of the distribution of data.
This function calls :meth:`matplotlib.pyplot.hist`, on each series in
the DataFrame, resulting in one histogram per column.

.. _histogram: <https://en.wikipedia.org/wiki/Histogram>

Parameters

data : DataFrame

The pandas object holding the data.

column : str or sequence, optional
If passed, will be used to limit data to a subset of columns.

by : object, optional
If passed, then used to form histograms for separate groups.

grid : bool, default True
Whether to show axis grid lines.

xlabelsize : int, default None
If specified changes the x-axis label size.

xrot : float, default None
Rotation of x axis labels. For example, a value of 90 displays the x labels rotated 90 degrees clockwise.

ylabelsize : int, default None
If specified changes the y-axis label size.

yrot : float, default None
Rotation of y axis labels. For example, a value of 90 displays the y labels rotated 90 degrees clockwise.

ax : Matplotlib axes object, default None
The axes to plot the histogram on.

sharex : bool, default True if ax is None else False
In case subplots=True, share x axis and set some x axis labels to invisible; defaults to True if ax is None otherwise False if an ax is passed in.
Note that passing in both an ax and sharex=True will alter all x axis labels for all subplots in a figure.

sharey : bool, default False
In case subplots=True, share y axis and set some y axis labels to invisible.

figsize : tuple, optional
The size in inches of the figure to create. Uses the value in `matplotlib.rcParams` by default.

layout : tuple, optional
Tuple of (rows, columns) for the layout of the histograms.

bins : int or sequence, default 10
Number of histogram bins to be used. If an integer is given, bins + 1 bin edges are calculated and returned. If bins is a sequence, gives bin edges, including left edge of first bin and right edge of last bin. In this case, bins is returned unmodified.

backend : str, default None
Backend to use instead of the backend specified in the option ``plotting.backend``. For instance, 'matplotlib'. Alternatively, to specify the ``plotting.backend`` for the whole session, set ``pd.options.plotting.backend``.

.. versionadded:: 1.0.0

legend : bool, default False
Whether to show the legend.

.. versionadded:: 1.1.0

****kwargs**
All other plotting keyword arguments to be passed to :meth:`matplotlib.pyplot.hist`.

Returns

matplotlib.AxesSubplot or numpy.ndarray of them

See Also

matplotlib.pyplot.hist : Plot a histogram using matplotlib.

Examples

This example draws a histogram based on the length and width of some animals, displayed in three bins

```
.. plot::
    :context: close-figs

    >>> df = pd.DataFrame({
    ...     'length': [1.5, 0.5, 1.2, 0.9, 3],
    ...     'width': [0.7, 0.2, 0.15, 0.2, 1.1]
    ...     }, index=['pig', 'rabbit', 'duck', 'chicken', 'horse'])
    >>> hist = df.hist(bins=3)
```

idxmax(self, axis: 'Axis' = 0, skipna: 'bool' = True) -> 'Series'
Return index of first occurrence of maximum over requested axis.

NA/null values are excluded.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0
The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-

wise.

skipna : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

Returns

Series

Indexes of maxima along the specified axis.

Raises

ValueError

* If the row/column is empty

See Also

Series.idxmax : Return index of the maximum element.

Notes

This method is the DataFrame version of ``ndarray.argmax``.

Examples

Consider a dataset containing food consumption in Argentina.

```
>>> df = pd.DataFrame({'consumption': [10.51, 103.11, 55.48],
...                    'co2_emissions': [37.2, 19.66, 1712]},
...                    index=['Pork', 'Wheat Products', 'Beef'])
```

```
>>> df
      consumption  co2_emissions
Pork           10.51           37.20
Wheat Products  103.11           19.66
Beef           55.48          1712.00
```

By default, it returns the index for the maximum value in each column.

```
>>> df.idxmax()
consumption    Wheat Products
co2_emissions           Beef
dtype: object
```

To return the index for the maximum value in each row, use

``axis="columns"``.

```
>>> df.idxmax(axis="columns")
Pork           co2_emissions
Wheat Products    consumption
Beef           co2_emissions
dtype: object
```

idxmin(self, axis: 'Axis' = 0, skipna: 'bool' = True) -> 'Series'
Return index of first occurrence of minimum over requested axis.

NA/null values are excluded.

Parameters

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-

wise.

skipna : bool, default True
Exclude NA/null values. If an entire row/column is NA, the result will be NA.

Returns

Series

Indexes of minima along the specified axis.

Raises

ValueError

* If the row/column is empty

See Also

Series.idxmin : Return index of the minimum element.

Notes

This method is the DataFrame version of ``ndarray.argmin``.

Examples

Consider a dataset containing food consumption in Argentina.

```
>>> df = pd.DataFrame({'consumption': [10.51, 103.11, 55.48],
...                     'co2_emissions': [37.2, 19.66, 1712]},
...                     index=['Pork', 'Wheat Products', 'Beef'])
```

```
>>> df
      consumption  co2_emissions
Pork           10.51           37.20
Wheat Products 103.11           19.66
Beef           55.48          1712.00
```

By default, it returns the index for the minimum value in each column.

```
>>> df.idxmin()
consumption      Pork
co2_emissions    Wheat Products
dtype: object
```

To return the index for the minimum value in each row, use
`axis="columns"`.

```
>>> df.idxmin(axis="columns")
Pork      consumption
Wheat Products  co2_emissions
Beef      consumption
dtype: object
```

```
info(self, verbose: 'bool | None' = None, buf: 'IO[str] | None' = None, max_cols:
'int | None' = None, memory_usage: 'bool | str | None' = None, show_counts: 'bool |
None' = None, null_counts: 'bool | None' = None) -> 'None'
```

Print a concise summary of a DataFrame.

This method prints information about a DataFrame including
the index dtype and columns, non-null values and memory usage.

Parameters

data : DataFrame

DataFrame to print information about.

verbose : bool, optional

Whether to print the full summary. By default, the setting in
`pandas.options.display.max_info_columns` is followed.

buf : writable buffer, defaults to sys.stdout

Where to send the output. By default, the output is printed to
sys.stdout. Pass a writable buffer if you need to further process
the output.

max_cols : int, optional

When to switch from the verbose to the truncated output. If the
DataFrame has more than `max_cols` columns, the truncated output
is used. By default, the setting in
`pandas.options.display.max_info_columns` is used.

memory_usage : bool, str, optional

Specifies whether total memory usage of the DataFrame
elements (including the index) should be displayed. By default,
this follows the `pandas.options.display.memory_usage` setting.

True always show memory usage. False never shows memory usage.

A value of 'deep' is equivalent to "True with deep introspection".

Memory usage is shown in human-readable units (base-2
representation). Without deep introspection a memory estimation is
made based in column dtype and number of rows assuming values
consume the same memory amount for corresponding dtypes. With deep
memory introspection, a real memory usage calculation is performed
at the cost of computational resources.

show_counts : bool, optional

Whether to show the non-null counts. By default, this is shown
only if the DataFrame is smaller than

`pandas.options.display.max_info_rows` and

`pandas.options.display.max_info_columns`. A value of True always
shows the counts, and False never shows the counts.

null_counts : bool, optional

.. deprecated:: 1.2.0

Use show_counts instead.

Returns

None

This method prints a summary of a DataFrame and returns None.

See Also

DataFrame.describe: Generate descriptive statistics of DataFrame
columns.

DataFrame.memory_usage: Memory usage of DataFrame columns.

Examples


```

-----
>>> int_values = [1, 2, 3, 4, 5]
>>> text_values = ['alpha', 'beta', 'gamma', 'delta', 'epsilon']
>>> float_values = [0.0, 0.25, 0.5, 0.75, 1.0]
>>> df = pd.DataFrame({"int_col": int_values, "text_col": text_values,
...                    "float_col": float_values})
>>> df
   int_col text_col float_col
0         1   alpha      0.00
1         2   beta      0.25
2         3  gamma      0.50
3         4  delta      0.75
4         5 epsilon      1.00

```

Prints information of all columns:

```

>>> df.info(verbose=True)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   int_col      5 non-null      int64
1   text_col     5 non-null      object
2   float_col    5 non-null      float64
dtypes: float64(1), int64(1), object(1)
memory usage: 248.0+ bytes

```

Prints a summary of columns count and its dtypes but not per column information:

```

>>> df.info(verbose=False)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5 entries, 0 to 4
Columns: 3 entries, int_col to float_col
dtypes: float64(1), int64(1), object(1)
memory usage: 248.0+ bytes

```

Pipe output of DataFrame.info to buffer instead of sys.stdout, get buffer content and writes to a text file:

```

>>> import io
>>> buffer = io.StringIO()
>>> df.info(buf=buffer)
>>> s = buffer.getvalue()
>>> with open("df_info.txt", "w",
...         encoding="utf-8") as f: # doctest: +SKIP
...     f.write(s)
260

```

The `memory_usage` parameter allows deep introspection mode, specially useful for big DataFrames and fine-tune memory optimization:

```

>>> random_strings_array = np.random.choice(['a', 'b', 'c'], 10 ** 6)
>>> df = pd.DataFrame({
...     'column_1': np.random.choice(['a', 'b', 'c'], 10 ** 6),
...     'column_2': np.random.choice(['a', 'b', 'c'], 10 ** 6),
...     'column_3': np.random.choice(['a', 'b', 'c'], 10 ** 6)
... })
>>> df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   column_1    1000000 non-null  object
1   column_2    1000000 non-null  object
2   column_3    1000000 non-null  object
dtypes: object(3)
memory usage: 22.9+ MB

>>> df.info(memory_usage='deep')
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000000 entries, 0 to 999999
Data columns (total 3 columns):
#   Column      Non-Null Count  Dtype
---  -
0   column_1    1000000 non-null  object
1   column_2    1000000 non-null  object
2   column_3    1000000 non-null  object
dtypes: object(3)
memory usage: 165.9 MB

```

insert(self, loc, column, value, allow_duplicates: 'bool' = False) -> 'None'
Insert column into DataFrame at specified location.

Raises a ValueError if `column` is already contained in the DataFrame,

unless `allow_duplicates` is set to True.

Parameters

loc : int
 Insertion index. Must verify $0 \leq \text{loc} \leq \text{len}(\text{columns})$.
column : str, number, or hashable object
 Label of the inserted column.
value : int, Series, or array-like
allow_duplicates : bool, optional

See Also

Index.insert : Insert new item by index.

Examples

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df
 col1 col2
0 1 3
1 2 4
>>> df.insert(1, "newcol", [99, 99])
>>> df
 col1 newcol col2
0 1 99 3
1 2 99 4
>>> df.insert(0, "col1", [100, 100], allow_duplicates=True)
>>> df
 col1 col1 newcol col2
0 100 1 99 3
1 100 2 99 4

Notice that pandas uses index alignment in case of `value` from type
`Series`:

```
>>> df.insert(0, "col0", pd.Series([5, 6], index=[1, 2]))  
>>> df  
   col0  col1  col1  newcol  col2  
0  NaN  100     1      99     3  
1  5.0  100     2      99     4
```

interpolate(self: 'DataFrame', method: 'str' = 'linear', axis: 'Axis' = 0, limit:
'int | None' = None, inplace: 'bool' = False, limit_direction: 'str | None' = None,
limit_area: 'str | None' = None, downcast: 'str | None' = None, **kwargs) ->
'DataFrame | None'

Fill NaN values using an interpolation method.

Please note that only ``method='linear'`` is supported for
DataFrame/Series with a MultiIndex.

Parameters

method : str, default 'linear'
 Interpolation technique to use. One of:

- * 'linear': Ignore the index and treat the values as equally spaced. This is the only method supported on MultiIndexes.
- * 'time': Works on daily and higher resolution data to interpolate given length of interval.
- * 'index', 'values': use the actual numerical values of the index.
- * 'pad': Fill in NaNs using existing values.
- * 'nearest', 'zero', 'slinear', 'quadratic', 'cubic', 'spline', 'barycentric', 'polynomial': Passed to `scipy.interpolate.interpld`. These methods use the numerical values of the index. Both 'polynomial' and 'spline' require that you also specify an `order` (int), e.g. ``df.interpolate(method='polynomial', order=5)``.
- * 'krogh', 'piecewise_polynomial', 'spline', 'pchip', 'akima', 'cubicspline': Wrappers around the SciPy interpolation methods of similar names. See `Notes`.
- * 'from_derivatives': Refers to `scipy.interpolate.BPoly.from_derivatives` which replaces 'piecewise_polynomial' interpolation method in scipy 0.18.

axis : {{0 or 'index', 1 or 'columns', None}}, default None
 Axis to interpolate along.

limit : int, optional
 Maximum number of consecutive NaNs to fill. Must be greater than 0.

inplace : bool, default False
 Update the data in place if possible.
limit_direction : {{'forward', 'backward', 'both'}}, Optional
 Consecutive NaNs will be filled in this direction.

If limit is specified:

```

        * If 'method' is 'pad' or 'ffill', 'limit_direction' must be
'forward'.
        * If 'method' is 'backfill' or 'bfill', 'limit_direction' must be
'backwards'.

    If 'limit' is not specified:
        * If 'method' is 'backfill' or 'bfill', the default is 'backward'
        * else the default is 'forward'

    .. versionchanged:: 1.1.0
        raises ValueError if `limit_direction` is 'forward' or 'both' and
        method is 'backfill' or 'bfill'.
        raises ValueError if `limit_direction` is 'backward' or 'both' and
        method is 'pad' or 'ffill'.

limit_area : {{None, 'inside', 'outside'}}, default None
    If limit is specified, consecutive NaNs will be filled with this
    restriction.

    * ``None``: No fill restriction.
    * 'inside': Only fill NaNs surrounded by valid values
      (interpolate).
    * 'outside': Only fill NaNs outside valid values (extrapolate).

downcast : optional, 'infer' or None, defaults to None
    Downcast dtypes if possible.
`**kwargs` : optional
    Keyword arguments to pass on to the interpolating function.

Returns
-----
Series or DataFrame or None
    Returns the same object type as the caller, interpolated at
    some or all ``NaN`` values or None if ``inplace=True``.

See Also
-----
fillna : Fill missing values using different methods.
scipy.interpolate.Akima1DInterpolator : Piecewise cubic polynomials
    (Akima interpolator).
scipy.interpolate.BPoly.from_derivatives : Piecewise polynomial in the
    Bernstein basis.
scipy.interpolate.interpld : Interpolate a 1-D function.
scipy.interpolate.KroghInterpolator : Interpolate polynomial (Krogh
    interpolator).
scipy.interpolate.PchipInterpolator : PCHIP 1-d monotonic cubic
    interpolation.
scipy.interpolate.CubicSpline : Cubic spline data interpolator.

Notes
-----
The 'krogh', 'piecewise_polynomial', 'spline', 'pchip' and 'akima'
methods are wrappers around the respective SciPy implementations of
similar names. These use the actual numerical values of the index.
For more information on their behavior, see the
`SciPy documentation
<https://docs.scipy.org/doc/scipy/reference/interpolate.html#univariate-
interpolation>`__
and `SciPy tutorial
<https://docs.scipy.org/doc/scipy/reference/tutorial/interpolate.html>`__.

Examples
-----
Filling in ``NaN`` in a :class:`~pandas.Series` via linear
interpolation.

>>> s = pd.Series([0, 1, np.nan, 3])
>>> s
0    0.0
1    1.0
2    NaN
3    3.0
dtype: float64
>>> s.interpolate()
0    0.0
1    1.0
2    2.0
3    3.0
dtype: float64

Filling in ``NaN`` in a Series by padding, but filling at most two
consecutive ``NaN`` at a time.

>>> s = pd.Series([np.nan, "single_one", np.nan,
...                "fill_two_more", np.nan, np.nan, np.nan,
...                4.71, np.nan])
>>> s

```

```

0      NaN
1    single_one
2      NaN
3    fill_two_more
4      NaN
5      NaN
6      NaN
7      4.71
8      NaN
dtype: object
>>> s.interpolate(method='pad', limit=2)
0      NaN
1    single_one
2    single_one
3    fill_two_more
4    fill_two_more
5    fill_two_more
6      NaN
7      4.71
8      4.71
dtype: object

```

Filling in ``NaN`` in a Series via polynomial interpolation or splines:
Both 'polynomial' and 'spline' methods require that you also specify
an ``order`` (int).

```

>>> s = pd.Series([0, 2, np.nan, 8])
>>> s.interpolate(method='polynomial', order=2)
0    0.000000
1    2.000000
2    4.666667
3    8.000000
dtype: float64

```

Fill the DataFrame forward (that is, going down) along each column
using linear interpolation.

Note how the last entry in column 'a' is interpolated differently,
because there is no entry after it to use for interpolation.
Note how the first entry in column 'b' remains ``NaN``, because there
is no entry before it to use for interpolation.

```

>>> df = pd.DataFrame([(0.0, np.nan, -1.0, 1.0),
...                     (np.nan, 2.0, np.nan, np.nan),
...                     (2.0, 3.0, np.nan, 9.0),
...                     (np.nan, 4.0, -4.0, 16.0)],
...                     columns=list('abcd'))
>>> df
   a    b    c    d
0  0.0 NaN -1.0  1.0
1 NaN  2.0 NaN  NaN
2  2.0  3.0 NaN  9.0
3 NaN  4.0 -4.0 16.0
>>> df.interpolate(method='linear', limit_direction='forward', axis=0)
   a    b    c    d
0  0.0 NaN -1.0  1.0
1  1.0  2.0 -2.0  5.0
2  2.0  3.0 -3.0  9.0
3  2.0  4.0 -4.0 16.0

```

Using polynomial interpolation.

```

>>> df['d'].interpolate(method='polynomial', order=2)
0    1.0
1    4.0
2    9.0
3   16.0
Name: d, dtype: float64

```

`isin(self, values) -> 'DataFrame'`

Whether each element in the DataFrame is contained in values.

Parameters

values : iterable, Series, DataFrame or dict

The result will only be true at a location if all the
labels match. If ``values`` is a Series, that's the index. If
``values`` is a dict, the keys must be the column names,
which must match. If ``values`` is a DataFrame,
then both the index and column labels must match.

Returns

DataFrame

DataFrame of booleans showing whether each element in the DataFrame
is contained in values.

See Also

DataFrame.eq: Equality test for DataFrame.

Series.isin: Equivalent method on Series.

Series.str.contains: Test if pattern or regex is contained within a string of a Series or Index.

Examples

```
>>> df = pd.DataFrame({'num_legs': [2, 4], 'num_wings': [2, 0]},
...                     index=['falcon', 'dog'])
```

```
>>> df
   num_legs  num_wings
falcon      2         2
dog         4         0
```

When ``values`` is a list check whether every value in the DataFrame is present in the list (which animals have 0 or 2 legs or wings)

```
>>> df.isin([0, 2])
   num_legs  num_wings
falcon    True      True
dog       False     True
```

When ``values`` is a dict, we can pass values to check for each column separately:

```
>>> df.isin({'num_wings': [0, 3]})
   num_legs  num_wings
falcon    False     False
dog       False     True
```

When ``values`` is a Series or DataFrame the index and column must match. Note that 'falcon' does not match based on the number of legs in df2.

```
>>> other = pd.DataFrame({'num_legs': [8, 2], 'num_wings': [0, 2]},
...                       index=['spider', 'falcon'])
>>> df.isin(other)
   num_legs  num_wings
falcon    True      True
dog       False     False
```

isna(self) -> 'DataFrame'
Detect missing values.

Return a boolean same-sized object indicating if the values are NA. NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values.

Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is an NA value.

See Also

DataFrame.isnull : Alias of isna.

DataFrame.notna : Boolean inverse of isna.

DataFrame.dropna : Omit axes labels with missing values.

isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                         born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                               pd.Timestamp('1940-04-25')],
...                         name=['Alfred', 'Batman', ''],
...                         toy=[None, 'Batmobile', 'Joker']))
>>> df
   age      born      name      toy
0  5.0      NaT  Alfred      None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25      Joker
```

```
>>> df.isna()
   age  born  name  toy
0 False  True False  True
1 False False False False
2  True False False False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64
```

```
>>> ser.isna()
0    False
1    False
2     True
dtype: bool
```

isnull(self) -> 'DataFrame'
Detect missing values.

Return a boolean same-sized object indicating if the values are NA.
NA values, such as None or :attr:`numpy.NaN`, gets mapped to True values.
Everything else gets mapped to False values. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``).

Returns

DataFrame
Mask of bool values for each element in DataFrame that indicates whether an element is an NA value.

See Also

DataFrame.isnull : Alias of isna.
DataFrame.notna : Boolean inverse of isna.
DataFrame.dropna : Omit axes labels with missing values.
isna : Top-level isna.

Examples

Show which entries in a DataFrame are NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                        born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                               pd.Timestamp('1940-04-25')],
...                        name=['Alfred', 'Batman', ''],
...                        toy=[None, 'Batmobile', 'Joker']))
>>> df
```

	age	born	name	toy
0	5.0	NaT	Alfred	None
1	6.0	1939-05-27	Batman	Batmobile
2	NaN	1940-04-25		Joker

```
>>> df.isna()
   age  born  name  toy
0  False  True  False  True
1  False  False  False  False
2   True  False  False  False
```

Show which entries in a Series are NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64
```

```
>>> ser.isna()
0    False
1    False
2     True
dtype: bool
```

items(self) -> 'Iterable[tuple[Hashable, Series]]'
Iterate over (column name, Series) pairs.

Iterates over the DataFrame columns, returning a tuple with the column name and the content as a Series.

Yields

label : object
The column names for the DataFrame being iterated over.
content : Series
The column entries belonging to each label, as a Series.

See Also

DataFrame.iterrows : Iterate over DataFrame rows as
(index, Series) pairs.

DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
of the values.

Examples

```
>>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],  
...                    'population': [1864, 22000, 80000]},  
...                    index=['panda', 'polar', 'koala'])  
>>> df
```

```
   species  population  
panda   bear      1864  
polar   bear     22000  
koala  marsupial  80000
```

```
>>> for label, content in df.items():  
...     print(f'label: {label}')  
...     print(f'content: {content}', sep='\n')  
...
```

label: species

content:

```
panda      bear  
polar      bear  
koala      marsupial
```

Name: species, dtype: object

label: population

content:

```
panda      1864  
polar      22000  
koala      80000
```

Name: population, dtype: int64

iteritems(self) -> 'Iterable[tuple[Hashable, Series]]'
Iterate over (column name, Series) pairs.

Iterates over the DataFrame columns, returning a tuple with
the column name and the content as a Series.

Yields

label : object

The column names for the DataFrame being iterated over.

content : Series

The column entries belonging to each label, as a Series.

See Also

DataFrame.iterrows : Iterate over DataFrame rows as
(index, Series) pairs.

DataFrame.itertuples : Iterate over DataFrame rows as namedtuples
of the values.

Examples

```
>>> df = pd.DataFrame({'species': ['bear', 'bear', 'marsupial'],  
...                    'population': [1864, 22000, 80000]},  
...                    index=['panda', 'polar', 'koala'])  
>>> df
```

```
   species  population  
panda   bear      1864  
polar   bear     22000  
koala  marsupial  80000
```

```
>>> for label, content in df.items():  
...     print(f'label: {label}')  
...     print(f'content: {content}', sep='\n')  
...
```

label: species

content:

```
panda      bear  
polar      bear  
koala      marsupial
```

Name: species, dtype: object

label: population

content:

```
panda      1864  
polar      22000  
koala      80000
```

Name: population, dtype: int64

iterrows(self) -> 'Iterable[tuple[Hashable, Series]]'
Iterate over DataFrame rows as (index, Series) pairs.

Yields

index : label or tuple of label

The index of the row. A tuple for a `MultiIndex`.
data : Series
The data of the row as a Series.

See Also

DataFrame.itertuples : Iterate over DataFrame rows as namedtuples of the values.

DataFrame.items : Iterate over (column name, Series) pairs.

Notes

- 1. Because ``iterrows`` returns a Series for each row, it does **not** preserve dtypes across the rows (dtypes are preserved across columns for DataFrames). For example,

```
>>> df = pd.DataFrame([[1, 1.5]], columns=['int', 'float'])
>>> row = next(df.iterrows())[1]
>>> row
int      1.0
float    1.5
Name: 0, dtype: float64
>>> print(row['int'].dtype)
float64
>>> print(df['int'].dtype)
int64
```

To preserve dtypes while iterating over the rows, it is better to use :meth:`itertuples` which returns namedtuples of the values and which is generally faster than ``iterrows``.

2. You should **never** modify something you are iterating over. This is not guaranteed to work in all cases. Depending on the data types, the iterator returns a copy and not a view, and writing to it will have no effect.

itertuples(self, index: 'bool' = True, name: 'str | None' = 'Pandas') -> Iterable[tuple[Any, ...]]
Iterate over DataFrame rows as namedtuples.

Parameters

index : bool, default True
If True, return the index as the first element of the tuple.
name : str or None, default "Pandas"
The name of the returned namedtuples or None to return regular tuples.

Returns

iterator
An object to iterate over namedtuples for each row in the DataFrame with the first field possibly being the index and following fields being the column values.

See Also

DataFrame.iterrows : Iterate over DataFrame rows as (index, Series) pairs.
DataFrame.items : Iterate over (column name, Series) pairs.

Notes

The column names will be renamed to positional names if they are invalid Python identifiers, repeated, or start with an underscore. On python versions < 3.7 regular tuples are returned for DataFrames with a large number of columns (>254).

Examples

>>> df = pd.DataFrame({'num_legs': [4, 2], 'num_wings': [0, 2]},
... index=['dog', 'hawk'])
>>> df

	num_legs	num_wings
dog	4	0
hawk	2	2

```
>>> for row in df.itertuples():
...     print(row)
...
Pandas(Index='dog', num_legs=4, num_wings=0)
Pandas(Index='hawk', num_legs=2, num_wings=2)
```

By setting the `index` parameter to False we can remove the index as the first element of the tuple:

```
>>> for row in df.itertuples(index=False):
...     print(row)
```



```

...
Pandas(num_legs=4, num_wings=0)
Pandas(num_legs=2, num_wings=2)

With the `name` parameter set we set a custom name for the yielded
namedtuples:

>>> for row in df.itertuples(name='Animal'):
...     print(row)
...
Animal(Index='dog', num_legs=4, num_wings=0)
Animal(Index='hawk', num_legs=2, num_wings=2)

join(self, other: 'FrameOrSeriesUnion', on: 'IndexLabel | None' = None, how:
'str' = 'left', lsuffix: 'str' = '', rsuffix: 'str' = '', sort: 'bool' = False) ->
'DataFrame'
    Join columns of another DataFrame.

    Join columns with `other` DataFrame either on index or on a key
    column. Efficiently join multiple DataFrame objects by index at once by
    passing a list.

    Parameters
    -----
    other : DataFrame, Series, or list of DataFrame
        Index should be similar to one of the columns in this one. If a
        Series is passed, its name attribute must be set, and that will be
        used as the column name in the resulting joined DataFrame.
    on : str, list of str, or array-like, optional
        Column or index level name(s) in the caller to join on the index
        in `other`, otherwise joins index-on-index. If multiple
        values given, the `other` DataFrame must have a MultiIndex. Can
        pass an array as the join key if it is not already contained in
        the calling DataFrame. Like an Excel VLOOKUP operation.
    how : {'left', 'right', 'outer', 'inner'}, default 'left'
        How to handle the operation of the two objects.

        * left: use calling frame's index (or column if on is specified)
        * right: use `other`'s index.
        * outer: form union of calling frame's index (or column if on is
          specified) with `other`'s index, and sort it.
          lexicographically.
        * inner: form intersection of calling frame's index (or column if
          on is specified) with `other`'s index, preserving the order
          of the calling's one.
    lsuffix : str, default ''
        Suffix to use from left frame's overlapping columns.
    rsuffix : str, default ''
        Suffix to use from right frame's overlapping columns.
    sort : bool, default False
        Order result DataFrame lexicographically by the join key. If False,
        the order of the join key depends on the join type (how keyword).

    Returns
    -----
    DataFrame
        A dataframe containing columns from both the caller and `other`.

    See Also
    -----
    DataFrame.merge : For column(s)-on-column(s) operations.

    Notes
    -----
    Parameters `on`, `lsuffix`, and `rsuffix` are not supported when
    passing a list of `DataFrame` objects.

    Support for specifying index levels as the `on` parameter was added
    in version 0.23.0.

    Examples
    -----
    >>> df = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3', 'K4', 'K5'],
...                        'A': ['A0', 'A1', 'A2', 'A3', 'A4', 'A5']})

    >>> df
       key  A
0    K0  A0
1    K1  A1
2    K2  A2
3    K3  A3
4    K4  A4
5    K5  A5

    >>> other = pd.DataFrame({'key': ['K0', 'K1', 'K2'],
...                           'B': ['B0', 'B1', 'B2']})

```

```
>>> other
   key  B
0  K0  B0
1  K1  B1
2  K2  B2
```

Join DataFrames using their indexes.

```
>>> df.join(other, lsuffix='_caller', rsuffix='_other')
   key_caller  A key_other  B
0          K0  A0          K0  B0
1          K1  A1          K1  B1
2          K2  A2          K2  B2
3          K3  A3          NaN NaN
4          K4  A4          NaN NaN
5          K5  A5          NaN NaN
```

If we want to join using the key columns, we need to set key to be the index in both `df` and `other`. The joined DataFrame will have key as its index.

```
>>> df.set_index('key').join(other.set_index('key'))
   key  A  B
key
K0    A0  B0
K1    A1  B1
K2    A2  B2
K3    A3  NaN
K4    A4  NaN
K5    A5  NaN
```

Another option to join using the key columns is to use the `on` parameter. DataFrame.join always uses `other`'s index but we can use any column in `df`. This method preserves the original DataFrame's index in the result.

```
>>> df.join(other.set_index('key'), on='key')
   key  A  B
0  K0  A0  B0
1  K1  A1  B1
2  K2  A2  B2
3  K3  A3  NaN
4  K4  A4  NaN
5  K5  A5  NaN
```

kurt(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
Return unbiased kurtosis over requested axis.

Kurtosis obtained using Fisher's definition of kurtosis (kurtosis of normal == 0.0). Normalized by N-1.

Parameters

```
-----
axis : {index (0), columns (1)}
    Axis for the function to be applied on.
skipna : bool, default True
    Exclude NA/null values when computing the result.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
numeric_only : bool, default None
    Include only float, int, boolean columns. If None, will attempt to use
    everything, then use only numeric data. Not implemented for Series.
**kwargs
    Additional keyword arguments to be passed to the function.
```

Returns

```
-----
Series or DataFrame (if level specified)
```

```
kurtosis = kurt(self, axis=None, skipna=None, level=None, numeric_only=None,
**kwargs)
```

```
le(self, other, axis='columns', level=None)
    Get Less than or equal to of dataframe and other, element-wise (binary
operator `le`).
```

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

```
-----
other : scalar, sequence, Series, or DataFrame
    Any single or multiple element data structure, or list-like object.
```

axis : {0 or 'index', 1 or 'columns'}, default 'columns'
Whether to compare by the index (0 or 'index') or columns
(1 or 'columns').
level : int or label
Broadcast across a level, matching Index values on the passed
MultiIndex level.

Returns

DataFrame of bool
Result of the comparison.

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.
DataFrame.ne : Compare DataFrames for inequality elementwise.
DataFrame.le : Compare DataFrames for less than inequality
or equality elementwise.
DataFrame.lt : Compare DataFrames for strictly less than
inequality elementwise.
DataFrame.ge : Compare DataFrames for greater than inequality
or equality elementwise.
DataFrame.gt : Compare DataFrames for strictly greater than
inequality elementwise.

Notes

Mismatched indices will be unioned together.
`NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

>>> df = pd.DataFrame({'cost': [250, 150, 100],
... 'revenue': [100, 250, 300]},
... index=['A', 'B', 'C'])
>>> df
 cost revenue
A 250 100
B 150 250
C 100 300

Comparison with a scalar, using either the operator or method:

```
>>> df == 100  
   cost  revenue  
A  False    True  
B  False    False  
C   True    False
```

```
>>> df.eq(100)  
   cost  revenue  
A  False    True  
B  False    False  
C   True    False
```

When `other` is a :class:`Series`, the columns of a DataFrame are aligned
with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])  
   cost  revenue  
A   True    True  
B   True   False  
C  False    True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')  
   cost  revenue  
A   True   False  
B   True    True  
C   True    True  
D   True    True
```

When comparing to an arbitrary sequence, the number of columns must
match the number elements in `other`:

```
>>> df == [250, 100]  
   cost  revenue  
A   True    True  
B  False   False  
C  False   False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')  
   cost  revenue  
A   True   False
```

```
B False    True
C  True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
```

```
>>> other
   revenue
A       300
B       250
C       100
D       150
```

```
>>> df.gt(other)
   cost  revenue
A  False   False
B  False   False
C  False    True
D  False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                              index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                     ['A', 'B', 'C', 'A', 'B', 'C']])
```

```
>>> df_multindex
   cost  revenue
Q1 A   250     100
   B   150     250
   C   100     300
Q2 A   150     200
   B   300     175
   C   220     225
```

```
>>> df.le(df_multindex, level=1)
```

```
   cost  revenue
Q1 A   True    True
   B   True    True
   C   True    True
Q2 A  False    True
   B   True   False
   C   True   False
```

```
lookup(self, row_labels: 'Sequence[IndexLabel]', col_labels:
'Sequence[IndexLabel]') -> 'np.ndarray'
```

Label-based "fancy indexing" function for DataFrame.

Given equal-length arrays of row and column labels, return an array of the values corresponding to each (row, col) pair.

.. deprecated:: 1.2.0

DataFrame.lookup is deprecated,
use DataFrame.melt and DataFrame.loc instead.

For further details see

:ref:`Looking up values by index/column labels <indexing.lookup>`.

Parameters

row_labels : sequence

The row labels to use for lookup.

col_labels : sequence

The column labels to use for lookup.

Returns

numpy.ndarray

The found values.

```
lt(self, other, axis='columns', level=None)
```

Get Less than of dataframe and other, element-wise (binary operator `lt`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<`, `<=`, `>`, `>=` with support to choose axis (rows or columns) and level for comparison.

Parameters

other : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns

(1 or 'columns').

level : int or label

Broadcast across a level, matching Index values on the passed

MultiIndex level.

Returns

DataFrame of bool

Result of the comparison.

See Also

DataFrame.eq : Compare DataFrames for equality elementwise.

DataFrame.ne : Compare DataFrames for inequality elementwise.

DataFrame.le : Compare DataFrames for less than inequality
or equality elementwise.

DataFrame.lt : Compare DataFrames for strictly less than
inequality elementwise.

DataFrame.ge : Compare DataFrames for greater than inequality
or equality elementwise.

DataFrame.gt : Compare DataFrames for strictly greater than
inequality elementwise.

Notes

Mismatched indices will be unioned together.

`NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                    'revenue': [100, 250, 300]},
...                   index=['A', 'B', 'C'])
>>> df
   cost  revenue
A   250     100
B   150     250
C   100     300
```

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
   cost  revenue
A  False     True
B  False    False
C   True    False
```

```
>>> df.eq(100)
   cost  revenue
A  False     True
B  False    False
C   True    False
```

When `other` is a `:class:`Series``, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
   cost  revenue
A   True     True
B   True    False
C  False     True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
   cost  revenue
A   True    False
B   True     True
C   True     True
D   True     True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
   cost  revenue
A   True     True
B  False    False
C  False    False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
   cost  revenue
A   True    False
B  False     True
C   True    False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                        index=['A', 'B', 'C', 'D'])
>>> other
   revenue
A       300
B       250
C       100
D       150
```

```
>>> df.gt(other)
      cost  revenue
A  False   False
B  False   False
C  False    True
D  False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                               'revenue': [100, 250, 300, 200, 175, 225]},
...                               index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                       ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250     100
   B    150     250
   C    100     300
Q2 A    150     200
   B    300     175
   C    220     225
```

```
>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A    True     True
   B    True     True
   C    True     True
Q2 A   False     True
   B    True    False
   C    True    False
```

`mad(self, axis=None, skipna=None, level=None)`
Return the mean absolute deviation of the values over the requested axis.

Parameters

`axis` : {index (0), columns (1)}
Axis for the function to be applied on.
`skipna` : bool, default None
Exclude NA/null values when computing the result.
`level` : int or level name, default None
If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

Returns

Series or DataFrame (if level specified)

`mask(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=<no_default>)`

Replace values where the condition is True.

Parameters

`cond` : bool Series/DataFrame, array-like, or callable
Where `'cond'` is False, keep the original value. Where True, replace with corresponding value from `'other'`.
If `'cond'` is callable, it is computed on the Series/DataFrame and should return boolean Series/DataFrame or array. The callable must not change input Series/DataFrame (though pandas doesn't check it).
`other` : scalar, Series/DataFrame, or callable
Entries where `'cond'` is True are replaced with corresponding value from `'other'`.
If `other` is callable, it is computed on the Series/DataFrame and should return scalar or Series/DataFrame. The callable must not change input Series/DataFrame (though pandas doesn't check it).
`inplace` : bool, default False
Whether to perform the operation in place on the data.
`axis` : int, default None
Alignment axis if needed.
`level` : int, default None
Alignment level if needed.
`errors` : str, {'raise', 'ignore'}, default 'raise'
Note that currently this parameter won't affect the results and will always coerce to a suitable dtype.

- 'raise' : allow exceptions to be raised.
- 'ignore' : suppress exceptions. On error return original object.

`try_cast : bool, default None`
Try to cast the result back to the input type (if possible).

.. deprecated:: 1.3.0
Manually cast back if necessary.

Returns

Same type as caller or None if `inplace=True`.

See Also

`:func:`DataFrame.where`` : Return an object of same shape as `self`.

Notes

The mask method is an application of the if-then idiom. For each element in the calling DataFrame, if `cond` is `False` the element is used; otherwise the corresponding element from the DataFrame `other` is used.

The signature for `:func:`DataFrame.where`` differs from `:func:`numpy.where``. Roughly `df1.where(m, df2)` is equivalent to `np.where(m, df1, df2)`.

For further details and examples see the `mask` documentation in `:ref:`indexing <indexing.where_mask>`.

Examples

```
>>> s = pd.Series(range(5))
>>> s.where(s > 0)
0    NaN
1    1.0
2    2.0
3    3.0
4    4.0
dtype: float64
>>> s.mask(s > 0)
0    0.0
1    NaN
2    NaN
3    NaN
4    NaN
dtype: float64

>>> s.where(s > 1, 10)
0    10
1    10
2     2
3     3
4     4
dtype: int64
>>> s.mask(s > 1, 10)
0     0
1     1
2    10
3    10
4    10
dtype: int64

>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])
>>> df
   A  B
0  0  1
1  2  3
2  4  5
3  6  7
4  8  9
>>> m = df % 3 == 0
>>> df.where(m, -df)
   A  B
0  0 -1
1 -2  3
2 -4 -5
3  6 -7
4 -8  9
>>> df.where(m, -df) == np.where(m, df, -df)
   A  B
0  True  True
1  True  True
2  True  True
3  True  True
4  True  True
>>> df.where(m, -df) == df.mask(~m, -df)
```

	A	B
0	True	True
1	True	True
2	True	True
3	True	True
4	True	True

```
max(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
    Return the maximum of the values over the requested axis.
```

If you want the `*index*` of the maximum, use ```idxmax```. This is the equivalent of the ```numpy.ndarray``` method ```argmax```.

Parameters

```
axis : {index (0), columns (1)}
    Axis for the function to be applied on.
skipna : bool, default True
    Exclude NA/null values when computing the result.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
numeric_only : bool, default None
    Include only float, int, boolean columns. If None, will attempt to use
    everything, then use only numeric data. Not implemented for Series.
**kwargs
    Additional keyword arguments to be passed to the function.
```

Returns

```
Series or DataFrame (if level specified)
```

See Also

```
Series.sum : Return the sum.
Series.min : Return the minimum.
Series.max : Return the maximum.
Series.idxmin : Return the index of the minimum.
Series.idxmax : Return the index of the maximum.
DataFrame.sum : Return the sum over the requested axis.
DataFrame.min : Return the minimum over the requested axis.
DataFrame.max : Return the maximum over the requested axis.
DataFrame.idxmin : Return the index of the minimum over the requested axis.
DataFrame.idxmax : Return the index of the maximum over the requested axis.
```

Examples

```
>>> idx = pd.MultiIndex.from_arrays([
...     ['warm', 'warm', 'cold', 'cold'],
...     ['dog', 'falcon', 'fish', 'spider']],
...     names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
>>> s
blooded animal
warm      dog      4
          falcon   2
cold      fish     0
          spider   8
Name: legs, dtype: int64

>>> s.max()
8
```

```
mean(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
    Return the mean of the values over the requested axis.
```

Parameters

```
axis : {index (0), columns (1)}
    Axis for the function to be applied on.
skipna : bool, default True
    Exclude NA/null values when computing the result.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
numeric_only : bool, default None
    Include only float, int, boolean columns. If None, will attempt to use
    everything, then use only numeric data. Not implemented for Series.
**kwargs
    Additional keyword arguments to be passed to the function.
```

Returns

```
Series or DataFrame (if level specified)
```

```
median(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
    Return the median of the values over the requested axis.
```



```

Parameters
-----
axis : {index (0), columns (1)}
    Axis for the function to be applied on.
skipna : bool, default True
    Exclude NA/null values when computing the result.
level : int or level name, default None
    If the axis is a MultiIndex (hierarchical), count along a
    particular level, collapsing into a Series.
numeric_only : bool, default None
    Include only float, int, boolean columns. If None, will attempt to use
    everything, then use only numeric data. Not implemented for Series.
**kwargs
    Additional keyword arguments to be passed to the function.

Returns
-----
Series or DataFrame (if level specified)

melt(self, id_vars=None, value_vars=None, var_name=None, value_name='value',
col_level: 'Level | None' = None, ignore_index: 'bool' = True) -> 'DataFrame'
    Unpivot a DataFrame from wide to long format, optionally leaving identifiers
set.

```

This function is useful to massage a DataFrame into a format where one or more columns are identifier variables ('id_vars'), while all other columns, considered measured variables ('value_vars'), are "unpivoted" to the row axis, leaving just two non-identifier columns, 'variable' and 'value'.

```

Parameters
-----
id_vars : tuple, list, or ndarray, optional
    Column(s) to use as identifier variables.
value_vars : tuple, list, or ndarray, optional
    Column(s) to unpivot. If not specified, uses all columns that
    are not set as `id_vars`.
var_name : scalar
    Name to use for the 'variable' column. If None it uses
    ``frame.columns.name`` or 'variable'.
value_name : scalar, default 'value'
    Name to use for the 'value' column.
col_level : int or str, optional
    If columns are a MultiIndex then use this level to melt.
ignore_index : bool, default True
    If True, original index is ignored. If False, the original index is
retained.
    Index labels will be repeated as necessary.

.. versionadded:: 1.1.0

```

```

Returns
-----
DataFrame
    Unpivoted DataFrame.

```

```

See Also
-----
melt : Identical method.
pivot_table : Create a spreadsheet-style pivot table as a DataFrame.
DataFrame.pivot : Return reshaped DataFrame organized
    by given index / column values.
DataFrame.explode : Explode a DataFrame from list-like
    columns to long format.

```

```

Examples
-----
>>> df = pd.DataFrame({'A': {0: 'a', 1: 'b', 2: 'c'},
...                     'B': {0: 1, 1: 3, 2: 5},
...                     'C': {0: 2, 1: 4, 2: 6}})
>>> df
   A  B  C
0  a  1  2
1  b  3  4
2  c  5  6

>>> df.melt(id_vars=['A'], value_vars=['B'])
   A variable  value
0  a         B      1
1  b         B      3
2  c         B      5

>>> df.melt(id_vars=['A'], value_vars=['B', 'C'])
   A variable  value
0  a         B      1
1  b         B      3

```

2	c	B	5
3	a	C	2
4	b	C	4
5	c	C	6

The names of 'variable' and 'value' columns can be customized:

```
>>> df.melt(id_vars=['A'], value_vars=['B'],
...         var_name='myVarname', value_name='myValname')
   A myVarname myValname
0  a          B         1
1  b          B         3
2  c          B         5
```

Original index values can be kept around:

```
>>> df.melt(id_vars=['A'], value_vars=['B', 'C'], ignore_index=False)
   A variable value
0  a          B     1
1  b          B     3
2  c          B     5
0  a          C     2
1  b          C     4
2  c          C     6
```

If you have multi-index columns:

```
>>> df.columns = [list('ABC'), list('DEF')]
>>> df
   A B C
   D E F
0  a 1 2
1  b 3 4
2  c 5 6

>>> df.melt(col_level=0, id_vars=['A'], value_vars=['B'])
   A variable value
0  a          B     1
1  b          B     3
2  c          B     5

>>> df.melt(id_vars=[('A', 'D')], value_vars=[('B', 'E')])
   (A, D) variable_0 variable_1 value
0      a          B           E     1
1      b          B           E     3
2      c          B           E     5
```

`memory_usage(self, index: 'bool' = True, deep: 'bool' = False) -> 'Series'`
Return the memory usage of each column in bytes.

The memory usage can optionally include the contribution of the index and elements of 'object' dtype.

This value is displayed in 'DataFrame.info' by default. This can be suppressed by setting 'pandas.options.display.memory_usage' to False.

Parameters

index : bool, default True
Specifies whether to include the memory usage of the DataFrame's index in returned Series. If 'index=True', the memory usage of the index is the first item in the output.

deep : bool, default False
If True, introspect the data deeply by interrogating 'object' dtypes for system-level memory consumption, and include it in the returned values.

Returns

Series
A Series whose index is the original column names and whose values is the memory usage of each column in bytes.

See Also

`numpy.ndarray.nbytes` : Total bytes consumed by the elements of an ndarray.

`Series.memory_usage` : Bytes consumed by a Series.

`Categorical` : Memory-efficient array for string values with many repeated values.

`DataFrame.info` : Concise summary of a DataFrame.

Examples

```
>>> dtypes = ['int64', 'float64', 'complex128', 'object', 'bool']
>>> data = dict([(t, np.ones(shape=5000, dtype=int).astype(t))
...              for t in dtypes])
```

```

>>> df = pd.DataFrame(data)
>>> df.head()
   int64  float64      complex128  object  bool
0      1      1.0      1.0+0.0j      1  True
1      1      1.0      1.0+0.0j      1  True
2      1      1.0      1.0+0.0j      1  True
3      1      1.0      1.0+0.0j      1  True
4      1      1.0      1.0+0.0j      1  True

>>> df.memory_usage()
Index          128
int64          40000
float64         40000
complex128      80000
object          40000
bool            5000
dtype: int64

>>> df.memory_usage(index=False)
int64          40000
float64         40000
complex128      80000
object          40000
bool            5000
dtype: int64

The memory footprint of `object` dtype columns is ignored by default:

>>> df.memory_usage(deep=True)
Index          128
int64          40000
float64         40000
complex128      80000
object         180000
bool            5000
dtype: int64

Use a Categorical for efficient storage of an object-dtype column with
many repeated values.

>>> df['object'].astype('category').memory_usage(deep=True)
5244

```

merge(self, right: 'FrameOrSeriesUnion', how: 'str' = 'inner', on: 'IndexLabel | None' = None, left_on: 'IndexLabel | None' = None, right_on: 'IndexLabel | None' = None, left_index: 'bool' = False, right_index: 'bool' = False, sort: 'bool' = False, suffixes: 'Suffixes' = ('_x', '_y'), copy: 'bool' = True, indicator: 'bool' = False, validate: 'str | None' = None) -> 'DataFrame'

Merge DataFrame or named Series objects with a database-style join.

A named Series object is treated as a DataFrame with a single named column.

The join is done on columns or indexes. If joining columns on columns, the DataFrame indexes *will be ignored*. Otherwise if joining indexes

on indexes or indexes on a column or columns, the index will be passed on. When performing a cross merge, no column specifications to merge on are allowed.

Parameters

right : DataFrame or named Series
Object to merge with.

how : {'left', 'right', 'outer', 'inner', 'cross'}, default 'inner'
Type of merge to be performed.

- * left: use only keys from left frame, similar to a SQL left outer join; preserve key order.
- * right: use only keys from right frame, similar to a SQL right outer join;
- * outer: use union of keys from both frames, similar to a SQL full outer join; sort keys lexicographically.
- * inner: use intersection of keys from both frames, similar to a SQL inner join; preserve the order of the left keys.
- * cross: creates the cartesian product from both frames, preserves the order of the left keys.

.. versionadded:: 1.2.0

on : label or list
Column or index level names to join on. These must be found in both DataFrames. If `on` is None and not merging on indexes then this defaults to the intersection of the columns in both DataFrames.

left_on : label or list, or array-like

```

        Column or index level names to join on in the left DataFrame. Can also
        be an array or list of arrays of the length of the left DataFrame.
        These arrays are treated as if they are columns.
    right_on : label or list, or array-like
        Column or index level names to join on in the right DataFrame. Can also
        be an array or list of arrays of the length of the right DataFrame.
        These arrays are treated as if they are columns.
    left_index : bool, default False
        Use the index from the left DataFrame as the join key(s). If it is a
        MultiIndex, the number of keys in the other DataFrame (either the index
        or a number of columns) must match the number of levels.
    right_index : bool, default False
        Use the index from the right DataFrame as the join key. Same caveats as
        left_index.
    sort : bool, default False
        Sort the join keys lexicographically in the result DataFrame. If False,
        the order of the join keys depends on the join type (how keyword).
    suffixes : list-like, default is ("_x", "_y")
        A length-2 sequence where each element is optionally a string
        indicating the suffix to add to overlapping column names in
        `left` and `right` respectively. Pass a value of `None` instead
        of a string to indicate that the column name from `left` or
        `right` should be left as-is, with no suffix. At least one of the
        values must not be None.
    copy : bool, default True
        If False, avoid copy if possible.
    indicator : bool or str, default False
        If True, adds a column to the output DataFrame called "_merge" with
        information on the source of each row. The column can be given a
different
        name by providing a string argument. The column will have a Categorical
        type with the value of "left_only" for observations whose merge key only
        appears in the left DataFrame, "right_only" for observations
        whose merge key only appears in the right DataFrame, and "both"
        if the observation's merge key is found in both DataFrames.

    validate : str, optional
        If specified, checks if merge is of specified type.

        * "one_to_one" or "1:1": check if merge keys are unique in both
          left and right datasets.
        * "one_to_many" or "1:m": check if merge keys are unique in left
          dataset.
        * "many_to_one" or "m:1": check if merge keys are unique in right
          dataset.
        * "many_to_many" or "m:m": allowed, but does not result in checks.

Returns
-----
DataFrame
    A DataFrame of the two merged objects.

See Also
-----
merge_ordered : Merge with optional filling/interpolation.
merge_asof : Merge on nearest keys.
DataFrame.join : Similar method using indices.

Notes
-----
Support for specifying index levels as the `on`, `left_on`, and
`right_on` parameters was added in version 0.23.0
Support for merging named Series objects was added in version 0.24.0

Examples
-----
>>> df1 = pd.DataFrame({'lkey': ['foo', 'bar', 'baz', 'foo'],
...                      'value': [1, 2, 3, 5]})
>>> df2 = pd.DataFrame({'rkey': ['foo', 'bar', 'baz', 'foo'],
...                      'value': [5, 6, 7, 8]})
>>> df1
   lkey value
0   foo     1
1   bar     2
2   baz     3
3   foo     5
>>> df2
   rkey value
0   foo     5
1   bar     6
2   baz     7
3   foo     8

Merge df1 and df2 on the lkey and rkey columns. The value columns have
the default suffixes, _x and _y, appended.

>>> df1.merge(df2, left_on='lkey', right_on='rkey')

```

	lkey	value_x	rkey	value_y
0	foo	1	foo	5
1	foo	1	foo	8
2	foo	5	foo	5
3	foo	5	foo	8
4	bar	2	bar	6
5	baz	3	baz	7

Merge DataFrames df1 and df2 with specified left and right suffixes appended to any overlapping columns.

```
>>> df1.merge(df2, left_on='lkey', right_on='rkey',
...           suffixes=('_left', '_right'))
   lkey  value_left rkey  value_right
0  foo           1  foo            5
1  foo           1  foo            8
2  foo           5  foo            5
3  foo           5  foo            8
4  bar           2  bar            6
5  baz           3  baz            7
```

Merge DataFrames df1 and df2, but raise an exception if the DataFrames have any overlapping columns.

```
>>> df1.merge(df2, left_on='lkey', right_on='rkey', suffixes=(False, False))
Traceback (most recent call last):
```

```
...
ValueError: columns overlap but no suffix specified:
Index(['value'], dtype='object')
```

```
>>> df1 = pd.DataFrame({'a': ['foo', 'bar'], 'b': [1, 2]})
>>> df2 = pd.DataFrame({'a': ['foo', 'baz'], 'c': [3, 4]})
```

```
>>> df1
   a  b
0  foo  1
1  bar  2
>>> df2
   a  c
0  foo  3
1  baz  4
```

```
>>> df1.merge(df2, how='inner', on='a')
   a  b  c
0  foo  1  3
```

```
>>> df1.merge(df2, how='left', on='a')
   a  b  c
0  foo  1  3.0
1  bar  2  NaN
```

```
>>> df1 = pd.DataFrame({'left': ['foo', 'bar']})
>>> df2 = pd.DataFrame({'right': [7, 8]})
>>> df1
   left
0  foo
1  bar
>>> df2
   right
0     7
1     8
```

```
>>> df1.merge(df2, how='cross')
   left  right
0  foo      7
1  foo      8
2  bar      7
3  bar      8
```

```
min(self, axis=None, skipna=None, level=None, numeric_only=None, **kwargs)
Return the minimum of the values over the requested axis.
```

If you want the `*index*` of the minimum, use ```idxmin```. This is the equivalent of the ```numpy.ndarray``` method ```argmin```.

Parameters

`axis` : {index (0), columns (1)}

Axis for the function to be applied on.

`skipna` : bool, default True

Exclude NA/null values when computing the result.

`level` : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

`numeric_only` : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

`**kwargs`

Additional keyword arguments to be passed to the function.

Returns

Series or DataFrame (if level specified)

See Also

Series.sum : Return the sum.

Series.min : Return the minimum.

Series.max : Return the maximum.

Series.idxmin : Return the index of the minimum.

Series.idxmax : Return the index of the maximum.

DataFrame.sum : Return the sum over the requested axis.

DataFrame.min : Return the minimum over the requested axis.

DataFrame.max : Return the maximum over the requested axis.

DataFrame.idxmin : Return the index of the minimum over the requested axis.

DataFrame.idxmax : Return the index of the maximum over the requested axis.

Examples

```
>>> idx = pd.MultiIndex.from_arrays([
...     ['warm', 'warm', 'cold', 'cold'],
...     ['dog', 'falcon', 'fish', 'spider']],
...     names=['blooded', 'animal'])
>>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
```

```
>>> s
blooded  animal
warm     dog      4
         falcon   2
cold     fish     0
         spider   8
Name: legs, dtype: int64
```

```
>>> s.min()
0
```

```
mod(self, other, axis='columns', level=None, fill_value=None)
Get Modulo of dataframe and other, element-wise (binary operator `mod`).
```

Equivalent to ``dataframe % other``, but with support to substitute a
fill_value

for missing data in one of the inputs. With reverse version, `rmod`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}

Whether to compare by the index (0 or 'index') or columns

(1 or 'columns'). For Series input, axis to match Series index on.

level : int or label

Broadcast across a level, matching Index values on the
passed MultiIndex level.

fill_value : float or None, default None

Fill existing missing (NaN) values, and any new element needed for
successful DataFrame alignment, with this value before computation.

If data in both corresponding DataFrame locations is missing
the result will be missing.

Returns

DataFrame

Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.

DataFrame.sub : Subtract DataFrames.

DataFrame.mul : Multiply DataFrames.

DataFrame.div : Divide DataFrames (float division).

DataFrame.truediv : Divide DataFrames (float division).

DataFrame.floordiv : Divide DataFrames (integer division).

DataFrame.mod : Calculate modulo (remainder after division).

DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
...                     'degrees': [360, 180, 360]},
```

```
... index=['circle', 'triangle', 'rectangle'])
```

```
>>> df
      angles  degrees
circle      0     360
triangle    3     180
rectangle   4     360
```

Add a scalar with operator version which return the same results.

```
>>> df + 1
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

```
>>> df.add(1)
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

Divide by constant with reverse version.

```
>>> df.div(10)
      angles  degrees
circle    0.0     36.0
triangle  0.3     18.0
rectangle 0.4     36.0
```

```
>>> df.rdiv(10)
      angles  degrees
circle      inf  0.027778
triangle  3.333333  0.055556
rectangle 2.500000  0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
      angles  degrees
circle     -1     358
triangle    2     178
rectangle   3     358
```

```
>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle     -1     358
triangle    2     178
rectangle   3     358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
      angles  degrees
circle     -1     359
triangle    2     179
rectangle   3     359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                        index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle    3
rectangle   4
```

```
>>> df * other
      angles  degrees
circle      0     NaN
triangle    9     NaN
rectangle  16     NaN
```

```
>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0     0.0
triangle    9     0.0
rectangle  16     0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                               'degrees': [360, 180, 360, 360, 540, 720]},
...                               index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                       ['circle', 'triangle', 'rectangle',
...                                        'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

```

            angles  degrees
A circle         0      360
  triangle       3      180
  rectangle       4      360
B square         4      360
  pentagon       5      540
  hexagon        6      720

>>> df.div(df_multindex, level=1, fill_value=0)
            angles  degrees
A circle         NaN      1.0
  triangle       1.0      1.0
  rectangle       1.0      1.0
B square         0.0      0.0
  pentagon       0.0      0.0
  hexagon        0.0      0.0

mode(self, axis: 'Axis' = 0, numeric_only: 'bool' = False, dropna: 'bool' = True)
-> 'DataFrame'
    Get the mode(s) of each element along the selected axis.

    The mode of a set of values is the value that appears most often.
    It can be multiple values.

    Parameters
    -----
    axis : {0 or 'index', 1 or 'columns'}, default 0
        The axis to iterate over while searching for the mode:

        * 0 or 'index' : get mode of each column
        * 1 or 'columns' : get mode of each row.

    numeric_only : bool, default False
        If True, only apply to numeric columns.
    dropna : bool, default True
        Don't consider counts of NaN/NaT.

    Returns
    -----
    DataFrame
        The modes of each column or row.

    See Also
    -----
    Series.mode : Return the highest frequency value in a Series.
    Series.value_counts : Return the counts of values in a Series.

    Examples
    -----
    >>> df = pd.DataFrame([('bird', 2, 2),
        ...                 ('mammal', 4, np.nan),
        ...                 ('arthropod', 8, 0),
        ...                 ('bird', 2, np.nan)],
        ...                 index=('falcon', 'horse', 'spider', 'ostrich'),
        ...                 columns=('species', 'legs', 'wings'))
    >>> df
           species  legs  wings
falcon      bird     2    2.0
horse      mammal     4    NaN
spider  arthropod     8    0.0
ostrich      bird     2    NaN

    By default, missing values are not considered, and the mode of wings
    are both 0 and 2. Because the resulting DataFrame has two rows,
    the second row of ``species`` and ``legs`` contains ``NaN``.

    >>> df.mode()
           species  legs  wings
0      bird     2.0    0.0
1      NaN     NaN    2.0

    Setting ``dropna=False`` ``NaN`` values are considered and they can be
    the mode (like for wings).

    >>> df.mode(dropna=False)
           species  legs  wings
0      bird     2    NaN

    Setting ``numeric_only=True``, only the mode of numeric columns is
    computed, and columns of other types are ignored.

    >>> df.mode(numeric_only=True)
           legs  wings
0      2.0    0.0
1      NaN    2.0

    To compute the mode over columns and not rows, use the axis parameter:

```



```

>>> df.mode(axis='columns', numeric_only=True)
      0      1
falcon  2.0  NaN
horse   4.0  NaN
spider  0.0  8.0
ostrich  2.0  NaN

mul(self, other, axis='columns', level=None, fill_value=None)
    Get Multiplication of dataframe and other, element-wise (binary operator
`mul`).

    Equivalent to ``dataframe * other``, but with support to substitute a
fill_value
    for missing data in one of the inputs. With reverse version, `rmul`.

    Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
    arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters
-----
other : scalar, sequence, Series, or DataFrame
    Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
    Whether to compare by the index (0 or 'index') or columns
    (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
    Broadcast across a level, matching Index values on the
    passed MultiIndex level.
fill_value : float or None, default None
    Fill existing missing (NaN) values, and any new element needed for
    successful DataFrame alignment, with this value before computation.
    If data in both corresponding DataFrame locations is missing
    the result will be missing.

Returns
-----
DataFrame
    Result of the arithmetic operation.

See Also
-----
DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes
-----
Mismatched indices will be unioned together.

Examples
-----
>>> df = pd.DataFrame({'angles': [0, 3, 4],
...                     'degrees': [360, 180, 360]},
...                     index=['circle', 'triangle', 'rectangle'])
>>> df
      angles  degrees
circle      0     360
triangle    3     180
rectangle   4     360

Add a scalar with operator version which return the same
results.

>>> df + 1
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361

>>> df.add(1)
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361

Divide by constant with reverse version.

>>> df.div(10)
      angles  degrees
circle    0.0     36.0
triangle  0.3     18.0

```

```
rectangle    0.4    36.0
```

```
>>> df.rdiv(10)
          angles  degrees
circle      inf  0.027778
triangle   3.333333  0.055556
rectangle  2.500000  0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
          angles  degrees
circle      -1    358
triangle     2    178
rectangle    3    358
```

```
>>> df.sub([1, 2], axis='columns')
          angles  degrees
circle      -1    358
triangle     2    178
rectangle    3    358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
          angles  degrees
circle      -1    359
triangle     2    179
rectangle    3    359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                       index=['circle', 'triangle', 'rectangle'])
>>> other
          angles
circle         0
triangle        3
rectangle       4
```

```
>>> df * other
          angles  degrees
circle         0     NaN
triangle        9     NaN
rectangle      16     NaN
```

```
>>> df.mul(other, fill_value=0)
          angles  degrees
circle         0     0.0
triangle        9     0.0
rectangle      16     0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                               'degrees': [360, 180, 360, 360, 540, 720]},
...                              index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                     ['circle', 'triangle', 'rectangle',
...                                      'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

```
          angles  degrees
A circle         0    360
  triangle        3    180
  rectangle        4    360
B square          4    360
  pentagon         5    540
  hexagon          6    720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
          angles  degrees
A circle      NaN     1.0
  triangle     1.0     1.0
  rectangle     1.0     1.0
B square       0.0     0.0
  pentagon     0.0     0.0
  hexagon      0.0     0.0
```

```
multiply = mul(self, other, axis='columns', level=None, fill_value=None)
```

```
ne(self, other, axis='columns', level=None)
```

Get Not equal to of dataframe and other, element-wise (binary operator `ne`).

Among flexible wrappers (`eq`, `ne`, `le`, `lt`, `ge`, `gt`) to comparison operators.

Equivalent to `==`, `!=`, `<=`, `<`, `>=`, `>` with support to choose axis (rows or columns) and level for comparison.

Parameters

`other` : scalar, sequence, Series, or DataFrame

Any single or multiple element data structure, or list-like object.

`axis` : {0 or 'index', 1 or 'columns'}, default 'columns'

Whether to compare by the index (0 or 'index') or columns (1 or 'columns').

`level` : int or label

Broadcast across a level, matching Index values on the passed MultiIndex level.

Returns

DataFrame of bool

Result of the comparison.

See Also

`DataFrame.eq` : Compare DataFrames for equality elementwise.

`DataFrame.ne` : Compare DataFrames for inequality elementwise.

`DataFrame.le` : Compare DataFrames for less than inequality or equality elementwise.

`DataFrame.lt` : Compare DataFrames for strictly less than inequality elementwise.

`DataFrame.ge` : Compare DataFrames for greater than inequality or equality elementwise.

`DataFrame.gt` : Compare DataFrames for strictly greater than inequality elementwise.

Notes

Mismatched indices will be unioned together.

`NaN` values are considered different (i.e. `NaN` != `NaN`).

Examples

```
>>> df = pd.DataFrame({'cost': [250, 150, 100],
...                    'revenue': [100, 250, 300]},
...                    index=['A', 'B', 'C'])
```

```
>>> df
   cost  revenue
A    250     100
B    150     250
C     100     300
```

Comparison with a scalar, using either the operator or method:

```
>>> df == 100
   cost  revenue
A  False     True
B  False     False
C   True     False
```

```
>>> df.eq(100)
   cost  revenue
A  False     True
B  False     False
C   True     False
```

When `other` is a `:class:`Series``, the columns of a DataFrame are aligned with the index of `other` and broadcast:

```
>>> df != pd.Series([100, 250], index=["cost", "revenue"])
   cost  revenue
A   True     True
B   True     False
C  False     True
```

Use the method to control the broadcast axis:

```
>>> df.ne(pd.Series([100, 300], index=["A", "D"]), axis='index')
   cost  revenue
A   True     False
B   True     True
C   True     True
D   True     True
```

When comparing to an arbitrary sequence, the number of columns must match the number elements in `other`:

```
>>> df == [250, 100]
   cost  revenue
A   True     True
B  False     False
C  False     False
```

Use the method to control the axis:

```
>>> df.eq([250, 250, 100], axis='index')
      cost  revenue
A   True   False
B   False   True
C    True   False
```

Compare to a DataFrame of different shape.

```
>>> other = pd.DataFrame({'revenue': [300, 250, 100, 150]},
...                       index=['A', 'B', 'C', 'D'])
>>> other
      revenue
A         300
B         250
C         100
D         150

>>> df.gt(other)
      cost  revenue
A   False   False
B   False   False
C   False   True
D   False   False
```

Compare to a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'cost': [250, 150, 100, 150, 300, 220],
...                              'revenue': [100, 250, 300, 200, 175, 225]},
...                              index=[['Q1', 'Q1', 'Q1', 'Q2', 'Q2', 'Q2'],
...                                      ['A', 'B', 'C', 'A', 'B', 'C']])
>>> df_multindex
      cost  revenue
Q1 A    250      100
   B    150      250
   C    100      300
Q2 A    150      200
   B    300      175
   C    220      225

>>> df.le(df_multindex, level=1)
      cost  revenue
Q1 A    True     True
   B    True     True
   C    True     True
Q2 A   False     True
   B    True     False
   C    True     False
```

`nlargest(self, n, columns, keep: 'str' = 'first') -> 'DataFrame'`
Return the first `n` rows ordered by `columns` in descending order.

Return the first `n` rows with the largest values in `columns`, in descending order. The columns that are not specified are returned as well, but not used for ordering.

This method is equivalent to
`df.sort_values(columns, ascending=False).head(n)`, but more performant.

Parameters

`n` : int

Number of rows to return.

`columns` : label or list of labels

Column label(s) to order by.

`keep` : {'first', 'last', 'all'}, default 'first'

Where there are duplicate values:

- 'first' : prioritize the first occurrence(s)
- 'last' : prioritize the last occurrence(s)
- 'all' : do not drop any duplicates, even it means selecting more than `n` items.

Returns

DataFrame

The first `n` rows ordered by the given columns in descending order.

See Also

`DataFrame.nsmallest` : Return the first `n` rows ordered by `columns` in ascending order.

`DataFrame.sort_values` : Sort DataFrame by the values.

`DataFrame.head` : Return the first `n` rows without re-ordering.

Notes

This function cannot be used with all column types. For example, when specifying columns with `object` or `category` dtypes, ``TypeError`` is raised.

Examples

```
>>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
...                                   434000, 434000, 337000, 11300,
...                                   11300, 11300],
...                    'GDP': [1937894, 2583560, 12011, 4520, 12128,
...                             17036, 182, 38, 311],
...                    'alpha-2': ["IT", "FR", "MT", "MV", "BN",
...                                 "IS", "NR", "TV", "AI"]},
...                    index=["Italy", "France", "Malta",
...                             "Maldives", "Brunei", "Iceland",
...                             "Nauru", "Tuvalu", "Anguilla"])
```

```
>>> df
      population    GDP alpha-2
Italy    59000000  1937894    IT
France    65000000  2583560    FR
Malta     434000   12011     MT
Maldives  434000   4520     MV
Brunei    434000   12128     BN
Iceland   337000   17036     IS
Nauru     11300    182      NR
Tuvalu    11300    38       TV
Anguilla  11300    311      AI
```

In the following example, we will use ``nlargest`` to select the three rows having the largest values in column "population".

```
>>> df.nlargest(3, 'population')
      population    GDP alpha-2
France    65000000  2583560    FR
Italy     59000000  1937894    IT
Malta      434000   12011     MT
```

When using ``keep='last'``, ties are resolved in reverse order:

```
>>> df.nlargest(3, 'population', keep='last')
      population    GDP alpha-2
France    65000000  2583560    FR
Italy     59000000  1937894    IT
Brunei    434000   12128     BN
```

When using ``keep='all'``, all duplicate items are maintained:

```
>>> df.nlargest(3, 'population', keep='all')
      population    GDP alpha-2
France    65000000  2583560    FR
Italy     59000000  1937894    IT
Malta      434000   12011     MT
Maldives  434000   4520     MV
Brunei    434000   12128     BN
```

To order by the largest values in column "population" and then "GDP", we can specify multiple columns like in the next example.

```
>>> df.nlargest(3, ['population', 'GDP'])
      population    GDP alpha-2
France    65000000  2583560    FR
Italy     59000000  1937894    IT
Brunei    434000   12128     BN
```

notna(self) -> 'DataFrame'

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA. Non-missing values get mapped to True. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``). NA values, such as None or :attr:`numpy.NaN`, get mapped to False values.

Returns

DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is not an NA value.

See Also

DataFrame.notnull : Alias of notna.

DataFrame.isna : Boolean inverse of notna.

DataFrame.dropna : Omit axes labels with missing values.

notna : Top-level notna.

Examples

Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                         born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                               pd.Timestamp('1940-04-25')],
...                         name=['Alfred', 'Batman', ''],
...                         toy=[None, 'Batmobile', 'Joker']))
>>> df
   age      born  name      toy
0  5.0      NaT  Alfred    None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25      Joker

>>> df.notna()
   age  born  name  toy
0  True False  True False
1  True  True  True  True
2 False  True  True  True
```

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
1    6.0
2    NaN
dtype: float64

>>> ser.notna()
0    True
1    True
2    False
dtype: bool
```

notnull(self) -> 'DataFrame'

Detect existing (non-missing) values.

Return a boolean same-sized object indicating if the values are not NA. Non-missing values get mapped to True. Characters such as empty strings ``''`` or :attr:`numpy.inf` are not considered NA values (unless you set ``pandas.options.mode.use_inf_as_na = True``). NA values, such as None or :attr:`numpy.NaN`, get mapped to False values.

Returns

DataFrame

Mask of bool values for each element in DataFrame that indicates whether an element is not an NA value.

See Also

DataFrame.notnull : Alias of notna.

DataFrame.isna : Boolean inverse of notna.

DataFrame.dropna : Omit axes labels with missing values.

notna : Top-level notna.

Examples

Show which entries in a DataFrame are not NA.

```
>>> df = pd.DataFrame(dict(age=[5, 6, np.NaN],
...                         born=[pd.NaT, pd.Timestamp('1939-05-27'),
...                               pd.Timestamp('1940-04-25')],
...                         name=['Alfred', 'Batman', ''],
...                         toy=[None, 'Batmobile', 'Joker']))
>>> df
   age      born  name      toy
0  5.0      NaT  Alfred    None
1  6.0  1939-05-27  Batman  Batmobile
2  NaN  1940-04-25      Joker

>>> df.notna()
   age  born  name  toy
0  True False  True False
1  True  True  True  True
2 False  True  True  True
```

Show which entries in a Series are not NA.

```
>>> ser = pd.Series([5, 6, np.NaN])
>>> ser
0    5.0
```

```
1    6.0
2    NaN
dtype: float64
```

```
>>> ser.notna()
0    True
1    True
2    False
dtype: bool
```

`nsmallest(self, n, columns, keep: 'str' = 'first') -> 'DataFrame'`
Return the first `n` rows ordered by `columns` in ascending order.

Return the first `n` rows with the smallest values in `columns`, in ascending order. The columns that are not specified are returned as well, but not used for ordering.

This method is equivalent to
``df.sort_values(columns, ascending=True).head(n)`` , but more performant.

Parameters

`n` : int
Number of items to retrieve.
`columns` : list or str
Column name or names to order by.
`keep` : {'first', 'last', 'all'}, default 'first'
Where there are duplicate values:

- ``first`` : take the first occurrence.
- ``last`` : take the last occurrence.
- ``all`` : do not drop any duplicates, even it means selecting more than `n` items.

Returns

DataFrame

See Also

`DataFrame.nlargest` : Return the first `n` rows ordered by `columns` in descending order.

`DataFrame.sort_values` : Sort DataFrame by the values.

`DataFrame.head` : Return the first `n` rows without re-ordering.

Examples

```
>>> df = pd.DataFrame({'population': [59000000, 65000000, 434000,
...                                   434000, 434000, 337000, 337000,
...                                   11300, 11300],
...                    'GDP': [1937894, 2583560, 12011, 4520, 12128,
...                             17036, 182, 38, 311],
...                    'alpha-2': ["IT", "FR", "MT", "MV", "BN",
...                                 "IS", "NR", "TV", "AI"]},
...                    index=["Italy", "France", "Malta",
...                             "Maldives", "Brunei", "Iceland",
...                             "Nauru", "Tuvalu", "Anguilla"])
>>> df
```

	population	GDP	alpha-2
Italy	59000000	1937894	IT
France	65000000	2583560	FR
Malta	434000	12011	MT
Maldives	434000	4520	MV
Brunei	434000	12128	BN
Iceland	337000	17036	IS
Nauru	337000	182	NR
Tuvalu	11300	38	TV
Anguilla	11300	311	AI

In the following example, we will use ``nsmallest`` to select the three rows having the smallest values in column "population".

```
>>> df.nsmallest(3, 'population')
   population  GDP alpha-2
Tuvalu      11300    38    TV
Anguilla     11300   311    AI
Iceland     337000 17036    IS
```

When using ``keep='last'``, ties are resolved in reverse order:

```
>>> df.nsmallest(3, 'population', keep='last')
   population  GDP alpha-2
Anguilla     11300   311    AI
Tuvalu       11300    38    TV
Nauru        337000  182    NR
```

When using ``keep='all'``, all duplicate items are maintained:

```
>>> df.nsmallest(3, 'population', keep='all')
      population  GDP alpha-2
Tuvalu      11300      38      TV
Anguilla    11300     311      AI
Iceland    337000  17036      IS
Nauru      337000     182      NR
```

To order by the smallest values in column "population" and then "GDP", we can specify multiple columns like in the next example.

```
>>> df.nsmallest(3, ['population', 'GDP'])
      population  GDP alpha-2
Tuvalu      11300      38      TV
Anguilla    11300     311      AI
Nauru      337000     182      NR
```

`nunique(self, axis: 'Axis' = 0, dropna: 'bool' = True) -> 'Series'`
Count number of distinct elements in specified axis.

Return Series with number of distinct elements. Can ignore NaN values.

Parameters

`axis` : {0 or 'index', 1 or 'columns'}, default 0
The axis to use. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
`dropna` : bool, default True
Don't include NaN in the counts.

Returns

Series

See Also

`Series.nunique`: Method `nunique` for Series.
`DataFrame.count`: Count non-NA cells for each column or row.

Examples

```
>>> df = pd.DataFrame({'A': [4, 5, 6], 'B': [4, 1, 1]})
>>> df.nunique()
A      3
B      2
dtype: int64

>>> df.nunique(axis=1)
0      1
1      2
2      2
dtype: int64
```

`pivot(self, index=None, columns=None, values=None) -> 'DataFrame'`
Return reshaped DataFrame organized by given index / column values.

Reshape data (produce a "pivot" table) based on column values. Uses unique values from specified `'index'` / `'columns'` to form axes of the resulting DataFrame. This function does not support data aggregation, multiple values will result in a MultiIndex in the columns. See the :ref:`User Guide <reshaping>` for more on reshaping.

Parameters

`index` : str or object or a list of str, optional
Column to use to make new frame's index. If None, uses existing index.

.. versionchanged:: 1.1.0
Also accept list of index names.

`columns` : str or object or a list of str
Column to use to make new frame's columns.

.. versionchanged:: 1.1.0
Also accept list of columns names.

`values` : str, object or a list of the previous, optional
Column(s) to use for populating new frame's values. If not specified, all remaining columns will be used and the result will have hierarchically indexed columns.

Returns

DataFrame

Returns reshaped DataFrame.

Raises

ValueError:

When there are any `index`, `columns` combinations with multiple values. `DataFrame.pivot_table` when you need to aggregate.

See Also

`DataFrame.pivot_table` : Generalization of pivot that can handle duplicate values for one index/column pair.

`DataFrame.unstack` : Pivot based on the index values instead of a column.

`wide_to_long` : Wide panel to long format. Less flexible but more user-friendly than melt.

Notes

For finer-tuned control, see hierarchical indexing documentation along with the related stack/unstack methods.

Examples

```
>>> df = pd.DataFrame({'foo': ['one', 'one', 'one', 'two', 'two',
...                             'two'],
...                    'bar': ['A', 'B', 'C', 'A', 'B', 'C'],
...                    'baz': [1, 2, 3, 4, 5, 6],
...                    'zoo': ['x', 'y', 'z', 'q', 'w', 't']})
```

```
>>> df
   foo  bar  baz  zoo
0  one   A    1    x
1  one   B    2    y
2  one   C    3    z
3  two   A    4    q
4  two   B    5    w
5  two   C    6    t
```

```
>>> df.pivot(index='foo', columns='bar', values='baz')
```

```
bar  A  B  C
foo
one  1  2  3
two  4  5  6
```

```
>>> df.pivot(index='foo', columns='bar')['baz']
```

```
bar  A  B  C
foo
one  1  2  3
two  4  5  6
```

```
>>> df.pivot(index='foo', columns='bar', values=['baz', 'zoo'])
```

```
      baz      zoo
bar  A  B  C  A  B  C
foo
one  1  2  3  x  y  z
two  4  5  6  q  w  t
```

You could also assign a list of column names or a list of index names.

```
>>> df = pd.DataFrame({
...     "lev1": [1, 1, 1, 2, 2, 2],
...     "lev2": [1, 1, 2, 1, 1, 2],
...     "lev3": [1, 2, 1, 2, 1, 2],
...     "lev4": [1, 2, 3, 4, 5, 6],
...     "values": [0, 1, 2, 3, 4, 5]})
```

```
>>> df
   lev1  lev2  lev3  lev4  values
0     1     1     1     1      0
1     1     1     2     2      1
2     1     2     1     3      2
3     2     1     2     4      3
4     2     1     1     5      4
5     2     2     2     6      5
```

```
>>> df.pivot(index="lev1", columns=["lev2", "lev3"], values="values")
```

```
lev2    1      2
lev3    1     2     1     2
lev1
1      0.0  1.0  2.0  NaN
2      4.0  3.0  NaN  5.0
```

```
>>> df.pivot(index=["lev1", "lev2"], columns=["lev3"], values="values")
```

```
lev3    1     2
lev1  lev2
1      1     0.0  1.0
      2     2.0  NaN
2      1     4.0  3.0
```

A `ValueError` is raised if there are any duplicates.

```
>>> df = pd.DataFrame({"foo": ['one', 'one', 'two', 'two'],
...                     "bar": ['A', 'A', 'B', 'C'],
...                     "baz": [1, 2, 3, 4]})
>>> df
   foo bar  baz
0  one  A    1
1  one  A    2
2  two  B    3
3  two  C    4
```

Notice that the first two rows are the same for our `'index'` and `'columns'` arguments.

```
>>> df.pivot(index='foo', columns='bar', values='baz')
Traceback (most recent call last):
...
ValueError: Index contains duplicate entries, cannot reshape
```

```
pivot_table(self, values=None, index=None, columns=None, aggfunc='mean',
fill_value=None, margins=False, dropna=True, margins_name='All', observed=False,
sort=True) -> 'DataFrame'
```

Create a spreadsheet-style pivot table as a `DataFrame`.

The levels in the pivot table will be stored in `MultiIndex` objects (hierarchical indexes) on the index and columns of the result `DataFrame`.

Parameters

```
-----
values : column to aggregate, optional
index : column, Grouper, array, or list of the previous
    If an array is passed, it must be the same length as the data. The
    list can contain any of the other types (except list).
    Keys to group by on the pivot table index. If an array is passed,
    it is being used as the same manner as column values.
columns : column, Grouper, array, or list of the previous
    If an array is passed, it must be the same length as the data. The
    list can contain any of the other types (except list).
    Keys to group by on the pivot table column. If an array is passed,
    it is being used as the same manner as column values.
aggfunc : function, list of functions, dict, default numpy.mean
    If list of functions passed, the resulting pivot table will have
    hierarchical columns whose top level are the function names
    (inferred from the function objects themselves)
    If dict is passed, the key is column to aggregate and value
    is function or list of functions.
fill_value : scalar, default None
    Value to replace missing values with (in the resulting pivot table,
    after aggregation).
margins : bool, default False
    Add all row / columns (e.g. for subtotal / grand totals).
dropna : bool, default True
    Do not include columns whose entries are all NaN.
margins_name : str, default 'All'
    Name of the row / column that will contain the totals
    when margins is True.
observed : bool, default False
    This only applies if any of the groupers are Categoricals.
    If True: only show observed values for categorical groupers.
    If False: show all values for categorical groupers.
```

.. versionchanged:: 0.25.0

```
sort : bool, default True
    Specifies if the result should be sorted.
```

.. versionadded:: 1.3.0

Returns

```
-----
DataFrame
```

An Excel style pivot table.

See Also

```
-----
DataFrame.pivot : Pivot without aggregation that can handle
    non-numeric data.
DataFrame.melt : Unpivot a DataFrame from wide to long format,
    optionally leaving identifiers set.
wide_to_long : Wide panel to long format. Less flexible but more
    user-friendly than melt.
```

Examples

```
-----
```



```

>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
...                     ('parrot', 'bird', 24.0),
...                     ('lion', 'mammal', 80.5),
...                     ('monkey', 'mammal', np.nan)],
...                     columns=('name', 'class', 'max_speed'))
>>> df
   name  class  max_speed
0  falcon   bird    389.0
1  parrot   bird     24.0
2   lion  mammal     80.5
3  monkey  mammal      NaN

>>> df.pop('class')
0    bird
1    bird
2  mammal
3  mammal
Name: class, dtype: object

>>> df
   name  max_speed
0  falcon    389.0
1  parrot    24.0
2   lion    80.5
3  monkey     NaN

```

`pow(self, other, axis='columns', level=None, fill_value=None)`
 Get Exponential power of dataframe and other, element-wise (binary operator ``pow``).

Equivalent to ``dataframe ** other``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``rpow``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

 Mismatched indices will be unioned together.

Examples

 >>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same results.

```
>>> df + 1
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

```
>>> df.add(1)
      angles  degrees
circle      1     361
triangle    4     181
rectangle   5     361
```

Divide by constant with reverse version.

```
>>> df.div(10)
      angles  degrees
circle    0.0    36.0
triangle  0.3    18.0
rectangle 0.4    36.0
```

```
>>> df.rdiv(10)
      angles  degrees
circle      inf  0.027778
triangle  3.333333  0.055556
rectangle 2.500000  0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
      angles  degrees
circle     -1    358
triangle    2    178
rectangle   3    358
```

```
>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle     -1    358
triangle    2    178
rectangle   3    358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
      angles  degrees
circle     -1    359
triangle    2    179
rectangle   3    359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                       index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle     3
rectangle    4
```

```
>>> df * other
      angles  degrees
circle      0      NaN
triangle     9      NaN
rectangle   16      NaN
```

```
>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0      0.0
triangle     9      0.0
rectangle   16      0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                              'degrees': [360, 180, 360, 360, 540, 720]},
...                              index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                      ['circle', 'triangle', 'rectangle',
...                                       'square', 'pentagon', 'hexagon']])
>>> df_multindex
      angles  degrees
A circle      0     360
  triangle    3     180
  rectangle    4     360
B square      4     360
  pentagon    5     540
  hexagon     6     720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
      angles  degrees
```

A circle	NaN	1.0
triangle	1.0	1.0
rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

```

    prod(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=0,
    **kwargs)
        Return the product of the values over the requested axis.

        Parameters
        -----
        axis : {index (0), columns (1)}
            Axis for the function to be applied on.
        skipna : bool, default True
            Exclude NA/null values when computing the result.
        level : int or level name, default None
            If the axis is a MultiIndex (hierarchical), count along a
            particular level, collapsing into a Series.
        numeric_only : bool, default None
            Include only float, int, boolean columns. If None, will attempt to use
            everything, then use only numeric data. Not implemented for Series.
        min_count : int, default 0
            The required number of valid values to perform the operation. If fewer
than
        ``min_count`` non-NA values are present the result will be NA.
    **kwargs
        Additional keyword arguments to be passed to the function.

    Returns
    -----
    Series or DataFrame (if level specified)

    See Also
    -----
    Series.sum : Return the sum.
    Series.min : Return the minimum.
    Series.max : Return the maximum.
    Series.idxmin : Return the index of the minimum.
    Series.idxmax : Return the index of the maximum.
    DataFrame.sum : Return the sum over the requested axis.
    DataFrame.min : Return the minimum over the requested axis.
    DataFrame.max : Return the maximum over the requested axis.
    DataFrame.idxmin : Return the index of the minimum over the requested axis.
    DataFrame.idxmax : Return the index of the maximum over the requested axis.

    Examples
    -----
    By default, the product of an empty or all-NA Series is ``1``

    >>> pd.Series([], dtype="float64").prod()
    1.0

    This can be controlled with the ``min_count`` parameter

    >>> pd.Series([], dtype="float64").prod(min_count=1)
    nan

    Thanks to the ``skipna`` parameter, ``min_count`` handles all-NA and
    empty series identically.

    >>> pd.Series([np.nan]).prod()
    1.0

    >>> pd.Series([np.nan]).prod(min_count=1)
    nan

    product = prod(self, axis=None, skipna=None, level=None, numeric_only=None,
min_count=0, **kwargs)

    quantile(self, q=0.5, axis: 'Axis' = 0, numeric_only: 'bool' = True,
interpolation: 'str' = 'linear')
        Return values at the given quantile over requested axis.

        Parameters
        -----
        q : float or array-like, default 0.5 (50% quantile)
            Value between 0 <= q <= 1, the quantile(s) to compute.
        axis : {0, 1, 'index', 'columns'}, default 0
            Equals 0 or 'index' for row-wise, 1 or 'columns' for column-wise.
        numeric_only : bool, default True
            If False, the quantile of datetime and timedelta data will be
            computed as well.
        interpolation : {'linear', 'lower', 'higher', 'midpoint', 'nearest'}
            This optional parameter specifies the interpolation method to use,
            when the desired quantile lies between two data points ``i`` and ``j``:

```

- * linear: $(j - i) * \text{fraction}$, where `fraction` is the fractional part of the index surrounded by `i` and `j`.
- * lower: `i`.
- * higher: `j`.
- * nearest: `i` or `j` whichever is nearest.
- * midpoint: $(i + j) / 2$.

Returns

Series or DataFrame

If `q` is an array, a DataFrame will be returned where the index is `q`, the columns are the columns of self, and the values are the quantiles.
If `q` is a float, a Series will be returned where the index is the columns of self and the values are the quantiles.

See Also

`core.window.Rolling.quantile`: Rolling quantile.
`numpy.percentile`: Numpy function to compute the percentile.

Examples

```
>>> df = pd.DataFrame(np.array([[1, 1], [2, 10], [3, 100], [4, 100]]),
...                    columns=['a', 'b'])
>>> df.quantile(.1)
a    1.3
b    3.7
Name: 0.1, dtype: float64
>>> df.quantile([.1, .5])
      a    b
0.1  1.3  3.7
0.5  2.5 55.0
```

Specifying `numeric_only=False` will also compute the quantile of datetime and timedelta data.

```
>>> df = pd.DataFrame({'A': [1, 2],
...                    'B': [pd.Timestamp('2010'),
...                          pd.Timestamp('2011')],
...                    'C': [pd.Timedelta('1 days'),
...                          pd.Timedelta('2 days')]})
>>> df.quantile(0.5, numeric_only=False)
A          1.5
B    2010-07-02 12:00:00
C          1 days 12:00:00
Name: 0.5, dtype: object
```

`query(self, expr: 'str', inplace: 'bool' = False, **kwargs)`
Query the columns of a DataFrame with a boolean expression.

Parameters

`expr` : str

The query string to evaluate.

You can refer to variables in the environment by prefixing them with an `@` character like `@a + b`.

You can refer to column names that are not valid Python variable names by surrounding them in backticks. Thus, column names containing spaces or punctuations (besides underscores) or starting with digits must be surrounded by backticks. (For example, a column named "Area (cm²)" would be referenced as ``Area (cm^2)``). Column names which are Python

keywords

(like "list", "for", "import", etc) cannot be used.

For example, if one of your columns is called `a`` and you want to sum it with `b``, your query should be ``a` + b``.

.. versionadded:: 0.25.0
Backtick quoting introduced.

.. versionadded:: 1.0.0
Expanding functionality of backtick quoting for more than only

spaces.

`inplace` : bool

Whether the query should modify the data in place or return a modified copy.

`**kwargs`

See the documentation for `:func:`eval`` for complete details on the keyword arguments accepted by `:meth:`DataFrame.query``.

```

Returns
-----
DataFrame or None
    DataFrame resulting from the provided query expression or
    None if ``inplace=True``.

See Also
-----
eval : Evaluate a string describing operations on
      DataFrame columns.
DataFrame.eval : Evaluate a string describing operations on
      DataFrame columns.

Notes
-----
The result of the evaluation of this expression is first passed to
:attr:`DataFrame.loc` and if that fails because of a
multidimensional key (e.g., a DataFrame) then the result will be passed
to :meth:`DataFrame.__getitem__`.

This method uses the top-level :func:`eval` function to
evaluate the passed query.

The :meth:`~pandas.DataFrame.query` method uses a slightly
modified Python syntax by default. For example, the ``&`` and ``|``
(bitwise) operators have the precedence of their boolean cousins,
:keyword:`and` and :keyword:`or`. This *is* syntactically valid Python,
however the semantics are different.

You can change the semantics of the expression by passing the keyword
argument ``parser='python'``. This enforces the same semantics as
evaluation in Python space. Likewise, you can pass ``engine='python'``
to evaluate an expression using Python itself as a backend. This is not
recommended as it is inefficient compared to using ``numexpr`` as the
engine.

The :attr:`DataFrame.index` and
:attr:`DataFrame.columns` attributes of the
:class:`~pandas.DataFrame` instance are placed in the query namespace
by default, which allows you to treat both the index and columns of the
frame as a column in the frame.
The identifier ``index`` is used for the frame index; you can also
use the name of the index to identify it in a query. Please note that
Python keywords may not be used as identifiers.

For further details and examples see the ``query`` documentation in
:ref:`indexing <indexing.query>`.

*Backtick quoted variables*

Backtick quoted variables are parsed as literal Python code and
are converted internally to a Python valid identifier.
This can lead to the following problems.

During parsing a number of disallowed characters inside the backtick
quoted string are replaced by strings that are allowed as a Python
identifier.
These characters include all operators in Python, the space character, the
question mark, the exclamation mark, the dollar sign, and the euro sign.
For other characters that fall outside the ASCII range (U+0001..U+007F)
and those that are not further specified in PEP 3131,
the query parser will raise an error.
This excludes whitespace different than the space character,
but also the hashtag (as it is used for comments) and the backtick
itself (backtick can also not be escaped).

In a special case, quotes that make a pair around a backtick can
confuse the parser.
For example, ``it's` > `that's`` will raise an error,
as it forms a quoted string (``'s > `that'``) with a backtick inside.

See also the Python documentation about lexical analysis
(https://docs.python.org/3/reference/lexical\_analysis.html)
in combination with the source code in
:mod:`pandas.core.computation.parsing`.

Examples
-----
>>> df = pd.DataFrame({'A': range(1, 6),
...                     'B': range(10, 0, -2),
...                     'C C': range(10, 5, -1)})
>>> df
   A  B  C C
0  1 10  10
1  2  8   9
2  3  6   8
3  4  4   7

```



```

4 5 2 6
>>> df.query('A > B')
   A  B  C  C
4 5 2  6

```

The previous expression is equivalent to

```

>>> df[df.A > df.B]
   A  B  C  C
4 5 2  6

```

For columns with spaces in their name, you can use backtick quoting.

```

>>> df.query('B == `C C`')
   A  B  C  C
0 1 10 10

```

The previous expression is equivalent to

```

>>> df[df.B == df['C C']]
   A  B  C  C
0 1 10 10

```

```

radd(self, other, axis='columns', level=None, fill_value=None)
    Get Addition of dataframe and other, element-wise (binary operator `radd`).

```

Equivalent to ``other + dataframe``, but with support to substitute a fill_value

for missing data in one of the inputs. With reverse version, `add`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.

level : int or label
Broadcast across a level, matching Index values on the passed MultiIndex level.

fill_value : float or None, default None
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

Returns

DataFrame
Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

```

>>> df = pd.DataFrame({'angles': [0, 3, 4],
...                     'degrees': [360, 180, 360]},
...                     index=['circle', 'triangle', 'rectangle'])
>>> df
      angles  degrees
circle      0     360
triangle    3     180
rectangle   4     360

```

Add a scalar with operator version which return the same results.

```

>>> df + 1
      angles  degrees
circle      1     361
triangle    4     181

```

```
rectangle      5      361
```

```
>>> df.add(1)
      angles  degrees
circle      1      361
triangle    4      181
rectangle    5      361
```

Divide by constant with reverse version.

```
>>> df.div(10)
      angles  degrees
circle      0.0     36.0
triangle    0.3     18.0
rectangle    0.4     36.0
```

```
>>> df.rdiv(10)
      angles  degrees
circle      inf     0.027778
triangle    3.333333  0.055556
rectangle    2.500000  0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
      angles  degrees
circle      -1     358
triangle     2     178
rectangle     3     358
```

```
>>> df.sub([1, 2], axis='columns')
      angles  degrees
circle      -1     358
triangle     2     178
rectangle     3     358
```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
...        axis='index')
      angles  degrees
circle      -1     359
triangle     2     179
rectangle     3     359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
...                       index=['circle', 'triangle', 'rectangle'])
>>> other
      angles
circle      0
triangle     3
rectangle     4
```

```
>>> df * other
      angles  degrees
circle      0      NaN
triangle     9      NaN
rectangle    16      NaN
```

```
>>> df.mul(other, fill_value=0)
      angles  degrees
circle      0      0.0
triangle     9      0.0
rectangle    16      0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
...                              'degrees': [360, 180, 360, 360, 540, 720]},
...                              index=[['A', 'A', 'A', 'B', 'B', 'B'],
...                                    ['circle', 'triangle', 'rectangle',
...                                     'square', 'pentagon', 'hexagon']])
>>> df_multindex
      angles  degrees
A circle      0      360
  triangle     3      180
  rectangle     4      360
B square      4      360
  pentagon     5      540
  hexagon      6      720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
      angles  degrees
A circle      NaN      1.0
  triangle     1.0      1.0
  rectangle     1.0      1.0
B square      0.0      0.0
```

```

    pentagon      0.0      0.0
    hexagon       0.0      0.0

    rdiv = rtruediv(self, other, axis='columns', level=None, fill_value=None)

    reindex(self, labels=None, index=None, columns=None, axis=None, method=None,
copy=True, level=None, fill_value=nan, limit=None, tolerance=None)
    Conform Series/DataFrame to new index with optional filling logic.

    Places NA/NaN in locations having no value in the previous index. A new
object
is produced unless the new index is equivalent to the current one and
``copy=False``.

    Parameters
    -----

    keywords for axes : array-like, optional
        New labels / index to conform to, should be specified using
        keywords. Preferably an Index object to avoid duplicating data.

    method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}
        Method to use for filling holes in reindexed DataFrame.
        Please note: this is only applicable to DataFrames/Series with a
        monotonically increasing/decreasing index.

        * None (default): don't fill gaps
        * pad / ffill: Propagate last valid observation forward to next
          valid.
        * backfill / bfill: Use next valid observation to fill gap.
        * nearest: Use nearest valid observations to fill gap.

    copy : bool, default True
        Return a new object, even if the passed indexes are the same.
    level : int or name
        Broadcast across a level, matching Index values on the
        passed MultiIndex level.
    fill_value : scalar, default np.NaN
        Value to use for missing values. Defaults to NaN, but can be any
        "compatible" value.
    limit : int, default None
        Maximum number of consecutive elements to forward or backward fill.
    tolerance : optional
        Maximum distance between original and new labels for inexact
        matches. The values of the index at the matching locations must
        satisfy the equation ``abs(index[indexer] - target) <= tolerance``.

        Tolerance may be a scalar value, which applies the same tolerance
        to all values, or list-like, which applies variable tolerance per
        element. List-like includes list, tuple, array, Series, and must be
        the same size as the index and its dtype must exactly match the
        index's type.

    Returns
    -----
    Series/DataFrame with changed index.

    See Also
    -----
    DataFrame.set_index : Set row labels.
    DataFrame.reset_index : Remove row labels or move them to new columns.
    DataFrame.reindex_like : Change to same indices as other DataFrame.

    Examples
    -----
    ``DataFrame.reindex`` supports two calling conventions

    * ``(index=index_labels, columns=column_labels, ...)``
    * ``(labels, axis={'index', 'columns'}, ...)``

    We *highly* recommend using keyword arguments to clarify your
    intent.

    Create a dataframe with some fictional data.

    >>> index = ['Firefox', 'Chrome', 'Safari', 'IE10', 'Konqueror']
    >>> df = pd.DataFrame({'http_status': [200, 200, 404, 404, 301],
    ...                   'response_time': [0.04, 0.02, 0.07, 0.08, 1.0]},
    ...                   index=index)
    >>> df

```

	http_status	response_time
Firefox	200	0.04
Chrome	200	0.02
Safari	404	0.07
IE10	404	0.08
Konqueror	301	1.00

Create a new index and reindex the dataframe. By default values in the new index that do not have corresponding records in the dataframe are assigned ``NaN``.

```
>>> new_index = ['Safari', 'Iceweasel', 'Comodo Dragon', 'IE10',
...              'Chrome']
>>> df.reindex(new_index)
           http_status  response_time
Safari              404             0.07
Iceweasel           NaN             NaN
Comodo Dragon       NaN             NaN
IE10                404             0.08
Chrome              200             0.02
```

We can fill in the missing values by passing a value to the keyword ``fill_value``. Because the index is not monotonically increasing or decreasing, we cannot use arguments to the keyword ``method`` to fill the ``NaN`` values.

```
>>> df.reindex(new_index, fill_value=0)
           http_status  response_time
Safari              404             0.07
Iceweasel           0             0.00
Comodo Dragon       0             0.00
IE10                404             0.08
Chrome              200             0.02

>>> df.reindex(new_index, fill_value='missing')
           http_status  response_time
Safari              404             0.07
Iceweasel          missing          missing
Comodo Dragon      missing          missing
IE10                404             0.08
Chrome              200             0.02
```

We can also reindex the columns.

```
>>> df.reindex(columns=['http_status', 'user_agent'])
           http_status  user_agent
Firefox            200          NaN
Chrome             200          NaN
Safari             404          NaN
IE10               404          NaN
Konqueror          301          NaN
```

Or we can use "axis-style" keyword arguments

```
>>> df.reindex(['http_status', 'user_agent'], axis="columns")
           http_status  user_agent
Firefox            200          NaN
Chrome             200          NaN
Safari             404          NaN
IE10               404          NaN
Konqueror          301          NaN
```

To further illustrate the filling functionality in ``reindex``, we will create a dataframe with a monotonically increasing index (for example, a sequence of dates).

```
>>> date_index = pd.date_range('1/1/2010', periods=6, freq='D')
>>> df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]},
...                    index=date_index)
>>> df2
           prices
2010-01-01    100.0
2010-01-02    101.0
2010-01-03     NaN
2010-01-04    100.0
2010-01-05     89.0
2010-01-06     88.0
```

Suppose we decide to expand the dataframe to cover a wider date range.

```
>>> date_index2 = pd.date_range('12/29/2009', periods=10, freq='D')
>>> df2.reindex(date_index2)
           prices
2009-12-29     NaN
2009-12-30     NaN
2009-12-31     NaN
2010-01-01    100.0
2010-01-02    101.0
2010-01-03     NaN
2010-01-04    100.0
2010-01-05     89.0
2010-01-06     88.0
```

2010-01-07 NaN

The index entries that did not have a value in the original data frame (for example, '2009-12-29') are by default filled with ``NaN``. If desired, we can fill in the missing values using one of several options.

For example, to back-propagate the last valid value to fill the ``NaN`` values, pass ``bfill`` as an argument to the ``method`` keyword.

```
>>> df2.reindex(date_index2, method='bfill')
```

```
prices
2009-12-29    100.0
2009-12-30    100.0
2009-12-31    100.0
2010-01-01    100.0
2010-01-02    101.0
2010-01-03     NaN
2010-01-04    100.0
2010-01-05     89.0
2010-01-06     88.0
2010-01-07     NaN
```

Please note that the ``NaN`` value present in the original dataframe (at index value 2010-01-03) will not be filled by any of the value propagation schemes. This is because filling while reindexing does not look at dataframe values, but only compares the original and desired indexes. If you do want to fill in the ``NaN`` values present in the original dataframe, use the ``fillna()`` method.

See the :ref:`user guide <basics.reindexing>` for more.

```
rename(self, mapper=None, index=None, columns=None, axis=None, copy=True,
inplace=False, level=None, errors='ignore')
```

Alter axes labels.

Function / dict values must be unique (1-to-1). Labels not contained in a dict / Series will be left as-is. Extra labels listed don't throw an error.

See the :ref:`user guide <basics.rename>` for more.

Parameters

mapper : dict-like or function

Dict-like or function transformations to apply to that axis' values. Use either ``mapper`` and ``axis`` to specify the axis to target with ``mapper``, or ``index`` and ``columns``.

index : dict-like or function

Alternative to specifying axis (``mapper, axis=0`` is equivalent to ``index=mapper``).

columns : dict-like or function

Alternative to specifying axis (``mapper, axis=1`` is equivalent to ``columns=mapper``).

axis : {0 or 'index', 1 or 'columns'}, default 0

Axis to target with ``mapper``. Can be either the axis name ('index', 'columns') or number (0, 1). The default is 'index'.

copy : bool, default True

Also copy underlying data.

inplace : bool, default False

Whether to return a new DataFrame. If True then value of copy is ignored.

level : int or level name, default None

In case of a MultiIndex, only rename labels in the specified level.

errors : {'ignore', 'raise'}, default 'ignore'

If 'raise', raise a `KeyError` when a dict-like `mapper`, `index`, or `columns` contains labels that are not present in the Index being transformed.

If 'ignore', existing keys will be renamed and extra keys will be ignored.

Returns

DataFrame or None

DataFrame with the renamed axis labels or None if ``inplace=True``.

Raises

KeyError

If any of the labels is not found in the selected axis and "errors='raise'".

See Also

DataFrame.rename_axis : Set the name of the axis.

Examples

```DataFrame.rename``` supports two calling conventions

```
* ``(index=index_mapper, columns=columns_mapper, ...)``
* ``(mapper, axis={'index', 'columns'}, ...)``
```

We *highly* recommend using keyword arguments to clarify your intent.

Rename columns using a mapping:

```
>>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
>>> df.rename(columns={"A": "a", "B": "c"})
 a c
0 1 4
1 2 5
2 3 6
```

Rename index using a mapping:

```
>>> df.rename(index={0: "x", 1: "y", 2: "z"})
 A B
x 1 4
y 2 5
z 3 6
```

Cast index labels to a different type:

```
>>> df.index
RangeIndex(start=0, stop=3, step=1)
>>> df.rename(index=str).index
Index(['0', '1', '2'], dtype='object')

>>> df.rename(columns={"A": "a", "B": "b", "C": "c"}, errors="raise")
Traceback (most recent call last):
KeyError: ['C'] not found in axis
```

Using axis-style parameters:

```
>>> df.rename(str.lower, axis='columns')
 a b
0 1 4
1 2 5
2 3 6

>>> df.rename({1: 2, 2: 4}, axis='index')
 A B
0 1 4
2 2 5
4 3 6
```

`reorder_levels(self, order: 'Sequence[Axis]', axis: 'Axis' = 0) -> 'DataFrame'`  
Rearrange index levels using input order. May not drop or duplicate levels.

## Parameters

-----

`order` : list of int or list of str  
List representing new level order. Reference level by number (position) or by key (label).  
`axis` : {0 or 'index', 1 or 'columns'}, default 0  
Where to reorder levels.

## Returns

-----

DataFrame

`replace(self, to_replace=None, value=None, inplace: 'bool' = False, limit=None, regex: 'bool' = False, method: 'str' = 'pad')`  
Replace values given in ``to_replace`` with ``value``.

Values of the DataFrame are replaced with other values dynamically.

This differs from updating with ```.loc``` or ```.iloc```, which require you to specify a location to update with some value.

## Parameters

-----

`to_replace` : str, regex, list, dict, Series, int, float, or None  
How to find the values that will be replaced.

\* numeric, str or regex:

- numeric: numeric values equal to ``to_replace`` will be replaced with ``value``
- str: string exactly matching ``to_replace`` will be replaced

```

 with `value`
 - regex: regexs matching `to_replace` will be replaced with
 `value`

* list of str, regex, or numeric:

 - First, if `to_replace` and `value` are both lists, they
 must be the same length.
 - Second, if ``regex=True`` then all of the strings in **both**
 lists will be interpreted as regexs otherwise they will match
 directly. This doesn't matter much for `value` since there
 are only a few possible substitution regexes you can use.
 - str, regex and numeric rules apply as above.

* dict:

 - Dicts can be used to specify different replacement values
 for different existing values. For example,
 ``{'a': 'b', 'y': 'z'}`` replaces the value 'a' with 'b' and
 'y' with 'z'. To use a dict in this way the `value`
 parameter should be `None`.
 - For a DataFrame a dict can specify that different values
 should be replaced in different columns. For example,
 ``{'a': 1, 'b': 'z'}`` looks for the value 1 in column 'a'
 and the value 'z' in column 'b' and replaces these values
 with whatever is specified in `value`. The `value` parameter
 should not be `None` in this case. You can treat this as a
 special case of passing two lists except that you are
 specifying the column to search in.
 - For a DataFrame nested dictionaries, e.g.,
 ``{'a': {'b': np.nan}}``, are read as follows: look in column
 'a' for the value 'b' and replace it with NaN. The `value`
 parameter should be `None` to use a nested dict in this
 way. You can nest regular expressions as well. Note that
 column names (the top-level dictionary keys in a nested
 dictionary) **cannot** be regular expressions.

* None:

 - This means that the `regex` argument must be a string,
 compiled regular expression, or list, dict, ndarray or
 Series of such elements. If `value` is also `None` then
 this **must** be a nested dictionary or Series.

See the examples section for examples of each of these.
value : scalar, dict, list, str, regex, default None
 Value to replace any values matching `to_replace` with.
 For a DataFrame a dict of values can be used to specify which
 value to use for each column (columns not in the dict will not be
 filled). Regular expressions, strings and lists or dicts of such
 objects are also allowed.

inplace : bool, default False
 If True, performs operation inplace and returns None.
limit : int, default None
 Maximum size gap to forward or backward fill.
regex : bool or same types as `to_replace`, default False
 Whether to interpret `to_replace` and/or `value` as regular
 expressions. If this is ``True`` then `to_replace` **must** be a
 string. Alternatively, this could be a regular expression or a
 list, dict, or array of regular expressions in which case
 `to_replace` must be `None`.
method : {'pad', 'ffill', 'bfill', 'None'}
 The method to use when for replacement, when `to_replace` is a
 scalar, list or tuple and `value` is `None`.

.. versionchanged:: 0.23.0
 Added to DataFrame.

Returns

DataFrame
 Object after replacement.

Raises

AssertionError
 * If `regex` is not a ``bool`` and `to_replace` is not
 ``None``.

TypeError
 * If `to_replace` is not a scalar, array-like, ``dict``, or ``None``
 * If `to_replace` is a ``dict`` and `value` is not a ``list``,
 ``dict``, ``ndarray``, or ``Series``
 * If `to_replace` is ``None`` and `regex` is not compilable
 into a regular expression or is a list, dict, ndarray, or
 Series.

```

\* When replacing multiple ``bool`` or ``datetime64`` objects and the arguments to ``to\_replace`` does not match the type of the value being replaced

ValueError

\* If a ``list`` or an ``ndarray`` is passed to ``to\_replace`` and ``value`` but they are not the same length.

See Also

-----  
DataFrame.fillna : Fill NA values.

DataFrame.where : Replace values based on boolean condition.

Series.str.replace : Simple string replacement.

Notes

-----  
\* Regex substitution is performed under the hood with ``re.sub``. The rules for substitution for ``re.sub`` are the same.  
\* Regular expressions will only substitute on strings, meaning you cannot provide, for example, a regular expression matching floating point numbers and expect the columns in your frame that have a numeric dtype to be matched. However, if those floating point numbers *are* strings, then you can do this.  
\* This method has *a lot* of options. You are encouraged to experiment and play with this method to gain intuition about how it works.  
\* When dict is used as the ``to\_replace`` value, it is like key(s) in the dict are the to\_replace part and value(s) in the dict are the value parameter.

Examples

-----  
\*\*Scalar ``to\_replace`` and ``value``\*\*

```
>>> s = pd.Series([0, 1, 2, 3, 4])
>>> s.replace(0, 5)
0 5
1 1
2 2
3 3
4 4
dtype: int64
```

```
>>> df = pd.DataFrame({'A': [0, 1, 2, 3, 4],
... 'B': [5, 6, 7, 8, 9],
... 'C': ['a', 'b', 'c', 'd', 'e']})
>>> df.replace(0, 5)
 A B C
0 5 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e
```

\*\*List-like ``to\_replace``\*\*

```
>>> df.replace([0, 1, 2, 3], 4)
 A B C
0 4 5 a
1 4 6 b
2 4 7 c
3 4 8 d
4 4 9 e
```

```
>>> df.replace([0, 1, 2, 3], [4, 3, 2, 1])
 A B C
0 4 5 a
1 3 6 b
2 2 7 c
3 1 8 d
4 4 9 e
```

```
>>> s.replace([1, 2], method='bfill')
0 0
1 3
2 3
3 3
4 4
dtype: int64
```

\*\*dict-like ``to\_replace``\*\*

```
>>> df.replace({0: 10, 1: 100})
 A B C
0 10 5 a
1 100 6 b
2 2 7 c
```



```

3 3 8 d
4 4 9 e

```

```
>>> df.replace({'A': 0, 'B': 5}, 100)
```

```

 A B C
0 100 100 a
1 1 6 b
2 2 7 c
3 3 8 d
4 4 9 e

```

```
>>> df.replace({'A': {0: 100, 4: 400}})
```

```

 A B C
0 100 5 a
1 1 6 b
2 2 7 c
3 3 8 d
4 400 9 e

```

**\*\*Regular expression `to\_replace`\*\***

```
>>> df = pd.DataFrame({'A': ['bat', 'foo', 'bait'],
... 'B': ['abc', 'bar', 'xyz']})
>>> df.replace(to_replace=r'^ba.$', value='new', regex=True)
```

```

 A B
0 new abc
1 foo new
2 bait xyz

```

```
>>> df.replace({'A': r'^ba.$'}, {'A': 'new'}, regex=True)
```

```

 A B
0 new abc
1 foo bar
2 bait xyz

```

```
>>> df.replace(regex=r'^ba.$', value='new')
```

```

 A B
0 new abc
1 foo new
2 bait xyz

```

```
>>> df.replace(regex={'r'^ba.$': 'new', 'foo': 'xyz'})
```

```

 A B
0 new abc
1 xyz new
2 bait xyz

```

```
>>> df.replace(regex=[r'^ba.$', 'foo'], value='new')
```

```

 A B
0 new abc
1 new new
2 bait xyz

```

Compare the behavior of ``s.replace({'a': None})`` and  
 ``s.replace('a', None)`` to understand the peculiarities  
 of the `to\_replace` parameter:

```
>>> s = pd.Series([10, 'a', 'a', 'b', 'a'])
```

When one uses a dict as the `to\_replace` value, it is like the  
 value(s) in the dict are equal to the `value` parameter.  
 ``s.replace({'a': None})`` is equivalent to  
 ``s.replace(to\_replace={'a': None}, value=None, method=None)``:

```
>>> s.replace({'a': None})
```

```

0 10
1 None
2 None
3 b
4 None
dtype: object

```

When ``value=None`` and `to\_replace` is a scalar, list or  
 tuple, `replace` uses the method parameter (default 'pad') to do the  
 replacement. So this is why the 'a' values are being replaced by 10  
 in rows 1 and 2 and 'b' in row 4 in this case.  
 The command ``s.replace('a', None)`` is actually equivalent to  
 ``s.replace(to\_replace='a', value=None, method='pad')``:

```
>>> s.replace('a', None)
```

```

0 10
1 10
2 10
3 b
4 b
dtype: object

```

```

| resample(self, rule, axis=0, closed: 'str | None' = None, label: 'str | None' =
None, convention: 'str' = 'start', kind: 'str | None' = None, loffset=None, base:
'int | None' = None, on=None, level=None, origin: 'str | TimestampConvertibleTypes' =
'start_day', offset: 'TimedeltaConvertibleTypes | None' = None) -> 'Resampler'
| Resample time-series data.

|
| Convenience method for frequency conversion and resampling of time series.
| The object must have a datetime-like index ('DatetimeIndex', 'PeriodIndex',
| or 'TimedeltaIndex'), or the caller must pass the label of a datetime-like
| series/index to the ``on``/``level`` keyword parameter.
|
| Parameters
| -----
| rule : DateOffset, Timedelta or str
| The offset string or object representing target conversion.
| axis : {0 or 'index', 1 or 'columns'}, default 0
| Which axis to use for up- or down-sampling. For 'Series' this
| will default to 0, i.e. along the rows. Must be
| 'DatetimeIndex', 'TimedeltaIndex' or 'PeriodIndex'.
| closed : {'right', 'left'}, default None
| Which side of bin interval is closed. The default is 'left'
| for all frequency offsets except for 'M', 'A', 'Q', 'BM',
| 'BA', 'BQ', and 'W' which all have a default of 'right'.
| label : {'right', 'left'}, default None
| Which bin edge label to label bucket with. The default is 'left'
| for all frequency offsets except for 'M', 'A', 'Q', 'BM',
| 'BA', 'BQ', and 'W' which all have a default of 'right'.
| convention : {'start', 'end', 's', 'e'}, default 'start'
| For 'PeriodIndex' only, controls whether to use the start or
| end of 'rule'.
| kind : {'timestamp', 'period'}, optional, default None
| Pass 'timestamp' to convert the resulting index to a
| 'DatetimeIndex' or 'period' to convert it to a 'PeriodIndex'.
| By default the input representation is retained.
| loffset : timedelta, default None
| Adjust the resampled time labels.
|
| .. deprecated:: 1.1.0
| You should add the loffset to the 'df.index' after the resample.
| See below.
|
| base : int, default 0
| For frequencies that evenly subdivide 1 day, the "origin" of the
| aggregated intervals. For example, for '5min' frequency, base could
| range from 0 through 4. Defaults to 0.
|
| .. deprecated:: 1.1.0
| The new arguments that you should use are 'offset' or 'origin'.
|
| on : str, optional
| For a DataFrame, column to use instead of index for resampling.
| Column must be datetime-like.
| level : str or int, optional
| For a MultiIndex, level (name or number) to use for
| resampling. 'level' must be datetime-like.
| origin : {'epoch', 'start', 'start_day', 'end', 'end_day'}, Timestamp
| or str, default 'start_day'
| The timestamp on which to adjust the grouping. The timezone of origin
| must match the timezone of the index.
| If a timestamp is not used, these values are also supported:
|
| - 'epoch': 'origin' is 1970-01-01
| - 'start': 'origin' is the first value of the timeseries
| - 'start_day': 'origin' is the first day at midnight of the timeseries
|
| .. versionadded:: 1.1.0
|
| - 'end': 'origin' is the last value of the timeseries
| - 'end_day': 'origin' is the ceiling midnight of the last day
|
| .. versionadded:: 1.3.0
|
| offset : Timedelta or str, default is None
| An offset timedelta added to the origin.
|
| .. versionadded:: 1.1.0
|
| Returns
| -----
| pandas.core.Resampler
| :class:`pandas.core.Resampler` object.
|
| See Also
| -----
| Series.resample : Resample a Series.
| DataFrame.resample : Resample a DataFrame.
| groupby : Group DataFrame by mapping, function, label, or list of labels.

```

asfreq : Reindex a DataFrame with the given frequency without grouping.

#### Notes

-----

See the `user guide

<[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/timeseries.html#resampling](https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html#resampling)>`\_\_  
for more.

To learn more about the offset strings, please see `this link

<[https://pandas.pydata.org/pandas-docs/stable/user\\_guide/timeseries.html#dateoffset-objects](https://pandas.pydata.org/pandas-docs/stable/user_guide/timeseries.html#dateoffset-objects)>`\_\_.

#### Examples

-----

Start by creating a series with 9 one minute timestamps.

```
>>> index = pd.date_range('1/1/2000', periods=9, freq='T')
>>> series = pd.Series(range(9), index=index)
>>> series
2000-01-01 00:00:00 0
2000-01-01 00:01:00 1
2000-01-01 00:02:00 2
2000-01-01 00:03:00 3
2000-01-01 00:04:00 4
2000-01-01 00:05:00 5
2000-01-01 00:06:00 6
2000-01-01 00:07:00 7
2000-01-01 00:08:00 8
Freq: T, dtype: int64
```

Downsample the series into 3 minute bins and sum the values of the timestamps falling into a bin.

```
>>> series.resample('3T').sum()
2000-01-01 00:00:00 3
2000-01-01 00:03:00 12
2000-01-01 00:06:00 21
Freq: 3T, dtype: int64
```

Downsample the series into 3 minute bins as above, but label each bin using the right edge instead of the left. Please note that the value in the bucket used as the label is not included in the bucket, which it labels. For example, in the original series the bucket ``2000-01-01 00:03:00`` contains the value 3, but the summed value in the resampled bucket with the label ``2000-01-01 00:03:00`` does not include 3 (if it did, the summed value would be 6, not 3). To include this value close the right side of the bin interval as illustrated in the example below this one.

```
>>> series.resample('3T', label='right').sum()
2000-01-01 00:03:00 3
2000-01-01 00:06:00 12
2000-01-01 00:09:00 21
Freq: 3T, dtype: int64
```

Downsample the series into 3 minute bins as above, but close the right side of the bin interval.

```
>>> series.resample('3T', label='right', closed='right').sum()
2000-01-01 00:00:00 0
2000-01-01 00:03:00 6
2000-01-01 00:06:00 15
2000-01-01 00:09:00 15
Freq: 3T, dtype: int64
```

Upsample the series into 30 second bins.

```
>>> series.resample('30S').asfreq()[0:5] # Select first 5 rows
2000-01-01 00:00:00 0.0
2000-01-01 00:00:30 NaN
2000-01-01 00:01:00 1.0
2000-01-01 00:01:30 NaN
2000-01-01 00:02:00 2.0
Freq: 30S, dtype: float64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``pad`` method.

```
>>> series.resample('30S').pad()[0:5]
2000-01-01 00:00:00 0
2000-01-01 00:00:30 0
2000-01-01 00:01:00 1
2000-01-01 00:01:30 1
2000-01-01 00:02:00 2
Freq: 30S, dtype: int64
```

Upsample the series into 30 second bins and fill the ``NaN`` values using the ``bfill`` method.

```
>>> series.resample('30S').bfill()[0:5]
2000-01-01 00:00:00 0
2000-01-01 00:00:30 1
2000-01-01 00:01:00 1
2000-01-01 00:01:30 2
2000-01-01 00:02:00 2
Freq: 30S, dtype: int64
```

Pass a custom function via ``apply``

```
>>> def custom_resampler(arraylike):
... return np.sum(arraylike) + 5
...
>>> series.resample('3T').apply(custom_resampler)
2000-01-01 00:00:00 8
2000-01-01 00:03:00 17
2000-01-01 00:06:00 26
Freq: 3T, dtype: int64
```

For a Series with a PeriodIndex, the keyword ``convention`` can be used to control whether to use the start or end of ``rule``.

Resample a year by quarter using 'start' ``convention``. Values are assigned to the first quarter of the period.

```
>>> s = pd.Series([1, 2], index=pd.period_range('2012-01-01',
... freq='A',
... periods=2))
>>> s
2012 1
2013 2
Freq: A-DEC, dtype: int64
>>> s.resample('Q', convention='start').asfreq()
2012Q1 1.0
2012Q2 NaN
2012Q3 NaN
2012Q4 NaN
2013Q1 2.0
2013Q2 NaN
2013Q3 NaN
2013Q4 NaN
Freq: Q-DEC, dtype: float64
```

Resample quarters by month using 'end' ``convention``. Values are assigned to the last month of the period.

```
>>> q = pd.Series([1, 2, 3, 4], index=pd.period_range('2018-01-01',
... freq='Q',
... periods=4))
>>> q
2018Q1 1
2018Q2 2
2018Q3 3
2018Q4 4
Freq: Q-DEC, dtype: int64
>>> q.resample('M', convention='end').asfreq()
2018-03 1.0
2018-04 NaN
2018-05 NaN
2018-06 2.0
2018-07 NaN
2018-08 NaN
2018-09 3.0
2018-10 NaN
2018-11 NaN
2018-12 4.0
Freq: M, dtype: float64
```

For DataFrame objects, the keyword ``on`` can be used to specify the column instead of the index for resampling.

```
>>> d = {'price': [10, 11, 9, 13, 14, 18, 17, 19],
... 'volume': [50, 60, 40, 100, 50, 100, 40, 50]}
>>> df = pd.DataFrame(d)
>>> df['week_starting'] = pd.date_range('01/01/2018',
... periods=8,
... freq='W')
>>> df
 price volume week_starting
0 10 50 2018-01-07
1 11 60 2018-01-14
2 9 40 2018-01-21
3 13 100 2018-01-28
4 14 50 2018-02-04
```

```

5 18 100 2018-02-11
6 17 40 2018-02-18
7 19 50 2018-02-25
>>> df.resample('M', on='week_starting').mean()
 price volume
week_starting
2018-01-31 10.75 62.5
2018-02-28 17.00 60.0

```

For a DataFrame with MultiIndex, the keyword `level` can be used to specify on which level the resampling needs to take place.

```

>>> days = pd.date_range('1/1/2000', periods=4, freq='D')
>>> d2 = {'price': [10, 11, 9, 13, 14, 18, 17, 19],
... 'volume': [50, 60, 40, 100, 50, 100, 40, 50]}
>>> df2 = pd.DataFrame(
... d2,
... index=pd.MultiIndex.from_product(
... [days, ['morning', 'afternoon']]
...)
...)
>>> df2
 price volume
2000-01-01 morning 10 50
 afternoon 11 60
2000-01-02 morning 9 40
 afternoon 13 100
2000-01-03 morning 14 50
 afternoon 18 100
2000-01-04 morning 17 40
 afternoon 19 50
>>> df2.resample('D', level=0).sum()
 price volume
2000-01-01 21 110
2000-01-02 22 140
2000-01-03 32 150
2000-01-04 36 90

```

If you want to adjust the start of the bins based on a fixed timestamp:

```

>>> start, end = '2000-10-01 23:30:00', '2000-10-02 00:30:00'
>>> rng = pd.date_range(start, end, freq='7min')
>>> ts = pd.Series(np.arange(len(rng)) * 3, index=rng)
>>> ts
2000-10-01 23:30:00 0
2000-10-01 23:37:00 3
2000-10-01 23:44:00 6
2000-10-01 23:51:00 9
2000-10-01 23:58:00 12
2000-10-02 00:05:00 15
2000-10-02 00:12:00 18
2000-10-02 00:19:00 21
2000-10-02 00:26:00 24
Freq: 7T, dtype: int64

>>> ts.resample('17min').sum()
2000-10-01 23:14:00 0
2000-10-01 23:31:00 9
2000-10-01 23:48:00 21
2000-10-02 00:05:00 54
2000-10-02 00:22:00 24
Freq: 17T, dtype: int64

>>> ts.resample('17min', origin='epoch').sum()
2000-10-01 23:18:00 0
2000-10-01 23:35:00 18
2000-10-01 23:52:00 27
2000-10-02 00:09:00 39
2000-10-02 00:26:00 24
Freq: 17T, dtype: int64

>>> ts.resample('17min', origin='2000-01-01').sum()
2000-10-01 23:24:00 3
2000-10-01 23:41:00 15
2000-10-01 23:58:00 45
2000-10-02 00:15:00 45
Freq: 17T, dtype: int64

```

If you want to adjust the start of the bins with an `offset` Timedelta, the following lines are equivalent:

```

>>> ts.resample('17min', origin='start').sum()
2000-10-01 23:30:00 9
2000-10-01 23:47:00 21
2000-10-02 00:04:00 54
2000-10-02 00:21:00 24

```

```
Freq: 17T, dtype: int64
```

```
>>> ts.resample('17min', offset='23h30min').sum()
2000-10-01 23:30:00 9
2000-10-01 23:47:00 21
2000-10-02 00:04:00 54
2000-10-02 00:21:00 24
Freq: 17T, dtype: int64
```

If you want to take the largest Timestamp as the end of the bins:

```
>>> ts.resample('17min', origin='end').sum()
2000-10-01 23:35:00 0
2000-10-01 23:52:00 18
2000-10-02 00:09:00 27
2000-10-02 00:26:00 63
Freq: 17T, dtype: int64
```

In contrast with the `'start_day'`, you can use `'end_day'` to take the ceiling midnight of the largest Timestamp as the end of the bins and drop the bins not containing data:

```
>>> ts.resample('17min', origin='end_day').sum()
2000-10-01 23:38:00 3
2000-10-01 23:55:00 15
2000-10-02 00:12:00 45
2000-10-02 00:29:00 45
Freq: 17T, dtype: int64
```

To replace the use of the deprecated `'base'` argument, you can now use `'offset'`, in this example it is equivalent to have `'base=2'`:

```
>>> ts.resample('17min', offset='2min').sum()
2000-10-01 23:16:00 0
2000-10-01 23:33:00 9
2000-10-01 23:50:00 36
2000-10-02 00:07:00 39
2000-10-02 00:24:00 24
Freq: 17T, dtype: int64
```

To replace the use of the deprecated `'loffset'` argument:

```
>>> from pandas.tseries.frequencies import to_offset
>>> loffset = '19min'
>>> ts_out = ts.resample('17min').sum()
>>> ts_out.index = ts_out.index + to_offset(loffset)
>>> ts_out
2000-10-01 23:33:00 0
2000-10-01 23:50:00 9
2000-10-02 00:07:00 21
2000-10-02 00:24:00 54
2000-10-02 00:41:00 24
Freq: 17T, dtype: int64
```

```
reset_index(self, level: 'Hashable | Sequence[Hashable] | None' = None, drop:
'bool' = False, inplace: 'bool' = False, col_level: 'Hashable' = 0, col_fill:
'Hashable' = '') -> 'DataFrame | None'
```

Reset the index, or a level of it.

Reset the index of the DataFrame, and use the default one instead.  
If the DataFrame has a MultiIndex, this method can remove one or more levels.

#### Parameters

-----

**level** : int, str, tuple, or list, default None

Only remove the given levels from the index. Removes all levels by default.

**drop** : bool, default False

Do not try to insert index into dataframe columns. This resets the index to the default integer index.

**inplace** : bool, default False

Modify the DataFrame in place (do not create a new object).

**col\_level** : int or str, default 0

If the columns have multiple levels, determines which level the labels are inserted into. By default it is inserted into the first level.

**col\_fill** : object, default ''

If the columns have multiple levels, determines how the other levels are named. If None then the index name is repeated.

#### Returns

-----  
DataFrame or None

DataFrame with the new index or None if `'inplace=True'`.

## See Also

-----  
DataFrame.set\_index : Opposite of reset\_index.  
DataFrame.reindex : Change to new indices or expand indices.  
DataFrame.reindex\_like : Change to same indices as other DataFrame.

## Examples

-----  
>>> df = pd.DataFrame([('bird', 389.0),  
... ('bird', 24.0),  
... ('mammal', 80.5),  
... ('mammal', np.nan)],  
... index=['falcon', 'parrot', 'lion', 'monkey'],  
... columns=('class', 'max\_speed'))  
>>> df

	class	max_speed
falcon	bird	389.0
parrot	bird	24.0
lion	mammal	80.5
monkey	mammal	NaN

When we reset the index, the old index is added as a column, and a new sequential index is used:

```
>>> df.reset_index()
 index class max_speed
0 falcon bird 389.0
1 parrot bird 24.0
2 lion mammal 80.5
3 monkey mammal NaN
```

We can use the `drop` parameter to avoid the old index being added as a column:

```
>>> df.reset_index(drop=True)
 class max_speed
0 bird 389.0
1 bird 24.0
2 mammal 80.5
3 mammal NaN
```

You can also use `reset\_index` with `MultiIndex`.

```
>>> index = pd.MultiIndex.from_tuples([('bird', 'falcon'),
... ('bird', 'parrot'),
... ('mammal', 'lion'),
... ('mammal', 'monkey')],
... names=['class', 'name'])
>>> columns = pd.MultiIndex.from_tuples([('speed', 'max'),
... ('species', 'type')])
>>> df = pd.DataFrame([(389.0, 'fly'),
... (24.0, 'fly'),
... (80.5, 'run'),
... (np.nan, 'jump')],
... index=index,
... columns=columns)
>>> df
```

		speed	species
		max	type
class	name		
bird	falcon	389.0	fly
	parrot	24.0	fly
mammal	lion	80.5	run
	monkey	NaN	jump

If the index has multiple levels, we can reset a subset of them:

```
>>> df.reset_index(level='class')
 class speed species
 max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

If we are not dropping the index, by default, it is placed in the top level. We can place it in another level:

```
>>> df.reset_index(level='class', col_level=1)
 class speed species
 max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

When the index is inserted under another level, we can specify under which one with the parameter `'col_fill'`:

```
>>> df.reset_index(level='class', col_level=1, col_fill='species')
 species speed species
 class max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

If we specify a nonexistent level for `'col_fill'`, it is created:

```
>>> df.reset_index(level='class', col_level=1, col_fill='genus')
 genus speed species
 class max type
name
falcon bird 389.0 fly
parrot bird 24.0 fly
lion mammal 80.5 run
monkey mammal NaN jump
```

`rfloordiv(self, other, axis='columns', level=None, fill_value=None)`  
Get Integer division of dataframe and other, element-wise (binary operator `'rfloordiv'`).

Equivalent to `'other // dataframe'`, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, `'floordiv'`.

Among flexible wrappers (`'add'`, `'sub'`, `'mul'`, `'div'`, `'mod'`, `'pow'`) to arithmetic operators: `'+'`, `'-'`, `'*'`, `'/'`, `'//'`, `'%'`, `'**'`.

#### Parameters

`other` : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.  
`axis` : {0 or 'index', 1 or 'columns'}  
Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.  
`level` : int or label  
Broadcast across a level, matching Index values on the passed MultiIndex level.  
`fill_value` : float or None, default None  
Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

DataFrame  
Result of the arithmetic operation.

#### See Also

`DataFrame.add` : Add DataFrames.  
`DataFrame.sub` : Subtract DataFrames.  
`DataFrame.mul` : Multiply DataFrames.  
`DataFrame.div` : Divide DataFrames (float division).  
`DataFrame.truediv` : Divide DataFrames (float division).  
`DataFrame.floordiv` : Divide DataFrames (integer division).  
`DataFrame.mod` : Calculate modulo (remainder after division).  
`DataFrame.pow` : Calculate exponential power.

#### Notes

Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360
```

Add a scalar with operator version which return the same results.

```
>>> df + 1
 angles degrees
```



circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
```

	angles	degrees
circle	-1	359
triangle	2	179
rectangle	3	359

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

		angles	degrees
A	circle	0	360
	triangle	3	180
	rectangle	4	360
B	square	4	360
	pentagon	5	540
	hexagon	6	720

```
>>> df.div(df_multindex, level=1, fill_value=0)
```

		angles	degrees
A	circle	NaN	1.0
	triangle	1.0	1.0

```

 rectangle 1.0 1.0
 B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

rmod(self, other, axis='columns', level=None, fill_value=None)
 Get Modulo of dataframe and other, element-wise (binary operator `rmod`).

 Equivalent to ``other % dataframe``, but with support to substitute a
fill_value
 for missing data in one of the inputs. With reverse version, `mod`.

 Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
 arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

>>> df.add(1)
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

Divide by constant with reverse version.

>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0

>>> df.rdiv(10)
 angles degrees

```

```
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778
```

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0
```

```
rmul(self, other, axis='columns', level=None, fill_value=None)
 Get Multiplication of dataframe and other, element-wise (binary operator
 `rmul`).
```

Equivalent to ``other \* dataframe``, but with support to substitute a  
fill\_value  
for missing data in one of the inputs. With reverse version, `mul`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to  
arithmetic operators: `+`, `-`, `\*`, `/`, `//`, `%`, `\*\*`.

Parameters

other : scalar, sequence, Series, or DataFrame  
Any single or multiple element data structure, or list-like object.

axis : {0 or 'index', 1 or 'columns'}  
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.  
 level : int or label  
 Broadcast across a level, matching Index values on the passed MultiIndex level.  
 fill\_value : float or None, default None  
 Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

#### Returns

#### DataFrame

Result of the arithmetic operation.

#### See Also

DataFrame.add : Add DataFrames.  
 DataFrame.sub : Subtract DataFrames.  
 DataFrame.mul : Multiply DataFrames.  
 DataFrame.div : Divide DataFrames (float division).  
 DataFrame.truediv : Divide DataFrames (float division).  
 DataFrame.floordiv : Divide DataFrames (integer division).  
 DataFrame.mod : Calculate modulo (remainder after division).  
 DataFrame.pow : Calculate exponential power.

#### Notes

Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

	angles	degrees
circle	0	360
triangle	3	180
rectangle	4	360

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
```

```
... axis='index')
... angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4
```

```
>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN
```

```
>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720
```

```
>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0
```

```
round(self, decimals: 'int | dict[IndexLabel, int] | Series' = 0, *args,
**kwargs) -> 'DataFrame'
```

Round a DataFrame to a variable number of decimal places.

Parameters

-----

**decimals** : int, dict, Series

Number of decimal places to round each column to. If an int is given, round each column to the same number of places. Otherwise dict and Series round to variable numbers of places. Column names should be in the keys if `decimals` is a dict-like, or in the index if `decimals` is a Series. Any columns not included in `decimals` will be left as is. Elements of `decimals` which are not columns of the input will be ignored.

**\*args**

Additional keywords have no effect but might be accepted for compatibility with numpy.

**\*\*kwargs**

Additional keywords have no effect but might be accepted for compatibility with numpy.

Returns

-----

**DataFrame**

A DataFrame with the affected columns rounded to the specified number of decimal places.

See Also

-----

**numpy.around** : Round a numpy array to the given number of decimals.  
**Series.round** : Round a Series to the given number of decimals.

Examples

```

>>> df = pd.DataFrame([(0.21, .32), (.01, .67), (.66, .03), (.21, .18)],
... columns=['dogs', 'cats'])
>>> df
 dogs cats
0 0.21 0.32
1 0.01 0.67
2 0.66 0.03
3 0.21 0.18

```

By providing an integer each column is rounded to the same number of decimal places

```

>>> df.round(1)
 dogs cats
0 0.2 0.3
1 0.0 0.7
2 0.7 0.0
3 0.2 0.2

```

With a dict, the number of places for specific columns can be specified with the column names as key and the number of decimal places as value

```

>>> df.round({'dogs': 1, 'cats': 0})
 dogs cats
0 0.2 0.0
1 0.0 1.0
2 0.7 0.0
3 0.2 0.0

```

Using a Series, the number of places for specific columns can be specified with the column names as index and the number of decimal places as value

```

>>> decimals = pd.Series([0, 1], index=['cats', 'dogs'])
>>> df.round(decimals)
 dogs cats
0 0.2 0.0
1 0.0 1.0
2 0.7 0.0
3 0.2 0.0

```

`rpow(self, other, axis='columns', level=None, fill_value=None)`  
Get Exponential power of dataframe and other, element-wise (binary operator ``rpow``).

Equivalent to ``other ** dataframe``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``pow``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

#### Parameters

```

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for successful DataFrame alignment, with this value before computation. If data in both corresponding DataFrame locations is missing the result will be missing.

```

#### Returns

```

DataFrame
 Result of the arithmetic operation.

```

#### See Also

```

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

```

#### Notes

-----  
Mismatched indices will be unioned together.

#### Examples

```
>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
```

	angles	degrees
circle	0	360
triangle	3	180
rectangle	4	360

Add a scalar with operator version which return the same results.

```
>>> df + 1
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
```

	angles	degrees
circle	-1	359
triangle	2	179
rectangle	3	359

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

Divide by a MultiIndex by level.

```

>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

rsub(self, other, axis='columns', level=None, fill_value=None)
 Get Subtraction of dataframe and other, element-wise (binary operator
 `rsub`).

 Equivalent to ``other - dataframe``, but with support to substitute a
 fill_value
 for missing data in one of the inputs. With reverse version, `sub`.

 Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
 arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

>>> df + 1
 angles degrees

```



circle	1	361
triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
```

	angles	degrees
circle	-1	359
triangle	2	179
rectangle	3	359

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

		angles	degrees
A	circle	0	360
	triangle	3	180
	rectangle	4	360
B	square	4	360
	pentagon	5	540
	hexagon	6	720

```
>>> df.div(df_multindex, level=1, fill_value=0)
```

		angles	degrees
A	circle	NaN	1.0
	triangle	1.0	1.0

rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

```

rtruediv(self, other, axis='columns', level=None, fill_value=None)
 Get Floating division of dataframe and other, element-wise (binary operator
`rtruediv`).

 Equivalent to ``other / dataframe``, but with support to substitute a
fill_value
for missing data in one of the inputs. With reverse version, `truediv`.

Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

>>> df.add(1)
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

Divide by constant with reverse version.

>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0

>>> df.rdiv(10)

```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
 angles
circle 0
triangle 3
rectangle 4

>>> df * other
 angles degrees
circle 0 NaN
triangle 9 NaN
rectangle 16 NaN

>>> df.mul(other, fill_value=0)
 angles degrees
circle 0 0.0
triangle 9 0.0
rectangle 16 0.0
```

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720

>>> df.div(df_multindex, level=1, fill_value=0)
 angles degrees
A circle NaN 1.0
 triangle 1.0 1.0
 rectangle 1.0 1.0
B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0
```

`select_dtypes(self, include=None, exclude=None) -> 'DataFrame'`

Return a subset of the DataFrame's columns based on the column dtypes.

Parameters

-----

`include, exclude` : scalar or list-like

A selection of dtypes or strings to be included/excluded. At least one of these parameters must be supplied.

Returns

-----

DataFrame

The subset of the frame including the dtypes in ``include`` and excluding the dtypes in ``exclude``.

## Raises

-----

### ValueError

- \* If both of ``include`` and ``exclude`` are empty
- \* If ``include`` and ``exclude`` have overlapping elements
- \* If any kind of string dtype is passed in.

## See Also

-----

DataFrame.dtypes: Return Series with the data type of each column.

## Notes

-----

- \* To select all *numeric* types, use ``np.number`` or ``'number'``
- \* To select strings you must use the ``object`` dtype, but note that this will return *all* object dtype columns
- \* See the *numpy dtype hierarchy*  
<<https://numpy.org/doc/stable/reference/arrays.scalars.html>>`\_
- \* To select datetimes, use ``np.datetime64``, ``'datetime'`` or ``'datetime64'``
- \* To select timedeltas, use ``np.timedelta64``, ``'timedelta'`` or ``'timedelta64'``
- \* To select Pandas categorical dtypes, use ``'category'``
- \* To select Pandas datetimetz dtypes, use ``'datetimetz'`` (new in 0.20.0) or ``'datetime64[ns, tz]'``

## Examples

-----

```
>>> df = pd.DataFrame({'a': [1, 2] * 3,
... 'b': [True, False] * 3,
... 'c': [1.0, 2.0] * 3})
```

```
>>> df
 a b c
0 1 True 1.0
1 2 False 2.0
2 1 True 1.0
3 2 False 2.0
4 1 True 1.0
5 2 False 2.0
```

```
>>> df.select_dtypes(include='bool')
```

```
 b
0 True
1 False
2 True
3 False
4 True
5 False
```

```
>>> df.select_dtypes(include=['float64'])
```

```
 c
0 1.0
1 2.0
2 1.0
3 2.0
4 1.0
5 2.0
```

```
>>> df.select_dtypes(exclude=['int64'])
```

```
 b c
0 True 1.0
1 False 2.0
2 True 1.0
3 False 2.0
4 True 1.0
5 False 2.0
```

```
sem(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
```

Return unbiased standard error of the mean over requested axis.

Normalized by N-1 by default. This can be changed using the ddof argument

## Parameters

-----

axis : {index (0), columns (1)}

skipna : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

ddof : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is N - ddof, where N represents the number of elements.

numeric\_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

#### Returns

Series or DataFrame (if level specified)

#### Notes

To have the same behaviour as ``numpy.std``, use ``ddof=0`` (instead of the default ``ddof=1``)

`set_axis(self, labels, axis: 'Axis' = 0, inplace: 'bool' = False)`  
Assign desired index to given axis.

Indexes for column or row labels can be changed by assigning a list-like or Index.

#### Parameters

`labels` : list-like, Index  
The values for the new index.

`axis` : {0 or 'index', 1 or 'columns'}, default 0  
The axis to update. The value 0 identifies the rows, and 1 identifies the

columns.

`inplace` : bool, default False  
Whether to return a new DataFrame instance.

#### Returns

`renamed` : DataFrame or None  
An object of type DataFrame or None if ``inplace=True``.

#### See Also

`DataFrame.rename_axis` : Alter the name of the index or columns.

#### Examples

```
>>> df = pd.DataFrame({"A": [1, 2, 3], "B": [4, 5, 6]})
```

Change the row labels.

```
>>> df.set_axis(['a', 'b', 'c'], axis='index')
 A B
a 1 4
b 2 5
c 3 6
```

Change the column labels.

```
>>> df.set_axis(['I', 'II'], axis='columns')
 I II
0 1 4
1 2 5
2 3 6
```

Now, update the labels inplace.

```
>>> df.set_axis(['i', 'ii'], axis='columns', inplace=True)
>>> df
 i ii
0 1 4
1 2 5
2 3 6
```

`set_index(self, keys, drop: 'bool' = True, append: 'bool' = False, inplace: 'bool' = False, verify_integrity: 'bool' = False)`  
Set the DataFrame index using existing columns.

Set the DataFrame index (row labels) using one or more existing columns or arrays (of the correct length). The index can replace the existing index or expand on it.

#### Parameters

`keys` : label or array-like or list of labels/arrays  
This parameter can be either a single column key, a single array of the same length as the calling DataFrame, or a list containing an arbitrary combination of column keys and arrays. Here, "array" encompasses `:class:`Series``, `:class:`Index``, ``np.ndarray``, and instances of `:class:`~collections.abc.Iterator``.

`drop` : bool, default True  
Delete columns to be used as the new index.

`append` : bool, default False

Whether to append columns to existing index.  
inplace : bool, default False  
If True, modifies the DataFrame in place (do not create a new object).  
verify\_integrity : bool, default False  
Check the new index for duplicates. Otherwise defer the check until necessary. Setting to False will improve the performance of this method.

#### Returns

-----  
DataFrame or None  
Changed row labels or None if ``inplace=True``.

#### See Also

-----  
DataFrame.reset\_index : Opposite of set\_index.  
DataFrame.reindex : Change to new indices or expand indices.  
DataFrame.reindex\_like : Change to same indices as other DataFrame.

#### Examples

-----  
>>> df = pd.DataFrame({'month': [1, 4, 7, 10],  
... 'year': [2012, 2014, 2013, 2014],  
... 'sale': [55, 40, 84, 31]})  
>>> df  
 month year sale  
0 1 2012 55  
1 4 2014 40  
2 7 2013 84  
3 10 2014 31

Set the index to become the 'month' column:

```
>>> df.set_index('month')
 year sale
month
1 2012 55
4 2014 40
7 2013 84
10 2014 31
```

Create a MultiIndex using columns 'year' and 'month':

```
>>> df.set_index(['year', 'month'])
 sale
year month
2012 1 55
2014 4 40
2013 7 84
2014 10 31
```

Create a MultiIndex using an Index and a column:

```
>>> df.set_index([pd.Index([1, 2, 3, 4]), 'year'])
 month sale
year
1 2012 1 55
2 2014 4 40
3 2013 7 84
4 2014 10 31
```

Create a MultiIndex using two Series:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> df.set_index([s, s**2])
 month year sale
1 1 1 2012 55
2 4 4 2014 40
3 9 7 2013 84
4 16 10 2014 31
```

```
shift(self, periods=1, freq: 'Frequency | None' = None, axis: 'Axis' = 0,
fill_value=<no_default>) -> 'DataFrame'
```

Shift index by desired number of periods with an optional time `freq`.

When `freq` is not passed, shift the index without realigning the data.  
If `freq` is passed (in this case, the index must be date or datetime,  
or it will raise a `NotImplementedError`), the index will be  
increased using the periods and the `freq`. `freq` can be inferred  
when specified as "infer" as long as either freq or inferred\_freq  
attribute is set in the index.

#### Parameters

-----  
periods : int  
Number of periods to shift. Can be positive or negative.  
freq : DateOffset, tseries.offsets, timedelta, or str, optional

Offset to use from the tseries module or time rule (e.g. 'EOM').  
 If `freq` is specified then the index values are shifted but the data is not realigned. That is, use `freq` if you would like to extend the index when shifting and preserve the original data.  
 If `freq` is specified as "infer" then it will be inferred from the freq or inferred\_freq attributes of the index. If neither of those attributes exist, a ValueError is thrown.

axis : {0 or 'index', 1 or 'columns', None}, default None

Shift direction.

fill\_value : object, optional  
 The scalar value to use for newly introduced missing values.  
 the default depends on the dtype of `self`.  
 For numeric data, ``np.nan`` is used.  
 For datetime, timedelta, or period data, etc. :attr:`NaT` is used.  
 For extension dtypes, ``self.dtype.na\_value`` is used.

.. versionchanged:: 1.1.0

#### Returns

-----

DataFrame

Copy of input object, shifted.

#### See Also

-----

Index.shift : Shift values of Index.

DatetimeIndex.shift : Shift values of DatetimeIndex.

PeriodIndex.shift : Shift values of PeriodIndex.

tshift : Shift the time index, using the index's frequency if available.

#### Examples

-----

```
>>> df = pd.DataFrame({"Col1": [10, 20, 15, 30, 45],
... "Col2": [13, 23, 18, 33, 48],
... "Col3": [17, 27, 22, 37, 52]},
... index=pd.date_range("2020-01-01", "2020-01-05"))
>>> df
```

	Col1	Col2	Col3
2020-01-01	10	13	17
2020-01-02	20	23	27
2020-01-03	15	18	22
2020-01-04	30	33	37
2020-01-05	45	48	52

```
>>> df.shift(periods=3)
 Col1 Col2 Col3
2020-01-01 NaN NaN NaN
2020-01-02 NaN NaN NaN
2020-01-03 NaN NaN NaN
2020-01-04 10.0 13.0 17.0
2020-01-05 20.0 23.0 27.0
```

```
>>> df.shift(periods=1, axis="columns")
 Col1 Col2 Col3
2020-01-01 NaN 10 13
2020-01-02 NaN 20 23
2020-01-03 NaN 15 18
2020-01-04 NaN 30 33
2020-01-05 NaN 45 48
```

```
>>> df.shift(periods=3, fill_value=0)
 Col1 Col2 Col3
2020-01-01 0 0 0
2020-01-02 0 0 0
2020-01-03 0 0 0
2020-01-04 10 13 17
2020-01-05 20 23 27
```

```
>>> df.shift(periods=3, freq="D")
 Col1 Col2 Col3
2020-01-04 10 13 17
2020-01-05 20 23 27
2020-01-06 15 18 22
2020-01-07 30 33 37
2020-01-08 45 48 52
```

```
>>> df.shift(periods=3, freq="infer")
 Col1 Col2 Col3
2020-01-04 10 13 17
2020-01-05 20 23 27
2020-01-06 15 18 22
2020-01-07 30 33 37
2020-01-08 45 48 52
```

skew(self, axis=None, skipna=None, level=None, numeric\_only=None, \*\*kwargs)  
 Return unbiased skew over requested axis.

Normalized by N-1.

#### Parameters

-----

axis : {index (0), columns (1)}

Axis for the function to be applied on.

skipna : bool, default True

Exclude NA/null values when computing the result.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

numeric\_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

\*\*kwargs

Additional keyword arguments to be passed to the function.

#### Returns

-----

Series or DataFrame (if level specified)

```
sort_index(self, axis: 'Axis' = 0, level: 'Level | None' = None, ascending: 'bool | int | Sequence[bool | int]' = True, inplace: 'bool' = False, kind: 'str' = 'quicksort', na_position: 'str' = 'last', sort_remaining: 'bool' = True, ignore_index: 'bool' = False, key: 'IndexKeyFunc' = None)
 Sort object by labels (along an axis).
```

Returns a new DataFrame sorted by label if `inplace` argument is ``False``, otherwise updates the original DataFrame and returns None.

#### Parameters

-----

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis along which to sort. The value 0 identifies the rows, and 1 identifies the columns.

level : int or level name or list of ints or list of level names

If not None, sort on values in specified index level(s).

ascending : bool or list-like of bools, default True

Sort ascending vs. descending. When the index is a MultiIndex the sort direction can be controlled for each level individually.

inplace : bool, default False

If True, perform operation in-place.

kind : {'quicksort', 'mergesort', 'heapsort', 'stable'}, default 'quicksort'

Choice of sorting algorithm. See also :func:`numpy.sort` for more information. `mergesort` and `stable` are the only stable algorithms. For DataFrames, this option is only applied when sorting on a single column or label.

na\_position : {'first', 'last'}, default 'last'

Puts NaNs at the beginning if `first`; `last` puts NaNs at the end. Not implemented for MultiIndex.

sort\_remaining : bool, default True

If True and sorting by level and index is multilevel, sort by other levels too (in order) after sorting by specified level.

ignore\_index : bool, default False

If True, the resulting axis will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.0.0

key : callable, optional

If not None, apply the key function to the index values before sorting. This is similar to the `key` argument in the builtin :meth:`sorted` function, with the notable difference that this `key` function should be \*vectorized\*. It should expect an ``Index`` and return an ``Index`` of the same shape. For MultiIndex inputs, the key is applied \*per level\*.

.. versionadded:: 1.1.0

#### Returns

-----

DataFrame or None

The original DataFrame sorted by the labels or None if ``inplace=True``.

#### See Also

-----

Series.sort\_index : Sort Series by the index.

DataFrame.sort\_values : Sort DataFrame by the value.

Series.sort\_values : Sort Series by the value.

#### Examples

-----

```
>>> df = pd.DataFrame([1, 2, 3, 4, 5], index=[100, 29, 234, 1, 150],
... columns=['A'])
```

```
>>> df.sort_index()
```

```
 A
```

```
1 4
```



```

29 2
100 1
150 5
234 3

```

By default, it sorts in ascending order, to sort in descending order, use ``ascending=False``

```

>>> df.sort_index(ascending=False)
 A
234 3
150 5
100 1
29 2
1 4

```

A key function can be specified which is applied to the index before sorting. For a ``MultiIndex`` this is applied to each level separately.

```

>>> df = pd.DataFrame({"a": [1, 2, 3, 4]}, index=['A', 'b', 'C', 'd'])
>>> df.sort_index(key=lambda x: x.str.lower())
 a
A 1
b 2
C 3
d 4

```

sort\_values(self, by, axis: 'Axis' = 0, ascending=True, inplace: 'bool' = False, kind: 'str' = 'quicksort', na\_position: 'str' = 'last', ignore\_index: 'bool' = False, key: 'ValueKeyFunc' = None)

Sort by the values along either axis.

Parameters

```

by : str or list of str
 Name or list of names to sort by.

 - if `axis` is 0 or `index` then `by` may contain index
 levels and/or column labels.
 - if `axis` is 1 or `columns` then `by` may contain column
 levels and/or index labels.
axis : {0 or 'index', 1 or 'columns'}, default 0
 Axis to be sorted.
ascending : bool or list of bool, default True
 Sort ascending vs. descending. Specify list for multiple sort
 orders. If this is a list of bools, must match the length of
 the by.
inplace : bool, default False
 If True, perform operation in-place.
kind : {'quicksort', 'mergesort', 'heapsort', 'stable'}, default 'quicksort'
 Choice of sorting algorithm. See also :func:`numpy.sort` for more
 information. `mergesort` and `stable` are the only stable algorithms.

```

For

```

 DataFrames, this option is only applied when sorting on a single
 column or label.
na_position : {'first', 'last'}, default 'last'
 Puts NaNs at the beginning if `first`; `last` puts NaNs at the
 end.
ignore_index : bool, default False
 If True, the resulting axis will be labeled 0, 1, ..., n - 1.

```

.. versionadded:: 1.0.0

```

key : callable, optional
 Apply the key function to the values
 before sorting. This is similar to the `key` argument in the
 builtin :meth:`sorted` function, with the notable difference that
 this `key` function should be *vectorized*. It should expect a
 ``Series`` and return a Series with the same shape as the input.
 It will be applied to each column in `by` independently.

```

.. versionadded:: 1.1.0

Returns

```

DataFrame or None
 DataFrame with sorted values or None if ``inplace=True``.

```

See Also

```

DataFrame.sort_index : Sort a DataFrame by the index.
Series.sort_values : Similar method for a Series.

```

Examples

```

>>> df = pd.DataFrame({
... 'coll': ['A', 'A', 'B', np.nan, 'D', 'C'],

```

```
... 'col2': [2, 1, 9, 8, 7, 4],
... 'col3': [0, 1, 9, 4, 2, 3],
... 'col4': ['a', 'B', 'c', 'D', 'e', 'F']
... })
>>> df
 col1 col2 col3 col4
0 A 2 0 a
1 A 1 1 B
2 B 9 9 c
3 NaN 8 4 D
4 D 7 2 e
5 C 4 3 F
```

Sort by col1

```
>>> df.sort_values(by=['col1'])
 col1 col2 col3 col4
0 A 2 0 a
1 A 1 1 B
2 B 9 9 c
5 C 4 3 F
4 D 7 2 e
3 NaN 8 4 D
```

Sort by multiple columns

```
>>> df.sort_values(by=['col1', 'col2'])
 col1 col2 col3 col4
1 A 1 1 B
0 A 2 0 a
2 B 9 9 c
5 C 4 3 F
4 D 7 2 e
3 NaN 8 4 D
```

Sort Descending

```
>>> df.sort_values(by='col1', ascending=False)
 col1 col2 col3 col4
4 D 7 2 e
5 C 4 3 F
2 B 9 9 c
0 A 2 0 a
1 A 1 1 B
3 NaN 8 4 D
```

Putting NAs first

```
>>> df.sort_values(by='col1', ascending=False, na_position='first')
 col1 col2 col3 col4
3 NaN 8 4 D
4 D 7 2 e
5 C 4 3 F
2 B 9 9 c
0 A 2 0 a
1 A 1 1 B
```

Sorting with a key function

```
>>> df.sort_values(by='col4', key=lambda col: col.str.lower())
 col1 col2 col3 col4
0 A 2 0 a
1 A 1 1 B
2 B 9 9 c
3 NaN 8 4 D
4 D 7 2 e
5 C 4 3 F
```

Natural sort with the key argument,  
using the `natsort` <<https://github.com/SethMMorton/natsort>> package.

```
>>> df = pd.DataFrame({
... "time": ['0hr', '128hr', '72hr', '48hr', '96hr'],
... "value": [10, 20, 30, 40, 50]
... })
>>> df
 time value
0 0hr 10
1 128hr 20
2 72hr 30
3 48hr 40
4 96hr 50
>>> from natsort import index_natsorted
>>> df.sort_values(
... by="time",
... key=lambda x: np.argsort(index_natsorted(df["time"]))
...)
```

	time	value
0	0hr	10
3	48hr	40
2	72hr	30
4	96hr	50
1	128hr	20

```
stack(self, level: 'Level' = -1, dropna: 'bool' = True)
 Stack the prescribed level(s) from columns to index.
```

Return a reshaped DataFrame or Series having a multi-level index with one or more new inner-most levels compared to the current DataFrame. The new inner-most levels are created by pivoting the columns of the current dataframe:

- if the columns have a single level, the output is a Series;
- if the columns have multiple levels, the new index level(s) is (are) taken from the prescribed level(s) and the output is a DataFrame.

#### Parameters

**level** : int, str, list, default -1  
 Level(s) to stack from the column axis onto the index axis, defined as one index or label, or a list of indices or labels.

**dropna** : bool, default True  
 Whether to drop rows in the resulting Frame/Series with missing values. Stacking a column level onto the index axis can create combinations of index and column values that are missing from the original dataframe. See Examples section.

#### Returns

DataFrame or Series  
 Stacked dataframe or series.

#### See Also

**DataFrame.unstack** : Unstack prescribed level(s) from index axis onto column axis.

**DataFrame.pivot** : Reshape dataframe from long format to wide format.

**DataFrame.pivot\_table** : Create a spreadsheet-style pivot table as a DataFrame.

#### Notes

The function is named by analogy with a collection of books being reorganized from being side by side on a horizontal position (the columns of the dataframe) to being stacked vertically on top of each other (in the index of the dataframe).

#### Examples

**\*\*Single level columns\*\***

```
>>> df_single_level_cols = pd.DataFrame([[0, 1], [2, 3]],
... index=['cat', 'dog'],
... columns=['weight', 'height'])
```

Stacking a dataframe with a single level column axis returns a Series:

```
>>> df_single_level_cols
 weight height
cat 0 1
dog 2 3
>>> df_single_level_cols.stack()
cat weight 0
 height 1
dog weight 2
 height 3
dtype: int64
```

**\*\*Multi level columns: simple case\*\***

```
>>> multicoll = pd.MultiIndex.from_tuples([('weight', 'kg'),
... ('weight', 'pounds')])
>>> df_multi_level_cols1 = pd.DataFrame([[1, 2], [2, 4]],
... index=['cat', 'dog'],
... columns=multicoll)
```

Stacking a dataframe with a multi-level column axis:

```
>>> df_multi_level_cols1
```

```

weight
kg pounds
cat 1 2
dog 2 4
>>> df_multi_level_cols1.stack()
weight
cat kg 1
pounds 2
dog kg 2
pounds 4

```

**\*\*Missing values\*\***

```

>>> multicol2 = pd.MultiIndex.from_tuples([('weight', 'kg'),
... ('height', 'm')])
>>> df_multi_level_cols2 = pd.DataFrame([[1.0, 2.0], [3.0, 4.0]],
... index=['cat', 'dog'],
... columns=multicol2)

```

It is common to have missing values when stacking a dataframe with multi-level columns, as the stacked dataframe typically has more values than the original dataframe. Missing values are filled with NaNs:

```

>>> df_multi_level_cols2
weight height
kg m
cat 1.0 2.0
dog 3.0 4.0
>>> df_multi_level_cols2.stack()
height weight
cat kg NaN 1.0
m 2.0 NaN
dog kg NaN 3.0
m 4.0 NaN

```

**\*\*Prescribing the level(s) to be stacked\*\***

The first parameter controls which level or levels are stacked:

```

>>> df_multi_level_cols2.stack(0)
kg m
cat height NaN 2.0
weight 1.0 NaN
dog height NaN 4.0
weight 3.0 NaN
>>> df_multi_level_cols2.stack([0, 1])
cat height m 2.0
weight kg 1.0
dog height m 4.0
weight kg 3.0
dtype: float64

```

**\*\*Dropping missing values\*\***

```

>>> df_multi_level_cols3 = pd.DataFrame([[None, 1.0], [2.0, 3.0]],
... index=['cat', 'dog'],
... columns=multicol2)

```

Note that rows where all values are missing are dropped by default but this behaviour can be controlled via the dropna keyword parameter:

```

>>> df_multi_level_cols3
weight height
kg m
cat NaN 1.0
dog 2.0 3.0
>>> df_multi_level_cols3.stack(dropna=False)
height weight
cat kg NaN NaN
m 1.0 NaN
dog kg NaN 2.0
m 3.0 NaN
>>> df_multi_level_cols3.stack(dropna=True)
height weight
cat m 1.0 NaN
dog kg NaN 2.0
m 3.0 NaN

```

```

std(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
Return sample standard deviation over requested axis.

Normalized by N-1 by default. This can be changed using the ddof argument

Parameters

```

```

axis : {index (0), columns (1)}
skipna : bool, default True
 Exclude NA/null values. If an entire row/column is NA, the result
 will be NA.
level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
ddof : int, default 1
 Delta Degrees of Freedom. The divisor used in calculations is N - ddof,
 where N represents the number of elements.
numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to use
 everything, then use only numeric data. Not implemented for Series.

Returns

Series or DataFrame (if level specified)

Notes

To have the same behaviour as `numpy.std`, use `ddof=0` (instead of the
default `ddof=1`)

sub(self, other, axis='columns', level=None, fill_value=None)
 Get Subtraction of dataframe and other, element-wise (binary operator `sub`).

 Equivalent to ``dataframe - other``, but with support to substitute a
fill_value
 for missing data in one of the inputs. With reverse version, `rsub`.

 Among flexible wrappers (`add`, `sub`, `mul`, `div`, `mod`, `pow`) to
 arithmetic operators: `+`, `-`, `*`, `/`, `//`, `%`, `**`.

Parameters

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

Returns

DataFrame
 Result of the arithmetic operation.

See Also

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

Notes

Mismatched indices will be unioned together.

Examples

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

Add a scalar with operator version which return the same
results.

>>> df + 1
 angles degrees
circle 1 361

```

triangle	4	181
rectangle	5	361

```
>>> df.add(1)
```

	angles	degrees
circle	1	361
triangle	4	181
rectangle	5	361

Divide by constant with reverse version.

```
>>> df.div(10)
```

	angles	degrees
circle	0.0	36.0
triangle	0.3	18.0
rectangle	0.4	36.0

```
>>> df.rdiv(10)
```

	angles	degrees
circle	inf	0.027778
triangle	3.333333	0.055556
rectangle	2.500000	0.027778

Subtract a list and Series by axis with operator version.

```
>>> df - [1, 2]
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub([1, 2], axis='columns')
```

	angles	degrees
circle	-1	358
triangle	2	178
rectangle	3	358

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
```

	angles	degrees
circle	-1	359
triangle	2	179
rectangle	3	359

Multiply a DataFrame of different shape with operator version.

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

Divide by a MultiIndex by level.

```
>>> df_multindex = pd.DataFrame({'angles': [0, 3, 4, 4, 5, 6],
... 'degrees': [360, 180, 360, 360, 540, 720]},
... index=[['A', 'A', 'A', 'B', 'B', 'B'],
... ['circle', 'triangle', 'rectangle',
... 'square', 'pentagon', 'hexagon']])
>>> df_multindex
```

	angles	degrees
A circle	0	360
triangle	3	180
rectangle	4	360
B square	4	360
pentagon	5	540
hexagon	6	720

```
>>> df.div(df_multindex, level=1, fill_value=0)
```

	angles	degrees
A circle	NaN	1.0
triangle	1.0	1.0
rectangle	1.0	1.0

```

 B square 0.0 0.0
 pentagon 0.0 0.0
 hexagon 0.0 0.0

subtract = sub(self, other, axis='columns', level=None, fill_value=None)

sum(self, axis=None, skipna=None, level=None, numeric_only=None, min_count=0,
**kwargs)
 Return the sum of the values over the requested axis.

 This is equivalent to the method ``numpy.sum``.

 Parameters

 axis : {index (0), columns (1)}
 Axis for the function to be applied on.
 skipna : bool, default True
 Exclude NA/null values when computing the result.
 level : int or level name, default None
 If the axis is a MultiIndex (hierarchical), count along a
 particular level, collapsing into a Series.
 numeric_only : bool, default None
 Include only float, int, boolean columns. If None, will attempt to use
 everything, then use only numeric data. Not implemented for Series.
 min_count : int, default 0
 The required number of valid values to perform the operation. If fewer
 than
 ``min_count`` non-NA values are present the result will be NA.
 **kwargs
 Additional keyword arguments to be passed to the function.

 Returns

 Series or DataFrame (if level specified)

 See Also

 Series.sum : Return the sum.
 Series.min : Return the minimum.
 Series.max : Return the maximum.
 Series.idxmin : Return the index of the minimum.
 Series.idxmax : Return the index of the maximum.
 DataFrame.sum : Return the sum over the requested axis.
 DataFrame.min : Return the minimum over the requested axis.
 DataFrame.max : Return the maximum over the requested axis.
 DataFrame.idxmin : Return the index of the minimum over the requested axis.
 DataFrame.idxmax : Return the index of the maximum over the requested axis.

 Examples

 >>> idx = pd.MultiIndex.from_arrays([
 ... ['warm', 'warm', 'cold', 'cold'],
 ... ['dog', 'falcon', 'fish', 'spider']],
 ... names=['blooded', 'animal'])
 >>> s = pd.Series([4, 2, 0, 8], name='legs', index=idx)
 >>> s
 blooded animal
 warm dog 4
 falcon 2
 cold fish 0
 spider 8
 Name: legs, dtype: int64

 >>> s.sum()
 14

 By default, the sum of an empty or all-NA Series is ``0``.

 >>> pd.Series([], dtype="float64").sum() # min_count=0 is the default
 0.0

 This can be controlled with the ``min_count`` parameter. For example, if
 you'd like the sum of an empty series to be NaN, pass ``min_count=1``.

 >>> pd.Series([], dtype="float64").sum(min_count=1)
 nan

 Thanks to the ``skipna`` parameter, ``min_count`` handles all-NA and
 empty series identically.

 >>> pd.Series([np.nan]).sum()
 0.0

 >>> pd.Series([np.nan]).sum(min_count=1)
 nan

 swaplevel(self, i: 'Axis' = -2, j: 'Axis' = -1, axis: 'Axis' = 0) -> 'DataFrame'

```

Swap levels i and j in a :class:`MultiIndex`.

Default is to swap the two innermost levels of the index.

#### Parameters

-----

i, j : int or str

Levels of the indices to be swapped. Can pass level name as string.

axis : {0 or 'index', 1 or 'columns'}, default 0

The axis to swap levels on. 0 or 'index' for row-wise, 1 or 'columns' for column-wise.

#### Returns

-----

DataFrame

DataFrame with levels swapped in MultiIndex.

#### Examples

-----

```
>>> df = pd.DataFrame(
... {"Grade": ["A", "B", "A", "C"]},
... index=[
... ["Final exam", "Final exam", "Coursework", "Coursework"],
... ["History", "Geography", "History", "Geography"],
... ["January", "February", "March", "April"],
...],
...)
>>> df
```

			Grade
Final exam	History	January	A
	Geography	February	B
Coursework	History	March	A
	Geography	April	C

In the following example, we will swap the levels of the indices. Here, we will swap the levels column-wise, but levels can be swapped

row-wise

in a similar manner. Note that column-wise is the default behaviour. By not supplying any arguments for i and j, we swap the last and

second to

last indices.

```
>>> df.swaplevel()
```

			Grade
Final exam	January	History	A
	February	Geography	B
Coursework	March	History	A
	April	Geography	C

By supplying one argument, we can choose which index to swap the last index with. We can for example swap the first index with the last one

as

follows.

```
>>> df.swaplevel(0)
```

			Grade
January	History	Final exam	A
February	Geography	Final exam	B
March	History	Coursework	A
April	Geography	Coursework	C

We can also define explicitly which indices we want to swap by

supplying values

for both i and j. Here, we for example swap the first and second

indices.

```
>>> df.swaplevel(0, 1)
```

			Grade
History	Final exam	January	A
Geography	Final exam	February	B
History	Coursework	March	A
Geography	Coursework	April	C

```
to_dict(self, orient: 'str' = 'dict', into=<class 'dict'>)
```

Convert the DataFrame to a dictionary.

The type of the key-value pairs can be customized with the parameters (see below).

#### Parameters

-----

orient : str {'dict', 'list', 'series', 'split', 'records', 'index'}

Determines the type of the values of the dictionary.

- 'dict' (default) : dict like {column -> {index -> value}}
- 'list' : dict like {column -> [values]}
- 'series' : dict like {column -> Series(values)}



- 'split' : dict like  
 {'index' -> [index], 'columns' -> [columns], 'data' -> [values]}
- 'records' : list like  
 [{column -> value}, ... , {column -> value}]
- 'index' : dict like {index -> {column -> value}}

Abbreviations are allowed. `s` indicates `series` and `sp` indicates `split`.

**into** : class, default dict  
 The collections.abc.Mapping subclass used for all Mappings in the return value. Can be the actual class or an empty instance of the mapping type you want. If you want a collections.defaultdict, you must pass it initialized.

#### Returns

dict, list or collections.abc.Mapping  
 Return a collections.abc.Mapping object representing the DataFrame. The resulting transformation depends on the `orient` parameter.

#### See Also

DataFrame.from\_dict: Create a DataFrame from a dictionary.  
 DataFrame.to\_json: Convert a DataFrame to JSON format.

#### Examples

```
>>> df = pd.DataFrame({'col1': [1, 2],
... 'col2': [0.5, 0.75]},
... index=['row1', 'row2'])
>>> df
 col1 col2
row1 1 0.50
row2 2 0.75
>>> df.to_dict()
{'col1': {'row1': 1, 'row2': 2}, 'col2': {'row1': 0.5, 'row2': 0.75}}
```

You can specify the return orientation.

```
>>> df.to_dict('series')
{'col1': row1 1
 row2 2
Name: col1, dtype: int64,
 'col2': row1 0.50
 row2 0.75
Name: col2, dtype: float64}

>>> df.to_dict('split')
{'index': ['row1', 'row2'], 'columns': ['col1', 'col2'],
 'data': [[1, 0.5], [2, 0.75]]}

>>> df.to_dict('records')
[{'col1': 1, 'col2': 0.5}, {'col1': 2, 'col2': 0.75}]

>>> df.to_dict('index')
{'row1': {'col1': 1, 'col2': 0.5}, 'row2': {'col1': 2, 'col2': 0.75}}
```

You can also specify the mapping type.

```
>>> from collections import OrderedDict, defaultdict
>>> df.to_dict(into=OrderedDict)
OrderedDict([('col1', OrderedDict([('row1', 1), ('row2', 2)])),
 ('col2', OrderedDict([('row1', 0.5), ('row2', 0.75)])])])
```

If you want a `defaultdict`, you need to initialize it:

```
>>> dd = defaultdict(list)
>>> df.to_dict('records', into=dd)
[defaultdict(<class 'list'>, {'col1': 1, 'col2': 0.5}),
 defaultdict(<class 'list'>, {'col1': 2, 'col2': 0.75})]
```

**to\_feather**(self, path: 'FilePathOrBuffer[AnyStr]', \*\*kwargs) -> 'None'  
 Write a DataFrame to the binary Feather format.

#### Parameters

**path** : str or file-like object  
 If a string, it will be used as Root Directory path.

**\*\*kwargs** :  
 Additional keywords passed to :func:`pyarrow.feather.write\_feather`. Starting with pyarrow 0.17, this includes the `compression`, `compression\_level`, `chunksize` and `version` keywords.

.. versionadded:: 1.1.0

**to\_gbq**(self, destination\_table: 'str', project\_id: 'str | None' = None,

```

chunksize: 'int | None' = None, reauth: 'bool' = False, if_exists: 'str' = 'fail',
auth_local_webserver: 'bool' = False, table_schema: 'list[dict[str, str]] | None' =
None, location: 'str | None' = None, progress_bar: 'bool' = True, credentials=None) -
> 'None'

 Write a DataFrame to a Google BigQuery table.

 This function requires the `pandas-gbq` package
 <https://pandas-gbq.readthedocs.io>`__`.

 See the `How to authenticate with Google BigQuery
 <https://pandas-gbq.readthedocs.io/en/latest/howto/authentication.html>`__
 guide for authentication instructions.

 Parameters

 destination_table : str
 Name of table to be written, in the form ``dataset.tablename``.
 project_id : str, optional
 Google BigQuery Account project ID. Optional when available from
 the environment.
 chunksize : int, optional
 Number of rows to be inserted in each chunk from the dataframe.
 Set to ``None`` to load the whole dataframe at once.
 reauth : bool, default False
 Force Google BigQuery to re-authenticate the user. This is useful
 if multiple accounts are used.
 if_exists : str, default 'fail'
 Behavior when the destination table exists. Value can be one of:

 ``'fail'``
 If table exists raise pandas_gbq.gbq.TableCreationError.
 ``'replace'``
 If table exists, drop it, recreate it, and insert data.
 ``'append'``
 If table exists, insert data. Create if does not exist.
 auth_local_webserver : bool, default False
 Use the `local webserver flow`_ instead of the `console flow`_
 when getting user credentials.

 .. _local webserver flow:
 https://google-auth-
 oauthlib.readthedocs.io/en/latest/reference/google_auth_oauthlib.flow.html#google_auth_
 h_oauthlib.flow.InstalledAppFlow.run_local_server
 .. _console flow:
 https://google-auth-
 oauthlib.readthedocs.io/en/latest/reference/google_auth_oauthlib.flow.html#google_auth_
 h_oauthlib.flow.InstalledAppFlow.run_console

 New in version 0.2.0 of pandas-gbq.
 table_schema : list of dicts, optional
 List of BigQuery table fields to which according DataFrame
 columns conform to, e.g. ``[{'name': 'col1', 'type':
 'STRING'},...]``. If schema is not provided, it will be
 generated according to dtypes of DataFrame columns. See
 BigQuery API documentation on available names of a field.

 New in version 0.3.1 of pandas-gbq.
 location : str, optional
 Location where the load job should run. See the `BigQuery locations
 documentation
 <https://cloud.google.com/bigquery/docs/dataset-locations>`__ for a
 list of available locations. The location must match that of the
 target dataset.

 New in version 0.5.0 of pandas-gbq.
 progress_bar : bool, default True
 Use the library `tqdm` to show the progress bar for the upload,
 chunk by chunk.

 New in version 0.5.0 of pandas-gbq.
 credentials : google.auth.credentials.Credentials, optional
 Credentials for accessing Google APIs. Use this parameter to
 override default credentials, such as to use Compute Engine
 :class:`google.auth.compute_engine.Credentials` or Service
 Account :class:`google.oauth2.service_account.Credentials`
 directly.

 New in version 0.8.0 of pandas-gbq.

 See Also

 pandas_gbq.to_gbq : This function in the pandas-gbq library.
 read_gbq : Read a DataFrame from Google BigQuery.

 to_html(self, buf: 'FilePathOrBuffer[str] | None' = None, columns: 'Sequence[str]
 | None' = None, col_space: 'ColspaceArgType | None' = None, header: 'bool |
 Sequence[str]' = True, index: 'bool' = True, na_rep: 'str' = 'NaN', formatters:

```

```
'FormattersType | None' = None, float_format: 'FloatFormatType | None' = None,
sparsify: 'bool | None' = None, index_names: 'bool' = True, justify: 'str | None' =
None, max_rows: 'int | None' = None, max_cols: 'int | None' = None, show_dimensions:
'bool | str' = False, decimal: 'str' = '.', bold_rows: 'bool' = True, classes: 'str |
list | tuple | None' = None, escape: 'bool' = True, notebook: 'bool' = False, border:
'int | None' = None, table_id: 'str | None' = None, render_links: 'bool' = False,
encoding: 'str | None' = None)
```

Render a DataFrame as an HTML table.

Parameters

```

buf : str, Path or StringIO-like, optional, default None
 Buffer to write to. If None, the output is returned as a string.
columns : sequence, optional, default None
 The subset of columns to write. Writes all columns by default.
col_space : str or int, list or dict of int or str, optional
 The minimum width of each column in CSS length units. An int is assumed
to be px units.
```

```
.. versionadded:: 0.25.0
 Ability to use str.
header : bool, optional
 Whether to print column labels, default True.
index : bool, optional, default True
 Whether to print index (row) labels.
na_rep : str, optional, default 'NaN'
 String representation of ``NaN`` to use.
formatters : list, tuple or dict of one-param. functions, optional
 Formatter functions to apply to columns' elements by position or
 name.
 The result of each function must be a unicode string.
 List/tuple must be of length equal to the number of columns.
float_format : one-parameter function, optional, default None
 Formatter function to apply to columns' elements if they are
 floats. This function must return a unicode string and will be
 applied only to the non-``NaN`` elements, with ``NaN`` being
 handled by ``na_rep``.
```

```
.. versionchanged:: 1.2.0

sparsify : bool, optional, default True
 Set to False for a DataFrame with a hierarchical index to print
 every multiindex key at each row.
index_names : bool, optional, default True
 Prints the names of the indexes.
justify : str, default None
 How to justify the column labels. If None uses the option from
 the print configuration (controlled by set_option), 'right' out
 of the box. Valid values are
```

```
* left
* right
* center
* justify
* justify-all
* start
* end
* inherit
* match-parent
* initial
* unset.
```

```
max_rows : int, optional
 Maximum number of rows to display in the console.
min_rows : int, optional
 The number of rows to display in the console in a truncated repr
 (when number of rows is above `max_rows`).
max_cols : int, optional
 Maximum number of columns to display in the console.
show_dimensions : bool, default False
 Display DataFrame dimensions (number of rows by number of columns).
decimal : str, default '.'
 Character recognized as decimal separator, e.g. ',' in Europe.
```

```
bold_rows : bool, default True
 Make the row labels bold in the output.
classes : str or list or tuple, default None
 CSS class(es) to apply to the resulting html table.
escape : bool, default True
 Convert the characters <, >, and & to HTML-safe sequences.
notebook : {True, False}, default False
 Whether the generated HTML is for IPython Notebook.
border : int
 A ``border=border`` attribute is included in the opening
 <table> tag. Default ``pd.options.display.html.border``.
encoding : str, default "utf-8"
 Set character encoding.
```

```

.. versionadded:: 1.0

table_id : str, optional
 A css id is included in the opening `<table>` tag if specified.
render_links : bool, default False
 Convert URLs to HTML links.

Returns

str or None
 If buf is None, returns the result as a string. Otherwise returns
 None.

See Also

to_string : Convert DataFrame to a string.

to_markdown(self, buf: 'IO[str] | str | None' = None, mode: 'str' = 'wt', index:
'bool' = True, storage_options: 'StorageOptions' = None, **kwargs) -> 'str | None'
 Print DataFrame in Markdown-friendly format.

.. versionadded:: 1.0.0

Parameters

buf : str, Path or StringIO-like, optional, default None
 Buffer to write to. If None, the output is returned as a string.
mode : str, optional
 Mode in which file is opened, "wt" by default.
index : bool, optional, default True
 Add index (row) labels.

.. versionadded:: 1.1.0
storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2.0

**kwargs
 These parameters will be passed to `tabulate
<https://pypi.org/project/tabulate>`_.

Returns

str
 DataFrame in Markdown-friendly format.

Notes

Requires the `tabulate <https://pypi.org/project/tabulate>`_ package.

Examples

>>> s = pd.Series(["elk", "pig", "dog", "quetzal"], name="animal")
>>> print(s.to_markdown())
| | animal |
|---|:-----|
| 0 | elk |
| 1 | pig |
| 2 | dog |
| 3 | quetzal |

Output markdown with a tabulate option.

>>> print(s.to_markdown(tablefmt="grid"))
+-----+
| | animal |
+====+
| 0 | elk |
+-----+
| 1 | pig |
+-----+
| 2 | dog |
+-----+
| 3 | quetzal |
+-----+

to_numpy(self, dtype: 'NpDtype | None' = None, copy: 'bool' = False, na_value=
'no_default') -> 'np.ndarray'
 Convert the DataFrame to a NumPy array.

By default, the dtype of the returned array will be the common NumPy
dtype of all types in the DataFrame. For example, if the dtypes are

```

``float16`` and ``float32``, the results dtype will be ``float32``. This may require copying data and coercing values, which may be expensive.

#### Parameters

-----  
dtype : str or numpy.dtype, optional  
The dtype to pass to :meth:`numpy.asarray`.  
copy : bool, default False  
Whether to ensure that the returned value is not a view on another array. Note that ``copy=False`` does not \*ensure\* that ``to\_numpy()`` is no-copy. Rather, ``copy=True`` ensure that a copy is made, even if not strictly necessary.  
na\_value : Any, optional  
The value to use for missing values. The default value depends on `dtype` and the dtypes of the DataFrame columns.

.. versionadded:: 1.1.0

#### Returns

-----  
numpy.ndarray

#### See Also

-----  
Series.to\_numpy : Similar method for Series.

#### Examples

-----  
>>> pd.DataFrame({"A": [1, 2], "B": [3, 4]}).to\_numpy()  
array([[1, 3],  
 [2, 4]])

With heterogeneous data, the lowest common type will have to be used.

>>> df = pd.DataFrame({"A": [1, 2], "B": [3.0, 4.5]})  
>>> df.to\_numpy()  
array([[1. , 3. ],  
 [2. , 4.5]])

For a mix of numeric and non-numeric types, the output array will have object dtype.

>>> df['C'] = pd.date\_range('2000', periods=2)  
>>> df.to\_numpy()  
array([[1, 3.0, Timestamp('2000-01-01 00:00:00')],  
 [2, 4.5, Timestamp('2000-01-02 00:00:00')]], dtype=object)

to\_parquet(self, path: 'FilePathOrBuffer | None' = None, engine: 'str' = 'auto', compression: 'str | None' = 'snappy', index: 'bool | None' = None, partition\_cols: 'list[str] | None' = None, storage\_options: 'StorageOptions' = None, \*\*kwargs) -> 'bytes | None'

Write a DataFrame to the binary parquet format.

This function writes the dataframe as a `parquet` file <<https://parquet.apache.org/>>`. You can choose different parquet backends, and have the option of compression. See :ref:`the user guide <io.parquet>` for more details.

#### Parameters

-----  
path : str or file-like object, default None  
If a string, it will be used as Root Directory path when writing a partitioned dataset. By file-like object, we refer to objects with a write() method, such as a file handle (e.g. via builtin open function) or io.BytesIO. The engine fastparquet does not accept file-like objects. If path is None, a bytes object is returned.

.. versionchanged:: 1.2.0

Previously this was "fname"

engine : {'auto', 'pyarrow', 'fastparquet'}, default 'auto'  
Parquet library to use. If 'auto', then the option ``io.parquet.engine`` is used. The default ``io.parquet.engine`` behavior is to try 'pyarrow', falling back to 'fastparquet' if 'pyarrow' is unavailable.  
compression : {'snappy', 'gzip', 'brotli', None}, default 'snappy'  
Name of the compression to use. Use ``None`` for no compression.  
index : bool, default None  
If ``True``, include the dataframe's index(es) in the file output. If ``False``, they will not be written to the file.  
If ``None``, similar to ``True`` the dataframe's index(es) will be saved. However, instead of being saved as values, the RangeIndex will be stored as a range in the metadata so it

```

 doesn't require much space and is faster. Other indexes will
 be included as columns in the file output.
partition_cols : list, optional, default None
 Column names by which to partition the dataset.
 Columns are partitioned in the order they are given.
 Must be None if path is not a string.
storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

 .. versionadded:: 1.2.0

**kwargs
 Additional arguments passed to the parquet library. See
 :ref:`pandas io <io.parquet>` for more details.

Returns

bytes if no path argument is provided else None

See Also

read_parquet : Read a parquet file.
DataFrame.to_csv : Write a csv file.
DataFrame.to_sql : Write to a sql table.
DataFrame.to_hdf : Write to hdf.

Notes

This function requires either the `fastparquet
<https://pypi.org/project/fastparquet>`_ or `pyarrow
<https://arrow.apache.org/docs/python/>`_ library.

Examples

>>> df = pd.DataFrame(data={'col1': [1, 2], 'col2': [3, 4]})
>>> df.to_parquet('df.parquet.gzip',
... compression='gzip') # doctest: +SKIP
>>> pd.read_parquet('df.parquet.gzip') # doctest: +SKIP
 col1 col2
0 1 3
1 2 4

If you want to get a buffer to the parquet content you can use a io.BytesIO
object, as long as you don't use partition_cols, which creates multiple
files.

>>> import io
>>> f = io.BytesIO()
>>> df.to_parquet(f)
>>> f.seek(0)
0
>>> content = f.read()

to_period(self, freq: 'Frequency | None' = None, axis: 'Axis' = 0, copy: 'bool' =
True) -> 'DataFrame'
 Convert DataFrame from DatetimeIndex to PeriodIndex.

 Convert DataFrame from DatetimeIndex to PeriodIndex with desired
 frequency (inferred from index if not passed).

Parameters

freq : str, default
 Frequency of the PeriodIndex.
axis : {0 or 'index', 1 or 'columns'}, default 0
 The axis to convert (the index by default).
copy : bool, default True
 If False then underlying input data is not copied.

Returns

DataFrame with PeriodIndex

to_records(self, index=True, column_dtypes=None, index_dtypes=None) ->
'np.recarray'
 Convert DataFrame to a NumPy record array.

 Index will be included as the first field of the record array if
 requested.

Parameters

index : bool, default True

```

```

 Include index in resulting record array, stored in 'index'
 field or using the index label, if set.
 column_dtypes : str, type, dict, default None
 If a string or type, the data type to store all columns. If
 a dictionary, a mapping of column names and indices (zero-indexed)
 to specific data types.
 index_dtypes : str, type, dict, default None
 If a string or type, the data type to store all index levels. If
 a dictionary, a mapping of index level names and indices
 (zero-indexed) to specific data types.

 This mapping is applied only if `index=True`.

Returns

numpy.recarray
 NumPy ndarray with the DataFrame labels as fields and each row
 of the DataFrame as entries.

See Also

DataFrame.from_records: Convert structured or record ndarray
 to DataFrame.
numpy.recarray: An ndarray that allows field access using
 attributes, analogous to typed columns in a
 spreadsheet.

Examples

>>> df = pd.DataFrame({'A': [1, 2], 'B': [0.5, 0.75]},
... index=['a', 'b'])
>>> df
 A B
a 1 0.50
b 2 0.75
>>> df.to_records()
rec.array([(('a', 1, 0.5), ('b', 2, 0.75))],
 dtype=[('index', 'O'), ('A', '<i8'), ('B', '<f8')])

If the DataFrame index has no label then the recarray field name
is set to 'index'. If the index has a label then this is used as the
field name:

>>> df.index = df.index.rename("I")
>>> df.to_records()
rec.array([(('a', 1, 0.5), ('b', 2, 0.75))],
 dtype=[('I', 'O'), ('A', '<i8'), ('B', '<f8')])

The index can be excluded from the record array:

>>> df.to_records(index=False)
rec.array([(1, 0.5), (2, 0.75)],
 dtype=[('A', '<i8'), ('B', '<f8')])

Data types can be specified for the columns:

>>> df.to_records(column_dtypes={"A": "int32"})
rec.array([(('a', 1, 0.5), ('b', 2, 0.75))],
 dtype=[('I', 'O'), ('A', '<i4'), ('B', '<f8')])

As well as for the index:

>>> df.to_records(index_dtypes="<S2")
rec.array([(('a', 1, 0.5), ('b', 2, 0.75))],
 dtype=[('I', 'S2'), ('A', '<i8'), ('B', '<f8')])

>>> index_dtypes = f"<S{df.index.str.len().max()}"
>>> df.to_records(index_dtypes=index_dtypes)
rec.array([(('a', 1, 0.5), ('b', 2, 0.75))],
 dtype=[('I', 'S1'), ('A', '<i8'), ('B', '<f8')])

to_stata(self, path: 'FilePathOrBuffer', convert_dates: 'dict[Hashable, str] |
None' = None, write_index: 'bool' = True, byteorder: 'str | None' = None, time_stamp:
'datetime.datetime | None' = None, data_label: 'str | None' = None, variable_labels:
'dict[Hashable, str] | None' = None, version: 'int | None' = 114, convert_strl:
'Sequence[Hashable] | None' = None, compression: 'CompressionOptions' = 'infer',
storage_options: 'StorageOptions' = None) -> 'None'
 Export DataFrame object to Stata dta format.

Writes the DataFrame to a Stata dataset file.
"dta" files contain a Stata dataset.

Parameters

path : str, buffer or path object
 String, path object (pathlib.Path or py_path.local.LocalPath) or
 object implementing a binary write() function. If using a buffer

```

```

then the buffer will not be automatically closed after the file
data has been written.

.. versionchanged:: 1.0.0

Previously this was "fname"

convert_dates : dict
 Dictionary mapping columns containing datetime types to stata
 internal format to use when writing the dates. Options are 'tc',
 'td', 'tm', 'tw', 'th', 'tq', 'ty'. Column can be either an integer
 or a name. Datetime columns that do not have a conversion type
 specified will be converted to 'tc'. Raises NotImplementedError if
 a datetime column has timezone information.
write_index : bool
 Write the index to Stata dataset.
byteorder : str
 Can be ">", "<", "little", or "big". default is `sys.byteorder`.
time_stamp : datetime
 A datetime to use as file creation date. Default is the current
 time.
data_label : str, optional
 A label for the data set. Must be 80 characters or smaller.
variable_labels : dict
 Dictionary containing columns as keys and variable labels as
 values. Each label must be 80 characters or smaller.
version : {114, 117, 118, 119, None}, default 114
 Version to use in the output dta file. Set to None to let pandas
 decide between 118 or 119 formats depending on the number of
 columns in the frame. Version 114 can be read by Stata 10 and
 later. Version 117 can be read by Stata 13 or later. Version 118
 is supported in Stata 14 and later. Version 119 is supported in
 Stata 15 and later. Version 114 limits string variables to 244
 characters or fewer while versions 117 and later allow strings
 with lengths up to 2,000,000 characters. Versions 118 and 119
 support Unicode characters, and version 119 supports more than
 32,767 variables.

 Version 119 should usually only be used when the number of
 variables exceeds the capacity of dta format 118. Exporting
 smaller datasets in format 119 may have unintended consequences,
 and, as of November 2020, Stata SE cannot read version 119 files.

.. versionchanged:: 1.0.0

 Added support for formats 118 and 119.

convert_strl : list, optional
 List of column names to convert to string columns to Stata StrL
 format. Only available if version is 117. Storing strings in the
 StrL format can produce smaller dta files if strings have more than
 8 characters and values are repeated.
compression : str or dict, default 'infer'
 For on-the-fly compression of the output dta. If string, specifies
 compression mode. If dict, value at key 'method' specifies
 compression mode. Compression mode must be one of {'infer', 'gzip',
 'bz2', 'zip', 'xz', None}. If compression mode is 'infer' and
 `fname` is path-like, then detect compression from the following
 extensions: '.gz', '.bz2', '.zip', or '.xz' (otherwise no
 compression). If dict and compression mode is one of {'zip',
 'gzip', 'bz2'}, or inferred as one of the above, other entries
 passed as additional compression options.

.. versionadded:: 1.1.0

storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2.0

Raises

NotImplementedError
 * If datetimes contain timezone information
 * Column dtype is not representable in Stata
ValueError
 * Columns listed in convert_dates are neither datetime64[ns]
 or datetime.datetime
 * Column listed in convert_dates is not in DataFrame
 * Categorical label contains more than 32,000 characters

See Also

```



```

read_stata : Import Stata data files.
io.stata.StataWriter : Low-level writer for Stata data files.
io.stata.StataWriter117 : Low-level writer for version 117 files.

Examples

>>> df = pd.DataFrame({'animal': ['falcon', 'parrot', 'falcon',
... 'parrot'],
... 'speed': [350, 18, 361, 15]})
>>> df.to_stata('animals.dta') # doctest: +SKIP

to_string(self, buf: 'FilePathOrBuffer[str] | None' = None, columns:
'Sequence[str] | None' = None, col_space: 'int | None' = None, header: 'bool |
Sequence[str]' = True, index: 'bool' = True, na_rep: 'str' = 'NaN', formatters:
'fmt.FormatterType | None' = None, float_format: 'fmt.FloatFormatType | None' =
None, sparsify: 'bool | None' = None, index_names: 'bool' = True, justify: 'str |
None' = None, max_rows: 'int | None' = None, min_rows: 'int | None' = None, max_cols:
'int | None' = None, show_dimensions: 'bool' = False, decimal: 'str' = '.',
line_width: 'int | None' = None, max_colwidth: 'int | None' = None, encoding: 'str |
None' = None) -> 'str | None'
 Render a DataFrame to a console-friendly tabular output.

Parameters

buf : str, Path or StringIO-like, optional, default None
 Buffer to write to. If None, the output is returned as a string.
columns : sequence, optional, default None
 The subset of columns to write. Writes all columns by default.
col_space : int, list or dict of int, optional
 The minimum width of each column.
header : bool or sequence, optional
 Write out the column names. If a list of strings is given, it is assumed
to be aliases for the column names.
index : bool, optional, default True
 Whether to print index (row) labels.
na_rep : str, optional, default 'NaN'
 String representation of ``NaN`` to use.
formatters : list, tuple or dict of one-param. functions, optional
 Formatter functions to apply to columns' elements by position or
 name.
 The result of each function must be a unicode string.
 List/tuple must be of length equal to the number of columns.
float_format : one-parameter function, optional, default None
 Formatter function to apply to columns' elements if they are
 floats. This function must return a unicode string and will be
 applied only to the non-``NaN`` elements, with ``NaN`` being
 handled by ``na_rep``.

 .. versionchanged:: 1.2.0

sparsify : bool, optional, default True
 Set to False for a DataFrame with a hierarchical index to print
 every multiindex key at each row.
index_names : bool, optional, default True
 Prints the names of the indexes.
justify : str, default None
 How to justify the column labels. If None uses the option from
 the print configuration (controlled by set_option), 'right' out
 of the box. Valid values are

 * left
 * right
 * center
 * justify
 * justify-all
 * start
 * end
 * inherit
 * match-parent
 * initial
 * unset.
max_rows : int, optional
 Maximum number of rows to display in the console.
min_rows : int, optional
 The number of rows to display in the console in a truncated repr
 (when number of rows is above `max_rows`).
max_cols : int, optional
 Maximum number of columns to display in the console.
show_dimensions : bool, default False
 Display DataFrame dimensions (number of rows by number of columns).
decimal : str, default '.'
 Character recognized as decimal separator, e.g. ',' in Europe.

line_width : int, optional
 Width to wrap a line in characters.
max_colwidth : int, optional
 Max width to truncate each column in characters. By default, no limit.

```

```

 .. versionadded:: 1.0.0
encoding : str, default "utf-8"
 Set character encoding.

 .. versionadded:: 1.0

Returns

str or None
 If buf is None, returns the result as a string. Otherwise returns
 None.

See Also

to_html : Convert DataFrame to HTML.

Examples

>>> d = {'col1': [1, 2, 3], 'col2': [4, 5, 6]}
>>> df = pd.DataFrame(d)
>>> print(df.to_string())
 col1 col2
0 1 4
1 2 5
2 3 6

to_timestamp(self, freq: 'Frequency | None' = None, how: 'str' = 'start', axis:
'Axis' = 0, copy: 'bool' = True) -> 'DataFrame'
 Cast to DatetimeIndex of timestamps, at *beginning* of period.

Parameters

freq : str, default frequency of PeriodIndex
 Desired frequency.
how : {'s', 'e', 'start', 'end'}
 Convention for converting period to timestamp; start of period
 vs. end.
axis : {0 or 'index', 1 or 'columns'}, default 0
 The axis to convert (the index by default).
copy : bool, default True
 If False then underlying input data is not copied.

Returns

DataFrame with DatetimeIndex

to_xml(self, path_or_buffer: 'FilePathOrBuffer | None' = None, index: 'bool' =
True, root_name: 'str | None' = 'data', row_name: 'str | None' = 'row', na_rep: 'str
| None' = None, attr_cols: 'str | list[str] | None' = None, elem_cols: 'str |
list[str] | None' = None, namespaces: 'dict[str | None, str] | None' = None, prefix:
'str | None' = None, encoding: 'str' = 'utf-8', xml_declaration: 'bool | None' =
True, pretty_print: 'bool | None' = True, parser: 'str | None' = 'lxml', stylesheet:
'FilePathOrBuffer | None' = None, compression: 'CompressionOptions' = 'infer',
storage_options: 'StorageOptions' = None) -> 'str | None'
 Render a DataFrame to an XML document.

 .. versionadded:: 1.3.0

Parameters

path_or_buffer : str, path object or file-like object, optional
 File to write output to. If None, the output is returned as a
 string.
index : bool, default True
 Whether to include index in XML document.
root_name : str, default 'data'
 The name of root element in XML document.
row_name : str, default 'row'
 The name of row element in XML document.
na_rep : str, optional
 Missing data representation.
attr_cols : list-like, optional
 List of columns to write as attributes in row element.
 Hierarchical columns will be flattened with underscore
 delimiting the different levels.
elem_cols : list-like, optional
 List of columns to write as children in row element. By default,
 all columns output as children of row element. Hierarchical
 columns will be flattened with underscore delimiting the
 different levels.
namespaces : dict, optional
 All namespaces to be defined in root element. Keys of dict
 should be prefix names and values of dict corresponding URIs.
 Default namespaces should be given empty string key. For
 example, ::

```

```

namespaces = {"": "https://example.com"}

prefix : str, optional
 Namespace prefix to be used for every element and/or attribute
 in document. This should be one of the keys in ``namespaces``
 dict.
encoding : str, default 'utf-8'
 Encoding of the resulting document.
xml_declaration : bool, default True
 Whether to include the XML declaration at start of document.
pretty_print : bool, default True
 Whether output should be pretty printed with indentation and
 line breaks.
parser : {'lxml', 'etree'}, default 'lxml'
 Parser module to use for building of tree. Only 'lxml' and
 'etree' are supported. With 'lxml', the ability to use XSLT
 stylesheet is supported.
stylesheet : str, path object or file-like object, optional
 A URL, file-like object, or a raw string containing an XSLT
 script used to transform the raw XML output. Script should use
 layout of elements and attributes from original output. This
 argument requires ``lxml`` to be installed. Only XSLT 1.0
 scripts and not later versions is currently supported.
compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default 'infer'
 For on-the-fly decompression of on-disk data. If 'infer', then use
 gzip, bz2, zip or xz if path_or_buffer is a string ending in
 '.gz', '.bz2', '.zip', or '.xz', respectively, and no decompression
 otherwise. If using 'zip', the ZIP file must contain only one data
 file to be read in. Set to None for no decompression.
storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

Returns

None or str
 If ``io`` is None, returns the resulting XML format as a
 string. Otherwise returns None.

See Also

to_json : Convert the pandas object to a JSON string.
to_html : Convert DataFrame to a html.

Examples

>>> df = pd.DataFrame({'shape': ['square', 'circle', 'triangle'],
... 'degrees': [360, 360, 180],
... 'sides': [4, np.nan, 3]})

>>> df.to_xml() # doctest: +SKIP
<?xml version='1.0' encoding='utf-8'?>
<data>
 <row>
 <index>0</index>
 <shape>square</shape>
 <degrees>360</degrees>
 <sides>4.0</sides>
 </row>
 <row>
 <index>1</index>
 <shape>circle</shape>
 <degrees>360</degrees>
 <sides/>
 </row>
 <row>
 <index>2</index>
 <shape>triangle</shape>
 <degrees>180</degrees>
 <sides>3.0</sides>
 </row>
</data>

>>> df.to_xml(attr_cols=[
... 'index', 'shape', 'degrees', 'sides'
...]) # doctest: +SKIP
<?xml version='1.0' encoding='utf-8'?>
<data>
 <row index="0" shape="square" degrees="360" sides="4.0"/>
 <row index="1" shape="circle" degrees="360"/>
 <row index="2" shape="triangle" degrees="180" sides="3.0"/>
</data>

>>> df.to_xml(namespaces={"doc": "https://example.com"},

```

```

... prefix="doc") # doctest: +SKIP
<?xml version='1.0' encoding='utf-8'?>
<doc:data xmlns:doc="https://example.com">
 <doc:row>
 <doc:index>0</doc:index>
 <doc:shape>square</doc:shape>
 <doc:degrees>360</doc:degrees>
 <doc:sides>4.0</doc:sides>
 </doc:row>
 <doc:row>
 <doc:index>1</doc:index>
 <doc:shape>circle</doc:shape>
 <doc:degrees>360</doc:degrees>
 <doc:sides/>
 </doc:row>
 <doc:row>
 <doc:index>2</doc:index>
 <doc:shape>triangle</doc:shape>
 <doc:degrees>180</doc:degrees>
 <doc:sides>3.0</doc:sides>
 </doc:row>
</doc:data>

```

```

transform(self, func: 'AggFuncType', axis: 'Axis' = 0, *args, **kwargs) ->
'DataFrame'

```

Call ``func`` on self producing a DataFrame with transformed values.

Produced DataFrame will have same axis length as self.

#### Parameters

**func** : function, str, list-like or dict-like  
 Function to use for transforming the data. If a function, must either work when passed a DataFrame or when passed to DataFrame.apply. If func is both list-like and dict-like, dict-like behavior takes precedence.

Accepted combinations are:

- function
- string function name
- list-like of functions and/or function names, e.g. ``[np.exp, 'sqrt']``
- dict-like of axis labels -> functions, function names or list-like of

such.

**axis** : {0 or 'index', 1 or 'columns'}, default 0  
 If 0 or 'index': apply function to each column.  
 If 1 or 'columns': apply function to each row.

**\*args**  
 Positional arguments to pass to `func`.

**\*\*kwargs**  
 Keyword arguments to pass to `func`.

#### Returns

**DataFrame**  
 A DataFrame that must have the same length as self.

#### Raises

**ValueError** : If the returned DataFrame has a different length than self.

#### See Also

**DataFrame.agg** : Only perform aggregating type operations.  
**DataFrame.apply** : Invoke function on a DataFrame.

#### Notes

Functions that mutate the passed object can produce unexpected behavior or errors and are not supported. See :ref:`gotchas.udf-mutation` for more details.

#### Examples

```

>>> df = pd.DataFrame({'A': range(3), 'B': range(1, 4)})
>>> df
 A B
0 0 1
1 1 2
2 2 3
>>> df.transform(lambda x: x + 1)
 A B
0 1 2
1 2 3
2 3 4

```

Even though the resulting DataFrame must have the same length as the input DataFrame, it is possible to provide several input functions:

```
>>> s = pd.Series(range(3))
>>> s
0 0
1 1
2 2
dtype: int64
>>> s.transform([np.sqrt, np.exp])
 sqrt exp
0 0.000000 1.000000
1 1.000000 2.718282
2 1.414214 7.389056
```

You can call transform on a GroupBy object:

```
>>> df = pd.DataFrame({
... "Date": [
... "2015-05-08", "2015-05-07", "2015-05-06", "2015-05-05",
... "2015-05-08", "2015-05-07", "2015-05-06", "2015-05-05"],
... "Data": [5, 8, 6, 1, 50, 100, 60, 120],
... })
>>> df
 Date Data
0 2015-05-08 5
1 2015-05-07 8
2 2015-05-06 6
3 2015-05-05 1
4 2015-05-08 50
5 2015-05-07 100
6 2015-05-06 60
7 2015-05-05 120
>>> df.groupby('Date')['Data'].transform('sum')
0 55
1 108
2 66
3 121
4 55
5 108
6 66
7 121
Name: Data, dtype: int64
```

```
>>> df = pd.DataFrame({
... "c": [1, 1, 1, 2, 2, 2, 2],
... "type": ["m", "n", "o", "m", "m", "n", "n"]
... })
>>> df
 c type
0 1 m
1 1 n
2 1 o
3 2 m
4 2 m
5 2 n
6 2 n
>>> df['size'] = df.groupby('c')['type'].transform(len)
>>> df
 c type size
0 1 m 3
1 1 n 3
2 1 o 3
3 2 m 4
4 2 m 4
5 2 n 4
6 2 n 4
```

`transpose(self, *args, copy: 'bool' = False) -> 'DataFrame'`  
Transpose index and columns.

Reflect the DataFrame over its main diagonal by writing rows as columns and vice-versa. The property `:attr:`.T`` is an accessor to the method `:meth:`transpose``.

#### Parameters

`*args` : tuple, optional  
Accepted for compatibility with NumPy.

`copy` : bool, default False  
Whether to copy the data after transposing, even for DataFrames with a single dtype.

Note that a copy is always required for mixed dtype DataFrames, or for DataFrames with any extension types.

#### Returns

-----  
DataFrame

The transposed DataFrame.

See Also

-----

`numpy.transpose` : Permute the dimensions of a given array.

Notes

-----

Transposing a DataFrame with mixed dtypes will result in a homogeneous DataFrame with the `object` dtype. In such a case, a copy of the data is always made.

Examples

-----

**\*\*Square DataFrame with homogeneous dtype\*\***

```
>>> d1 = {'col1': [1, 2], 'col2': [3, 4]}
>>> df1 = pd.DataFrame(data=d1)
>>> df1
 col1 col2
0 1 3
1 2 4

>>> df1_transposed = df1.T # or df1.transpose()
>>> df1_transposed
 0 1
col1 1 2
col2 3 4
```

When the dtype is homogeneous in the original DataFrame, we get a transposed DataFrame with the same dtype:

```
>>> df1.dtypes
col1 int64
col2 int64
dtype: object
>>> df1_transposed.dtypes
0 int64
1 int64
dtype: object
```

**\*\*Non-square DataFrame with mixed dtypes\*\***

```
>>> d2 = {'name': ['Alice', 'Bob'],
... 'score': [9.5, 8],
... 'employed': [False, True],
... 'kids': [0, 0]}
>>> df2 = pd.DataFrame(data=d2)
>>> df2
 name score employed kids
0 Alice 9.5 False 0
1 Bob 8.0 True 0

>>> df2_transposed = df2.T # or df2.transpose()
>>> df2_transposed
 0 1
name Alice Bob
score 9.5 8.0
employed False True
kids 0 0
```

When the DataFrame has mixed dtypes, we get a transposed DataFrame with the `object` dtype:

```
>>> df2.dtypes
name object
score float64
employed bool
kids int64
dtype: object
>>> df2_transposed.dtypes
0 object
1 object
dtype: object
```

`truediv(self, other, axis='columns', level=None, fill_value=None)`

Get Floating division of dataframe and other, element-wise (binary operator ``truediv``).

Equivalent to ``dataframe / other``, but with support to substitute a `fill_value` for missing data in one of the inputs. With reverse version, ``rtruediv``.

Among flexible wrappers (``add``, ``sub``, ``mul``, ``div``, ``mod``, ``pow``) to arithmetic operators: ``+``, ``-``, ``*``, ``/``, ``//``, ``%``, ``**``.

Parameters

```

other : scalar, sequence, Series, or DataFrame
 Any single or multiple element data structure, or list-like object.
axis : {0 or 'index', 1 or 'columns'}
 Whether to compare by the index (0 or 'index') or columns
 (1 or 'columns'). For Series input, axis to match Series index on.
level : int or label
 Broadcast across a level, matching Index values on the
 passed MultiIndex level.
fill_value : float or None, default None
 Fill existing missing (NaN) values, and any new element needed for
 successful DataFrame alignment, with this value before computation.
 If data in both corresponding DataFrame locations is missing
 the result will be missing.

```

#### Returns

```

DataFrame
 Result of the arithmetic operation.

```

#### See Also

```

DataFrame.add : Add DataFrames.
DataFrame.sub : Subtract DataFrames.
DataFrame.mul : Multiply DataFrames.
DataFrame.div : Divide DataFrames (float division).
DataFrame.truediv : Divide DataFrames (float division).
DataFrame.floordiv : Divide DataFrames (integer division).
DataFrame.mod : Calculate modulo (remainder after division).
DataFrame.pow : Calculate exponential power.

```

#### Notes

```

Mismatched indices will be unioned together.

```

#### Examples

```

>>> df = pd.DataFrame({'angles': [0, 3, 4],
... 'degrees': [360, 180, 360]},
... index=['circle', 'triangle', 'rectangle'])
>>> df
 angles degrees
circle 0 360
triangle 3 180
rectangle 4 360

```

Add a scalar with operator version which return the same results.

```

>>> df + 1
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

```

```

>>> df.add(1)
 angles degrees
circle 1 361
triangle 4 181
rectangle 5 361

```

Divide by constant with reverse version.

```

>>> df.div(10)
 angles degrees
circle 0.0 36.0
triangle 0.3 18.0
rectangle 0.4 36.0

```

```

>>> df.rdiv(10)
 angles degrees
circle inf 0.027778
triangle 3.333333 0.055556
rectangle 2.500000 0.027778

```

Subtract a list and Series by axis with operator version.

```

>>> df - [1, 2]
 angles degrees
circle -1 358
triangle 2 178
rectangle 3 358

```

```

>>> df.sub([1, 2], axis='columns')
 angles degrees
circle -1 358
triangle 2 178

```

```
>>> df.sub(pd.Series([1, 1, 1], index=['circle', 'triangle', 'rectangle']),
... axis='index')
 angles degrees
circle -1 359
triangle 2 179
rectangle 3 359
```

```
>>> other = pd.DataFrame({'angles': [0, 3, 4]},
... index=['circle', 'triangle', 'rectangle'])
>>> other
```

	angles
circle	0
triangle	3
rectangle	4

```
>>> df * other
```

	angles	degrees
circle	0	NaN
triangle	9	NaN
rectangle	16	NaN

```
>>> df.mul(other, fill_value=0)
```

	angles	degrees
circle	0	0.0
triangle	9	0.0
rectangle	16	0.0

[illegible]

```
>>> df_multindex
 angles degrees
A circle 0 360
 triangle 3 180
 rectangle 4 360
B square 4 360
 pentagon 5 540
 hexagon 6 720
```

	angles	degrees
A circle	NaN	1.0
triangle	1.0	1.0
rectangle	1.0	1.0
B square	0.0	0.0
pentagon	0.0	0.0
hexagon	0.0	0.0

level

If the index is not a MultiIndex, the output will be a Series (the analogue of stack when the columns are not a MultiIndex).

```

level : int, str, or list of these, default -1 (last level)
 Level(s) of index to unstack, can pass level name.
fill_value : int, str or dict
 Replace NaN with this value if the unstack produces missing values.

```

## Series or DataFrame

```
DataFrame.pivot : Pivot a table based on column values.
DataFrame.stack : Pivot a level of the column labels (inverse operation
 from `unstack`).
```

[illegible]



```
>>> s = pd.Series(np.arange(1.0, 5.0), index=index)
>>> s
one a 1.0
 b 2.0
two a 3.0
 b 4.0
dtype: float64
```

```
>>> s.unstack(level=-1)
 a b
one 1.0 2.0
two 3.0 4.0
```

```
>>> s.unstack(level=0)
 one two
a 1.0 3.0
b 2.0 4.0
```

```
>>> df = s.unstack(level=0)
>>> df.unstack()
one a 1.0
 b 2.0
two a 3.0
 b 4.0
dtype: float64
```

update(self, other, join: 'str' = 'left', overwrite: 'bool' = True, filter\_func=None, errors: 'str' = 'ignore') -> 'None'  
 Modify in place using non-NA values from another DataFrame.

Aligns on indices. There is no return value.

#### Parameters

other : DataFrame, or object coercible into a DataFrame  
 Should have at least one matching index/column label with the original DataFrame. If a Series is passed, its name attribute must be set, and that will be used as the column name to align with the original DataFrame.

join : {'left'}, default 'left'  
 Only left join is implemented, keeping the index and columns of the original object.

overwrite : bool, default True  
 How to handle non-NA values for overlapping keys:

- \* True: overwrite original DataFrame's values with values from 'other'.
- \* False: only update values that are NA in the original DataFrame.

filter\_func : callable(1d-array) -> bool 1d-array, optional  
 Can choose to replace values other than NA. Return True for values that should be updated.

errors : {'raise', 'ignore'}, default 'ignore'  
 If 'raise', will raise a ValueError if the DataFrame and 'other' both contain non-NA data in the same place.

#### Returns

None : method directly changes calling object

#### Raises

ValueError  
 \* When 'errors='raise' and there's overlapping non-NA data.  
 \* When 'errors' is not either 'ignore' or 'raise'

NotImplementedError  
 \* If 'join != 'left'

#### See Also

dict.update : Similar method for dictionaries.  
 DataFrame.merge : For column(s)-on-column(s) operations.

#### Examples

```
>>> df = pd.DataFrame({'A': [1, 2, 3],
... 'B': [400, 500, 600]})
>>> new_df = pd.DataFrame({'B': [4, 5, 6],
... 'C': [7, 8, 9]})
>>> df.update(new_df)
>>> df
 A B
0 1 4
1 2 5
2 3 6
```

The DataFrame's length does not increase as a result of the update, only values at matching index/column labels are updated.

```
>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_df = pd.DataFrame({'B': ['d', 'e', 'f', 'g', 'h', 'i']})
>>> df.update(new_df)
>>> df
 A B
0 a d
1 b e
2 c f
```

For Series, its name attribute must be set.

```
>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_column = pd.Series(['d', 'e'], name='B', index=[0, 2])
>>> df.update(new_column)
>>> df
 A B
0 a d
1 b y
2 c e
>>> df = pd.DataFrame({'A': ['a', 'b', 'c'],
... 'B': ['x', 'y', 'z']})
>>> new_df = pd.DataFrame({'B': ['d', 'e']}, index=[1, 2])
>>> df.update(new_df)
>>> df
 A B
0 a x
1 b d
2 c e
```

If `other` contains NaNs the corresponding values are not updated in the original dataframe.

```
>>> df = pd.DataFrame({'A': [1, 2, 3],
... 'B': [400, 500, 600]})
>>> new_df = pd.DataFrame({'B': [4, np.nan, 6]})
>>> df.update(new_df)
>>> df
 A B
0 1 4.0
1 2 500.0
2 3 6.0
```

```
value_counts(self, subset: 'Sequence[Hashable] | None' = None, normalize: 'bool'
= False, sort: 'bool' = True, ascending: 'bool' = False, dropna: 'bool' = True)
 Return a Series containing counts of unique rows in the DataFrame.
```

.. versionadded:: 1.1.0

Parameters

subset : list-like, optional  
Columns to use when counting unique combinations.  
normalize : bool, default False  
Return proportions rather than frequencies.  
sort : bool, default True  
Sort by frequencies.  
ascending : bool, default False  
Sort in ascending order.  
dropna : bool, default True  
Don't include counts of rows that contain NA values.

.. versionadded:: 1.3.0

Returns

Series

See Also

Series.value\_counts: Equivalent method on Series.

Notes

The returned Series will have a MultiIndex with one level per input column. By default, rows that contain any NA values are omitted from the result. By default, the resulting Series will be in descending order so that the first element is the most frequently-occurring row.

Examples

```
>>> df = pd.DataFrame({'num_legs': [2, 4, 4, 6],
... 'num_wings': [2, 0, 0, 0]},
```

```
... index=['falcon', 'dog', 'cat', 'ant'])
>>> df
 num_legs num_wings
falcon 2 2
dog 4 0
cat 4 0
ant 6 0
```

```
>>> df.value_counts()
num_legs num_wings
4 0 2
2 2 1
6 0 1
dtype: int64
```

```
>>> df.value_counts(sort=False)
num_legs num_wings
2 2 1
4 0 2
6 0 1
dtype: int64
```

```
>>> df.value_counts(ascending=True)
num_legs num_wings
2 2 1
6 0 1
4 0 2
dtype: int64
```

```
>>> df.value_counts(normalize=True)
num_legs num_wings
4 0 0.50
2 2 0.25
6 0 0.25
dtype: float64
```

With `dropna` set to `False` we can also count rows with NA values.

```
>>> df = pd.DataFrame({'first_name': ['John', 'Anne', 'John', 'Beth'],
... 'middle_name': ['Smith', pd.NA, pd.NA, 'Louise']})
>>> df
 first_name middle_name
0 John Smith
1 Anne <NA>
2 John <NA>
3 Beth Louise
```

```
>>> df.value_counts()
first_name middle_name
Beth Louise 1
John Smith 1
dtype: int64
```

```
>>> df.value_counts(dropna=False)
first_name middle_name
Anne NaN 1
Beth Louise 1
John Smith 1
 NaN 1
dtype: int64
```

```
var(self, axis=None, skipna=None, level=None, ddof=1, numeric_only=None,
**kwargs)
```

Return unbiased variance over requested axis.

Normalized by N-1 by default. This can be changed using the ddof argument

Parameters

-----  
axis : {index (0), columns (1)}

skipna : bool, default True

Exclude NA/null values. If an entire row/column is NA, the result will be NA.

level : int or level name, default None

If the axis is a MultiIndex (hierarchical), count along a particular level, collapsing into a Series.

ddof : int, default 1

Delta Degrees of Freedom. The divisor used in calculations is N - ddof, where N represents the number of elements.

numeric\_only : bool, default None

Include only float, int, boolean columns. If None, will attempt to use everything, then use only numeric data. Not implemented for Series.

Returns

-----  
Series or DataFrame (if level specified)

## Notes

-----

To have the same behaviour as `'numpy.std'`, use `'ddof=0'` (instead of the default `'ddof=1'`)

`where(self, cond, other=nan, inplace=False, axis=None, level=None, errors='raise', try_cast=<no_default>)`

Replace values where the condition is False.

## Parameters

-----

`cond` : bool Series/DataFrame, array-like, or callable

Where `'cond'` is True, keep the original value. Where

False, replace with corresponding value from `'other'`.

If `'cond'` is callable, it is computed on the Series/DataFrame and

should return boolean Series/DataFrame or array. The callable must

not change input Series/DataFrame (though pandas doesn't check it).

`other` : scalar, Series/DataFrame, or callable

Entries where `'cond'` is False are replaced with

corresponding value from `'other'`.

If `other` is callable, it is computed on the Series/DataFrame and

should return scalar or Series/DataFrame. The callable must not

change input Series/DataFrame (though pandas doesn't check it).

`inplace` : bool, default False

Whether to perform the operation in place on the data.

`axis` : int, default None

Alignment axis if needed.

`level` : int, default None

Alignment level if needed.

`errors` : str, {'raise', 'ignore'}, default 'raise'

Note that currently this parameter won't affect

the results and will always coerce to a suitable dtype.

- 'raise' : allow exceptions to be raised.

- 'ignore' : suppress exceptions. On error return original object.

`try_cast` : bool, default None

Try to cast the result back to the input type (if possible).

.. deprecated:: 1.3.0

Manually cast back if necessary.

## Returns

-----

Same type as caller or None if `'inplace=True'`.

## See Also

-----

`:func:'DataFrame.mask'` : Return an object of same shape as `self`.

## Notes

-----

The `where` method is an application of the if-then idiom. For each element in the calling DataFrame, if `'cond'` is `'True'` the element is used; otherwise the corresponding element from the DataFrame `'other'` is used.

The signature for `:func:'DataFrame.where'` differs from

`:func:'numpy.where'`. Roughly `'df1.where(m, df2)'` is equivalent to `'np.where(m, df1, df2)'`.

For further details and examples see the `'where'` documentation in `:ref:'indexing <indexing.where_mask>'`.

## Examples

-----

```
>>> s = pd.Series(range(5))
```

```
>>> s.where(s > 0)
```

```
0 NaN
```

```
1 1.0
```

```
2 2.0
```

```
3 3.0
```

```
4 4.0
```

```
dtype: float64
```

```
>>> s.mask(s > 0)
```

```
0 0.0
```

```
1 NaN
```

```
2 NaN
```

```
3 NaN
```

```
4 NaN
```

```
dtype: float64
```

```
>>> s.where(s > 1, 10)
```

```
0 10
```

```
1 10
```

```
2 2
```

```

3 3
4 4
dtype: int64
>>> s.mask(s > 1, 10)
0 0
1 1
2 10
3 10
4 10
dtype: int64

>>> df = pd.DataFrame(np.arange(10).reshape(-1, 2), columns=['A', 'B'])
>>> df
 A B
0 0 1
1 2 3
2 4 5
3 6 7
4 8 9
>>> m = df % 3 == 0
>>> df.where(m, -df)
 A B
0 0 -1
1 -2 3
2 -4 -5
3 6 -7
4 -8 9
>>> df.where(m, -df) == np.where(m, df, -df)
 A B
0 True True
1 True True
2 True True
3 True True
4 True True
>>> df.where(m, -df) == df.mask(~m, -df)
 A B
0 True True
1 True True
2 True True
3 True True
4 True True

```

-----  
Class methods defined here:

`from_dict(data, orient: 'str' = 'columns', dtype: 'Dtype | None' = None, columns=None) -> 'DataFrame' from builtins.type`

Construct DataFrame from dict of array-like or dicts.

Creates DataFrame object from dictionary by columns or by index allowing dtype specification.

Parameters

-----

`data` : dict

Of the form {field : array-like} or {field : dict}.

`orient` : {'columns', 'index'}, default 'columns'

The "orientation" of the data. If the keys of the passed dict should be the columns of the resulting DataFrame, pass 'columns' (default). Otherwise if the keys should be rows, pass 'index'.

`dtype` : dtype, default None

Data type to force, otherwise infer.

`columns` : list, default None

Column labels to use when ``orient='index'``. Raises a ValueError if used with ``orient='columns'``.

Returns

-----

DataFrame

See Also

-----

`DataFrame.from_records` : DataFrame from structured ndarray, sequence of tuples or dicts, or DataFrame.

`DataFrame` : DataFrame object creation using constructor.

Examples

-----

By default the keys of the dict become the DataFrame columns:

```

>>> data = {'col_1': [3, 2, 1, 0], 'col_2': ['a', 'b', 'c', 'd']}
>>> pd.DataFrame.from_dict(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d

```

Specify ``orient='index'`` to create the DataFrame using dictionary keys as rows:

```
>>> data = {'row_1': [3, 2, 1, 0], 'row_2': ['a', 'b', 'c', 'd']}
>>> pd.DataFrame.from_dict(data, orient='index')
 0 1 2 3
row_1 3 2 1 0
row_2 a b c d
```

When using the 'index' orientation, the column names can be specified manually:

```
>>> pd.DataFrame.from_dict(data, orient='index',
... columns=['A', 'B', 'C', 'D'])
 A B C D
row_1 3 2 1 0
row_2 a b c d
```

`from_records(data, index=None, exclude=None, columns=None, coerce_float: 'bool' = False, nrows: 'int | None' = None) -> 'DataFrame'` from `builtins.type`  
Convert structured or record ndarray to DataFrame.

Creates a DataFrame object from a structured ndarray, sequence of tuples or dicts, or DataFrame.

#### Parameters

`data` : structured ndarray, sequence of tuples or dicts, or DataFrame  
Structured input data.  
`index` : str, list of fields, array-like  
Field of array to use as the index, alternately a specific set of input labels to use.  
`exclude` : sequence, default None  
Columns or fields to exclude.  
`columns` : sequence, default None  
Column names to use. If the passed data do not have names associated with them, this argument provides names for the columns. Otherwise this argument indicates the order of the columns in the result (any names not found in the data will become all-NA columns).  
`coerce_float` : bool, default False  
Attempt to convert values of non-string, non-numeric objects (like decimal.Decimal) to floating point, useful for SQL result sets.  
`nrows` : int, default None  
Number of rows to read if data is an iterator.

#### Returns

DataFrame

#### See Also

`DataFrame.from_dict` : DataFrame from dict of array-like or dicts.  
`DataFrame` : DataFrame object creation using constructor.

#### Examples

Data can be provided as a structured ndarray:

```
>>> data = np.array([(3, 'a'), (2, 'b'), (1, 'c'), (0, 'd')],
... dtype=[('col_1', 'i4'), ('col_2', 'U1')])
>>> pd.DataFrame.from_records(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

Data can be provided as a list of dicts:

```
>>> data = [{'col_1': 3, 'col_2': 'a'},
... {'col_1': 2, 'col_2': 'b'},
... {'col_1': 1, 'col_2': 'c'},
... {'col_1': 0, 'col_2': 'd'}]
>>> pd.DataFrame.from_records(data)
 col_1 col_2
0 3 a
1 2 b
2 1 c
3 0 d
```

Data can be provided as a list of tuples with corresponding columns:

```
>>> data = [(3, 'a'), (2, 'b'), (1, 'c'), (0, 'd')]
>>> pd.DataFrame.from_records(data, columns=['col_1', 'col_2'])
 col_1 col_2
```

0	3	a
1	2	b
2	1	c
3	0	d

-----  
Data descriptors defined here:

T

axes

Return a list representing the axes of the DataFrame.

It has the row axis labels and column axis labels as the only members.  
They are returned in that order.

Examples

```

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.axes
[RangeIndex(start=0, stop=2, step=1), Index(['col1', 'col2'],
dtype='object')]

```

columns

The column labels of the DataFrame.

index

The index (row labels) of the DataFrame.

shape

Return a tuple representing the dimensionality of the DataFrame.

See Also

-----  
ndarray.shape : Tuple of array dimensions.

Examples

```

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.shape
(2, 2)

>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4],
... 'col3': [5, 6]})
>>> df.shape
(2, 3)

```

style

Returns a Styler object.

Contains methods for building a styled HTML representation of the DataFrame.

See Also

-----  
io.formats.style.Styler : Helps style a DataFrame or Series according to the data with HTML and CSS.

values

Return a Numpy representation of the DataFrame.

.. warning::

We recommend using :meth:`DataFrame.to\_numpy` instead.

Only the values in the DataFrame will be returned, the axes labels will be removed.

Returns

-----  
numpy.ndarray  
The values of the DataFrame.

See Also

-----  
DataFrame.to\_numpy : Recommended alternative to this method.  
DataFrame.index : Retrieve the index labels.  
DataFrame.columns : Retrieving the column names.

Notes

-----  
The dtype will be a lower-common-denominator dtype (implicit upcasting); that is to say if the dtypes (even of numeric types) are mixed, the one that accommodates all will be chosen. Use this with care if you are not dealing with the blocks.

e.g. If the dtypes are float16 and float32, dtype will be upcast to float32. If dtypes are int32 and uint8, dtype will be upcast to

int32. By :func:`numpy.find\_common\_type` convention, mixing int64 and uint64 will result in a float64 dtype.

#### Examples

A DataFrame where all columns are the same type (e.g., int64) results in an array of the same type.

```
>>> df = pd.DataFrame({'age': [3, 29],
... 'height': [94, 170],
... 'weight': [31, 115]})
>>> df
 age height weight
0 3 94 31
1 29 170 115
>>> df.dtypes
age int64
height int64
weight int64
dtype: object
>>> df.values
array([[3, 94, 31],
 [29, 170, 115]])
```

A DataFrame with mixed type columns (e.g., str/object, int64, float32) results in an ndarray of the broadest type that accommodates these mixed types (e.g., object).

```
>>> df2 = pd.DataFrame([('parrot', 24.0, 'second'),
... ('lion', 80.5, 1),
... ('monkey', np.nan, None)],
... columns=('name', 'max_speed', 'rank'))
>>> df2.dtypes
name object
max_speed float64
rank object
dtype: object
>>> df2.values
array([('parrot', 24.0, 'second'],
 ['lion', 80.5, 1],
 ['monkey', nan, None]], dtype=object)
```

-----  
Data and other attributes defined here:

`__annotations__` = {'\_AXIS\_TO\_AXIS\_NUMBER': 'dict[Axis, int]', '\_access...

`plot` = <class 'pandas.plotting.\_core.PlotAccessor'>  
Make plots of Series or DataFrame.

Uses the backend specified by the option ``plotting.backend``. By default, matplotlib is used.

#### Parameters

-----  
`data` : Series or DataFrame  
The object for which the method is called.  
`x` : label or position, default None  
Only used if data is a DataFrame.  
`y` : label, position or list of label, positions, default None  
Allows plotting of one column versus another. Only used if data is a DataFrame.  
`kind` : str  
The kind of plot to produce:

- 'line' : line plot (default)
- 'bar' : vertical bar plot
- 'barh' : horizontal bar plot
- 'hist' : histogram
- 'box' : boxplot
- 'kde' : Kernel Density Estimation plot
- 'density' : same as 'kde'
- 'area' : area plot
- 'pie' : pie plot
- 'scatter' : scatter plot (DataFrame only)
- 'hexbin' : hexbin plot (DataFrame only)

`ax` : matplotlib axes object, default None  
An axes of the current figure.  
`subplots` : bool, default False  
Make separate subplots for each column.  
`sharex` : bool, default True if ax is None else False  
In case ``subplots=True``, share x axis and set some x axis labels to invisible; defaults to True if ax is None otherwise False if an ax is passed in; Be aware, that passing in both an ax and ``sharex=True`` will alter all x axis labels for all axis in a figure.  
`sharey` : bool, default False  
In case ``subplots=True``, share y axis and set some y axis labels to



invisible.

layout : tuple, optional  
(rows, columns) for the layout of subplots.

figsize : a tuple (width, height) in inches  
Size of a figure object.

use\_index : bool, default True  
Use index as ticks for x axis.

title : str or list  
Title to use for the plot. If a string is passed, print the string at the top of the figure. If a list is passed and 'subplots' is True, print each item in the list above the corresponding subplot.

grid : bool, default None (matlab style default)  
Axis grid lines.

legend : bool or {'reverse'}  
Place legend on axis subplots.

style : list or dict  
The matplotlib line style per column.

logx : bool or 'sym', default False  
Use log scaling or symlog scaling on x axis.  
.. versionchanged:: 0.25.0

logy : bool or 'sym' default False  
Use log scaling or symlog scaling on y axis.  
.. versionchanged:: 0.25.0

loglog : bool or 'sym', default False  
Use log scaling or symlog scaling on both x and y axes.  
.. versionchanged:: 0.25.0

xticks : sequence  
Values to use for the xticks.

yticks : sequence  
Values to use for the yticks.

xlim : 2-tuple/list  
Set the x limits of the current axes.

ylim : 2-tuple/list  
Set the y limits of the current axes.

xlabel : label, optional  
Name to use for the xlabel on x-axis. Default uses index name as xlabel,

or the  
x-column name for planar plots.  
.. versionadded:: 1.1.0  
.. versionchanged:: 1.2.0  
Now applicable to planar plots ('scatter', 'hexbin').

ylabel : label, optional  
Name to use for the ylabel on y-axis. Default will show no ylabel, or the y-column name for planar plots.  
.. versionadded:: 1.1.0  
.. versionchanged:: 1.2.0  
Now applicable to planar plots ('scatter', 'hexbin').

rot : int, default None  
Rotation for ticks (xticks for vertical, yticks for horizontal plots).

fontsize : int, default None  
Font size for xticks and yticks.

colormap : str or matplotlib colormap object, default None  
Colormap to select colors from. If string, load colormap with that name from matplotlib.

colorbar : bool, optional  
If True, plot colorbar (only relevant for 'scatter' and 'hexbin' plots).

position : float  
Specify relative alignments for bar plot layout.  
From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5 (center).

table : bool, Series or DataFrame, default False  
If True, draw a table using the data in the DataFrame and the data will be transposed to meet matplotlib's default layout.  
If a Series or DataFrame is passed, use passed data to draw a table.

yerr : DataFrame, Series, array-like, dict and str  
See :ref:`Plotting with Error Bars <visualization.errorbars>` for detail.

xerr : DataFrame, Series, array-like, dict and str  
Equivalent to yerr.

stacked : bool, default False in line and bar plots, and True in area plot  
If True, create stacked plot.

sort\_columns : bool, default False  
Sort column names to determine plot ordering.

```

secondary_y : bool or sequence, default False
 Whether to plot on the secondary y-axis if a list/tuple, which
 columns to plot on secondary y-axis.
mark_right : bool, default True
 When using a secondary_y axis, automatically mark the column
 labels with "(right)" in the legend.
include_bool : bool, default is False
 If True, boolean values can be plotted.
backend : str, default None
 Backend to use instead of the backend specified in the option
 ``plotting.backend``. For instance, 'matplotlib'. Alternatively, to
 specify the ``plotting.backend`` for the whole session, set
 ``pd.options.plotting.backend``.

.. versionadded:: 1.0.0

**kwargs
 Options to pass to matplotlib plotting method.

Returns

:class:`matplotlib.axes.Axes` or numpy.ndarray of them
 If the backend is not the default matplotlib one, the return value
 will be the object returned by the backend.

Notes

- See matplotlib documentation online for more on this subject
- If `kind` = 'bar' or 'barh', you can specify relative alignments
 for bar plot layout by `position` keyword.
 From 0 (left/bottom-end) to 1 (right/top-end). Default is 0.5
 (center)

sparse = <class 'pandas.core.arrays.sparse.accessor.SparseFrameAccesso...
 DataFrame accessor for sparse data.

.. versionadded:: 0.25.0

Methods inherited from pandas.core.generic.NDFrame:

__abs__(self: 'FrameOrSeries') -> 'FrameOrSeries'
__array__(self, dtype: 'NpDtype | None' = None) -> 'np.ndarray'
__array_ufunc__(self, ufunc: 'np.ufunc', method: 'str', *inputs: 'Any', **kwargs:
'Any')
__array_wrap__(self, result: 'np.ndarray', context: 'tuple[Callable, tuple[Any,
...], int] | None' = None)
 Gets called after a ufunc and other functions.

Parameters

result: np.ndarray
 The result of the ufunc or other function called on the NumPy array
 returned by __array__
context: tuple of (func, tuple, int)
 This parameter is returned by ufuncs as a 3-element tuple: (name of the
 ufunc, arguments of the ufunc, domain of the ufunc), but is not set by
 other numpy functions.q

Notes

Series implements __array_ufunc__ so this not called for ufunc on Series.

__bool__ = __nonzero__(self)
__contains__(self, key) -> 'bool_t'
 True if the key is in the info axis
__copy__(self: 'FrameOrSeries', deep: 'bool_t' = True) -> 'FrameOrSeries'
__deepcopy__(self: 'FrameOrSeries', memo=None) -> 'FrameOrSeries'
Parameters

memo, default None
 Standard signature. Unused

__delitem__(self, key) -> 'None'
 Delete item

__finalize__(self: 'FrameOrSeries', other, method: 'str | None' = None, **kwargs)
-> 'FrameOrSeries'
 Propagate metadata from other to self.

Parameters

```

```

other : the object from which to get the attributes that we are going
 to propagate
method : str, optional
 A passed method name providing context on where ``__finalize__``
 was called.

.. warning::

 The value passed as `method` are not currently considered
 stable across pandas releases.

__getattr__(self, name: 'str')
 After regular attribute access, try looking up the name
 This allows simpler access to columns for interactive use.

__getstate__(self) -> 'dict[str, Any]'

__iadd__(self, other)

__iand__(self, other)

__ifloordiv__(self, other)

__imod__(self, other)

__imul__(self, other)

__invert__(self)

__ior__(self, other)

__ipow__(self, other)

__isub__(self, other)

__iter__(self)
 Iterate over info axis.

 Returns

 iterator
 Info axis as iterator.

__itruediv__(self, other)

__ixor__(self, other)

__neg__(self)

__nonzero__(self)

__pos__(self)

__round__(self: 'FrameOrSeries', decimals: 'int' = 0) -> 'FrameOrSeries'

__setattr__(self, name: 'str', value) -> 'None'
 After regular attribute access, try setting the name
 This allows simpler access to columns for interactive use.

__setstate__(self, state)

abs(self: 'FrameOrSeries') -> 'FrameOrSeries'
 Return a Series/DataFrame with absolute numeric value of each element.

 This function only applies to elements that are all numeric.

 Returns

 abs
 Series/DataFrame containing the absolute value of each element.

 See Also

 numpy.absolute : Calculate the absolute value element-wise.

 Notes

 For ``complex`` inputs, ``1.2 + 1j``, the absolute value is
 :math:\sqrt{a^2 + b^2}``.

 Examples

 Absolute numeric values in a Series.

 >>> s = pd.Series([-1.10, 2, -3.33, 4])
 >>> s.abs()

```

```

0 1.10
1 2.00
2 3.33
3 4.00
dtype: float64

```

Absolute numeric values in a Series with complex numbers.

```

>>> s = pd.Series([1.2 + 1j])
>>> s.abs()
0 1.56205
dtype: float64

```

Absolute numeric values in a Series with a Timedelta element.

```

>>> s = pd.Series([pd.Timedelta('1 days')])
>>> s.abs()
0 1 days
dtype: timedelta64[ns]

```

Select rows with data closest to certain value using argsort (from [StackOverflow <https://stackoverflow.com/a/17758115>](https://stackoverflow.com/a/17758115)).

```

>>> df = pd.DataFrame({
... 'a': [4, 5, 6, 7],
... 'b': [10, 20, 30, 40],
... 'c': [100, 50, -30, -50]
... })
>>> df
 a b c
0 4 10 100
1 5 20 50
2 6 30 -30
3 7 40 -50
>>> df.loc[(df.c - 43).abs().argsort()]
 a b c
1 5 20 50
0 4 10 100
2 6 30 -30
3 7 40 -50

```

`add_prefix(self: 'FrameOrSeries', prefix: 'str') -> 'FrameOrSeries'`  
Prefix labels with string ``prefix``.

For Series, the row labels are prefixed.  
For DataFrame, the column labels are prefixed.

Parameters

`prefix : str`  
The string to add before each label.

Returns

Series or DataFrame  
New Series or DataFrame with updated labels.

See Also

`Series.add_suffix`: Suffix row labels with string ``suffix``.  
`DataFrame.add_suffix`: Suffix column labels with string ``suffix``.

Examples

```

>>> s = pd.Series([1, 2, 3, 4])
>>> s
0 1
1 2
2 3
3 4
dtype: int64

```

```

>>> s.add_prefix('item_')
item_0 1
item_1 2
item_2 3
item_3 4
dtype: int64

```

```

>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
>>> df
 A B
0 1 3
1 2 4
2 3 5
3 4 6

```

```
>>> df.add_prefix('col_')
 col_A col_B
0 1 3
1 2 4
2 3 5
3 4 6
```

`add_suffix(self: 'FrameOrSeries', suffix: 'str') -> 'FrameOrSeries'`  
 Suffix labels with string `suffix`.

For Series, the row labels are suffixed.  
 For DataFrame, the column labels are suffixed.

Parameters

-----  
 suffix : str  
 The string to add after each label.

Returns

-----  
 Series or DataFrame  
 New Series or DataFrame with updated labels.

See Also

-----  
 Series.add\_prefix: Prefix row labels with string `prefix`.  
 DataFrame.add\_prefix: Prefix column labels with string `prefix`.

Examples

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s
0 1
1 2
2 3
3 4
dtype: int64
```

```
>>> s.add_suffix('_item')
0_item 1
1_item 2
2_item 3
3_item 4
dtype: int64
```

```
>>> df = pd.DataFrame({'A': [1, 2, 3, 4], 'B': [3, 4, 5, 6]})
>>> df
 A B
0 1 3
1 2 4
2 3 5
3 4 6
```

```
>>> df.add_suffix('_col')
 A_col B_col
0 1 3
1 2 4
2 3 5
3 4 6
```

`asof(self, where, subset=None)`  
 Return the last row(s) without any NaNs before `where`.

The last row (for each element in `where`, if list) without any NaN is taken.  
 In case of a :class:`~pandas.DataFrame`, the last row without NaN considering only the subset of columns (if not `None`)

If there is no good value, NaN is returned for a Series or a Series of NaN values for a DataFrame

Parameters

-----  
 where : date or array-like of dates  
 Date(s) before which the last row(s) are returned.  
 subset : str or array-like of str, default `None`  
 For DataFrame, if not `None`, only use these columns to check for NaNs.

Returns

-----  
 scalar, Series, or DataFrame

The return can be:

- \* scalar : when `self` is a Series and `where` is a scalar
- \* Series: when `self` is a Series and `where` is an array-like,

or when `self` is a DataFrame and `where` is a scalar  
\* DataFrame : when `self` is a DataFrame and `where` is an array-like

Return scalar, Series, or DataFrame.

See Also

merge\_asof : Perform an asof merge. Similar to left join.

Notes

Dates are assumed to be sorted. Raises if this is not the case.

Examples

A Series and a scalar `where`.

```
>>> s = pd.Series([1, 2, np.nan, 4], index=[10, 20, 30, 40])
>>> s
10 1.0
20 2.0
30 NaN
40 4.0
dtype: float64
```

```
>>> s.asof(20)
2.0
```

For a sequence `where`, a Series is returned. The first value is NaN, because the first element of `where` is before the first index value.

```
>>> s.asof([5, 20])
5 NaN
20 2.0
dtype: float64
```

Missing values are not considered. The following is ``2.0``, not NaN, even though NaN is at the index location for ``30``.

```
>>> s.asof(30)
2.0
```

Take all columns into consideration

```
>>> df = pd.DataFrame({'a': [10, 20, 30, 40, 50],
... 'b': [None, None, None, None, 500]},
... index=pd.DatetimeIndex(['2018-02-27 09:01:00',
... '2018-02-27 09:02:00',
... '2018-02-27 09:03:00',
... '2018-02-27 09:04:00',
... '2018-02-27 09:05:00']))
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
... '2018-02-27 09:04:30']))
...
 a b
2018-02-27 09:03:30 NaN NaN
2018-02-27 09:04:30 NaN NaN
```

Take a single column into consideration

```
>>> df.asof(pd.DatetimeIndex(['2018-02-27 09:03:30',
... '2018-02-27 09:04:30']),
... subset=['a'])
...
 a b
2018-02-27 09:03:30 30.0 NaN
2018-02-27 09:04:30 40.0 NaN
```

```
astype(self: 'FrameOrSeries', dtype, copy: 'bool_t' = True, errors: 'str' =
'raise') -> 'FrameOrSeries'
Cast a pandas object to a specified dtype ``dtype``.
```

Parameters

dtype : data type, or dict of column name -> data type

Use a numpy.dtype or Python type to cast entire pandas object to the same type. Alternatively, use {col: dtype, ...}, where col is a column label and dtype is a numpy.dtype or Python type to cast one or more of the DataFrame's columns to column-specific types.

copy : bool, default True

Return a copy when ``copy=True`` (be very careful setting ``copy=False`` as changes to values then may propagate to other pandas objects).

errors : {'raise', 'ignore'}, default 'raise'

Control raising of exceptions on invalid data for provided dtype.

- ``raise`` : allow exceptions to be raised

- ``ignore`` : suppress exceptions. On error return original object.

#### Returns

-----  
casted : same type as caller

#### See Also

-----  
to\_datetime : Convert argument to datetime.  
to\_timedelta : Convert argument to timedelta.  
to\_numeric : Convert argument to a numeric type.  
numpy.ndarray.astype : Cast a numpy array to a specified type.

#### Notes

-----  
.. deprecated:: 1.3.0  
  
Using ``astype`` to convert from timezone-naive dtype to  
timezone-aware dtype is deprecated and will raise in a  
future version. Use :meth:`Series.dt.tz\_localize` instead.

#### Examples

-----  
Create a DataFrame:

```
>>> d = {'col1': [1, 2], 'col2': [3, 4]}
>>> df = pd.DataFrame(data=d)
>>> df.dtypes
col1 int64
col2 int64
dtype: object
```

Cast all columns to int32:

```
>>> df.astype('int32').dtypes
col1 int32
col2 int32
dtype: object
```

Cast col1 to int32 using a dictionary:

```
>>> df.astype({'col1': 'int32'}).dtypes
col1 int32
col2 int64
dtype: object
```

Create a series:

```
>>> ser = pd.Series([1, 2], dtype='int32')
>>> ser
0 1
1 2
dtype: int32
>>> ser.astype('int64')
0 1
1 2
dtype: int64
```

Convert to categorical type:

```
>>> ser.astype('category')
0 1
1 2
dtype: category
Categories (2, int64): [1, 2]
```

Convert to ordered categorical type with custom ordering:

```
>>> from pandas.api.types import CategoricalDtype
>>> cat_dtype = CategoricalDtype(
... categories=[2, 1], ordered=True)
>>> ser.astype(cat_dtype)
0 1
1 2
dtype: category
Categories (2, int64): [2 < 1]
```

Note that using ``copy=False`` and changing data on a new  
pandas object may propagate changes:

```
>>> s1 = pd.Series([1, 2])
>>> s2 = s1.astype('int64', copy=False)
>>> s2[0] = 10
>>> s1 # note that s1[0] has changed too
0 10
1 2
dtype: int64
```

```

Create a series of dates:

>>> ser_date = pd.Series(pd.date_range('20200101', periods=3))
>>> ser_date
0 2020-01-01
1 2020-01-02
2 2020-01-03
dtype: datetime64[ns]

at_time(self: 'FrameOrSeries', time, asof: 'bool_t' = False, axis=None) ->
'FrameOrSeries'
 Select values at particular time of day (e.g., 9:30AM).

Parameters

time : datetime.time or str
axis : {0 or 'index', 1 or 'columns'}, default 0

Returns

Series or DataFrame

Raises

TypeError
 If the index is not a :class:`DatetimeIndex`

See Also

between_time : Select values between particular times of the day.
first : Select initial periods of time series based on a date offset.
last : Select final periods of time series based on a date offset.
DatetimeIndex.indexer_at_time : Get just the index locations for
 values at particular time of the day.

Examples

>>> i = pd.date_range('2018-04-09', periods=4, freq='12H')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
 A
2018-04-09 00:00:00 1
2018-04-09 12:00:00 2
2018-04-10 00:00:00 3
2018-04-10 12:00:00 4

>>> ts.at_time('12:00')
 A
2018-04-09 12:00:00 2
2018-04-10 12:00:00 4

backfill = bfill(self: 'FrameOrSeries', axis: 'None | Axis' = None, inplace:
'bool_t' = False, limit: 'None | int' = None, downcast=None) -> 'FrameOrSeries |
None'
 Synonym for :meth:`DataFrame.fillna` with ``method='bfill'``.

Returns

Series/DataFrame or None
 Object with missing values filled or None if ``inplace=True``.

between_time(self: 'FrameOrSeries', start_time, end_time, include_start: 'bool_t'
= True, include_end: 'bool_t' = True, axis=None) -> 'FrameOrSeries'
 Select values between particular times of the day (e.g., 9:00-9:30 AM).

By setting ``start_time`` to be later than ``end_time``,
you can get the times that are *not* between the two times.

Parameters

start_time : datetime.time or str
 Initial time as a time filter limit.
end_time : datetime.time or str
 End time as a time filter limit.
include_start : bool, default True
 Whether the start time needs to be included in the result.
include_end : bool, default True
 Whether the end time needs to be included in the result.
axis : {0 or 'index', 1 or 'columns'}, default 0
 Determine range time on index or columns value.

Returns

Series or DataFrame
 Data from the original object filtered to the specified dates range.

```



```

Raises

TypeError
 If the index is not a :class:`DatetimeIndex`

See Also

at_time : Select values at a particular time of the day.
first : Select initial periods of time series based on a date offset.
last : Select final periods of time series based on a date offset.
DatetimeIndex.indexer_between_time : Get just the index locations for
 values between particular times of the day.

Examples

>>> i = pd.date_range('2018-04-09', periods=4, freq='1D20min')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
 A
2018-04-09 00:00:00 1
2018-04-10 00:20:00 2
2018-04-11 00:40:00 3
2018-04-12 01:00:00 4

>>> ts.between_time('0:15', '0:45')
 A
2018-04-10 00:20:00 2
2018-04-11 00:40:00 3

You get the times that are *not* between two times by setting
``start_time`` later than ``end_time``:

>>> ts.between_time('0:45', '0:15')
 A
2018-04-09 00:00:00 1
2018-04-12 01:00:00 4

bool(self)
Return the bool of a single element Series or DataFrame.

This must be a boolean scalar value, either True or False. It will raise a
ValueError if the Series or DataFrame does not have exactly 1 element, or
that element is not boolean (integer values 0 and 1 will also raise an exception).

Returns

bool
 The value in the Series or DataFrame.

See Also

Series.astype : Change the data type of a Series, including to boolean.
DataFrame.astype : Change the data type of a DataFrame, including to boolean.
numpy.bool_ : NumPy boolean data type, used by pandas for boolean values.

Examples

The method will only work for single element objects with a boolean value:

>>> pd.Series([True]).bool()
True
>>> pd.Series([False]).bool()
False

>>> pd.DataFrame({'col': [True]}).bool()
True
>>> pd.DataFrame({'col': [False]}).bool()
False

convert_dtypes(self: 'FrameOrSeries', infer_objects: 'bool_t' = True,
convert_string: 'bool_t' = True, convert_integer: 'bool_t' = True, convert_boolean:
'bool_t' = True, convert_floating: 'bool_t' = True) -> 'FrameOrSeries'
 Convert columns to best possible dtypes using dtypes supporting ``pd.NA``.

 .. versionadded:: 1.0.0

Parameters

infer_objects : bool, default True
 Whether object dtypes should be converted to the best possible types.
convert_string : bool, default True
 Whether object dtypes should be converted to ``StringDtype()``.
convert_integer : bool, default True
 Whether, if possible, conversion can be done to integer extension types.
convert_boolean : bool, defaults True
 Whether object dtypes should be converted to ``BooleanDtypes()``.

```

`convert_floating` : bool, defaults True  
Whether, if possible, conversion can be done to floating extension types.  
If `'convert_integer'` is also True, preference will be given to integer dtypes if the floats can be faithfully casted to integers.

.. versionadded:: 1.2.0

Returns

-----

Series or DataFrame

Copy of input object with new dtype.

See Also

-----

`infer_objects` : Infer dtypes of objects.

`to_datetime` : Convert argument to datetime.

`to_timedelta` : Convert argument to timedelta.

`to_numeric` : Convert argument to a numeric type.

Notes

-----

By default, `'convert_dtypes'` will attempt to convert a Series (or each Series in a DataFrame) to dtypes that support `'pd.NA'`. By using the options `'convert_string'`, `'convert_integer'`, `'convert_boolean'` and `'convert_floating'`, it is possible to turn off individual conversions to `'StringDtype'`, the integer extension types, `'BooleanDtype'` or floating extension types, respectively.

For object-dtyped columns, if `'infer_objects'` is `'True'`, use the inference

rules as during normal Series/DataFrame construction. Then, if possible, convert to `'StringDtype'`, `'BooleanDtype'` or an appropriate integer or floating extension type, otherwise leave as `'object'`.

If the dtype is integer, convert to an appropriate integer extension type.

If the dtype is numeric, and consists of all integers, convert to an appropriate integer extension type. Otherwise, convert to an appropriate floating extension type.

.. versionchanged:: 1.2

Starting with pandas 1.2, this method also converts float columns to the nullable floating extension type.

In the future, as new dtypes are added that support `'pd.NA'`, the results of this method will change to support those new dtypes.

Examples

-----

```
>>> df = pd.DataFrame(
... {
... "a": pd.Series([1, 2, 3], dtype=np.dtype("int32")),
... "b": pd.Series(["x", "y", "z"], dtype=np.dtype("O")),
... "c": pd.Series([True, False, np.nan], dtype=np.dtype("O")),
... "d": pd.Series(["h", "i", np.nan], dtype=np.dtype("O")),
... "e": pd.Series([10, np.nan, 20], dtype=np.dtype("float")),
... "f": pd.Series([np.nan, 100.5, 200], dtype=np.dtype("float")),
... }
...)
```

Start with a DataFrame with default dtypes.

```
>>> df
 a b c d e f
0 1 x True h 10.0 NaN
1 2 y False i NaN 100.5
2 3 z NaN NaN 20.0 200.0
```

```
>>> df.dtypes
```

```
a int32
b object
c object
d object
e float64
f float64
dtype: object
```

Convert the DataFrame to use best possible dtypes.

```
>>> dfn = df.convert_dtypes()
```

```
>>> dfn
```

```
 a b c d e f
0 1 x True h 10 <NA>
1 2 y False i <NA> 100.5
2 3 z <NA> <NA> 20 200.0
```

```
>>> dfn.dtypes
```

```
a Int32
b string
c boolean
d string
e Int64
f Float64
dtype: object
```

Start with a Series of strings and missing data represented by ``np.nan``.

```
>>> s = pd.Series(["a", "b", np.nan])
>>> s
0 a
1 b
2 NaN
dtype: object
```

Obtain a Series with dtype ``StringDtype``.

```
>>> s.convert_dtypes()
0 a
1 b
2 <NA>
dtype: string
```

`copy(self: 'FrameOrSeries', deep: 'bool_t' = True) -> 'FrameOrSeries'`  
Make a copy of this object's indices and data.

When ``deep=True`` (default), a new object will be created with a copy of the calling object's data and indices. Modifications to the data or indices of the copy will not be reflected in the original object (see notes below).

When ``deep=False``, a new object will be created without copying the calling object's data or index (only references to the data and index are copied). Any changes to the data of the original will be reflected in the shallow copy (and vice versa).

#### Parameters

`deep` : bool, default True  
Make a deep copy, including a copy of the data and the indices.  
With ``deep=False`` neither the indices nor the data are copied.

#### Returns

`copy` : Series or DataFrame  
Object type matches caller.

#### Notes

When ``deep=True``, data is copied but actual Python objects will not be copied recursively, only the reference to the object. This is in contrast to `copy.deepcopy` in the Standard Library, which recursively copies object data (see examples below).

While ``Index`` objects are copied when ``deep=True``, the underlying numpy array is not copied for performance reasons. Since ``Index`` is immutable, the underlying data can be safely shared and a copy is not needed.

#### Examples

```
>>> s = pd.Series([1, 2], index=["a", "b"])
>>> s
a 1
b 2
dtype: int64
```

```
>>> s_copy = s.copy()
>>> s_copy
a 1
b 2
dtype: int64
```

**\*\*Shallow copy versus default (deep) copy:\*\***

```
>>> s = pd.Series([1, 2], index=["a", "b"])
>>> deep = s.copy()
>>> shallow = s.copy(deep=False)
```

Shallow copy shares data and index with original.

```
>>> s is shallow
False
>>> s.values is shallow.values and s.index is shallow.index
True
```

Deep copy has own copy of data and index.

```
>>> s is deep
False
>>> s.values is deep.values or s.index is deep.index
False
```

Updates to the data shared by shallow copy and original is reflected in both; deep copy remains unchanged.

```
>>> s[0] = 3
>>> shallow[1] = 4
>>> s
a 3
b 4
dtype: int64
>>> shallow
a 3
b 4
dtype: int64
>>> deep
a 1
b 2
dtype: int64
```

Note that when copying an object containing Python objects, a deep copy will copy the data, but will not do so recursively. Updating a nested data object will be reflected in the deep copy.

```
>>> s = pd.Series([[1, 2], [3, 4]])
>>> deep = s.copy()
>>> s[0][0] = 10
>>> s
0 [10, 2]
1 [3, 4]
dtype: object
>>> deep
0 [10, 2]
1 [3, 4]
dtype: object
```

```
describe(self: 'FrameOrSeries', percentiles=None, include=None, exclude=None,
datetime_is_numeric=False) -> 'FrameOrSeries'
Generate descriptive statistics.
```

Descriptive statistics include those that summarize the central tendency, dispersion and shape of a dataset's distribution, excluding ``NaN`` values.

Analyzes both numeric and object series, as well as ``DataFrame`` column sets of mixed data types. The output will vary depending on what is provided. Refer to the notes below for more detail.

#### Parameters

percentiles : list-like of numbers, optional

The percentiles to include in the output. All should fall between 0 and 1. The default is ``[.25, .5, .75]``, which returns the 25th, 50th, and 75th percentiles.

include : 'all', list-like of dtypes or None (default), optional

A white list of data types to include in the result. Ignored for ``Series``. Here are the options:

- 'all' : All columns of the input will be included in the output.
- A list-like of dtypes : Limits the results to the provided data types.  
To limit the result to numeric types submit ``numpy.number``. To limit it instead to object columns submit the ``numpy.object`` data type. Strings can also be used in the style of ``select\_dtypes`` (e.g. ``df.describe(include=['O'])``). To select pandas categorical columns, use ``category``
- None (default) : The result will include all numeric columns.

exclude : list-like of dtypes or None (default), optional,  
A black list of data types to omit from the result. Ignored for ``Series``. Here are the options:

- A list-like of dtypes : Excludes the provided data types from the result. To exclude numeric types submit ``numpy.number``. To exclude object columns submit the data type ``numpy.object``. Strings can also be used in the style of ``select\_dtypes`` (e.g. ``df.describe(include=['O'])``). To exclude pandas categorical columns, use ``category``
- None (default) : The result will exclude nothing.

`datetime_is_numeric` : bool, default False  
Whether to treat datetime dtypes as numeric. This affects statistics calculated for the column. For DataFrame input, this also controls whether datetime columns are included by default.

.. versionadded:: 1.1.0

#### Returns

-----

Series or DataFrame

Summary statistics of the Series or Dataframe provided.

#### See Also

-----

`DataFrame.count`: Count number of non-NA/null observations.

`DataFrame.max`: Maximum of the values in the object.

`DataFrame.min`: Minimum of the values in the object.

`DataFrame.mean`: Mean of the values.

`DataFrame.std`: Standard deviation of the observations.

`DataFrame.select_dtypes`: Subset of a DataFrame including/excluding columns based on their dtype.

#### Notes

-----

For numeric data, the result's index will include `count`, `mean`, `std`, `min`, `max` as well as lower, `50` and upper percentiles. By default the lower percentile is `25` and the upper percentile is `75`. The `50` percentile is the same as the median.

For object data (e.g. strings or timestamps), the result's index will include `count`, `unique`, `top`, and `freq`. The `top` is the most common value. The `freq` is the most common value's frequency. Timestamps also include the `first` and `last` items.

If multiple object values have the highest count, then the `count` and `top` results will be arbitrarily chosen from among those with the highest count.

For mixed data types provided via a `DataFrame`, the default is to return only an analysis of numeric columns. If the dataframe consists only of object and categorical data without any numeric columns, the default is to return an analysis of both the object and categorical columns. If `include='all'` is provided as an option, the result will include a union of attributes of each type.

The `include` and `exclude` parameters can be used to limit which columns in a `DataFrame` are analyzed for the output. The parameters are ignored when analyzing a `Series`.

#### Examples

-----

Describing a numeric `Series`.

```
>>> s = pd.Series([1, 2, 3])
>>> s.describe()
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
dtype: float64
```

Describing a categorical `Series`.

```
>>> s = pd.Series(['a', 'a', 'b', 'c'])
>>> s.describe()
count 4
unique 3
top a
freq 2
dtype: object
```

Describing a timestamp `Series`.

```
>>> s = pd.Series([
... np.datetime64("2000-01-01"),
... np.datetime64("2010-01-01"),
... np.datetime64("2010-01-01")
...])
>>> s.describe(datetime_is_numeric=True)
count 3
mean 2006-09-01 08:00:00
min 2000-01-01 00:00:00
```

```
25% 2004-12-31 12:00:00
50% 2010-01-01 00:00:00
75% 2010-01-01 00:00:00
max 2010-01-01 00:00:00
dtype: object
```

Describing a ``DataFrame``. By default only numeric fields are returned.

```
>>> df = pd.DataFrame({'categorical': pd.Categorical(['d','e','f']),
... 'numeric': [1, 2, 3],
... 'object': ['a', 'b', 'c']}
... })
>>> df.describe()
 numeric
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
```

Describing all columns of a ``DataFrame`` regardless of data type.

```
>>> df.describe(include='all') # doctest: +SKIP
 categorical numeric object
count 3 3.0 3
unique 3 NaN 3
top f NaN a
freq 1 NaN 1
mean NaN 2.0 NaN
std NaN 1.0 NaN
min NaN 1.0 NaN
25% NaN 1.5 NaN
50% NaN 2.0 NaN
75% NaN 2.5 NaN
max NaN 3.0 NaN
```

Describing a column from a ``DataFrame`` by accessing it as an attribute.

```
>>> df.numeric.describe()
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
Name: numeric, dtype: float64
```

Including only numeric columns in a ``DataFrame`` description.

```
>>> df.describe(include=[np.number])
 numeric
count 3.0
mean 2.0
std 1.0
min 1.0
25% 1.5
50% 2.0
75% 2.5
max 3.0
```

Including only string columns in a ``DataFrame`` description.

```
>>> df.describe(include=[object]) # doctest: +SKIP
 object
count 3
unique 3
top a
freq 1
```

Including only categorical columns from a ``DataFrame`` description.

```
>>> df.describe(include=['category'])
 categorical
count 3
unique 3
top d
freq 1
```

Excluding numeric columns from a ``DataFrame`` description.

```
>>> df.describe(exclude=[np.number]) # doctest: +SKIP
categorical object
count 3 3
unique 3 3
top f a
freq 1 1
```

Excluding object columns from a ``DataFrame`` description.

```
>>> df.describe(exclude=[object]) # doctest: +SKIP
categorical numeric
count 3 3.0
unique 3 NaN
top f NaN
freq 1 NaN
mean NaN 2.0
std NaN 1.0
min NaN 1.0
25% NaN 1.5
50% NaN 2.0
75% NaN 2.5
max NaN 3.0
```

```
droplevel(self: 'FrameOrSeries', level, axis=0) -> 'FrameOrSeries'
Return Series/DataFrame with requested index / column level(s) removed.
```

Parameters

```

level : int, str, or list-like
 If a string is given, must be the name of a level
 If list-like, elements must be names or positional indexes
 of levels.
```

```
axis : {0 or 'index', 1 or 'columns'}, default 0
 Axis along which the level(s) is removed:
```

```
 * 0 or 'index': remove level(s) in column.
 * 1 or 'columns': remove level(s) in row.
```

Returns

```

Series/DataFrame
 Series/DataFrame with requested index / column level(s) removed.
```

Examples

```

>>> df = pd.DataFrame([
... [1, 2, 3, 4],
... [5, 6, 7, 8],
... [9, 10, 11, 12]
...]).set_index([0, 1]).rename_axis(['a', 'b'])
```

```
>>> df.columns = pd.MultiIndex.from_tuples([
... ('c', 'e'), ('d', 'f')
...], names=['level_1', 'level_2'])
```

```
>>> df
level_1 c d
level_2 e f
a b
1 2 3 4
5 6 7 8
9 10 11 12
```

```
>>> df.droplevel('a')
level_1 c d
level_2 e f
b
2 3 4
6 7 8
10 11 12
```

```
>>> df.droplevel('level_2', axis=1)
level_1 c d
a b
1 2 3 4
5 6 7 8
9 10 11 12
```

```
equals(self, other: 'object') -> 'bool_t'
Test whether two objects contain the same elements.
```

This function allows two Series or DataFrames to be compared against each other to see if they have the same shape and elements. NaNs in the same location are considered equal.

The row/column index do not need to have the same type, as long

as the values are considered equal. Corresponding columns must be of the same dtype.

#### Parameters

**other** : Series or DataFrame

The other Series or DataFrame to be compared with the first.

#### Returns

bool

True if all elements are the same in both objects, False otherwise.

#### See Also

**Series.eq** : Compare two Series objects of the same length and return a Series where each element is True if the element in each Series is equal, False otherwise.

**DataFrame.eq** : Compare two DataFrame objects of the same shape and return a DataFrame where each element is True if the respective element in each DataFrame is equal, False otherwise.

**testing.assert\_series\_equal** : Raises an AssertionError if left and right are not equal. Provides an easy interface to ignore inequality in dtypes, indexes and precision among others.

**testing.assert\_frame\_equal** : Like assert\_series\_equal, but targets DataFrames.

**numpy.array\_equal** : Return True if two arrays have the same shape and elements, False otherwise.

#### Examples

```
>>> df = pd.DataFrame({'1': [10], '2': [20]})
>>> df
 1 2
0 10 20
```

DataFrames df and exactly\_equal have the same types and values for their elements and column labels, which will return True.

```
>>> exactly_equal = pd.DataFrame({'1': [10], '2': [20]})
>>> exactly_equal
 1 2
0 10 20
>>> df.equals(exactly_equal)
True
```

DataFrames df and different\_column\_type have the same element types and values, but have different types for the column labels, which will still return True.

```
>>> different_column_type = pd.DataFrame({'1.0': [10], '2.0': [20]})
>>> different_column_type
 1.0 2.0
0 10 20
>>> df.equals(different_column_type)
True
```

DataFrames df and different\_data\_type have different types for the same values for their elements, and will return False even though their column labels are the same values and types.

```
>>> different_data_type = pd.DataFrame({'1': [10.0], '2': [20.0]})
>>> different_data_type
 1 2
0 10.0 20.0
>>> df.equals(different_data_type)
False
```

```
ewm(self, com: 'float | None' = None, span: 'float | None' = None, halflife:
'float | TimedeltaConvertibleTypes | None' = None, alpha: 'float | None' = None,
min_periods: 'int | None' = 0, adjust: 'bool_t' = True, ignore_na: 'bool_t' = False,
axis: 'Axis' = 0, times: 'str | np.ndarray | FrameOrSeries | None' = None) ->
'ExponentialMovingWindow'
```

Provide exponential weighted (EW) functions.

Available EW functions: ``mean()``, ``var()``, ``std()``, ``corr()``, ``cov()``.

Exactly one parameter: ``com``, ``span``, ``halflife``, or ``alpha`` must be provided.

#### Parameters

**com** : float, optional

Specify decay in terms of center of mass,

:math:\alpha = 1 / (1 + com), for :math:com \geq 0.



```

span : float, optional
 Specify decay in terms of span,
 :math:\alpha = 2 / (\text{span} + 1), for :math:\text{span} \geq 1.
halflife : float, str, timedelta, optional
 Specify decay in terms of half-life,
 :math:\alpha = 1 - \exp\left(-\ln(2) / \text{halflife}\right), for
 :math:\text{halflife} > 0.

 If ``times`` is specified, the time unit (str or timedelta) over which an
 observation decays to half its value. Only applicable to ``mean()``
 and halflife value will not apply to the other functions.

 .. versionadded:: 1.1.0

alpha : float, optional
 Specify smoothing factor :math:\alpha directly,
 :math:0 < \alpha \leq 1.
min_periods : int, default 0
 Minimum number of observations in window required to have a value
 (otherwise result is NA).
adjust : bool, default True
 Divide by decaying adjustment factor in beginning periods to account
 for imbalance in relative weightings (viewing EWMA as a moving average).

 - When ``adjust=True`` (default), the EW function is calculated using
weights
series
 :math:w_i = (1 - \alpha)^i. For example, the EW moving average of the
 series
 [:math:x_0, x_1, \dots, x_t] would be:

 .. math::
 y_t = \frac{x_t + (1 - \alpha)x_{t-1} + (1 - \alpha)^2 x_{t-2} + \dots
+ (1 - \alpha)^t x_0}{1 + (1 - \alpha) + (1 - \alpha)^2 + \dots + (1 - \alpha)^t}

 - When ``adjust=False``, the exponentially weighted function is
calculated
 recursively:

 .. math::
 \begin{split}
y_0 &= x_0 \\
y_t &= (1 - \alpha) y_{t-1} + \alpha x_t,
\end{split}

ignore_na : bool, default False
 Ignore missing values when calculating weights; specify ``True`` to
reproduce
pre-0.15.0 behavior.

 - When ``ignore_na=False`` (default), weights are based on absolute
positions.
 For example, the weights of :math:x_0 and :math:x_2 used in
calculating
 the final weighted average of [:math>x_0, \text{None}, x_2] are
 :math:(1 - \alpha)^2 and :math:1 if ``adjust=True``, and
 :math:(1 - \alpha)^2 and :math:\alpha if ``adjust=False``.

 - When ``ignore_na=True`` (reproducing pre-0.15.0 behavior), weights are
based
 on relative positions. For example, the weights of :math>x_0 and
:math:x_2
 used in calculating the final weighted average of
 [:math>x_0, \text{None}, x_2] are :math:1 - \alpha and :math:1 if
 ``adjust=True``, and :math:1 - \alpha and :math:\alpha if
 ``adjust=False``.
axis : {0, 1}, default 0
 The axis to use. The value 0 identifies the rows, and 1
 identifies the columns.
times : str, np.ndarray, Series, default None

 .. versionadded:: 1.1.0

 Times corresponding to the observations. Must be monotonically increasing
and
 ``datetime64[ns]`` dtype.

 If str, the name of the column in the DataFrame representing the times.

 If 1-D array like, a sequence with the same shape as the observations.

 Only applicable to ``mean()``.

Returns

DataFrame
 A Window sub-classed for the particular operation.

```

See Also

-----

rolling : Provides rolling window calculations.  
expanding : Provides expanding transformations.

Notes

-----

More details can be found at:

:ref:`Exponentially weighted windows <window.exponentially\_weighted>`.

Examples

-----

```
>>> df = pd.DataFrame({'B': [0, 1, 2, np.nan, 4]})
```

```
>>> df
```

```
 B
```

```
0 0.0
```

```
1 1.0
```

```
2 2.0
```

```
3 NaN
```

```
4 4.0
```

```
>>> df.ewm(com=0.5).mean()
```

```
 B
```

```
0 0.000000
```

```
1 0.750000
```

```
2 1.615385
```

```
3 1.615385
```

```
4 3.670213
```

Specifying ``times`` with a `timedelta` ``half-life`` when computing mean.

```
>>> times = ['2020-01-01', '2020-01-03', '2020-01-10', '2020-01-15', '2020-01-17']
```

```
>>> df.ewm(half-life='4 days', times=pd.DatetimeIndex(times)).mean()
```

```
 B
```

```
0 0.000000
```

```
1 0.585786
```

```
2 1.523889
```

```
3 1.523889
```

```
4 3.233686
```

```
expanding(self, min_periods: 'int' = 1, center: 'bool_t | None' = None, axis: 'Axis' = 0, method: 'str' = 'single') -> 'Expanding'
```

Provide expanding transformations.

Parameters

-----

min\_periods : int, default 1

Minimum number of observations in window required to have a value (otherwise result is NA).

center : bool, default False

Set the labels at the center of the window.

axis : int or str, default 0

method : str {'single', 'table'}, default 'single'

Execute the rolling operation per single column or row (``'single'``) or over the entire object (``'table'``).

This argument is only implemented when specifying ``engine='numba'`` in the method call.

.. versionadded:: 1.3.0

Returns

-----

a Window sub-classed for the particular operation

See Also

-----

rolling : Provides rolling window calculations.

ewm : Provides exponential weighted functions.

Notes

-----

By default, the result is set to the right edge of the window. This can be changed to the center of the window by setting ``center=True``.

Examples

-----

```
>>> df = pd.DataFrame({"B": [0, 1, 2, np.nan, 4]})
```

```
>>> df
```

```
 B
```

```
0 0.0
```

```
1 1.0
```

```
2 2.0
```

```
3 NaN
```

```
4 4.0
```

```
>>> df.expanding(2).sum()
 B
0 NaN
1 1.0
2 3.0
3 3.0
4 7.0
```

```
filter(self: 'FrameOrSeries', items=None, like: 'str | None' = None, regex: 'str
None' = None, axis=None) -> 'FrameOrSeries'
```

Subset the dataframe rows or columns according to the specified index labels.

Note that this routine does not filter a dataframe on its contents. The filter is applied to the labels of the index.

#### Parameters

-----

items : list-like

Keep labels from axis which are in items.

like : str

Keep labels from axis for which "like in label == True".

regex : str (regular expression)

Keep labels from axis for which re.search(regex, label) == True.

axis : {0 or 'index', 1 or 'columns', None}, default None

The axis to filter on, expressed either as an index (int) or axis name (str). By default this is the info axis, 'index' for Series, 'columns' for DataFrame.

#### Returns

-----

same type as input object

#### See Also

-----

DataFrame.loc : Access a group of rows and columns by label(s) or a boolean array.

#### Notes

-----

The ``items``, ``like``, and ``regex`` parameters are enforced to be mutually exclusive.

``axis`` defaults to the info axis that is used when indexing with ``[]``.

#### Examples

-----

```
>>> df = pd.DataFrame(np.array([[1, 2, 3], [4, 5, 6]]),
... index=['mouse', 'rabbit'],
... columns=['one', 'two', 'three'])
>>> df
```

```
 one two three
mouse 1 2 3
rabbit 4 5 6
```

```
>>> # select columns by name
```

```
>>> df.filter(items=['one', 'three'])
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select columns by regular expression
```

```
>>> df.filter(regex='e$', axis=1)
```

```
 one three
mouse 1 3
rabbit 4 6
```

```
>>> # select rows containing 'bbi'
```

```
>>> df.filter(like='bbi', axis=0)
```

```
 one two three
rabbit 4 5 6
```

```
first(self: 'FrameOrSeries', offset) -> 'FrameOrSeries'
```

Select initial periods of time series data based on a date offset.

When having a DataFrame with dates as index, this function can select the first few rows based on a date offset.

#### Parameters

-----

offset : str, DateOffset or dateutil.relativedelta

The offset length of the data that will be selected. For instance, '1M' will display all the rows having their index within the first month.

#### Returns

```

Series or DataFrame
 A subset of the caller.

Raises

TypeError
 If the index is not a :class:`DatetimeIndex`

See Also

last : Select final periods of time series based on a date offset.
at_time : Select values at a particular time of the day.
between_time : Select values between particular times of the day.

Examples

>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
 A
2018-04-09 1
2018-04-11 2
2018-04-13 3
2018-04-15 4

Get the rows for the first 3 days:

>>> ts.first('3D')
 A
2018-04-09 1
2018-04-11 2

Notice the data for 3 first calendar days were returned, not the first
3 days observed in the dataset, and therefore data for 2018-04-13 was
not returned.

first_valid_index(self) -> 'Hashable | None'
 Return index for first non-NA value or None, if no NA value is found.

Returns

scalar : type of index

Notes

If all elements are non-NA/null, returns None.
Also returns None for empty Series/DataFrame.

get(self, key, default=None)
 Get item from object for given key (ex: DataFrame column).

 Returns default value if not found.

Parameters

key : object

Returns

value : same type as items contained in object

head(self: 'FrameOrSeries', n: 'int' = 5) -> 'FrameOrSeries'
 Return the first `n` rows.

 This function returns the first `n` rows for the object based
 on position. It is useful for quickly testing if your object
 has the right type of data in it.

 For negative values of `n`, this function returns all rows except
 the last `n` rows, equivalent to ``df[:-n]``.

Parameters

n : int, default 5
 Number of rows to select.

Returns

same type as caller
 The first `n` rows of the caller object.

See Also

DataFrame.tail: Returns the last `n` rows.

Examples

```

```

>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon', 'lion',
... 'monkey', 'parrot', 'shark', 'whale', 'zebra']})
>>> df
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
6 shark
7 whale
8 zebra

```

Viewing the first 5 lines

```

>>> df.head()
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey

```

Viewing the first `n` lines (three in this case)

```

>>> df.head(3)
 animal
0 alligator
1 bee
2 falcon

```

For negative values of `n`

```

>>> df.head(-3)
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot

```

`infer_objects(self: 'FrameOrSeries') -> 'FrameOrSeries'`  
 Attempt to infer better dtypes for object columns.

Attempts soft conversion of object-dtyped columns, leaving non-object and unconvertible columns unchanged. The inference rules are the same as during normal Series/DataFrame construction.

Returns

-----  
 converted : same type as input object

See Also

-----  
`to_datetime` : Convert argument to datetime.  
`to_timedelta` : Convert argument to timedelta.  
`to_numeric` : Convert argument to numeric type.  
`convert_dtypes` : Convert argument to best possible dtype.

Examples

```

>>> df = pd.DataFrame({"A": ["a", 1, 2, 3]})
>>> df = df.iloc[1:]
>>> df
 A
1 1
2 2
3 3

```

```

>>> df.dtypes
A object
dtype: object

```

```

>>> df.infer_objects().dtypes
A int64
dtype: object

```

`keys(self)`  
 Get the 'info axis' (see Indexing for more).

This is index for Series, columns for DataFrame.

Returns

```

Index
 Info axis.

last(self: 'FrameOrSeries', offset) -> 'FrameOrSeries'
 Select final periods of time series data based on a date offset.

 For a DataFrame with a sorted DatetimeIndex, this function
 selects the last few rows based on a date offset.

Parameters

offset : str, DateOffset, dateutil.relativedelta
 The offset length of the data that will be selected. For instance,
 '3D' will display all the rows having their index within the last 3 days.

Returns

Series or DataFrame
 A subset of the caller.

Raises

TypeError
 If the index is not a :class:`DatetimeIndex`

See Also

first : Select initial periods of time series based on a date offset.
at_time : Select values at a particular time of the day.
between_time : Select values between particular times of the day.

Examples

>>> i = pd.date_range('2018-04-09', periods=4, freq='2D')
>>> ts = pd.DataFrame({'A': [1, 2, 3, 4]}, index=i)
>>> ts
 A
2018-04-09 1
2018-04-11 2
2018-04-13 3
2018-04-15 4

Get the rows for the last 3 days:

>>> ts.last('3D')
 A
2018-04-13 3
2018-04-15 4

Notice the data for 3 last calendar days were returned, not the last
3 observed days in the dataset, and therefore data for 2018-04-11 was
not returned.

last_valid_index(self) -> 'Hashable | None'
 Return index for last non-NA value or None, if no NA value is found.

Returns

scalar : type of index

Notes

If all elements are non-NA/null, returns None.
Also returns None for empty Series/DataFrame.

pad = ffill(self: 'FrameOrSeries', axis: 'None | Axis' = None, inplace: 'bool_t'
= False, limit: 'None | int' = None, downcast=None) -> 'FrameOrSeries | None'
 Synonym for :meth:`DataFrame.fillna` with ``method='ffill'``.

Returns

Series/DataFrame or None
 Object with missing values filled or None if ``inplace=True``.

pct_change(self: 'FrameOrSeries', periods=1, fill_method='pad', limit=None,
freq=None, **kwargs) -> 'FrameOrSeries'
 Percentage change between the current and a prior element.

 Computes the percentage change from the immediately previous row by
 default. This is useful in comparing the percentage of change in a time
 series of elements.

Parameters

periods : int, default 1
 Periods to shift for forming percent change.

```

```

fill_method : str, default 'pad'
 How to handle NAs before computing percent changes.
limit : int, default None
 The number of consecutive NAs to fill before stopping.
freq : DateOffset, timedelta, or str, optional
 Increment to use from time series API (e.g. 'M' or BDay()).
**kwargs
 Additional keyword arguments are passed into
 `DataFrame.shift` or `Series.shift`.

Returns

chg : Series or DataFrame
 The same type as the calling object.

See Also

Series.diff : Compute the difference of two elements in a Series.
DataFrame.diff : Compute the difference of two elements in a DataFrame.
Series.shift : Shift the index by some number of periods.
DataFrame.shift : Shift the index by some number of periods.

Examples

Series

>>> s = pd.Series([90, 91, 85])
>>> s
0 90
1 91
2 85
dtype: int64

>>> s.pct_change()
0 NaN
1 0.011111
2 -0.065934
dtype: float64

>>> s.pct_change(periods=2)
0 NaN
1 NaN
2 -0.055556
dtype: float64

See the percentage change in a Series where filling NAs with last
valid observation forward to next valid.

>>> s = pd.Series([90, 91, None, 85])
>>> s
0 90.0
1 91.0
2 NaN
3 85.0
dtype: float64

>>> s.pct_change(fill_method='ffill')
0 NaN
1 0.011111
2 0.000000
3 -0.065934
dtype: float64

DataFrame

Percentage change in French franc, Deutsche Mark, and Italian lira from
1980-01-01 to 1980-03-01.

>>> df = pd.DataFrame({
... 'FR': [4.0405, 4.0963, 4.3149],
... 'GR': [1.7246, 1.7482, 1.8519],
... 'IT': [804.74, 810.01, 860.13]},
... index=['1980-01-01', '1980-02-01', '1980-03-01'])
>>> df
 FR GR IT
1980-01-01 4.0405 1.7246 804.74
1980-02-01 4.0963 1.7482 810.01
1980-03-01 4.3149 1.8519 860.13

>>> df.pct_change()
 FR GR IT
1980-01-01 NaN NaN NaN
1980-02-01 0.013810 0.013684 0.006549
1980-03-01 0.053365 0.059318 0.061876

Percentage of change in GOOG and APPL stock volume. Shows computing
the percentage change between columns.

```

```

>>> df = pd.DataFrame({
... '2016': [1769950, 30586265],
... '2015': [1500923, 40912316],
... '2014': [1371819, 41403351]},
... index=['GOOG', 'APPL'])
>>> df
 2016 2015 2014
GOOG 1769950 1500923 1371819
APPL 30586265 40912316 41403351

>>> df.pct_change(axis='columns', periods=-1)
 2016 2015 2014
GOOG 0.179241 0.094112 NaN
APPL -0.252395 -0.011860 NaN

```

pipe(self, func: 'Callable[..., T] | tuple[Callable[..., T], str]', \*args, \*\*kwargs) -> 'T'

Apply func(self, \\*args, \\*\\*kwargs).

Parameters

-----

func : function

Function to apply to the Series/DataFrame. \\*args\\*, and \\*\\*kwargs\\* are passed into \\*func\\*. Alternatively a \\*(callable, data\_keyword)\\* tuple where \\*data\_keyword\\* is a string indicating the keyword of \\*callable\\* that expects the Series/DataFrame.

args : iterable, optional

Positional arguments passed into \\*func\\*.

kwargs : mapping, optional

A dictionary of keyword arguments passed into \\*func\\*.

Returns

-----

object : the return type of \\*func\\*.

See Also

-----

DataFrame.apply : Apply a function along input axis of DataFrame.

DataFrame.applymap : Apply a function elementwise on a whole DataFrame.

Series.map : Apply a mapping correspondence on a :class:`~pandas.Series`.

Notes

-----

Use \\*.pipe\\* when chaining together functions that expect Series, DataFrames or GroupBy objects. Instead of writing

```

>>> func(g(h(df), arg1=a), arg2=b, arg3=c) # doctest: +SKIP

```

You can write

```

>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe(func, arg2=b, arg3=c)
...) # doctest: +SKIP

```

If you have a function that takes the data as (say) the second argument, pass a tuple indicating which keyword expects the data. For example, suppose \\*f\\* takes its data as \\*arg2\\*:

```

>>> (df.pipe(h)
... .pipe(g, arg1=a)
... .pipe((func, 'arg2'), arg1=a, arg3=c)
...) # doctest: +SKIP

```

rank(self: 'FrameOrSeries', axis=0, method: 'str' = 'average', numeric\_only: 'bool\_t | None' = None, na\_option: 'str' = 'keep', ascending: 'bool\_t' = True, pct: 'bool\_t' = False) -> 'FrameOrSeries'

Compute numerical data ranks (1 through n) along axis.

By default, equal values are assigned a rank that is the average of the ranks of those values.

Parameters

-----

axis : {0 or 'index', 1 or 'columns'}, default 0

Index to direct ranking.

method : {'average', 'min', 'max', 'first', 'dense'}, default 'average'

How to rank the group of records that have the same value (i.e. ties):

- \* average: average rank of the group
- \* min: lowest rank in the group
- \* max: highest rank in the group
- \* first: ranks assigned in order they appear in the array
- \* dense: like 'min', but rank always increases by 1 between groups.



numeric\_only : bool, optional  
 For DataFrame objects, rank only numeric columns if set to True.  
 na\_option : {'keep', 'top', 'bottom'}, default 'keep'  
 How to rank NaN values:

- \* keep: assign NaN rank to NaN values
- \* top: assign lowest rank to NaN values
- \* bottom: assign highest rank to NaN values

ascending : bool, default True  
 Whether or not the elements should be ranked in ascending order.  
 pct : bool, default False  
 Whether or not to display the returned rankings in percentile form.

#### Returns

-----  
 same type as caller  
 Return a Series or DataFrame with data ranks as values.

#### See Also

-----  
 core.groupby.GroupBy.rank : Rank of values within each group.

#### Examples

-----  
 >>> df = pd.DataFrame(data={'Animal': ['cat', 'penguin', 'dog',  
 ... 'spider', 'snake'],  
 ... 'Number\_legs': [4, 2, 4, 8, np.nan]})  
 >>> df

	Animal	Number_legs
0	cat	4.0
1	penguin	2.0
2	dog	4.0
3	spider	8.0
4	snake	NaN

The following example shows how the method behaves with the above parameters:

- \* default\_rank: this is the default behaviour obtained without using any parameter.
- \* max\_rank: setting ``method = 'max'`` the records that have the same values are ranked using the highest rank (e.g.: since 'cat' and 'dog' are both in the 2nd and 3rd position, rank 3 is assigned.)
- \* NA\_bottom: choosing ``na\_option = 'bottom'``, if there are records with NaN values they are placed at the bottom of the ranking.
- \* pct\_rank: when setting ``pct = True``, the ranking is expressed as percentile rank.

```
>>> df['default_rank'] = df['Number_legs'].rank()
>>> df['max_rank'] = df['Number_legs'].rank(method='max')
>>> df['NA_bottom'] = df['Number_legs'].rank(na_option='bottom')
>>> df['pct_rank'] = df['Number_legs'].rank(pct=True)
>>> df
```

	Animal	Number_legs	default_rank	max_rank	NA_bottom	pct_rank
0	cat	4.0	2.5	3.0	2.5	0.625
1	penguin	2.0	1.0	1.0	1.0	0.250
2	dog	4.0	2.5	3.0	2.5	0.625
3	spider	8.0	4.0	4.0	4.0	1.000
4	snake	NaN	NaN	NaN	5.0	NaN

```
reindex_like(self: 'FrameOrSeries', other, method: 'str | None' = None, copy:
bool_t' = True, limit=None, tolerance=None) -> 'FrameOrSeries'
Return an object with matching indices as other object.
```

Conform the object to the same index on all axes. Optional filling logic, placing NaN in locations having no value in the previous index. A new object is produced unless the new index is equivalent to the current one and copy=False.

#### Parameters

-----  
 other : Object of the same data type  
 Its row and column indices are used to define the new indices of this object.  
 method : {None, 'backfill'/'bfill', 'pad'/'ffill', 'nearest'}  
 Method to use for filling holes in reindexed DataFrame.  
 Please note: this is only applicable to DataFrames/Series with a monotonically increasing/decreasing index.

- \* None (default): don't fill gaps
- \* pad / ffill: propagate last valid observation forward to next valid
- \* backfill / bfill: use next valid observation to fill gap
- \* nearest: use nearest valid observations to fill gap.

copy : bool, default True  
 Return a new object, even if the passed indexes are the same.  
 limit : int, default None  
 Maximum number of consecutive labels to fill for inexact matches.  
 tolerance : optional  
 Maximum distance between original and new labels for inexact matches. The values of the index at the matching locations must satisfy the equation `abs(index[indexer] - target) <= tolerance`.  
  
 Tolerance may be a scalar value, which applies the same tolerance to all values, or list-like, which applies variable tolerance per element. List-like includes list, tuple, array, Series, and must be the same size as the index and its dtype must exactly match the index's type.

#### Returns

-----

Series or DataFrame

Same type as caller, but with changed indices on each axis.

#### See Also

-----

DataFrame.set\_index : Set row labels.

DataFrame.reset\_index : Remove row labels or move them to new columns.

DataFrame.reindex : Change to new indices or expand indices.

#### Notes

-----

Same as calling

``.reindex(index=other.index, columns=other.columns,...)``.

#### Examples

-----

```
>>> df1 = pd.DataFrame([[24.3, 75.7, 'high'],
... [31, 87.8, 'high'],
... [22, 71.6, 'medium'],
... [35, 95, 'medium']],
... columns=['temp_celsius', 'temp_fahrenheit',
... 'windspeed'],
... index=pd.date_range(start='2014-02-12',
... end='2014-02-15', freq='D'))
```

```
>>> df1
 temp_celsius temp_fahrenheit windspeed
2014-02-12 24.3 75.7 high
2014-02-13 31.0 87.8 high
2014-02-14 22.0 71.6 medium
2014-02-15 35.0 95.0 medium
```

```
>>> df2 = pd.DataFrame([[28, 'low'],
... [30, 'low'],
... [35.1, 'medium']],
... columns=['temp_celsius', 'windspeed'],
... index=pd.DatetimeIndex(['2014-02-12', '2014-02-13',
... '2014-02-15']))
```

```
>>> df2
 temp_celsius windspeed
2014-02-12 28.0 low
2014-02-13 30.0 low
2014-02-15 35.1 medium
```

```
>>> df2.reindex_like(df1)
 temp_celsius temp_fahrenheit windspeed
2014-02-12 28.0 NaN low
2014-02-13 30.0 NaN low
2014-02-14 NaN NaN NaN
2014-02-15 35.1 NaN medium
```

```
rename_axis(self, mapper=None, index=None, columns=None, axis=None, copy=True,
inplace=False)
```

Set the name of the axis for the index or columns.

#### Parameters

-----

mapper : scalar, list-like, optional

Value to set the axis name attribute.

index, columns : scalar, list-like, dict-like or function, optional

A scalar, list-like, dict-like or functions transformations to apply to that axis' values.

Note that the `columns` parameter is not allowed if the object is a Series. This parameter only apply for DataFrame type objects.

Use either `mapper` and `axis` to specify the axis to target with `mapper`, or `index`

and/or ``columns``.  
axis : {0 or 'index', 1 or 'columns'}, default 0  
The axis to rename.  
copy : bool, default True  
Also copy underlying data.  
inplace : bool, default False  
Modifies the object directly, instead of creating a new Series  
or DataFrame.

#### Returns

Series, DataFrame, or None  
The same type as the caller or None if ``inplace=True``.

#### See Also

Series.rename : Alter Series index labels or name.  
DataFrame.rename : Alter DataFrame index labels or name.  
Index.rename : Set new names on index.

#### Notes

``DataFrame.rename\_axis`` supports two calling conventions

```
* ``(index=index_mapper, columns=columns_mapper, ...)``
* ``(mapper, axis={'index', 'columns'}, ...)``
```

The first calling convention will only modify the names of  
the index and/or the names of the Index object that is the columns.  
In this case, the parameter ``copy`` is ignored.

The second calling convention will modify the names of the  
corresponding index if mapper is a list or a scalar.  
However, if mapper is dict-like or a function, it will use the  
deprecated behavior of modifying the axis \*labels\*.

We *highly* recommend using keyword arguments to clarify your  
intent.

#### Examples

**Series**

```
>>> s = pd.Series(["dog", "cat", "monkey"])
>>> s
0 dog
1 cat
2 monkey
dtype: object
>>> s.rename_axis("animal")
animal
0 dog
1 cat
2 monkey
dtype: object
```

**DataFrame**

```
>>> df = pd.DataFrame({"num_legs": [4, 4, 2],
... "num_arms": [0, 0, 2]},
... ["dog", "cat", "monkey"])
>>> df
 num_legs num_arms
dog 4 0
cat 4 0
monkey 2 2
>>> df = df.rename_axis("animal")
>>> df
 num_legs num_arms
animal
dog 4 0
cat 4 0
monkey 2 2
>>> df = df.rename_axis("limbs", axis="columns")
>>> df
limbs num_legs num_arms
animal
dog 4 0
cat 4 0
monkey 2 2
```

**MultiIndex**

```
>>> df.index = pd.MultiIndex.from_product(['mammal'],
... ['dog', 'cat', 'monkey']),
... names=['type', 'name'])
>>> df
```

limbs		num_legs	num_arms
type	name		
mammal	dog	4	0
	cat	4	0
	monkey	2	2

```
>>> df.rename_axis(index={'type': 'class'})
```

limbs		num_legs	num_arms
class	name		
mammal	dog	4	0
	cat	4	0
	monkey	2	2

```
>>> df.rename_axis(columns=str.upper)
```

LIMBS		num_legs	num_arms
type	name		
mammal	dog	4	0
	cat	4	0
	monkey	2	2

```
rolling(self, window: 'int | timedelta | BaseOffset | BaseIndexer', min_periods:
'int | None' = None, center: 'bool_t' = False, win_type: 'str | None' = None, on:
'str | None' = None, axis: 'Axis' = 0, closed: 'str | None' = None, method: 'str' =
'single')
```

Provide rolling window calculations.

Parameters

-----

window : int, offset, or BaseIndexer subclass

Size of the moving window. This is the number of observations used for calculating the statistic. Each window will be a fixed size.

If its an offset then this will be the time period of each window. Each window will be a variable sized based on the observations included in the time-period. This is only valid for datetimelike indexes.

If a BaseIndexer subclass is passed, calculates the window boundaries based on the defined ``get\_window\_bounds`` method. Additional rolling keyword arguments, namely ``min\_periods``, ``center``, and ``closed`` will be passed to ``get\_window\_bounds``.

min\_periods : int, default None

Minimum number of observations in window required to have a value (otherwise result is NA). For a window that is specified by an offset, ``min\_periods`` will default to 1. Otherwise, ``min\_periods`` will default to the size of the window.

center : bool, default False

Set the labels at the center of the window.

win\_type : str, default None

Provide a window type. If ``None``, all points are evenly weighted. See the notes below for further information.

on : str, optional

For a DataFrame, a datetime-like column or Index level on which to calculate the rolling window, rather than the DataFrame's index. Provided integer column is ignored and excluded from result since an integer index is not used to calculate the rolling window.

axis : int or str, default 0

closed : str, default None

Make the interval closed on the 'right', 'left', 'both' or 'neither' endpoints. Defaults to 'right'.

.. versionchanged:: 1.2.0

The closed parameter with fixed windows is now supported.

method : str {'single', 'table'}, default 'single'

Execute the rolling operation per single column or row (``'single'``) or over the entire object (``'table'``).

This argument is only implemented when specifying ``engine='numba'`` in the method call.

.. versionadded:: 1.3.0

Returns

-----

a Window or Rolling sub-classed for the particular operation

See Also

-----

expanding : Provides expanding transformations.

ewm : Provides exponential weighted functions.

Notes

-----

By default, the result is set to the right edge of the window. This can be changed to the center of the window by setting ``center=True``.

To learn more about the offsets & frequency strings, please see [this link](#)



```
>>> df
 B
2013-01-01 09:00:00 0.0
2013-01-01 09:00:02 1.0
2013-01-01 09:00:03 2.0
2013-01-01 09:00:05 NaN
2013-01-01 09:00:06 4.0
```

Contrasting to an integer rolling window, this will roll a variable length window corresponding to the time period. The default for min\_periods is 1.

```
>>> df.rolling('2s').sum()
 B
2013-01-01 09:00:00 0.0
2013-01-01 09:00:02 1.0
2013-01-01 09:00:03 3.0
2013-01-01 09:00:05 NaN
2013-01-01 09:00:06 4.0
```

```
sample(self: 'FrameOrSeries', n=None, frac: 'float | None' = None, replace:
'bool_t' = False, weights=None, random_state=None, axis: 'Axis | None' = None,
ignore_index: 'bool_t' = False) -> 'FrameOrSeries'
```

Return a random sample of items from an axis of object.

You can use `random\_state` for reproducibility.

Parameters

```

n : int, optional
 Number of items from axis to return. Cannot be used with `frac`.
 Default = 1 if `frac` = None.
frac : float, optional
 Fraction of axis items to return. Cannot be used with `n`.
replace : bool, default False
 Allow or disallow sampling of the same row more than once.
weights : str or ndarray-like, optional
 Default 'None' results in equal probability weighting.
 If passed a Series, will align with target object on index. Index
 values in weights not found in sampled object will be ignored and
 index values in sampled object not in weights will be assigned
 weights of zero.
 If called on a DataFrame, will accept the name of a column
 when axis = 0.
 Unless weights are a Series, weights must be same length as axis
 being sampled.
 If weights do not sum to 1, they will be normalized to sum to 1.
 Missing values in the weights column will be treated as zero.
 Infinite values not allowed.
random_state : int, array-like, BitGenerator, np.random.RandomState, optional
 If int, array-like, or BitGenerator (NumPy>=1.17), seed for
 random number generator
 If np.random.RandomState, use as numpy RandomState object.

.. versionchanged:: 1.1.0

 array-like and BitGenerator (for NumPy>=1.17) object now passed to
 np.random.RandomState() as seed

axis : {0 or 'index', 1 or 'columns', None}, default None
 Axis to sample. Accepts axis number or name. Default is stat axis
 for given data type (0 for Series and DataFrames).
ignore_index : bool, default False
 If True, the resulting index will be labeled 0, 1, ..., n - 1.

.. versionadded:: 1.3.0
```

Returns

```

Series or DataFrame
 A new object of same type as caller containing `n` items randomly
 sampled from the caller object.
```

See Also

```

DataFrameGroupBy.sample: Generates random samples from each group of a
 DataFrame object.
SeriesGroupBy.sample: Generates random samples from each group of a
 Series object.
numpy.random.choice: Generates a random sample from a given 1-D numpy
 array.
```

Notes

```

If `frac` > 1, `replacement` should be set to `True`.
```

Examples

```

>>> df = pd.DataFrame({'num_legs': [2, 4, 8, 0],
... 'num_wings': [2, 0, 0, 0],
... 'num_specimen_seen': [10, 2, 1, 8]},
... index=['falcon', 'dog', 'spider', 'fish'])
>>> df
 num_legs num_wings num_specimen_seen
falcon 2 2 10
dog 4 0 2
spider 8 0 1
fish 0 0 8

```

Extract 3 random elements from the ``Series`` ``df['num\_legs']``:  
Note that we use `random\_state` to ensure the reproducibility of the examples.

```

>>> df['num_legs'].sample(n=3, random_state=1)
fish 0
spider 8
falcon 2
Name: num_legs, dtype: int64

```

A random 50% sample of the ``DataFrame`` with replacement:

```

>>> df.sample(frac=0.5, replace=True, random_state=1)
 num_legs num_wings num_specimen_seen
dog 4 0 2
fish 0 0 8

```

An upsample sample of the ``DataFrame`` with replacement:  
Note that `replace` parameter has to be `True` for `frac` parameter > 1.

```

>>> df.sample(frac=2, replace=True, random_state=1)
 num_legs num_wings num_specimen_seen
dog 4 0 2
fish 0 0 8
falcon 2 2 10
falcon 2 2 10
fish 0 0 8
dog 4 0 2
fish 0 0 8
dog 4 0 2

```

Using a DataFrame column as weights. Rows with larger value in the `num\_specimen\_seen` column are more likely to be sampled.

```

>>> df.sample(n=2, weights='num_specimen_seen', random_state=1)
 num_legs num_wings num_specimen_seen
falcon 2 2 10
fish 0 0 8

```

```

set_flags(self: 'FrameOrSeries', *, copy: 'bool_t' = False,
allows_duplicate_labels: 'bool_t | None' = None) -> 'FrameOrSeries'
Return a new object with updated flags.

```

#### Parameters

```

allows_duplicate_labels : bool, optional
 Whether the returned object allows duplicate labels.

```

#### Returns

```

Series or DataFrame
 The same type as the caller.

```

#### See Also

```

DataFrame.attrs : Global metadata applying to this dataset.
DataFrame.flags : Global flags applying to this object.

```

#### Notes

This method returns a new object that's a view on the same data as the input. Mutating the input or the output values will be reflected in the other.

This method is intended to be used in method chains.

"Flags" differ from "metadata". Flags reflect properties of the pandas object (the Series or DataFrame). Metadata refer to properties of the dataset, and should be stored in :attr:`DataFrame.attrs`.

#### Examples

```

>>> df = pd.DataFrame({"A": [1, 2]})
>>> df.flags.allows_duplicate_labels
True

```

```

>>> df2 = df.set_flags(allows_duplicate_labels=False)
>>> df2.flags.allows_duplicate_labels
False

```

slice\_shift(self: 'FrameOrSeries', periods: 'int' = 1, axis=0) -> 'FrameOrSeries'

Equivalent to `shift` without copying data.  
The shifted data will not include the dropped periods and the shifted axis will be smaller than the original.

.. deprecated:: 1.2.0  
slice\_shift is deprecated,  
use DataFrame/Series.shift instead.

Parameters  
-----  
periods : int  
Number of periods to move, can be positive or negative.

Returns  
-----  
shifted : same type as caller

Notes  
-----  
While the `slice\_shift` is faster than `shift`, you may pay for it later during alignment.

squeeze(self, axis=None)

Squeeze 1 dimensional axis objects into scalars.

Series or DataFrames with a single element are squeezed to a scalar.  
DataFrames with a single column or a single row are squeezed to a Series. Otherwise the object is unchanged.

This method is most useful when you don't know if your object is a Series or DataFrame, but you do know it has just a single column. In that case you can safely call `squeeze` to ensure you have a Series.

Parameters  
-----  
axis : {0 or 'index', 1 or 'columns', None}, default None  
A specific axis to squeeze. By default, all length-1 axes are squeezed.

Returns  
-----  
DataFrame, Series, or scalar  
The projection after squeezing `axis` or all the axes.

See Also  
-----  
Series.iloc : Integer-location based indexing for selecting scalars.  
DataFrame.iloc : Integer-location based indexing for selecting Series.  
Series.to\_frame : Inverse of DataFrame.squeeze for a single-column DataFrame.

Examples  
-----

```

>>> primes = pd.Series([2, 3, 5, 7])

```

Slicing might produce a Series with a single value:

```

>>> even_primes = primes[primes % 2 == 0]
>>> even_primes
0 2
dtype: int64

>>> even_primes.squeeze()
2

```

Squeezing objects with more than one value in every axis does nothing:

```

>>> odd_primes = primes[primes % 2 == 1]
>>> odd_primes
1 3
2 5
3 7
dtype: int64

>>> odd_primes.squeeze()
1 3
2 5
3 7
dtype: int64

```

Squeezing is even more effective when used with DataFrames.



```
>>> df = pd.DataFrame([[1, 2], [3, 4]], columns=['a', 'b'])
>>> df
 a b
0 1 2
1 3 4
```

Slicing a single column will produce a DataFrame with the columns having only one value:

```
>>> df_a = df[['a']]
>>> df_a
 a
0 1
1 3
```

So the columns can be squeezed down, resulting in a Series:

```
>>> df_a.squeeze('columns')
0 1
1 3
Name: a, dtype: int64
```

Slicing a single row from a single column will produce a single scalar DataFrame:

```
>>> df_0a = df.loc[df.index < 1, ['a']]
>>> df_0a
 a
0 1
```

Squeezing the rows produces a single scalar Series:

```
>>> df_0a.squeeze('rows')
a 1
Name: 0, dtype: int64
```

Squeezing all axes will project directly into a scalar:

```
>>> df_0a.squeeze()
1
```

swapaxes(self: 'FrameOrSeries', axis1, axis2, copy=True) -> 'FrameOrSeries'  
Interchange axes and swap values axes appropriately.

Returns  
-----  
y : same as input

tail(self: 'FrameOrSeries', n: 'int' = 5) -> 'FrameOrSeries'  
Return the last `n` rows.

This function returns last `n` rows from the object based on position. It is useful for quickly verifying data, for example, after sorting or appending rows.

For negative values of `n`, this function returns all rows except the first `n` rows, equivalent to ``df[n:]``.

Parameters  
-----  
n : int, default 5  
Number of rows to select.

Returns  
-----  
type of caller  
The last `n` rows of the caller object.

See Also  
-----  
DataFrame.head : The first `n` rows of the caller object.

Examples  
-----  
>>> df = pd.DataFrame({'animal': ['alligator', 'bee', 'falcon', 'lion',  
... 'monkey', 'parrot', 'shark', 'whale', 'zebra']})  
>>> df
 animal
0 alligator
1 bee
2 falcon
3 lion
4 monkey
5 parrot
6 shark
7 whale

```
8 zebra
```

Viewing the last 5 lines

```
>>> df.tail()
 animal
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

Viewing the last `n` lines (three in this case)

```
>>> df.tail(3)
 animal
6 shark
7 whale
8 zebra
```

For negative values of `n`

```
>>> df.tail(-3)
 animal
3 lion
4 monkey
5 parrot
6 shark
7 whale
8 zebra
```

```
take(self: 'FrameOrSeries', indices, axis=0, is_copy: 'bool_t | None' = None,
**kwargs) -> 'FrameOrSeries'
```

Return the elements in the given *\*positional\** indices along an axis.

This means that we are not indexing according to actual values in the index attribute of the object. We are indexing according to the actual position of the element in the object.

Parameters

-----

**indices** : array-like

An array of ints indicating which positions to take.

**axis** : {0 or 'index', 1 or 'columns', None}, default 0

The axis on which to select elements. ``0`` means that we are selecting rows, ``1`` means that we are selecting columns.

**is\_copy** : bool

Before pandas 1.0, ``is\_copy=False`` can be specified to ensure that the return value is an actual copy. Starting with pandas 1.0, ``take`` always returns a copy, and the keyword is therefore deprecated.

.. deprecated:: 1.0.0

**\*\*kwargs**

For compatibility with :meth:`numpy.take`. Has no effect on the output.

Returns

-----

**taken** : same type as caller

An array-like containing the elements taken from the object.

See Also

-----

**DataFrame.loc** : Select a subset of a DataFrame by labels.

**DataFrame.iloc** : Select a subset of a DataFrame by positions.

**numpy.take** : Take elements from an array along an axis.

Examples

-----

```
>>> df = pd.DataFrame([('falcon', 'bird', 389.0),
... ('parrot', 'bird', 24.0),
... ('lion', 'mammal', 80.5),
... ('monkey', 'mammal', np.nan)],
... columns=['name', 'class', 'max_speed'],
... index=[0, 2, 3, 1])
>>> df
 name class max_speed
0 falcon bird 389.0
2 parrot bird 24.0
3 lion mammal 80.5
1 monkey mammal NaN
```

Take elements at positions 0 and 3 along the axis 0 (default).

Note how the actual indices selected (0 and 1) do not correspond to our selected indices 0 and 3. That's because we are selecting the 0th

and 3rd rows, not rows whose indices equal 0 and 3.

```
>>> df.take([0, 3])
 name class max_speed
0 falcon bird 389.0
1 monkey mammal NaN
```

Take elements at indices 1 and 2 along the axis 1 (column selection).

```
>>> df.take([1, 2], axis=1)
 class max_speed
0 bird 389.0
2 bird 24.0
3 mammal 80.5
1 mammal NaN
```

We may take elements using negative integers for positive indices, starting from the end of the object, just like with Python lists.

```
>>> df.take([-1, -2])
 name class max_speed
1 monkey mammal NaN
3 lion mammal 80.5
```

to\_clipboard(self, excel: 'bool\_t' = True, sep: 'str | None' = None, \*\*kwargs) -> 'None'

Copy object to the system clipboard.

Write a text representation of object to the system clipboard.  
This can be pasted into Excel, for example.

#### Parameters

-----

excel : bool, default True

Produce output in a csv format for easy pasting into excel.

- True, use the provided separator for csv pasting.
- False, write a string representation of the object to the clipboard.

sep : str, default ``\t``  
Field delimiter.

\*\*kwargs

These parameters will be passed to DataFrame.to\_csv.

#### See Also

-----

DataFrame.to\_csv : Write a DataFrame to a comma-separated values (csv) file.

read\_clipboard : Read text from clipboard and pass to read\_table.

#### Notes

-----

Requirements for your platform.

- Linux : `xclip`, or `xsel` (with `PyQt4` modules)
- Windows : none
- OS X : none

#### Examples

-----

Copy the contents of a DataFrame to the clipboard.

```
>>> df = pd.DataFrame([[1, 2, 3], [4, 5, 6]], columns=['A', 'B', 'C'])
```

```
>>> df.to_clipboard(sep=',') # doctest: +SKIP
... # Wrote the following to the system clipboard:
... # ,A,B,C
... # 0,1,2,3
... # 1,4,5,6
```

We can omit the index by passing the keyword `index` and setting it to false.

```
>>> df.to_clipboard(sep=',', index=False) # doctest: +SKIP
... # Wrote the following to the system clipboard:
... # A,B,C
... # 1,2,3
... # 4,5,6
```

to\_csv(self, path\_or\_buf: 'FilePathOrBuffer[AnyStr] | None' = None, sep: 'str' = ',', na\_rep: 'str' = '', float\_format: 'str | None' = None, columns: 'Sequence[Hashable] | None' = None, header: 'bool\_t | list[str]' = True, index: 'bool\_t' = True, index\_label: 'IndexLabel | None' = None, mode: 'str' = 'w', encoding: 'str | None' = None, compression: 'CompressionOptions' = 'infer', quoting: 'int | None' = None, quotechar: 'str' = '"', line\_terminator: 'str | None' = None, chunksize: 'int | None' = None, date\_format: 'str | None' = None, doublequote: 'bool\_t' = True, escapechar: 'str | None' = None, decimal: 'str' = '.', errors: 'str'

```
= 'strict', storage_options: 'StorageOptions' = None) -> 'str | None'
 Write object to a comma-separated values (csv) file.

Parameters

path_or_buf : str or file handle, default None
 File path or object, if None is provided the result is returned as
 a string. If a non-binary file object is passed, it should be opened
 with `newline=''`, disabling universal newlines. If a binary
 file object is passed, `mode` might need to contain a `b`'.

.. versionchanged:: 1.2.0

 Support for binary file objects was introduced.

sep : str, default ','
 String of length 1. Field delimiter for the output file.
na_rep : str, default ''
 Missing data representation.
float_format : str, default None
 Format string for floating point numbers.
columns : sequence, optional
 Columns to write.
header : bool or list of str, default True
 Write out the column names. If a list of strings is given it is
 assumed to be aliases for the column names.
index : bool, default True
 Write row names (index).
index_label : str or sequence, or False, default None
 Column label for index column(s) if desired. If None is given, and
 `header` and `index` are True, then the index names are used. A
 sequence should be given if the object uses MultiIndex. If
 False do not print fields for index names. Use index_label=False
 for easier importing in R.
mode : str
 Python write mode, default 'w'.
encoding : str, optional
 A string representing the encoding to use in the output file,
 defaults to 'utf-8'. `encoding` is not supported if `path_or_buf`
 is a non-binary file object.
compression : str or dict, default 'infer'
 If str, represents compression mode. If dict, value at 'method' is
 the compression mode. Compression mode may be any of the following
 possible values: {'infer', 'gzip', 'bz2', 'zip', 'xz', None}. If
 compression mode is 'infer' and `path_or_buf` is path-like, then
 detect compression mode from the following extensions: '.gz',
 '.bz2', '.zip' or '.xz'. (otherwise no compression). If dict given
 and mode is one of {'zip', 'gzip', 'bz2'}, or inferred as
 one of the above, other entries passed as
 additional compression options.

.. versionchanged:: 1.0.0

 May now be a dict with key 'method' as compression mode
 and other entries as additional compression options if
 compression mode is 'zip'.

.. versionchanged:: 1.1.0

 Passing compression options as keys in dict is
 supported for compression modes 'gzip' and 'bz2'
 as well as 'zip'.

.. versionchanged:: 1.2.0

 Compression is supported for binary file objects.

.. versionchanged:: 1.2.0

 Previous versions forwarded dict entries for 'gzip' to
 `gzip.open` instead of `gzip.GzipFile` which prevented
 setting `mtime`.

quoting : optional constant from csv module
 Defaults to csv.QUOTE_MINIMAL. If you have set a `float_format`
 then floats are converted to strings and thus csv.QUOTE_NONNUMERIC
 will treat them as non-numeric.
quotechar : str, default '"'
 String of length 1. Character used to quote fields.
line_terminator : str, optional
 The newline character or character sequence to use in the output
 file. Defaults to `os.linesep`, which depends on the OS in which
 this method is called (`\n` for linux, `\\r\\n` for Windows, i.e.).
chunksize : int or None
 Rows to write at a time.
date_format : str, default None
 Format string for datetime objects.
```

```

doublequote : bool, default True
 Control quoting of `quotechar` inside a field.
escapechar : str, default None
 String of length 1. Character used to escape `sep` and `quotechar`
 when appropriate.
decimal : str, default '.'
 Character recognized as decimal separator. E.g. use ',' for
 European data.
errors : str, default 'strict'
 Specifies how encoding and decoding errors are to be handled.
 See the errors argument for :func:`open` for a full list
 of options.

.. versionadded:: 1.1.0

storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2.0

Returns

None or str
 If path_or_buf is None, returns the resulting csv format as a
 string. Otherwise returns None.

See Also

read_csv : Load a CSV file into a DataFrame.
to_excel : Write DataFrame to an Excel file.

Examples

>>> df = pd.DataFrame({'name': ['Raphael', 'Donatello'],
... 'mask': ['red', 'purple'],
... 'weapon': ['sai', 'bo staff']})
>>> df.to_csv(index=False)
'name,mask,weapon\nRaphael,red,sai\nDonatello,purple,bo staff\n'

Create 'out.zip' containing 'out.csv'

>>> compression_opts = dict(method='zip',
... archive_name='out.csv') # doctest: +SKIP
>>> df.to_csv('out.zip', index=False,
... compression=compression_opts) # doctest: +SKIP

to_excel(self, excel_writer, sheet_name: 'str' = 'Sheet1', na_rep: 'str' = '',
float_format: 'str | None' = None, columns=None, header=True, index=True,
index_label=None, startrow=0, startcol=0, engine=None, merge_cells=True,
encoding=None, inf_rep='inf', verbose=True, freeze_panes=None, storage_options:
'StorageOptions' = None) -> 'None'
 Write object to an Excel sheet.

To write a single object to an Excel .xlsx file it is only necessary to
specify a target file name. To write to multiple sheets it is necessary to
create an `ExcelWriter` object with a target file name, and specify a sheet
in the file to write to.

Multiple sheets may be written to by specifying unique `sheet_name`.
With all data written to the file it is necessary to save the changes.
Note that creating an `ExcelWriter` object with a file name that already
exists will result in the contents of the existing file being erased.

Parameters

excel_writer : path-like, file-like, or ExcelWriter object
 File path or existing ExcelWriter.
sheet_name : str, default 'Sheet1'
 Name of sheet which will contain DataFrame.
na_rep : str, default ''
 Missing data representation.
float_format : str, optional
 Format string for floating point numbers. For example
 ``float_format="%.2f"`` will format 0.1234 to 0.12.
columns : sequence or list of str, optional
 Columns to write.
header : bool or list of str, default True
 Write out the column names. If a list of string is given it is
 assumed to be aliases for the column names.
index : bool, default True
 Write row names (index).
index_label : str or sequence, optional
 Column label for index column(s) if desired. If not specified, and

```

```

 `header` and `index` are True, then the index names are used. A
 sequence should be given if the DataFrame uses MultiIndex.
startrow : int, default 0
 Upper left cell row to dump data frame.
startcol : int, default 0
 Upper left cell column to dump data frame.
engine : str, optional
 Write engine to use, 'openpyxl' or 'xlsxwriter'. You can also set this
 via the options ``io.excel.xlsx.writer``, ``io.excel.xls.writer``, and
 ``io.excel.xlsm.writer``.

.. deprecated:: 1.2.0

 As the `xlwt` <https://pypi.org/project/xlwt/>__ package is no longer
 maintained, the ``xlwt`` engine will be removed in a future version
 of pandas.

merge_cells : bool, default True
 Write MultiIndex and Hierarchical Rows as merged cells.
encoding : str, optional
 Encoding of the resulting excel file. Only necessary for xlwt,
 other writers support unicode natively.
inf_rep : str, default 'inf'
 Representation for infinity (there is no native representation for
 infinity in Excel).
verbose : bool, default True
 Display more information in the error logs.
freeze_panes : tuple of int (length 2), optional
 Specifies the one-based bottommost row and rightmost column that
 is to be frozen.
storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2.0

```

#### See Also

```

to_csv : Write DataFrame to a comma-separated values (csv) file.
ExcelWriter : Class for writing DataFrame objects into excel sheets.
read_excel : Read an Excel file into a pandas DataFrame.
read_csv : Read a comma-separated values (csv) file into DataFrame.

```

#### Notes

```

For compatibility with :meth:`~DataFrame.to_csv`,
to_excel serializes lists and dicts to strings before writing.

```

Once a workbook has been saved it is not possible to write further data without rewriting the whole workbook.

#### Examples

Create, write to and save a workbook:

```

>>> df1 = pd.DataFrame([['a', 'b'], ['c', 'd']],
... index=['row 1', 'row 2'],
... columns=['col 1', 'col 2'])
>>> df1.to_excel("output.xlsx") # doctest: +SKIP

```

To specify the sheet name:

```

>>> df1.to_excel("output.xlsx",
... sheet_name='Sheet_name_1') # doctest: +SKIP

```

If you wish to write to more than one sheet in the workbook, it is necessary to specify an ExcelWriter object:

```

>>> df2 = df1.copy()
>>> with pd.ExcelWriter('output.xlsx') as writer: # doctest: +SKIP
... df1.to_excel(writer, sheet_name='Sheet_name_1')
... df2.to_excel(writer, sheet_name='Sheet_name_2')

```

ExcelWriter can also be used to append to an existing Excel file:

```

>>> with pd.ExcelWriter('output.xlsx',
... mode='a') as writer: # doctest: +SKIP
... df.to_excel(writer, sheet_name='Sheet_name_3')

```

To set the library that is used to write the Excel file, you can pass the `engine` keyword (the default engine is automatically chosen depending on the file extension):

```
>>> df1.to_excel('output1.xlsx', engine='xlsxwriter') # doctest: +SKIP
```

to\_hdf(self, path\_or\_buf, key: 'str', mode: 'str' = 'a', complevel: 'int | None' = None, complib: 'str | None' = None, append: 'bool\_t' = False, format: 'str | None' = None, index: 'bool\_t' = True, min\_itemsize: 'int | dict[str, int] | None' = None, nan\_rep=None, dropna: 'bool\_t | None' = None, data\_columns: 'bool\_t | list[str] | None' = None, errors: 'str' = 'strict', encoding: 'str' = 'UTF-8') -> 'None'

Write the contained data to an HDF5 file using HDFStore.

Hierarchical Data Format (HDF) is self-describing, allowing an application to interpret the structure and contents of a file with no outside information. One HDF file can hold a mix of related objects which can be accessed as a group or as individual objects.

In order to add another DataFrame or Series to an existing HDF file please use append mode and a different a key.

.. warning::

One can store a subclass of ``DataFrame`` or ``Series`` to HDF5, but the type of the subclass is lost upon storing.

For more information see the :ref:`user guide <io.hdf5>`.

Parameters

-----

path\_or\_buf : str or pandas.HDFStore  
File path or HDFStore object.

key : str  
Identifier for the group in the store.

mode : {'a', 'w', 'r+'}, default 'a'  
Mode to open file:

- 'w': write, a new file is created (an existing file with the same name would be deleted).
- 'a': append, an existing file is opened for reading and writing, and if the file does not exist it is created.
- 'r+': similar to 'a', but the file must already exist.

complevel : {0-9}, optional  
Specifies a compression level for data.  
A value of 0 disables compression.

complib : {'zlib', 'lzo', 'bzip2', 'blosc'}, default 'zlib'  
Specifies the compression library to be used.  
As of v0.20.2 these additional compressors for Blosc are supported (default if no compressor specified: 'blosc:blosclz'):  
{'blosc:blosclz', 'blosc:lz4', 'blosc:lz4hc', 'blosc:snappy', 'blosc:zlib', 'blosc:zstd'}.

Specifying a compression library which is not available issues a ValueError.

append : bool, default False  
For Table formats, append the input data to the existing.

format : {'fixed', 'table', None}, default 'fixed'  
Possible values:

- 'fixed': Fixed format. Fast writing/reading. Not-appendable, nor searchable.
- 'table': Table format. Write as a PyTables Table structure which may perform worse but allow more flexible operations like searching / selecting subsets of the data.
- If None, pd.get\_option('io.hdf.default\_format') is checked, followed by fallback to "fixed"

errors : str, default 'strict'  
Specifies how encoding and decoding errors are to be handled.  
See the errors argument for :func:`open` for a full list of options.

encoding : str, default "UTF-8"

min\_itemsize : dict or int, optional  
Map column names to minimum string sizes for columns.

nan\_rep : Any, optional  
How to represent null values as str.  
Not allowed with append=True.

data\_columns : list of columns or True, optional  
List of columns to create as indexed data columns for on-disk queries, or True to use all columns. By default only the axes of the object are indexed. See :ref:`io.hdf5-query-data-columns`. Applicable only to format='table'.

See Also

-----

read\_hdf : Read from HDF file.  
DataFrame.to\_parquet : Write a DataFrame to the binary parquet format.  
DataFrame.to\_sql : Write to a SQL table.  
DataFrame.to\_feather : Write out feather-format for DataFrames.  
DataFrame.to\_csv : Write out to a csv file.

Examples

-----

```
>>> df = pd.DataFrame({'A': [1, 2, 3], 'B': [4, 5, 6]},
... index=['a', 'b', 'c'])
>>> df.to_hdf('data.h5', key='df', mode='w')
```

We can add another object to the same file:

```
>>> s = pd.Series([1, 2, 3, 4])
>>> s.to_hdf('data.h5', key='s')
```

Reading from HDF file:

```
>>> pd.read_hdf('data.h5', 'df')
A B
a 1 4
b 2 5
c 3 6
>>> pd.read_hdf('data.h5', 's')
0 1
1 2
2 3
3 4
dtype: int64
```

Deleting file with data:

```
>>> import os
>>> os.remove('data.h5')
```

```
to_json(self, path_or_buf: 'FilePathOrBuffer | None' = None, orient: 'str | None'
= None, date_format: 'str | None' = None, double_precision: 'int' = 10, force_ascii:
'bool_t' = True, date_unit: 'str' = 'ms', default_handler: 'Callable[[Any],
JSONSerializable] | None' = None, lines: 'bool_t' = False, compression:
'CompressionOptions' = 'infer', index: 'bool_t' = True, indent: 'int | None' = None,
storage_options: 'StorageOptions' = None) -> 'str | None'
```

Convert the object to a JSON string.

Note NaN's and None will be converted to null and datetime objects will be converted to UNIX timestamps.

Parameters

-----

path\_or\_buf : str or file handle, optional

File path or object. If not specified, the result is returned as a string.

orient : str

Indication of expected JSON string format.

\* Series:

- default is 'index'
- allowed values are: {'split', 'records', 'index', 'table'}.

\* DataFrame:

- default is 'columns'
- allowed values are: {'split', 'records', 'index', 'columns', 'values', 'table'}.

\* The format of the JSON string:

- 'split' : dict like {'index' -> [index], 'columns' -> [columns], 'data' -> [values]}
- 'records' : list like [{column -> value}, ... , {column -> value}]
- 'index' : dict like {index -> {column -> value}}
- 'columns' : dict like {column -> {index -> value}}
- 'values' : just the values array
- 'table' : dict like {'schema': {schema}, 'data': {data}}

Describing the data, where data component is like

```
`orient='records'``.
```

date\_format : {None, 'epoch', 'iso'}

Type of date conversion. 'epoch' = epoch milliseconds, 'iso' = ISO8601. The default depends on the 'orient'. For 'orient='table'', the default is 'iso'. For all other orients, the default is 'epoch'.

double\_precision : int, default 10

The number of decimal places to use when encoding floating point values.

force\_ascii : bool, default True

Force encoded string to be ASCII.

date\_unit : str, default 'ms' (milliseconds)

The time unit to encode to, governs timestamp and ISO8601 precision. One of 's', 'ms', 'us', 'ns' for second, millisecond, microsecond, and nanosecond respectively.

default\_handler : callable, default None

Handler to call if object cannot otherwise be converted to a



```

 suitable format for JSON. Should receive a single argument which is
 the object to convert and return a serialisable object.
lines : bool, default False
 If 'orient' is 'records' write out line-delimited json format. Will
 throw ValueError if incorrect 'orient' since others are not
 list-like.

compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}

 A string representing the compression to use in the output file,
 only used when the first argument is a filename. By default, the
 compression is inferred from the filename.
index : bool, default True
 Whether to include the index values in the JSON string. Not
 including the index ('`index=False`') is only supported when
 orient is 'split' or 'table'.
indent : int, optional
 Length of whitespace used to indent each record.

.. versionadded:: 1.0.0

storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to
 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

.. versionadded:: 1.2.0

Returns

None or str
 If path_or_buf is None, returns the resulting json format as a
 string. Otherwise returns None.

See Also

read_json : Convert a JSON string to pandas object.

Notes

The behavior of ``indent=0`` varies from the stdlib, which does not
indent the output but does insert newlines. Currently, ``indent=0``
and the default ``indent=None`` are equivalent in pandas, though this
may change in a future release.

``orient='table'`` contains a 'pandas_version' field under 'schema'.
This stores the version of 'pandas' used in the latest revision of the
schema.

Examples

>>> import json
>>> df = pd.DataFrame(
... [{"a", "b"}, {"c", "d"}],
... index=["row 1", "row 2"],
... columns=["col 1", "col 2"],
...)

>>> result = df.to_json(orient="split")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
 "columns": [
 "col 1",
 "col 2"
],
 "index": [
 "row 1",
 "row 2"
],
 "data": [
 [
 "a",
 "b"
],
 [
 "c",
 "d"
]
]
}


```

Encoding/decoding a Dataframe using ``'records'`` formatted JSON.  
 Note that index labels are not preserved with this encoding.

```
>>> result = df.to_json(orient="records")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
[
 {
 "col 1": "a",
 "col 2": "b"
 },
 {
 "col 1": "c",
 "col 2": "d"
 }
]
```

Encoding/decoding a Dataframe using ``'index'`` formatted JSON:

```
>>> result = df.to_json(orient="index")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
 "row 1": {
 "col 1": "a",
 "col 2": "b"
 },
 "row 2": {
 "col 1": "c",
 "col 2": "d"
 }
}
```

Encoding/decoding a Dataframe using ``'columns'`` formatted JSON:

```
>>> result = df.to_json(orient="columns")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
 "col 1": {
 "row 1": "a",
 "row 2": "c"
 },
 "col 2": {
 "row 1": "b",
 "row 2": "d"
 }
}
```

Encoding/decoding a Dataframe using ``'values'`` formatted JSON:

```
>>> result = df.to_json(orient="values")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
[
 [
 "a",
 "b"
],
 [
 "c",
 "d"
]
]
```

Encoding with Table Schema:

```
>>> result = df.to_json(orient="table")
>>> parsed = json.loads(result)
>>> json.dumps(parsed, indent=4) # doctest: +SKIP
{
 "schema": {
 "fields": [
 {
 "name": "index",
 "type": "string"
 },
 {
 "name": "col 1",
 "type": "string"
 },
 {
 "name": "col 2",
 "type": "string"
 }
],
 "primaryKey": [
 "index"
],
 "pandas_version": "0.20.0"
 }
}
```

```

 },
 "data": [
 {
 "index": "row 1",
 "col 1": "a",
 "col 2": "b"
 },
 {
 "index": "row 2",
 "col 1": "c",
 "col 2": "d"
 }
]
}

```

`to_latex(self, buf=None, columns=None, col_space=None, header=True, index=True, na_rep='NaN', formatters=None, float_format=None, sparsify=None, index_names=True, bold_rows=False, column_format=None, longtable=None, escape=None, encoding=None, decimal='.', multicolumn=None, multicolumn_format=None, multirow=None, caption=None, label=None, position=None)`

Render object to a LaTeX tabular, longtable, or nested table/tabular.

Requires `\\usepackage{booktabs}`. The output can be copy/pasted into a main LaTeX document or read from an external file with `\\input{table.tex}`.

.. versionchanged:: 1.0.0

Added caption and label arguments.

.. versionchanged:: 1.2.0

Added position argument, changed meaning of caption argument.

#### Parameters

`buf` : str, Path or StringIO-like, optional, default None  
Buffer to write to. If None, the output is returned as a string.

`columns` : list of label, optional  
The subset of columns to write. Writes all columns by default.

`col_space` : int, optional  
The minimum width of each column.

`header` : bool or list of str, default True  
Write out the column names. If a list of strings is given, it is assumed to be aliases for the column names.

`index` : bool, default True  
Write row names (index).

`na_rep` : str, default 'NaN'  
Missing data representation.

`formatters` : list of functions or dict of {str: function}, optional  
Formatter functions to apply to columns' elements by position or name. The result of each function must be a unicode string. List must be of length equal to the number of columns.

`float_format` : one-parameter function or str, optional, default None  
Formatter for floating point numbers. For example `\\float_format="%0.2f"` and `\\float_format="{:0.2f}".format` will both result in 0.1234 being formatted as 0.12.

`sparsify` : bool, optional  
Set to False for a DataFrame with a hierarchical index to print every multiindex key at each row. By default, the value will be read from the config module.

`index_names` : bool, default True  
Prints the names of the indexes.

`bold_rows` : bool, default False  
Make the row labels bold in the output.

`column_format` : str, optional  
The columns format as specified in 'LaTeX table format' <<https://en.wikibooks.org/wiki/LaTeX/Tables>>\_\_ e.g. 'rcl' for 3 columns. By default, 'l' will be used for all columns except columns of numbers, which default to 'r'.

`longtable` : bool, optional  
By default, the value will be read from the pandas config module. Use a longtable environment instead of tabular. Requires adding a `\\usepackage{longtable}` to your LaTeX preamble.

`escape` : bool, optional  
By default, the value will be read from the pandas config module. When set to False prevents from escaping latex special characters in column names.

`encoding` : str, optional  
A string representing the encoding to use in the output file, defaults to 'utf-8'.

`decimal` : str, default '.'  
Character recognized as decimal separator, e.g. ',' in Europe.

`multicolumn` : bool, default True  
Use `\\multicolumn` to enhance MultiIndex columns. The default will be read from the config module.

`multicolumn_format` : str, default 'l'  
The alignment for multicolumns, similar to `\\column_format`. The default will be read from the config module.

```

multirow : bool, default False
 Use \multirow to enhance MultiIndex rows. Requires adding a
 \usepackage{multirow} to your LaTeX preamble. Will print
 centered labels (instead of top-aligned) across the contained
 rows, separating groups via clines. The default will be read
 from the pandas config module.
caption : str or tuple, optional
 Tuple (full_caption, short_caption),
 which results in ``\caption[short_caption]{full_caption}``;
 if a single string is passed, no short caption will be set.

 .. versionadded:: 1.0.0

 .. versionchanged:: 1.2.0
 Optionally allow caption to be a tuple ``(full_caption,
short_caption)``.

label : str, optional
 The LaTeX label to be placed inside ``\label{}`` in the output.
 This is used with ``\ref{}`` in the main ``.tex`` file.

 .. versionadded:: 1.0.0
position : str, optional
 The LaTeX positional argument for tables, to be placed after
 ``\begin{}`` in the output.

 .. versionadded:: 1.2.0

 Returns

 str or None
 If buf is None, returns the result as a string. Otherwise returns
 None.

See Also

DataFrame.to_string : Render a DataFrame to a console-friendly
 tabular output.
DataFrame.to_html : Render a DataFrame as an HTML table.

Examples

>>> df = pd.DataFrame(dict(name=['Raphael', 'Donatello'],
... mask=['red', 'purple'],
... weapon=['sai', 'bo staff']))
>>> print(df.to_latex(index=False)) # doctest: +NORMALIZE_WHITESPACE
\begin{tabular}{lll}
\toprule
 name & mask & weapon \\
\midrule
 Raphael & red & sai \\
 Donatello & purple & bo staff \\
\bottomrule
\end{tabular}

 to_pickle(self, path, compression: 'CompressionOptions' = 'infer', protocol:
'int' = 4, storage_options: 'StorageOptions' = None) -> 'None'
 Pickle (serialize) object to file.

Parameters

path : str
 File path where the pickled object will be stored.
compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default
'infer'
 A string representing the compression to use in the output file. By
 default, infers from the file extension in specified path.
 Compression mode may be any of the following possible
 values: {'infer', 'gzip', 'bz2', 'zip', 'xz', None}. If compression
 mode is 'infer' and path_or_buf is path-like, then detect
 compression mode from the following extensions:
 '.gz', '.bz2', '.zip' or '.xz'. (otherwise no compression).
 If dict given and mode is 'zip' or inferred as 'zip', other entries
 passed as additional compression options.
protocol : int
 Int which indicates which protocol should be used by the pickler,
 default HIGHEST_PROTOCOL (see [1]_ paragraph 12.1.2). The possible
 values are 0, 1, 2, 3, 4, 5. A negative value for the protocol
 parameter is equivalent to setting its value to HIGHEST_PROTOCOL.

 .. [1] https://docs.python.org/3/library/pickle.html.

storage_options : dict, optional
 Extra options that make sense for a particular storage connection, e.g.
 host, port, username, password, etc. For HTTP(S) URLs the key-value pairs
 are forwarded to ``urllib`` as header options. For other URLs (e.g.
 starting with "s3://", and "gcs://") the key-value pairs are forwarded to

```

```

 ``fsspec``. Please see ``fsspec`` and ``urllib`` for more details.

 .. versionadded:: 1.2.0

See Also

read_pickle : Load pickled pandas object (or any object) from file.
DataFrame.to_hdf : Write DataFrame to an HDF5 file.
DataFrame.to_sql : Write DataFrame to a SQL database.
DataFrame.to_parquet : Write a DataFrame to the binary parquet format.

Examples

>>> original_df = pd.DataFrame({"foo": range(5), "bar": range(5, 10)})
>>> original_df
 foo bar
0 0 5
1 1 6
2 2 7
3 3 8
4 4 9
>>> original_df.to_pickle("./dummy.pkl")

>>> unpickled_df = pd.read_pickle("./dummy.pkl")
>>> unpickled_df
 foo bar
0 0 5
1 1 6
2 2 7
3 3 8
4 4 9

>>> import os
>>> os.remove("./dummy.pkl")

to_sql(self, name: 'str', con, schema=None, if_exists: 'str' = 'fail', index:
'bool_t' = True, index_label=None, chunksize=None, dtype: 'DtypeArg | None' = None,
method=None) -> 'None'
 Write records stored in a DataFrame to a SQL database.

 Databases supported by SQLAlchemy [1]_ are supported. Tables can be
 newly created, appended to, or overwritten.

 Parameters

 name : str
 Name of SQL table.
 con : sqlalchemy.engine.(Engine or Connection) or sqlite3.Connection
 Using SQLAlchemy makes it possible to use any DB supported by that
 library. Legacy support is provided for sqlite3.Connection objects. The
user
 is responsible for engine disposal and connection closure for the
SQLAlchemy
 connectable See `here
<https://docs.sqlalchemy.org/en/13/core/connections.html>`_.

 schema : str, optional
 Specify the schema (if database flavor supports this). If None, use
 default schema.
 if_exists : {'fail', 'replace', 'append'}, default 'fail'
 How to behave if the table already exists.

 * fail: Raise a ValueError.
 * replace: Drop the table before inserting new values.
 * append: Insert new values to the existing table.

 index : bool, default True
 Write DataFrame index as a column. Uses `index_label` as the column
 name in the table.
 index_label : str or sequence, default None
 Column label for index column(s). If None is given (default) and
 `index` is True, then the index names are used.
 A sequence should be given if the DataFrame uses MultiIndex.
 chunksize : int, optional
 Specify the number of rows in each batch to be written at a time.
 By default, all rows will be written at once.
 dtype : dict or scalar, optional
 Specifying the datatype for columns. If a dictionary is used, the
 keys should be the column names and the values should be the
 SQLAlchemy types or strings for the sqlite3 legacy mode. If a
 scalar is provided, it will be applied to all columns.
 method : {None, 'multi', callable}, optional
 Controls the SQL insertion clause used:

 * None : Uses standard SQL ``INSERT`` clause (one per row).
 * 'multi': Pass multiple values in a single ``INSERT`` clause.
 * callable with signature ``(pd_table, conn, keys, data_iter)``.

```

Details and a sample callable implementation can be found in the section :ref:`insert method <io.sql.method>`.

#### Raises

-----

##### ValueError

When the table already exists and `if\_exists` is 'fail' (the default).

#### See Also

-----

`read_sql` : Read a DataFrame from a table.

#### Notes

-----

Timezone aware datetime columns will be written as ``Timestamp with timezone`` type with SQLAlchemy if supported by the database. Otherwise, the datetimes will be stored as timezone unaware timestamps local to the original timezone.

#### References

-----

- .. [1] <https://docs.sqlalchemy.org>
- .. [2] <https://www.python.org/dev/peps/pep-0249/>

#### Examples

-----

Create an in-memory SQLite database.

```
>>> from sqlalchemy import create_engine
>>> engine = create_engine('sqlite://', echo=False)
```

Create a table from scratch with 3 rows.

```
>>> df = pd.DataFrame({'name' : ['User 1', 'User 2', 'User 3']})
>>> df
 name
0 User 1
1 User 2
2 User 3
```

```
>>> df.to_sql('users', con=engine)
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3')]
```

An ``sqlalchemy.engine.Connection`` can also be passed to ``con``:

```
>>> with engine.begin() as connection:
... df1 = pd.DataFrame({'name' : ['User 4', 'User 5']})
... df1.to_sql('users', con=connection, if_exists='append')
```

This is allowed to support operations that require that the same DBAPI connection is used for the entire operation.

```
>>> df2 = pd.DataFrame({'name' : ['User 6', 'User 7']})
>>> df2.to_sql('users', con=engine, if_exists='append')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 1'), (1, 'User 2'), (2, 'User 3'),
 (0, 'User 4'), (1, 'User 5'), (0, 'User 6'),
 (1, 'User 7')]
```

Overwrite the table with just ``df2``.

```
>>> df2.to_sql('users', con=engine, if_exists='replace',
... index_label='id')
>>> engine.execute("SELECT * FROM users").fetchall()
[(0, 'User 6'), (1, 'User 7')]
```

Specify the dtype (especially useful for integers with missing values). Notice that while pandas is forced to store the data as floating point, the database supports nullable integers. When fetching the data with Python, we get back integer scalars.

```
>>> df = pd.DataFrame({"A": [1, None, 2]})
>>> df
 A
0 1.0
1 NaN
2 2.0
```

```
>>> from sqlalchemy.types import Integer
>>> df.to_sql('integers', con=engine, index=False,
... dtype={"A": Integer()})
>>> engine.execute("SELECT * FROM integers").fetchall()
[(1,), (None,), (2,)]
```

```

to_xarray(self)
 Return an xarray object from the pandas object.

Returns

xarray.DataArray or xarray.Dataset
 Data in the pandas structure converted to Dataset if the object is
 a DataFrame, or a DataArray if the object is a Series.

See Also

DataFrame.to_hdf : Write DataFrame to an HDF5 file.
DataFrame.to_parquet : Write a DataFrame to the binary parquet format.

Notes

See the `xarray docs <https://xarray.pydata.org/en/stable/>`__

Examples

>>> df = pd.DataFrame([('falcon', 'bird', 389.0, 2),
... ('parrot', 'bird', 24.0, 2),
... ('lion', 'mammal', 80.5, 4),
... ('monkey', 'mammal', np.nan, 4)],
... columns=['name', 'class', 'max_speed',
... 'num_legs'])
>>> df
 name class max_speed num_legs
0 falcon bird 389.0 2
1 parrot bird 24.0 2
2 lion mammal 80.5 4
3 monkey mammal NaN 4

>>> df.to_xarray()
<xarray.Dataset>
Dimensions: (index: 4)
Coordinates:
 * index (index) int64 0 1 2 3
Data variables:
 name (index) object 'falcon' 'parrot' 'lion' 'monkey'
 class (index) object 'bird' 'bird' 'mammal' 'mammal'
 max_speed (index) float64 389.0 24.0 80.5 nan
 num_legs (index) int64 2 2 4 4

>>> df['max_speed'].to_xarray()
<xarray.DataArray 'max_speed' (index: 4)>
array([389. , 24. , 80.5, nan])
Coordinates:
 * index (index) int64 0 1 2 3

>>> dates = pd.to_datetime(['2018-01-01', '2018-01-01',
... '2018-01-02', '2018-01-02'])
>>> df_multiindex = pd.DataFrame({'date': dates,
... 'animal': ['falcon', 'parrot',
... 'falcon', 'parrot'],
... 'speed': [350, 18, 361, 15]})
>>> df_multiindex = df_multiindex.set_index(['date', 'animal'])

>>> df_multiindex
 speed
date animal
2018-01-01 falcon 350
 parrot 18
2018-01-02 falcon 361
 parrot 15

>>> df_multiindex.to_xarray()
<xarray.Dataset>
Dimensions: (animal: 2, date: 2)
Coordinates:
 * date (date) datetime64[ns] 2018-01-01 2018-01-02
 * animal (animal) object 'falcon' 'parrot'
Data variables:
 speed (date, animal) int64 350 18 361 15

truncate(self: 'FrameOrSeries', before=None, after=None, axis=None, copy:
'bool_t' = True) -> 'FrameOrSeries'
 Truncate a Series or DataFrame before and after some index value.

 This is a useful shorthand for boolean indexing based on index
 values above or below certain thresholds.

Parameters

before : date, str, int
 Truncate all rows before this index value.

```

after : date, str, int  
Truncate all rows after this index value.  
axis : {0 or 'index', 1 or 'columns'}, optional  
Axis to truncate. Truncates the index (rows) by default.  
copy : bool, default is True,  
Return a copy of the truncated section.

#### Returns

-----

type of caller

The truncated Series or DataFrame.

#### See Also

-----

DataFrame.loc : Select a subset of a DataFrame by label.

DataFrame.iloc : Select a subset of a DataFrame by position.

#### Notes

-----

If the index being truncated contains only datetime values,  
'before' and 'after' may be specified as strings instead of  
Timestamps.

#### Examples

-----

```
>>> df = pd.DataFrame({'A': ['a', 'b', 'c', 'd', 'e'],
... 'B': ['f', 'g', 'h', 'i', 'j'],
... 'C': ['k', 'l', 'm', 'n', 'o']},
... index=[1, 2, 3, 4, 5])
```

```
>>> df
 A B C
1 a f k
2 b g l
3 c h m
4 d i n
5 e j o
```

```
>>> df.truncate(before=2, after=4)
```

```
 A B C
2 b g l
3 c h m
4 d i n
```

The columns of a DataFrame can be truncated.

```
>>> df.truncate(before="A", after="B", axis="columns")
```

```
 A B
1 a f
2 b g
3 c h
4 d i
5 e j
```

For Series, only rows can be truncated.

```
>>> df['A'].truncate(before=2, after=4)
```

```
2 b
3 c
4 d
Name: A, dtype: object
```

The index values in ``truncate`` can be datetimes or string  
dates.

```
>>> dates = pd.date_range('2016-01-01', '2016-02-01', freq='s')
>>> df = pd.DataFrame(index=dates, data={'A': 1})
>>> df.tail()
```

```
 A
2016-01-31 23:59:56 1
2016-01-31 23:59:57 1
2016-01-31 23:59:58 1
2016-01-31 23:59:59 1
2016-02-01 00:00:00 1
```

```
>>> df.truncate(before=pd.Timestamp('2016-01-05'),
... after=pd.Timestamp('2016-01-10')).tail()
```

```
 A
2016-01-09 23:59:56 1
2016-01-09 23:59:57 1
2016-01-09 23:59:58 1
2016-01-09 23:59:59 1
2016-01-10 00:00:00 1
```

Because the index is a DatetimeIndex containing only dates, we can  
specify 'before' and 'after' as strings. They will be coerced to  
Timestamps before truncation.



```

>>> df.truncate('2016-01-05', '2016-01-10').tail()
 A
2016-01-09 23:59:56 1
2016-01-09 23:59:57 1
2016-01-09 23:59:58 1
2016-01-09 23:59:59 1
2016-01-10 00:00:00 1

Note that ``truncate`` assumes a 0 value for any unspecified time
component (midnight). This differs from partial string slicing, which
returns any partially matching dates.

>>> df.loc['2016-01-05':'2016-01-10', :].tail()
 A
2016-01-10 23:59:55 1
2016-01-10 23:59:56 1
2016-01-10 23:59:57 1
2016-01-10 23:59:58 1
2016-01-10 23:59:59 1

tshift(self: 'FrameOrSeries', periods: 'int' = 1, freq=None, axis: 'Axis' = 0) ->
'FrameOrSeries'
 Shift the time index, using the index's frequency if available.

 .. deprecated:: 1.1.0
 Use ``shift`` instead.

 Parameters

 periods : int
 Number of periods to move, can be positive or negative.
 freq : DateOffset, timedelta, or str, default None
 Increment to use from the tseries module
 or time rule expressed as a string (e.g. 'EOM').
 axis : {0 or 'index', 1 or 'columns', None}, default 0
 Corresponds to the axis that contains the Index.

 Returns

 shifted : Series/DataFrame

 Notes

 If freq is not specified then tries to use the freq or inferred_freq
 attributes of the index. If neither of those attributes exist, a
 ValueError is thrown

tz_convert(self: 'FrameOrSeries', tz, axis=0, level=None, copy: 'bool_t' = True)
-> 'FrameOrSeries'
 Convert tz-aware axis to target time zone.

 Parameters

 tz : str or tzinfo object
 axis : the axis to convert
 level : int, str, default None
 If axis is a MultiIndex, convert a specific level. Otherwise
 must be None.
 copy : bool, default True
 Also make a copy of the underlying data.

 Returns

 {klass}
 Object with time zone converted axis.

 Raises

 TypeError
 If the axis is tz-naive.

tz_localize(self: 'FrameOrSeries', tz, axis=0, level=None, copy: 'bool_t' = True,
ambiguous='raise', nonexistent: 'str' = 'raise') -> 'FrameOrSeries'
 Localize tz-naive index of a Series or DataFrame to target time zone.

 This operation localizes the Index. To localize the values in a
 timezone-naive Series, use :meth:`Series.dt.tz_localize`.

 Parameters

 tz : str or tzinfo
 axis : the axis to localize
 level : int, str, default None
 If axis ia a MultiIndex, localize a specific level. Otherwise
 must be None.
 copy : bool, default True
 Also make a copy of the underlying data.

```

`ambiguous` : 'infer', bool-ndarray, 'NaT', default 'raise'

When clocks moved backward due to DST, ambiguous times may arise. For example in Central European Time (UTC+01), when going from 03:00 DST to 02:00 non-DST, 02:30:00 local time occurs both at 00:30:00 UTC and at 01:30:00 UTC. In such a situation, the `'ambiguous'` parameter dictates how ambiguous times should be handled.

- 'infer' will attempt to infer fall dst-transition hours based on order
- bool-ndarray where True signifies a DST time, False designates a non-DST time (note that this flag is only applicable for ambiguous times)
- 'NaT' will return NaT where there are ambiguous times
- 'raise' will raise an `AmbiguousTimeError` if there are ambiguous times.

`nonexistent` : str, default 'raise'

A nonexistent time does not exist in a particular timezone where clocks moved forward due to DST. Valid values are:

- 'shift\_forward' will shift the nonexistent time forward to the closest existing time
- 'shift\_backward' will shift the nonexistent time backward to the closest existing time
- 'NaT' will return NaT where there are nonexistent times
- `timedelta` objects will shift nonexistent times by the `timedelta`
- 'raise' will raise an `NonExistentTimeError` if there are nonexistent times.

Returns

-----

Series or DataFrame

Same type as the input.

Raises

-----

`TypeError`

If the `TimeSeries` is tz-aware and `tz` is not `None`.

Examples

-----

Localize local times:

```
>>> s = pd.Series([1],
... index=pd.DatetimeIndex(['2018-09-15 01:30:00']))
>>> s.tz_localize('CET')
2018-09-15 01:30:00+02:00 1
dtype: int64
```

Be careful with DST changes. When there is sequential data, pandas can infer the DST time:

```
>>> s = pd.Series(range(7),
... index=pd.DatetimeIndex(['2018-10-28 01:30:00',
... '2018-10-28 02:00:00',
... '2018-10-28 02:30:00',
... '2018-10-28 02:00:00',
... '2018-10-28 02:30:00',
... '2018-10-28 03:00:00',
... '2018-10-28 03:30:00']))
>>> s.tz_localize('CET', ambiguous='infer')
2018-10-28 01:30:00+02:00 0
2018-10-28 02:00:00+02:00 1
2018-10-28 02:30:00+02:00 2
2018-10-28 02:00:00+01:00 3
2018-10-28 02:30:00+01:00 4
2018-10-28 03:00:00+01:00 5
2018-10-28 03:30:00+01:00 6
dtype: int64
```

In some cases, inferring the DST is impossible. In such cases, you can pass an ndarray to the `ambiguous` parameter to set the DST explicitly

```
>>> s = pd.Series(range(3),
... index=pd.DatetimeIndex(['2018-10-28 01:20:00',
... '2018-10-28 02:36:00',
... '2018-10-28 03:46:00']))
>>> s.tz_localize('CET', ambiguous=np.array([True, True, False]))
2018-10-28 01:20:00+02:00 0
2018-10-28 02:36:00+02:00 1
2018-10-28 03:46:00+01:00 2
dtype: int64
```

If the DST transition causes nonexistent times, you can shift these dates forward or backward with a `timedelta` object or `'shift_forward'` or `'shift_backward'`.

```
>>> s = pd.Series(range(2),
... index=pd.DatetimeIndex(['2015-03-29 02:30:00',
... '2015-03-29 03:30:00']))
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_forward')
2015-03-29 03:00:00+02:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent='shift_backward')
2015-03-29 01:59:59.999999999+01:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
>>> s.tz_localize('Europe/Warsaw', nonexistent=pd.Timedelta('1H'))
2015-03-29 03:30:00+02:00 0
2015-03-29 03:30:00+02:00 1
dtype: int64
```

`xs(self, key, axis=0, level=None, drop_level: 'bool_t' = True)`  
Return cross-section from the Series/DataFrame.

This method takes a `key` argument to select data at a particular level of a MultiIndex.

#### Parameters

-----  
**key** : label or tuple of label  
Label contained in the index, or partially in a MultiIndex.  
**axis** : {0 or 'index', 1 or 'columns'}, default 0  
Axis to retrieve cross-section on.  
**level** : object, defaults to first n levels (n=1 or len(key))  
In case of a key partially contained in a MultiIndex, indicate which levels are used. Levels can be referred by label or position.  
**drop\_level** : bool, default True  
If False, returns object with same levels as self.

#### Returns

-----  
Series or DataFrame  
Cross-section from the original Series or DataFrame corresponding to the selected index levels.

#### See Also

-----  
**DataFrame.loc** : Access a group of rows and columns by label(s) or a boolean array.  
**DataFrame.iloc** : Purely integer-location based indexing for selection by position.

#### Notes

-----  
`xs` can not be used to set values.

MultiIndex Slicers is a generic way to get/set values on any level or levels.

It is a superset of `xs` functionality, see :ref:`MultiIndex Slicers <advanced.mi\_slicers>`.

#### Examples

```
>>> d = {'num_legs': [4, 4, 2, 2],
... 'num_wings': [0, 0, 2, 2],
... 'class': ['mammal', 'mammal', 'mammal', 'bird'],
... 'animal': ['cat', 'dog', 'bat', 'penguin'],
... 'locomotion': ['walks', 'walks', 'flies', 'walks']}
>>> df = pd.DataFrame(data=d)
>>> df = df.set_index(['class', 'animal', 'locomotion'])
>>> df
```

			num_legs	num_wings
class	animal	locomotion		
mammal	cat	walks	4	0
	dog	walks	4	0
	bat	flies	2	2
bird	penguin	walks	2	2

Get values at specified index

```
>>> df.xs('mammal')
 num_legs num_wings
animal locomotion
cat walks 4 0
dog walks 4 0
bat flies 2 2
```

Get values at several indexes

```
>>> df.xs(['mammal', 'dog'])
 num_legs num_wings
locomotion
```

```
walks 4 0

Get values at specified index and level

>>> df.xs('cat', level=1)
 num_legs num_wings
class locomotion
mammal walks 4 0

Get values at several indexes and levels

>>> df.xs(('bird', 'walks'),
... level=[0, 'locomotion'])
 num_legs num_wings
animal
penguin 2 2

Get values at specified column and axis

>>> df.xs('num_wings', axis=1)
class animal locomotion
mammal cat walks 0
 dog walks 0
 bat flies 2
bird penguin walks 2
Name: num_wings, dtype: int64
```

-----  
Data descriptors inherited from pandas.core.generic.NDFrame:

attrs

Dictionary of global attributes of this dataset.

.. warning::

attrs is experimental and may change without warning.

See Also

-----  
DataFrame.flags : Global flags applying to this object.

dtypes

Return the dtypes in the DataFrame.

This returns a Series with the data type of each column.  
The result's index is the original DataFrame's columns. Columns  
with mixed types are stored with the ``object`` dtype. See  
:ref:`the User Guide <basics.dtypes>` for more.

Returns

-----  
pandas.Series

The data type of each column.

Examples

```
>>> df = pd.DataFrame({'float': [1.0],
... 'int': [1],
... 'datetime': [pd.Timestamp('20180310')],
... 'string': ['foo']})
>>> df.dtypes
float float64
int int64
datetime datetime64[ns]
string object
dtype: object
```

empty

Indicator whether DataFrame is empty.

True if DataFrame is entirely empty (no items), meaning any of the  
axes are of length 0.

Returns

-----  
bool

If DataFrame is empty, return True, if not return False.

See Also

-----  
Series.dropna : Return series without null values.  
DataFrame.dropna : Return DataFrame with labels on given axis omitted  
where (all or any) data are missing.

Notes

-----  
If DataFrame contains only NaNs, it is still not considered empty. See

the example below.

#### Examples

-----

An example of an actual empty DataFrame. Notice the index is empty:

```
>>> df_empty = pd.DataFrame({'A' : []})
>>> df_empty
Empty DataFrame
Columns: [A]
Index: []
>>> df_empty.empty
True
```

If we only have NaNs in our DataFrame, it is not considered empty! We will need to drop the NaNs to make the DataFrame empty:

```
>>> df = pd.DataFrame({'A' : [np.nan]})
>>> df
 A
0 NaN
>>> df.empty
False
>>> df.dropna().empty
True
```

#### flags

Get the properties associated with this pandas object.

The available flags are

\* :attr:`Flags.allows\_duplicate\_labels`

#### See Also

-----

Flags : Flags that apply to pandas objects.

DataFrame.attrs : Global metadata applying to this dataset.

#### Notes

-----

"Flags" differ from "metadata". Flags reflect properties of the pandas object (the Series or DataFrame). Metadata refer to properties of the dataset, and should be stored in :attr:`DataFrame.attrs`.

#### Examples

-----

```
>>> df = pd.DataFrame({"A": [1, 2]})
>>> df.flags
<Flags(allows_duplicate_labels=True)>
```

Flags can be get or set using ``.``

```
>>> df.flags.allows_duplicate_labels
True
>>> df.flags.allows_duplicate_labels = False
```

Or by slicing with a key

```
>>> df.flags["allows_duplicate_labels"]
False
>>> df.flags["allows_duplicate_labels"] = True
```

#### ndim

Return an int representing the number of axes / array dimensions.

Return 1 if Series. Otherwise return 2 if DataFrame.

#### See Also

-----

ndarray.ndim : Number of array dimensions.

#### Examples

-----

```
>>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
>>> s.ndim
1
```

```
>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
>>> df.ndim
2
```

#### size

Return an int representing the number of elements in this object.

Return the number of rows if Series. Otherwise return the number of rows times number of columns if DataFrame.

#### See Also

-----

`ndarray.size` : Number of elements in the array.

#### Examples

-----

```
>>> s = pd.Series({'a': 1, 'b': 2, 'c': 3})
```

```
>>> s.size
```

```
3
```

```
>>> df = pd.DataFrame({'col1': [1, 2], 'col2': [3, 4]})
```

```
>>> df.size
```

```
4
```

-----  
Data and other attributes inherited from `pandas.core.generic.NDFrame`:

```
__array_priority__ = 1000
```

-----  
Methods inherited from `pandas.core.base.PandasObject`:

```
__sizeof__(self) -> 'int'
```

Generates the total memory usage for an object that returns  
either a value or Series of values

-----  
Methods inherited from `pandas.core.accessor.DirNamesMixin`:

```
__dir__(self) -> 'list[str]'
```

Provide method name lookup and completion.

#### Notes

-----

Only provide 'public' methods.

-----  
Data descriptors inherited from `pandas.core.accessor.DirNamesMixin`:

```
__dict__
```

dictionary for instance variables (if defined)

```
__weakref__
```

list of weak references to the object (if defined)

-----  
Data descriptors inherited from `pandas.core.indexing.IndexingMixin`:

`at`

Access a single value for a row/column label pair.

Similar to `loc`, in that both provide label-based lookups. Use  
`at` if you only need to get or set a single value in a DataFrame  
or Series.

#### Raises

-----

`KeyError`

If 'label' does not exist in DataFrame.

#### See Also

-----

`DataFrame.iat` : Access a single value for a row/column pair by integer  
position.

`DataFrame.loc` : Access a group of rows and columns by label(s).

`Series.at` : Access a single value using a label.

#### Examples

-----

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
... index=[4, 5, 6], columns=['A', 'B', 'C'])
```

```
>>> df
```

```
 A B C
4 0 2 3
5 0 4 1
6 10 20 30
```

Get value at specified row/column pair

```
>>> df.at[4, 'B']
```

```
2
```

Set value at specified row/column pair

```
>>> df.at[4, 'B'] = 10
```

```
>>> df.at[4, 'B']
```

```
10
```

Get value within a Series

```
>>> df.loc[5].at['B']
4
```

iat

Access a single value for a row/column pair by integer position.

Similar to ``iloc``, in that both provide integer-based lookups. Use ``iat`` if you only need to get or set a single value in a DataFrame or Series.

Raises

-----

IndexError

When integer position is out of bounds.

See Also

-----

DataFrame.at : Access a single value for a row/column label pair.

DataFrame.loc : Access a group of rows and columns by label(s).

DataFrame.iloc : Access a group of rows and columns by integer position(s).

Examples

-----

```
>>> df = pd.DataFrame([[0, 2, 3], [0, 4, 1], [10, 20, 30]],
... columns=['A', 'B', 'C'])
```

```
>>> df
```

```
 A B C
0 0 2 3
1 0 4 1
2 10 20 30
```

Get value at specified row/column pair

```
>>> df.iat[1, 2]
1
```

Set value at specified row/column pair

```
>>> df.iat[1, 2] = 10
>>> df.iat[1, 2]
10
```

Get value within a series

```
>>> df.loc[0].iat[1]
2
```

iloc

Purely integer-location based indexing for selection by position.

``.iloc[]`` is primarily integer position based (from ``0`` to ``length-1`` of the axis), but may also be used with a boolean array.

Allowed inputs are:

- An integer, e.g. ``5``.
- A list or array of integers, e.g. ``[4, 3, 0]``.
- A slice object with ints, e.g. ``1:7``.
- A boolean array.
- A ``callable`` function with one argument (the calling Series or DataFrame) and that returns valid output for indexing (one of the above). This is useful in method chains, when you don't have a reference to the calling object, but would like to base your selection on some value.

``.iloc`` will raise ``IndexError`` if a requested indexer is out-of-bounds, except \*slice\* indexers which allow out-of-bounds indexing (this conforms with python/numpy \*slice\* semantics).

See more at :ref:`Selection by Position <indexing.integer>`.

See Also

-----

DataFrame.iat : Fast integer location scalar accessor.

DataFrame.loc : Purely label-location based indexer for selection by label.

Series.iloc : Purely integer-location based indexing for selection by position.

Examples

-----

```
>>> mydict = [{'a': 1, 'b': 2, 'c': 3, 'd': 4},
... {'a': 100, 'b': 200, 'c': 300, 'd': 400},
... {'a': 1000, 'b': 2000, 'c': 3000, 'd': 4000}]
>>> df = pd.DataFrame(mydict)
```

```
>>> df
 a b c d
0 1 2 3 4
1 100 200 300 400
2 1000 2000 3000 4000
```

**\*\*Indexing just the rows\*\***

With a scalar integer.

```
>>> type(df.iloc[0])
<class 'pandas.core.series.Series'>
>>> df.iloc[0]
a 1
b 2
c 3
d 4
Name: 0, dtype: int64
```

With a list of integers.

```
>>> df.iloc[[0]]
 a b c d
0 1 2 3 4
>>> type(df.iloc[[0]])
<class 'pandas.core.frame.DataFrame'>
```

```
>>> df.iloc[[0, 1]]
 a b c d
0 1 2 3 4
1 100 200 300 400
```

With a `slice` object.

```
>>> df.iloc[:3]
 a b c d
0 1 2 3 4
1 100 200 300 400
2 1000 2000 3000 4000
```

With a boolean mask the same length as the index.

```
>>> df.iloc[[True, False, True]]
 a b c d
0 1 2 3 4
2 1000 2000 3000 4000
```

With a callable, useful in method chains. The `x` passed to the ``lambda`` is the DataFrame being sliced. This selects the rows whose index label even.

```
>>> df.iloc[lambda x: x.index % 2 == 0]
 a b c d
0 1 2 3 4
2 1000 2000 3000 4000
```

**\*\*Indexing both axes\*\***

You can mix the indexer types for the index and columns. Use ``:``` to select the entire axis.

With scalar integers.

```
>>> df.iloc[0, 1]
2
```

With lists of integers.

```
>>> df.iloc[[0, 2], [1, 3]]
 b d
0 2 4
2 2000 4000
```

With `slice` objects.

```
>>> df.iloc[1:3, 0:3]
 a b c
1 100 200 300
2 1000 2000 3000
```

With a boolean array whose length matches the columns.

```
>>> df.iloc[:, [True, False, True, False]]
 a c
0 1 3
1 100 300
2 1000 3000
```



With a callable function that expects the Series or DataFrame.

```
>>> df.iloc[:, lambda df: [0, 2]]
 a c
0 1 3
1 100 300
2 1000 3000
```

loc

Access a group of rows and columns by label(s) or a boolean array.

`df.loc[]` is primarily label based, but may also be used with a boolean array.

Allowed inputs are:

- A single label, e.g. `df.loc[5]` or `df.loc['a']`, (note that `df.loc[5]` is interpreted as a \*label\* of the index, and **never** as an integer position along the index).
  - A list or array of labels, e.g. `df.loc['a', 'b', 'c']`.
  - A slice object with labels, e.g. `df.loc['a':'f']`.
- .. warning:: Note that contrary to usual python slices, **both** the start and the stop are included
- A boolean array of the same length as the axis being sliced, e.g. `df.loc[[True, False, True]]`.
  - An alignable boolean Series. The index of the key will be aligned before masking.
  - An alignable Index. The Index of the returned selection will be the input.
  - A `callable` function with one argument (the calling Series or DataFrame) and that returns valid output for indexing (one of the above)

See more at :ref:`Selection by Label <indexing.label>`.

Raises

-----

KeyError

If any items are not found.

IndexingError

If an indexed key is passed and its index is unalignable to the frame

index.

See Also

-----

DataFrame.at : Access a single value for a row/column label pair.

DataFrame.iloc : Access group of rows and columns by integer position(s).

DataFrame.xs : Returns a cross-section (row(s) or column(s)) from the Series/DataFrame.

Series.loc : Access group of values using labels.

Examples

-----

**\*\*Getting values\*\***

```
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
... index=['cobra', 'viper', 'sidewinder'],
... columns=['max_speed', 'shield'])
>>> df
```

```
 max_speed shield
cobra 1 2
viper 4 5
sidewinder 7 8
```

Single label. Note this returns the row as a Series.

```
>>> df.loc['viper']
max_speed 4
shield 5
Name: viper, dtype: int64
```

List of labels. Note using `df.loc[]` returns a DataFrame.

```
>>> df.loc[['viper', 'sidewinder']]
 max_speed shield
viper 4 5
sidewinder 7 8
```

Single label for row and column

```
>>> df.loc['cobra', 'shield']
2
```

Slice with labels for row and single label for column. As mentioned above, note that both the start and stop of the slice are included.

```
>>> df.loc['cobra':'viper', 'max_speed']
cobra 1
viper 4
Name: max_speed, dtype: int64
```

Boolean list with the same length as the row axis

```
>>> df.loc[[False, False, True]]
 max_speed shield
sidewinder 7 8
```

Alignable boolean Series:

```
>>> df.loc[pd.Series([False, True, False],
... index=['viper', 'sidewinder', 'cobra'])]
 max_speed shield
sidewinder 7 8
```

Index (same behavior as ``df.reindex``)

```
>>> df.loc[pd.Index(["cobra", "viper"], name="foo")]
 max_speed shield
foo
cobra 1 2
viper 4 5
```

Conditional that returns a boolean Series

```
>>> df.loc[df['shield'] > 6]
 max_speed shield
sidewinder 7 8
```

Conditional that returns a boolean Series with column labels specified

```
>>> df.loc[df['shield'] > 6, ['max_speed']]
 max_speed
sidewinder 7
```

Callable that returns a boolean Series

```
>>> df.loc[lambda df: df['shield'] == 8]
 max_speed shield
sidewinder 7 8
```

**\*\*Setting values\*\***

Set value for all items matching the list of labels

```
>>> df.loc[['viper', 'sidewinder'], ['shield']] = 50
>>> df
 max_speed shield
cobra 1 2
viper 4 50
sidewinder 7 50
```

Set value for an entire row

```
>>> df.loc['cobra'] = 10
>>> df
 max_speed shield
cobra 10 10
viper 4 50
sidewinder 7 50
```

Set value for an entire column

```
>>> df.loc[:, 'max_speed'] = 30
>>> df
 max_speed shield
cobra 30 10
viper 30 50
sidewinder 30 50
```

Set value for rows matching callable condition

```
>>> df.loc[df['shield'] > 35] = 0
>>> df
 max_speed shield
cobra 30 10
viper 0 0
sidewinder 0 0
```

**\*\*Getting values on a DataFrame with an index that has integer labels\*\***

Another example using integers for the index

```
>>> df = pd.DataFrame([[1, 2], [4, 5], [7, 8]],
```

```
... index=[7, 8, 9], columns=['max_speed', 'shield'])
>>> df
 max_speed shield
7 1 2
8 4 5
9 7 8
```

Slice with integer labels for rows. As mentioned above, note that both the start and stop of the slice are included.

```
>>> df.loc[7:9]
 max_speed shield
7 1 2
8 4 5
9 7 8
```

**\*\*Getting values with a MultiIndex\*\***

A number of examples using a DataFrame with a MultiIndex

```
>>> tuples = [
... ('cobra', 'mark i'), ('cobra', 'mark ii'),
... ('sidewinder', 'mark i'), ('sidewinder', 'mark ii'),
... ('viper', 'mark ii'), ('viper', 'mark iii')
...]
>>> index = pd.MultiIndex.from_tuples(tuples)
>>> values = [[12, 2], [0, 4], [10, 20],
... [1, 4], [7, 1], [16, 36]]
>>> df = pd.DataFrame(values, columns=['max_speed', 'shield'], index=index)
>>> df
 max_speed shield
cobra mark i 12 2
 mark ii 0 4
sidewinder mark i 10 20
 mark ii 1 4
viper mark ii 7 1
 mark iii 16 36
```

Single label. Note this returns a DataFrame with a single index.

```
>>> df.loc['cobra']
 max_speed shield
mark i 12 2
mark ii 0 4
```

Single index tuple. Note this returns a Series.

```
>>> df.loc[('cobra', 'mark ii')]
max_speed 0
shield 4
Name: (cobra, mark ii), dtype: int64
```

Single label for row and column. Similar to passing in a tuple, this returns a Series.

```
>>> df.loc['cobra', 'mark i']
max_speed 12
shield 2
Name: (cobra, mark i), dtype: int64
```

Single tuple. Note using ``[]`` returns a DataFrame.

```
>>> df.loc[('cobra', 'mark ii')]
 max_speed shield
cobra mark ii 0 4
```

Single tuple for the index with a single label for the column

```
>>> df.loc[('cobra', 'mark i'), 'shield']
2
```

Slice from index tuple to single label

```
>>> df.loc[('cobra', 'mark i'):'viper']
 max_speed shield
cobra mark i 12 2
 mark ii 0 4
sidewinder mark i 10 20
 mark ii 1 4
viper mark ii 7 1
 mark iii 16 36
```

Slice from index tuple to index tuple

```
>>> df.loc[('cobra', 'mark i'):(viper', 'mark ii')]
 max_speed shield
cobra mark i 12 2
```

	mark ii	0	4
sidewinder	mark i	10	20
	mark ii	1	4
viper	mark ii	7	1

-----  
Methods inherited from pandas.core.arraylike.OpsMixin:

```

__add__(self, other)
__and__(self, other)
__eq__(self, other)
 Return self==value.
__floordiv__(self, other)
__ge__(self, other)
 Return self>=value.
__gt__(self, other)
 Return self>value.
__le__(self, other)
 Return self<=value.
__lt__(self, other)
 Return self<value.
__mod__(self, other)
__mul__(self, other)
__ne__(self, other)
 Return self!=value.
__or__(self, other)
__pow__(self, other)
__radd__(self, other)
__rand__(self, other)
__rfloordiv__(self, other)
__rmod__(self, other)
__rmul__(self, other)
__ror__(self, other)
__rpow__(self, other)
__rsub__(self, other)
__rtruediv__(self, other)
__rxor__(self, other)
__sub__(self, other)
__truediv__(self, other)
__xor__(self, other)

```

-----  
Data and other attributes inherited from pandas.core.arraylike.OpsMixin:

```
__hash__ = None
```

### Try it yourself!

Make a list of the shape of all of the tables on the syllabus Achievements page.

```

achievements_url =
'https://rhodyprog4ds.github.io/BrownFall21/syllabus/achievements.html'

```

```
[(14, 3), (15, 5), (15, 15), (15, 6)]
```

This solution uses a list comprehension which allows us to compress a loop. It's equivalent to the following with a for loop

```
[(14, 3), (15, 5), (15, 15), (15, 6)]
```

## 5.6. Lambdas and Dictionaries for switching

What if we want to print out the first column for the dataFrame if it has more than 3 columns and the whole thing if it has 3 or less columns?

Two ways of writing a function

```
typical way
def first_col_f(d):
 return d[d.columns[0]]

first_col_l = lambda d: d[d.columns[0]]

first_col_f(help_df) == first_col_l(help_df)
```

```
0 True
1 True
2 True
3 True
4 True
5 True
Name: Day, dtype: bool
```

We can put functions in dictionaries, or even define a lambda right in the dictionary.

```
df_display = {True: lambda d: d[d.columns[0]],
 False: lambda d: d}

for df in pd.read_html(achievements_url):
 _, n_cols = df.shape
 print(df_display[n_cols>3](df))
```

### Note

Python does have [ternary operators](#) but the dictionary is a more common way to achieve this and goal and more common patterns are better for readability

```

 Unnamed: 0_level_0
 week
0 1
1 2
2 3
3 4
4 5
5 6
6 7
7 8
8 9
9 10
10 11
11 12
12 13
13 14
 topics \
 Unnamed: 1_level_1
 [admin, python review]
 Loading data, Python review
 Exploratory Data Analysis
 Data Cleaning
 Databases, Merging DataFrames
 Modeling, Naive Bayes, classification performa...
 decision trees, cross validation
 Regression
 Clustering
 SVM, parameter tuning
 KNN, Model comparison
 Text Analysis
 Images Analysis
 Deep Learning

```

```

 skills
 Unnamed: 2_level_1
 process
0 [access, prepare, summarize]
1 [summarize, visualize]
2 [prepare, summarize, visualize]
3 [access, construct, summarize]
4 [classification, evaluate]
5 [classification, evaluate]
6 [regression, evaluate]
7 [clustering, evaluate]
8 [optimize, tools]
9 [compare, tools]
10 [unstructured]
11 [unstructured, tools]
12 [tools, compare]
13

```

```

0 python
1 process
2 access
3 construct
4 summarize
5 visualize
6 prepare
7 classification
8 regression
9 clustering
10 evaluate
11 optimize
12 compare
13 unstructured
14 workflow

```

Name: (Unnamed: 0\_level\_0, keyword), dtype: object

```

0 python
1 process
2 access
3 construct
4 summarize
5 visualize
6 prepare
7 classification
8 regression
9 clustering
10 evaluate
11 optimize
12 compare
13 unstructured
14 workflow

```

Name: (Unnamed: 0\_level\_0, keyword), dtype: object

```

0 python
1 process
2 access
3 construct
4 summarize
5 visualize
6 prepare
7 classification
8 regression
9 clustering
10 evaluate
11 optimize
12 compare
13 unstructured
14 workflow

```

Name: (Unnamed: 0\_level\_0, keyword), dtype: object

### Try it Yourself!

What does the `_` do?

## 5.7. Questions after class

### Note

add a question with a pull request; earn 1-2 ram tokens for submitting a question with the answer (with sources)

## 5.8. More Practice

1. What `type` is the shape of a `pandas.DataFrame`?
2. use a list comprehension to create a list that you could use as column names for data that consists of `N` measurements. Set `N=5` for now, but you suspect that the number might change.
1. create a list of items with different types, then Create a dictionary with the types as keys using a dictionary comprehension. Dictionary comprehensions are similar to list comprehensions, in their form.

```
File "/tmp/ipykernel_1694/1953963119.py", line 1
 about_prof_brown_dict = {type(fact):fact} for fact in about_prof_brown_list
 ^
SyntaxError: invalid syntax
```

1. Create a lambda function to print return the first 2 rows of a data frame

# 1. Portfolio Setup, Data Science, and Python

Due: 2020-09-12

## 1.1. Objective & Evaluation

This assignment is an opportunity to earn level 2 achievements for the `process` and `python` and confirm that you have all of your tools setup, including your portfolio.

## 1.2. To Do

### Important

If you have trouble, check the GitHub FAQ on the left before e-mailing

```
```{warning}
If you have trouble with the (*)d steps, don't worry, we can help work around these later. To help
us out, document the errors as bugs on your repository.
```
```

Your task is to:

1. Install required software from the Tools & Resource page
2. Create your portfolio, by [accepting the assignment](#)
3. Learn about your portfolio from the README file on your repository.
4. edit `_config.yml` to set your name as author and change the logo if you wish
5. Fill in `about/index.md` with information about yourself(not evaluated, but useful) and your own definition of data science (graded for **level 1 process**)

### Note

If you get stuck on any of this after accepting the assignment and creating a repository, you can create an issue on your repository, describing what you're stuck on and tag us: `@rhodypro4dg/fall21instructors`

To do this click Issues at the top, the green "New Issue" button and then type away.

6. (\*) Install some additional python packages with: `pip install pip install -r requirements.txt` (this is a python operation, so use anaconda prompt on Windows, if the pip version doesn't work, try it with conda: `conda install --file requirements.txt`) form inside the portfolio folder
7. (\*) Configure precommit to help keep your repo clean with `pre-commit install`. If this step doesn't work, see the portfolio README under "Using your Jupyter Book Portfolio"
8. Add a Jupyter notebook called `grading.ipynb` to the `about` folder and write a function that computes a grade for this course, with the following docstring. Include:
  - a Markdown cell with a heading
  - your function called `compute_grade`
  - three calls to your function that verify it returns the correct value for different number of badges that produce at three different letter grades.
  - a basic function that uses conditionals in python will earn **level 1 python**
  - to earn **level 2 python** use pythonic code to write a loop that tests your function's correctness, by iterating over a list or dictionary. Remember you will have many chances to earn level 2 achievement in python
9. Add the line `- file: about/grading` in your `_toc.yml` file.

### ! Important

remember to add, commit, and push your changes so we can see them

```
'''
 Computes a grade for CSC/DSP310 from numbers of achievements at each level

 Parameters:

 num_level1 : int
 number of level 1 achievements earned
 num_level2 : int
 number of level 2 achievements earned
 num_level3 : int
 number of level 3 achievements earned

 Returns:

 letter_grade : string
 letter grade with possible modifier (+/-)
'''
```

Here are some sample tests you could run to confirm that your function works correctly:

```
assert compute_grade(15,15,15) == 'A'
assert compute_grade(15,15,13) == 'A-'
assert compute_grade(15,14,14) == 'B-'
assert compute_grade(14,14,14) == 'C-'
assert compute_grade(4,3,1) == 'D'
assert compute_grade(15,15,6) == 'B+'
```

### ⚠ Warning

your function can have a different name than `compute_grade`, but make sure it's your function name, with those parameter values in your tests.

### i Note

when the value of the expression after `assert` is `True`, it will look like nothing happened. `assert` is used for testing

## 1.3. Submission Instructions

Create a Jupyter Notebook with your function in your portfolio folder commit and push the changes.

In your browser, view the `gh-pages` branch to see your compiled submission, as `portfolio.pdf` or by viewing your website.

There will be a pull request on your repository that is made by GitHub classroom, [request a review](#) from @rhodypro4dg/fall21instructors.

## 2. Assignment 2: Practicing Python and Accessing Data

due : 2020-09-21



## 2.1. Objective & Evaluation

This assignment is an opportunity to earn level 1 and 2 achievements in **python** and **access** and begin working toward level 1 for **summarize**. You can also earn level 1 for **process**.

In this assignment, you'll practice/ review python skills by manipulating datasets and extracting

| Task                                                                                   | Skills (max level) |
|----------------------------------------------------------------------------------------|--------------------|
| identify possible uses for data in a data science pipeline                             | [process (1)]      |
| load data from one file format                                                         | [ access (1)]      |
| load data from at least two of (.csv, .tsv, .dat, database, .json)                     | [access (2)]       |
| compare the data formats                                                               | [ access (2)]      |
| complete the assignment in python                                                      | python (1)]        |
| use python data types (eg dictionaries) to prepare information about datasets          | [python (2)]       |
| use informative variable names, pythonic iteration, and other common PEP 8 conventions | [python (2)]       |
| display DataFrame properties                                                           | [summarize (1)]    |

Table 2.1 practice python by manipulating data files, load datasets of different types

First, [accept the assignment](#) . It contains a notebook with some template structure (and will set you up for grading).

## 2.2. Find Datasets

Find 3 datasets of interest to you that are provided in at least two different file formats. Choose datasets that are not too big, so that they do not take more than a few second to load. At least one dataset, must have non numerical (eg string or boolean) data in at least 1 column.

In your notebook, create a markdown cell for each notebook that includes:

- heading of the dataset's name
- a link to where someone can learn about the dataset
- a 1-2 sentence summary of what the dataset contains and why it was collected
- 1-2 questions you would like to answer with that dataset.

## 2.3. Store them for loading

Create a list of dictionaries in **datasets.py**, so that there is one dictionary for each dataset with the url, a name, and what function should be used to load the data into a **pandas.DataFrame**.

## 2.4. Make a dataset about your datasets

Import the list fomr the **datasets** module you created in the step above. Then terate over the list of dictionaries, and:

1. save it to a local csv using the short name you provided for the dataset as the file name, without writing the index column to the file.
2. record attributes about the dataset as in the table below in a list of lists:
3. Use that to create a DataFrame with the following columns:

### Tip

Urls are strings. The **string** class in python has a lot of helpful methods for manipulating strings, like [split](#).

|               |                                              |
|---------------|----------------------------------------------|
| name          | a short name for the dataset                 |
| source        | a url to where you found the data            |
| num_rows      | number of rows in the dataset                |
| num_columns   | number of columns in the dataset             |
| num_numerical | number of numerical variables in the dataset |

Table 2.2 Meta Data Description of the DataFrame to build

## 2.5. Manipulate your datasets

For one dataset that includes nonnumerical data:

- display the heading and the last 4 rows
- make and display a new data frame with only the numerical columns (select these programmatically)

For any other dataset:

- display the heading and the first three rows
- display the datatype for each column
- Are there any variables where pandas may have read in the data as a datatype that's not what you expect (eg a numerical column mistaken for strings)? If so, investigate and try to figure out why.

For the third dataset:

- display the first 3 odd rows (eg 1,3,5) of the data for two columns of your choice

## 2.6. Exploring data files

For each dataset, in a separate section of your notebook titled **When things go wrong**:

- try reading in data with the wrong `read_` function and make notes about what happens.
- was the format that the data was provided in a good format? why or why not?
- try to read in the `.csv` file that's included in the template repository (), use the error messages you get to try to fix the file manually (any text editor, including jupyter can edit a `.csv`), making notes about what changes you made in a markdown cell.

## 2.7. Thinking ahead

### ⚠ Warning

This section is not required, but is intended to help you get started thinking about ideas for your portfolio. If you complete it, we'll give your feedback to help shape your ideas to get to level 3 achievements. If you want to focus only on level 2 at this moment in time, feel free to skip this part.

1. When might you prefer one datatype over another?
2. How does PEP 8 standard code help you be collaborative?
3. Learn about [Datasheets for Datasets](#) eg this [google scholar result](#) How could something like this impact your work as a datascientist?

## Portfolio Dates and Key Facts

```

NameError Traceback (most recent call last)
/tmp/ipykernel_1719/311973014.py in <module>
 5
 6 rubric_df = pd.concat([rubric_df,
----> 7 assignment_dummies,
 8 portfolio_dummies],axis=1)
 9

NameError: name 'assignment_dummies' is not defined

```

This section of the site has a set of portfolio prompts and this page has instructions for portfolio submissions.

Starting in week 3 it is recommended that you spend some time each week working on items for your portfolio, that way when it's time to submit you only have a little bit to add before submission. The portfolio is your only chance to earn Level 3 achievements, however, if you have not earned a level 2 for any of the skills in a given check, you could earn level 2 then instead. The prompts provide a starting point, but remember that to earn achievements, you'll be evaluated by the rubric. You can see the full rubric for all portfolios in the [syllabus](#). Your portfolio is also an opportunity to be creative, explore things, and answer your own questions that we haven't answered in class to dig deeper on the topics we're covering. Use the feedback you get on assignments to inspire your portfolio.

Each submission should include an introduction and a number of 'chapters'. The grade will be based on both that you demonstrate skills through your chapters that are inspired by the prompts and that your summary demonstrates that you *know* you learned the skills. See the [formatting tips](#) for advice on how to structure files.

On each chapter(for a file) of your portfolio, you should identify which skills by their keyword, you are applying.

You can view a (fake) example [in this repository](#) as a [pdf](#) or as a [rendered website](#)

## Current

### Check 1

The first portfolio check will be due October 15

```

NameError Traceback (most recent call last)
/tmp/ipykernel_1719/2819809228.py in <module>
----> 1 portfolio_df[portfolio_df['P1']==1]

NameError: name 'portfolio_df' is not defined

```

## Upcoming Checks

Check 2: November 12 Check 3: December 5 Check 4: December 20

## Formatting Tips

### ⚠ Warning

This is all based on you having accepted the portfolio assignment on github and having a cloned copy of the template. If you are not enrolled or the initial assignment has not been issued, you can view [the template on GitHub](#)

Your portfolio is a [jupyter book](#). This means a few things:

- it uses [myst markdown](#)
- it will run and compile Jupyter notebooks

This page will cover a few basic tips.

## Organization

The summary of for the **part** or whole submission, should match the skills to the chapters. Which prompt you're addressing is not important, the prompts are a *starting point* not the end goal of your portfolio.

## Data Files

Also note that for your portfolio to build, you will have to:

- include the data files in the repository and use a relative path OR
- load via url

using a full local path(eg that starts with `///file:`) **will not work** and will render your portfolio unreadable.

## Structure of plain markdown

Use a heading like this:

```
Heading of page
Heading 2
Heading 3
```

in the file and it will appear in the sidebar.

You can also make text *italic* or **bold** with either `*asterics*` or `__underscores__` with `_one` for *italic* or `**two` for **bold\*\*** in either case

## File Naming

It is best practice to name files without spaces. Each **chapter** or file should have a descriptive file name (`with_no_spaces`) and descriptive title for it.

## Syncing markdown and ipynb files

If you have the precommit hook working, git will call a script and convert your notebook files from the ipynb format (which is json like) to Myst Markdown, which is more plain text with some header information. The markdown format works better with version control, largely because it doesn't contain the outputs.

If you don't get the precommit hook working, but you do get jupyter installed, you can set each file to sync.

## Adding annotations with formatting or margin notes

You can either install [jupyter](#) and convert locally or upload /push a notebook to your repository and let GitHub convert. Then edit the .md file with a [text editor](#) of your choice. You can run by uploading if you don't have jupyter installed, or locally if you have installed jupyter or jupyterbook.

In your .md file use backticks to mark [special content blocks](#)

```
```{note}
Here is a note!
```
```

```
```{warning}
Here is a warning!
```
```

```
```{tip}
Here is a tip!
```
```

```
```{margin}
Here is a margin note!
```
```

For a complete list of options, see [the sphinx-book-theme documentation](#).

## Links

Markdown syntax for links

```
[text to show](path/or/url)
```

## Configurations

Things like the menus and links at the top are controlled as [settings](#), in `_config.yml`. The following are some things that you might change in your configuration file.

### Show errors and continue

To show errors and continue running the rest, add the following to your configuration file:

```
Execution settings
execute:
 allow_errors : true
```

## Using additional packages

You'll have to add any additional packages you use (beyond pandas and seaborn) to the `requirements.txt` file in your portfolio.

## FAQ

This section will grow as questions are asked and new content is introduced to the site. You can submit questions:

- via e-mail to Dr. Brown (brownsarahm) or Beibhinn (beibhinn)
- via Prismia.chat during class
- by creating an [issue](#)

## Syllabus FAQ

How much does assignment x, class participation, or a portfolio check weigh in my grade?



Can I submit this assignment late if ...?



## Git and GitHub FAQ

The content I added to my portfolio isn't in the pdf



My command line says I cannot use a password



My .ipynb file isn't showing in the staging area or didn't push



My portfolio won't compile



Help! I accidentally merged the Feedback Pull Request before my assignment was graded

## Common Debugging Issues

Key Error

<bound method

## Glossary

### DataFrame

a data structure provided by pandas for tabular data in python.

### **git**

a version control tool; it's a fully open source and always free tool, that can be hosted by anyone or used without a host, locally only.

### **GitHub**

a hosting service for git repositories

### **interpreter**

the translator from human readable python code to something the computer can run. An interpreted language means you can work with python interactively

### **kernel**

in the jupyter environment, [the kernel](#) is a language specific computational engine

### **PEP 8**

[Python Enhancement Proposal 8](#), the Style Guide for Python Code

### **repository**

a project folder with tracking information in it in the form of a .git file

### **TraceBack**

an error message in python that traces back from the line of code that had caused the exception back through all of the functions that called other functions to reach that line. This is sometimes call tracing back through the stack

## General Tips and Resources

This section is for materials that are not specific to this course, but are likely useful. They are not generally required readings or installs, but are options or advice I provide frequently.


### on email

- [how to e-mail professors](#)

## How to Study in this class

This is a programming intensive course and it's about data science. This course is designed to help you learn how to program for data science and in the process build general skills in programming and using data to understand the world. Learning two things at once is more complex. In this page, I break down how I expect learning to work for this class.

Remember the goal is to avoid this:



## Why this way?

Learning to program requires iterative practice. It does not require memorizing all of the specific commands, but instead learning the basic patterns.

Using reference materials frequently is a built in part of programming, most languages have built in help as a part of the language for this reason.

A new book that might be of interest if you find programming classes hard is [the Programmers Brain](#). As of 2020-09-07, it is available for free by clicking on chapters at that linked table of contents section.

### **i** Where are your help tools?

In Python and Jupyter notebooks, what help tools do you have?

## Learning in class

### **i** Important

My goal is to use class time so that you can be successful with *minimal frustration* while working outside of class time.

Programming requires both practical skills and abstract concepts. During class time, we will cover the practical aspects and introduce the basic concepts. You will get to see the basic practical details and real examples of debugging during class sessions. Learning to debug something you've never encountered before and setting up your programming environment, for example, are *high frustration* activities, when you're learning, because you don't know what you don't know. On the other hand, diving deeper into options and more complex applications of what you have already seen in class, while challenging, is something I'm confident that you can all be successful at with minimal frustration once you've seen basic ideas in class. My goal is that you can repeat the patterns and processes we use in class outside of class to complete assignments, while acknowledging that you will definitely have to look things up and read documentation outside of class.

Each class will open with some time to review what was covered in the last session before adding new material.

To get the most out of class sessions, you should have a laptop with you. During class you should be following along with Dr. Brown, typing and running the same code. You'll answer questions on Prismia chat, when you do so, you should try running necessary code to answer those questions. If you encounter errors, share them via prismia chat so that we can see and help you.

## After class

After class, you should practice with the concepts introduced.

This means reviewing the notes: both yours from class and the annotated notes posted to the course website.

When you review the notes, you should be adding comments on tricky aspects of the code and narrative text between code blocks in markdown cells. While you review your notes and the annotated course notes, you should also read the documentation for new modules, libraries, or functions introduced that day.

In the annotated notes, there will often be extra questions or ideas on how to extend and practice the concepts. Try these out.

If you find anything hard to understand or unclear, write it down to bring to class the next day.

There will be additional drills posted to Prismia chat for you to try.

## Assignments

In assignments, you will be asked to practice with specific concepts at an intermediate level. Assignments will apply the concepts from class with minimal extensions. You will probably need to use help functions and read documentation to complete assignments, but mostly to look up things you saw in class and make minor variations. Most of what you need for assignments will be in the class notes, which is another reason to read them after class.

## Portfolios


In portfolios, your goal is to extend and apply the concepts taught in class and practiced in assignments to solve more realistic problems. You may also reflect on your learning in order to demonstrate deep understanding. These will require significant reading beyond what we cover in class.

## Getting Help with Programming

### Asking Questions

JULIA EVANS  
@b0rk

### asking good questions



One of my favorite resources that describes how to ask good questions is [this blog post](#) by Julia Evans, a developer who writes comics about the things she learns in the course of her work and publisher of [wizard zines](#).

## Describing what you have so far

Stackoverflow is a common place for programmers to post and answer questions. As such, they have written a good [guide on creating a minimal, reproducible example](#).

Creating a minimal reproducible example may even help you debug your own code, but if it does not, it will definitely make it easier for another person to understand what you have, what your goal is, and what's working.

## Understanding Errors

### Note

A fun version of this is [rubber duck debugging](#)



Error messages from the compiler are not always straight forward.

The [TraceBack](#) can be a really long list of errors that seem like they are not even from your code. It will trace back to all of the places that the error occurred. It is often about how you called the functions from a library, but the compiler cannot tell that.

To understand what the traceback is, how to read one, and common examples, see [this post on Real Python](#).

One thing to try, is [friendly.traceback](#) a python package that is designed to make that error message text more clear and help you figure out what to do next.

## References on Python

- [Course Text](#)

## Cheatsheet

Patterns and examples of how to use common tips in class

### Axes

First build a small dataset that's just enough to display

```
data = [[1,0],[5,4],[1,4]]
df = pd.DataFrame(data = data,
 columns = ['A','B'])

df
```

```
 A B
0 1 0
1 5 4
2 1 4
```

```
df.sum(axis=0)
```

```
A 7
B 8
dtype: int64
```

```
df.sum(axis=1)
```

```
0 1
1 9
2 5
dtype: int64
```

```
df.apply(sum,axis=0)
```

```
A 7
B 8
dtype: int64
```

```
df.apply(sum,axis=1)
```

```
0 1
1 9
2 5
dtype: int64
```

```
df['A'][1]
```

5

```
df.iloc[0][1]
```

0

## Data Sources

### Best for loading directly into a notebook

- [Tidy Tuesday](#) inside the folder for each year there is a README file with list of the datasets. These are .csv files
- [Json Datasets](#)
- [National Center for Education Statistics Digest 2019](#) These data tables are available for download as excel and visible on the page.
- Lots of wikipedia pages have tables in them.

### Requires some more work

- [Stackoverflow Developer Survey](#). This data comes with readme info all packaged together in a .zip. You'll need to unzip it first.
- [Google Dataset Search](#)
- [Kaggle](#) most Kaggle datasets will require you to download and unzip them first and then you can copy them into your repo folder.
- [UCI Data Repository](#).

## Databases

- [SQLite Databases](#)

If you have others please share by creating a pull request or issue on this repo (from the GitHub logo at the top right, [suggest edit](#)).

---

By Professor Sarah M Brown

© Copyright 2021.